Writing R Extensions

Version 3.3.1 (2016-06-21)

This manual is for R, version 3.3.1 (2016-06-21).

Copyright © 1999–2016 R Core Team

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions, except that this permission notice may be stated in a translation approved by the R Core Team.

Table of Contents

A	cknow	$\operatorname{ledgements}$	1
1	Crea	ting R packages	2
		kage structure	
	1.1.1	The DESCRIPTION file	
	1.1.2	Licensing	
	1.1.3	Package Dependencies	
	1.1	L.3.1 Suggested packages	
	1.1.4	The INDEX file	
	1.1.5	Package subdirectories	
	1.1.6	Data in packages	15
	1.1.7	Non-R scripts in packages	
	1.1.8	Specifying URLs	17
	1.2 Cor	afigure and cleanup	17
	1.2.1	Using Makevars	20
	1.2	2.1.1 OpenMP support	23
	1.2	2.1.2 Using pthreads	
	1.2	2.1.3 Compiling in sub-directories	
	1.2.2	Configure example	
	1.2.3	Using F95 code	
	1.2.4	Using C++11 code	
	1.3 $Ch\epsilon$	ecking and building packages	
	1.3.1	Checking packages	
	1.3.2	Building package tarballs	
	1.3.3	Building binary packages	
		ting package vignettes	
	1.4.1	Encodings and vignettes	
	1.4.2	Non-Sweave vignettes	
		kage namespaces	
	1.5.1	Specifying imports and exports	
	1.5.2	Registering S3 methods	
	1.5.3	Load hooks	
	1.5.4	useDynLib	
	1.5.5	An example	
	1.5.6	Namespaces with S4 classes and methods	
		ting portable packages	
	1.6.1	PDF size	
	1.6.2	Check timing	
	1.6.3	Encoding issues	
	1.6.4 $1.6.5$	Binary distribution	
		gnostic messages	
		gnostic messages	
	1.8.1	C-level messages	
	1.8.2	R messages	
	1.8.3 1.9 CIT	Preparing translations	
		ackage types	
	1.1U 1 a	ULIMIGU UY DUU	1

	1.10.1 Frontend	. 57
	1.11 Services	. 57
_		
2	Writing R documentation files	
	2.1 Rd format	
	2.1.1 Documenting functions	
	2.1.2 Documenting data sets	
	2.1.3 Documenting S4 classes and methods	
	2.1.4 Documenting packages	
	2.2 Sectioning	
	2.3 Marking text	
	2.4 Lists and tables 2.5 Cross-references.	
	2.6 Mathematics	
	2.7 Figures	
	2.8 Insertions	
	2.9 Indices	
	2.10 Platform-specific documentation	
	2.11 Conditional text	
	2.12 Dynamic pages	
	2.13 User-defined macros	
	2.14 Encoding	74
	2.15 Processing documentation files	74
	2.16 Editing Rd files	75
3	Tidying and profiling R code	76
	3.1 Tidying R code	76
	3.2 Profiling R code for speed	
	3.3 Profiling R code for memory use	78
	3.3.1 Memory statistics from Rprof	78
	3.3.2 Tracking memory allocations	79
	3.3.3 Tracing copies of an object	
	3.4 Profiling compiled code	
	3.4.1 Linux	
	3.4.1.1 sprof	
	3.4.1.2 oprofile and operf	
	3.4.2 Solaris	
	3.4.3 OS X	83
4		0.4
4	8 8 8 8 8	
	4.1 Browsing	
	4.2 Debugging R code	
	4.3 Checking memory access	
	4.3.1 Using gctorture	
	4.3.2 Using valgrind	
	4.3.3 Using the Address Sanitizer	
	4.3.3.1 Using the Leak Sanitizer	
	4.3.4 Using the Undefined Behaviour Sanitizer	
	4.3.5 Other analyses with 'clang'	
	4.3.6 Using 'Dr. Memory'	
	4.4 Debugging compiled code	
	4.4.1 Finding entry points in dynamically loaded code	

	4.4.2 Inspecting R objects when debugging	97
5	System and foreign language interfaces	99
	5.1 Operating system access	99
	5.2 Interface functions .C and .Fortran	99
	5.3 dyn.load and dyn.unload	101
	5.4 Registering native routines	102
	5.4.1 Speed considerations	
	5.4.2 Linking to native routines in other packages	
	5.5 Creating shared objects	
	5.6 Interfacing C++ code	
	5.7 Fortran I/O	
	5.8 Linking to other packages	
	5.8.1 Unix-alikes	
	5.8.2 Windows	
	5.9 Handling R objects in C	
	5.9.1 Handling the effects of garbage collection	
	5.9.2 Allocating storage	
	5.9.3 Details of R types	
	5.9.4 Attributes	
	5.9.5 Classes	
	5.9.6 Handling lists	
	5.9.7 Handling character data	
	5.9.9 Some convenience functions	
	5.9.9.1 Semi-internal convenience functions	
	5.9.10 Named objects and copying	
	5.10 Interface functions .Call and .External	
	5.10.1 Calling .Call	
	5.10.2 Calling .External	
	5.10.3 Missing and special values	
	5.11 Evaluating R expressions from C	
	5.11.1 Zero-finding	
	5.11.2 Calculating numerical derivatives	
	5.12 Parsing R code from C	
	5.12.1 Accessing source references	
	5.13 External pointers and weak references	
	5.13.1 An example	131
	5.14 Vector accessor functions	132
	5.15 Character encoding issues	
6	The R API: entry points for C code	13/
U	6.1 Memory allocation	
	6.1.1 Transient storage allocation	
	6.1.2 User-controlled memory	
	6.2 Error handling	
	6.2.1 Error handling from FORTRAN	
	6.3 Random number generation	
	6.4 Missing and IEEE special values	
	6.5 Printing	
	6.5.1 Printing from FORTRAN	
	6.6 Calling C from FORTRAN and vice versa	
	6.7 Numerical analysis subroutines	

	6.7.1 Distribution functions	139
	6.7.2 Mathematical functions	
	6.7.3 Numerical Utilities	
	6.7.4 Mathematical constants	142
6	Optimization	143
6	Integration	144
6	0 Utility functions	
6	1 Re-encoding	147
6	2 Allowing interrupts	147
6	3 Platform and version information	
6	4 Inlining C functions	148
6	5 Controlling visibility	
6	6 Using these functions in your own C code	
6	7 Organization of header files	150
7	Generic functions and methods Adding new generics	
_		·
8	Linking GUIs and other front-ends to R	
8 8	_	
	Embedding R under Unix-alikes	
	Embedding R under Unix-alikes	
	Embedding R under Unix-alikes. 8.1.1 Compiling against the R library. 8.1.2 Setting R callbacks. 8.1.3 Registering symbols.	
	Embedding R under Unix-alikes. 8.1.1 Compiling against the R library. 8.1.2 Setting R callbacks. 8.1.3 Registering symbols. 8.1.4 Meshing event loops.	
8	Embedding R under Unix-alikes 8.1.1 Compiling against the R library 8.1.2 Setting R callbacks 8.1.3 Registering symbols 8.1.4 Meshing event loops 8.1.5 Threading issues	
8	Embedding R under Unix-alikes 8.1.1 Compiling against the R library 8.1.2 Setting R callbacks 8.1.3 Registering symbols 8.1.4 Meshing event loops 8.1.5 Threading issues Embedding R under Windows	
8	Embedding R under Unix-alikes 8.1.1 Compiling against the R library 8.1.2 Setting R callbacks 8.1.3 Registering symbols 8.1.4 Meshing event loops 8.1.5 Threading issues Embedding R under Windows 8.2.1 Using (D)COM	
8	Embedding R under Unix-alikes 8.1.1 Compiling against the R library 8.1.2 Setting R callbacks 8.1.3 Registering symbols 8.1.4 Meshing event loops 8.1.5 Threading issues Embedding R under Windows 8.2.1 Using (D)COM 8.2.2 Calling R.dll directly	
8	Embedding R under Unix-alikes 8.1.1 Compiling against the R library 8.1.2 Setting R callbacks 8.1.3 Registering symbols 8.1.4 Meshing event loops 8.1.5 Threading issues Embedding R under Windows 8.2.1 Using (D)COM	
8	Embedding R under Unix-alikes 8.1.1 Compiling against the R library 8.1.2 Setting R callbacks 8.1.3 Registering symbols 8.1.4 Meshing event loops 8.1.5 Threading issues Embedding R under Windows 8.2.1 Using (D)COM 8.2.2 Calling R.dll directly	

Acknowledgements

The contributions to early versions of this manual by Saikat DebRoy (who wrote the first draft of a guide to using .Call and .External) and Adrian Trapletti (who provided information on the C++ interface) are gratefully acknowledged.

1 Creating R packages

Packages provide a mechanism for loading optional code, data and documentation as needed. The R distribution itself includes about 30 packages.

In the following, we assume that you know the library() command, including its lib.loc argument, and we also assume basic knowledge of the R CMD INSTALL utility. Otherwise, please look at R's help pages on

?library
?INSTALL

before reading on.

For packages which contain code to be compiled, a computing environment including a number of tools is assumed; the "R Installation and Administration" manual describes what is needed for each OS.

Once a source package is created, it must be installed by the command R CMD INSTALL. See Section "Add-on-packages" in R Installation and Administration.

Other types of extensions are supported (but rare): See Section 1.10 [Package types], page 57.

Some notes on terminology complete this introduction. These will help with the reading of this manual, and also in describing concepts accurately when asking for help.

A package is a directory of files which extend R, a source package (the master files of a package), or a tarball containing the files of a source package, or an installed package, the result of running R CMD INSTALL on a source package. On some platforms (notably OS X and Windows) there are also binary packages, a zip file or tarball containing the files of an installed package which can be unpacked rather than installing from sources.

A package is **not**¹ a *library*. The latter is used in two senses in R documentation.

- A directory into which packages are installed, e.g. /usr/lib/R/library: in that sense it is sometimes referred to as a *library directory* or *library tree* (since the library is a directory which contains packages as directories, which themselves contain directories).
- That used by the operating system, as a shared, dynamic or static library or (especially on Windows) a DLL, where the second L stands for 'library'. Installed packages may contain compiled code in what is known on Unix-alikes as a shared object and on Windows as a DLL. The concept of a shared library (dynamic library on OS X) as a collection of compiled code to which a package might link is also used, especially for R itself on some platforms. On most platforms these concepts are interchangeable (shared objects and DLLs can both be loaded into the R process and be linked against), but OS X distinguishes between shared objects (extension .so) and dynamic libraries (extension .dylib).

There are a number of well-defined operations on source packages.

- The most common is *installation* which takes a source package and installs it in a library using R CMD INSTALL or install.packages.
- Source packages can be *built*. This involves taking a source directory and creating a tarball ready for distribution, including cleaning it up and creating PDF documentation from any *vignettes* it may contain. Source packages (and most often tarballs) can be *checked*, when a test installation is done and tested (including running its examples); also, the contents of the package are tested in various ways for consistency and portability.
- Compilation is not a correct term for a package. Installing a source package which contains C, C++ or Fortran code will involve compiling that code. There is also the possibility of 'byte' compiling the R code in a package (using the facilities of package compiler): already

¹ although this is a persistent mis-usage. It seems to stem from S, whose analogues of R's packages were officially known as *library sections* and later as *chapters*, but almost always referred to as *libraries*.

base and recommended packages are normally byte-compiled and this can be specified for other packages. So *compiling* a package may come to mean byte-compiling its R code.

• It used to be unambiguous to talk about *loading* an installed package using library(), but since the advent of package namespaces this has been less clear: people now often talk about *loading* the package's namespace and then *attaching* the package so it becomes visible on the search path. Function library performs both steps, but a package's namespace can be loaded without the package being attached (for example by calls like splines::ns).

The concept of *lazy loading* of code or data is mentioned at several points. This is part of the installation, always selected for R code but optional for data. When used the R objects of the package are created at installation time and stored in a database in the R directory of the installed package, being loaded into the session at first use. This makes the R session start up faster and use less (virtual) memory. (For technical details, see Section "Lazy loading" in R Internals.)

CRAN is a network of WWW sites holding the R distributions and contributed code, especially R packages. Users of R are encouraged to join in the collaborative project and to submit their own packages to CRAN: current instructions are linked from https://CRAN.R-project.org/banner.shtml#submitting.

1.1 Package structure

The sources of an R package consists of a subdirectory containing a files DESCRIPTION and NAMESPACE, and the subdirectories R, data, demo, exec, inst, man, po, src, tests, tools and vignettes (some of which can be missing, but which should not be empty). The package subdirectory may also contain files INDEX, configure, cleanup, LICENSE, LICENCE and NEWS. Other files such as INSTALL (for non-standard installation instructions), README/README.md², or ChangeLog will be ignored by R, but may be useful to end users. The utility R CMD build may add files in a build directory (but this should not be used for other purposes).

Except where specifically mentioned,³ packages should not contain Unix-style 'hidden' files/directories (that is, those whose name starts with a dot).

The DESCRIPTION and INDEX files are described in the subsections below. The NAMESPACE file is described in the section on Section 1.5 [Package namespaces], page 39.

The optional files configure and cleanup are (Bourne) shell scripts which are, respectively, executed before and (if option --clean was given) after installation on Unix-alikes, see Section 1.2 [Configure and cleanup], page 17. The analogues on Windows are configure.win and cleanup.win.

For the conventions for files NEWS and ChangeLog in the GNU project see https://www.gnu.org/prep/standards/standards.html#Documentation.

The package subdirectory should be given the same name as the package. Because some file systems (e.g., those on Windows and by default on OS X) are not case-sensitive, to maintain portability it is strongly recommended that case distinctions not be used to distinguish different packages. For example, if you have a package named foo, do not also create a package named foo.

To ensure that file names are valid across file systems and supported operating systems, the ASCII control characters as well as the characters '"', '*', ':', '/', '<', '>', '?', '\', and '|' are not allowed in file names. In addition, files with names 'con', 'prn', 'aux', 'clock\$', 'nul', 'com1' to

² This seems to be commonly used for a file in 'markdown' format. Be aware that most users of R will not know that, nor know how to view such a file: platforms such as OS X and Windows do not have a default viewer set in their file associations. The CRAN package web pages render such files in HTML: the converter used expects the file to be encoded in UTF-8.

 $^{^3}$ currently, top-level files .Rbuildignore and .Rinstignore, and vignettes/.install_extras.

'com9', and 'lpt1' to 'lpt9' after conversion to lower case and stripping possible "extensions" (e.g., 'lpt5.foo.bar'), are disallowed. Also, file names in the same directory must not differ only by case (see the previous paragraph). In addition, the basenames of '.Rd' files may be used in URLs and so must be ASCII and not contain %. For maximal portability filenames should only contain only ASCII characters not excluded already (that is A-Za-z0-9._!#\$%&+,;=@^(){}'[] — we exclude space as many utilities do not accept spaces in file paths): non-English alphabetic characters cannot be guaranteed to be supported in all locales. It would be good practice to avoid the shell metacharacters (){}'[]\$\sigma': \sigma' is also used as part of '8.3' filenames on Windows. In addition, packages are normally distributed as tarballs, and these have a limit on path lengths: for maximal portability 100 bytes.

A source package if possible should not contain binary executable files: they are not portable, and a security risk if they are of the appropriate architecture. R CMD check will warn about them⁴ unless they are listed (one filepath per line) in a file BinaryFiles at the top level of the package. Note that CRAN will not accept submissions containing binary files even if they are listed.

The R function package.skeleton can help to create the structure for a new package: see its help page for details.

1.1.1 The DESCRIPTION file

The DESCRIPTION file contains basic information about the package in the following format:

```
Package: pkgname
Version: 0.5-1
Date: 2015-01-01
Title: My First Collection of Functions
Authors@R: c(person("Joe", "Developer", role = c("aut", "cre"),
     email = "Joe.Developer@some.domain.net"),
      person("Pat", "Developer", role = "aut"),
person("A.", "User", role = "ctb",
     email = "A.User@whereever.net"))
Author: Joe Developer [aut, cre],
  Pat Developer [aut],
  A. User [ctb]
Maintainer: Joe Developer <Joe.Developer@some.domain.net>
Depends: R (>= 3.1.0), nlme
Suggests: MASS
Description: A (one paragraph) description of what
  the package does and why it may be useful.
License: GPL (>= 2)
URL: https://www.r-project.org, http://www.another.url
BugReports: https://pkgname.bugtracker.url
```

The format is that of a version of a 'Debian Control File' (see the help for 'read.dcf' and https://www.debian.org/doc/debian-policy/ch-controlfields.html: R does not require encoding in UTF-8 and does not support comments starting with '#'). Fields start with an ASCII name immediately followed by a colon: the value starts after the colon and a space. Continuation lines (for example, for descriptions longer than one line) start with a space or tab. Field names are case-sensitive: all those used by R are capitalized.

For maximal portability, the DESCRIPTION file should be written entirely in ASCII — if this is not possible it must contain an 'Encoding' field (see below).

Several optional fields take *logical values*: these can be specified as 'yes', 'true', 'no' or 'false': capitalized values are also accepted.

The 'Package', 'Version', 'License', 'Description', 'Title', 'Author', and 'Maintainer' fields are mandatory, all other fields are optional. Fields 'Author' and 'Maintainer' can be

⁴ false positives are possible, but only a handful have been seen so far.

auto-generated from 'Authors@R', and may be omitted if the latter is provided: however if they are not ASCII we recommend that they are provided.

The mandatory 'Package' field gives the name of the package. This should contain only (ASCII) letters, numbers and dot, have at least two characters and start with a letter and not end in a dot. If it needs explaining, this should be done in the 'Description' field (and not the 'Title' field).

The mandatory 'Version' field gives the version of the package. This is a sequence of at least two (and usually three) non-negative integers separated by single '.' or '-' characters. The canonical form is as shown in the example, and a version such as '0.01' or '0.01.0' will be handled as if it were '0.1-0'. It is **not** a decimal number, so for example 0.9 < 0.75 since 9 < 75.

The mandatory 'License' field is discussed in the next subsection.

The mandatory 'Title' field should give a *short* description of the package. Some package listings may truncate the title to 65 characters. It should use *title case* (that is, use capitals for the principal words: tools::toTitleCase can help you with this), not use any markup, not have any continuation lines, and not end in a period (unless part of . . .). Do not repeat the package name: it is often used prefixed by the name. Refer to other packages and external software in single quotes, and to book titles (and similar) in double quotes.

The mandatory 'Description' field should give a *comprehensive* description of what the package does. One can use several (complete) sentences, but only one paragraph. It should be intelligible to all the intended readership (e.g. for a CRAN package to all CRAN users). It is good practice not to start with the package name, 'This package' or similar. As with the 'Title' field, double quotes should be used for quotations (including titles of books and articles), and single quotes for non-English usage, including names of other packages and external software. This field should also be used for explaining the package name if necessary. URLs should be enclosed in angle brackets, e.g. 'https://www.r-project.org': see also Section 1.1.8 [Specifying URLs], page 17.

The mandatory 'Author' field describes who wrote the package. It is a plain text field intended for human readers, but not for automatic processing (such as extracting the email addresses of all listed contributors: for that use 'Authors@R'). Note that all significant contributors must be included: if you wrote an R wrapper for the work of others included in the src directory, you are not the sole (and maybe not even the main) author.

The mandatory 'Maintainer' field should give a *single* name followed by a *valid* (RFC 2822) email address in angle brackets. It should not end in a period or comma. This field is what is reported by the maintainer function and used by bug.report. For a CRAN package it should be a *person*, not a mailing list and not a corporate entity: do ensure that it is valid and will remain valid for the lifetime of the package.

Note that the *display name* (the part before the address in angle brackets) should be enclosed in double quotes if it contains non-alphanumeric characters such as comma or period. (The current standard, RFC 5322, allows periods but RFC 2822 did not.)

Both 'Author' and 'Maintainer' fields can be omitted if a suitable 'Authors@R' field is given. This field can be used to provide a refined and machine-readable description of the package "authors" (in particular specifying their precise roles), via suitable R code. It should create an object of class "person', by either a call to person or a series of calls (one per "author") concatenated by c()): see the example DESCRIPTION file above. The roles can include '"aut"' (author) for full authors, '"cre"' (creator) for the package maintainer, and '"ctb"' (contributor) for other contributors, '"cph"' (copyright holder), among others. See 'person for more information. Note that no role is assumed by default. Auto-generated package citation information

takes advantage of this specification. The 'Author' and 'Maintainer' fields are auto-generated from it if needed when building⁵ or installing.

An optional 'Copyright' field can be used where the copyright holder(s) are not the authors. If necessary, this can refer to an installed file: the convention is to use file inst/COPYRIGHTS.

The optional 'Date' field gives the *release date* of the current version of the package. It is strongly recommended⁶ to use the 'yyyy-mm-dd' format conforming to the ISO 8601 standard.

The 'Depends', 'Imports', 'Suggests', 'Enhances', 'LinkingTo' and 'Additional_repositories' fields are discussed in a later subsection.

Dependencies external to the R system should be listed in the 'SystemRequirements' field, possibly amplified in a separate README file.

The 'URL' field may give a list of URLs separated by commas or whitespace, for example the homepage of the author or a page where additional material describing the software can be found. These URLs are converted to active hyperlinks in CRAN package listings. See Section 1.1.8 [Specifying URLs], page 17.

The 'BugReports' field may contain a single URL to which bug reports about the package should be submitted. This URL will be used by bug.report instead of sending an email to the maintainer.

Base and recommended packages (i.e., packages contained in the R source distribution or available from CRAN and recommended to be included in every binary distribution of R) have a 'Priority' field with value 'base' or 'recommended', respectively. These priorities must not be used by other packages.

A 'Collate' field can be used for controlling the collation order for the R code files in a package when these are processed for package installation. The default is to collate according to the 'C' locale. If present, the collate specification must list *all* R code files in the package (taking possible OS-specific subdirectories into account, see Section 1.1.5 [Package subdirectories], page 12) as a whitespace separated list of file paths relative to the R subdirectory. Paths containing white space or quotes need to be quoted. An OS-specific collation field ('Collate.unix' or 'Collate.windows') will be used in preference to 'Collate'.

The 'LazyData' logical field controls whether the R datasets use lazy-loading. A 'LazyLoad' field was used in versions prior to 2.14.0, but now is ignored.

The 'KeepSource' logical field controls if the package code is sourced using keep.source = TRUE or FALSE: it might be needed exceptionally for a package designed to always be used with keep.source = TRUE.

The 'ByteCompile' logical field controls if the package code is to be byte-compiled on installation: the default is currently not to, so this may be useful for a package known to benefit particularly from byte-compilation (which can take quite a long time and increases the installed size of the package). It is used for the recommended packages, as they are byte-compiled when R is installed and for consistency should be byte-compiled when updated. This can be overridden by installing with flag --no-byte-compile.

The 'ZipData' logical field was used to control whether the automatic Windows build would zip up the data directory or not prior to R 2.13.0: it is now ignored.

The 'Biarch' logical field is used on Windows to select the INSTALL option --force-biarch for this package. (Introduced in R 3.0.0.)

The 'BuildVignettes' logical field can be set to a false value to stop R CMD build from attempting to build the vignettes, as well as preventing R CMD check from testing this. This

 $^{^{5}}$ at least if this is done in a locale which matches the package encoding.

 $^{^{6}\,}$ and required by CRAN, so checked by R CMD check --as-cran.

 $^{^{7}}$ But it is checked for Open Source packages by R CMD check --as-cran.

should only be used exceptionally, for example if the PDFs include large figures which are not part of the package sources (and hence only in packages which do not have an Open Source license).

The 'VignetteBuilder' field names (in a comma-separated list) packages that provide an engine for building vignettes. These may include the current package, or ones listed in 'Depends', 'Suggests' or 'Imports'. The utils package is always implicitly appended. See Section 1.4.2 [Non-Sweave vignettes], page 38, for details.

If the DESCRIPTION file is not entirely in ASCII it should contain an 'Encoding' field specifying an encoding. This is used as the encoding of the DESCRIPTION file itself and of the R and NAMESPACE files, and as the default encoding of .Rd files. The examples are assumed to be in this encoding when running R CMD check, and it is used for the encoding of the CITATION file. Only encoding names latin1, latin2 and UTF-8 are known to be portable. (Do not specify an encoding unless one is actually needed: doing so makes the package less portable. If a package has a specified encoding, you should run R CMD build etc in a locale using that encoding.)

The 'NeedsCompilation' field should be set to "yes" if the package contains code which to be compiled, otherwise "no" (when the package could be installed from source on any platform without additional tools). This is used by install.packages(type = "both") in R >= 2.15.2 on platforms where binary packages are the norm: it is normally set by R CMD build or the repository assuming compilation is required if and only if the package has a src directory.

The 'OS_type' field specifies the OS(es) for which the package is intended. If present, it should be one of unix or windows, and indicates that the package can only be installed on a platform with '.Platform\$OS.type' having that value.

The 'Type' field specifies the type of the package: see Section 1.10 [Package types], page 57.

One can add subject classifications for the content of the package using the fields 'Classification/ACM' or 'Classification/ACM-2012' (using the Computing Classification System of the Association for Computing Machinery, http://www.acm.org/about/class/; the former refers to the 1998 version), 'Classification/JEL' (the Journal of Economic Literature Classification System, https://www.aeaweb.org/econlit/jelCodes.php, or 'Classification/MSC' or 'Classification/MSC-2010' (the Mathematics Subject Classification of the American Mathematical Society, http://www.ams.org/msc/; the former refers to the 2000 version). The subject classifications should be comma-separated lists of the respective classification codes, e.g., 'Classification/ACM: G.4, H.2.8, I.5.1'.

A 'Language' field can be used to indicate if the package documentation is not in English: this should be a comma-separated list of standard (not private use or grandfathered) IETF language tags as currently defined by RFC 5646 (https://tools.ietf.org/html/rfc5646, see also https://en.wikipedia.org/wiki/IETF_language_tag), i.e., use language subtags which in essence are 2-letter ISO 639-1 (https://en.wikipedia.org/wiki/ISO_639-3) language codes.

An 'RdMacros' field can be used to hold a comma-separated list of packages from which the current package will import Rd macro definitions. These will be imported after the system macros, in the order listed in the 'RdMacros' field, before any macro definitions in the current package are loaded. Macro definitions in individual .Rd files in the man directory are loaded last, and are local to later parts of that file. In case of duplicates, the last loaded definition will be used Both R CMD Rd2pdf and R CMD Rdconv have an optional flag --RdMacros=pkglist. The option is also a comma-separated list of package names, and has priority over the value given in DESCRIPTION. Packages using Rd macros should depend on R 3.2.0 or later.

Note: There should be no 'Built' or 'Packaged' fields, as these are added by the package management tools.

⁸ Duplicate definitions may trigger a warning: see Section 2.13 [User-defined macros], page 73.

There is no restriction on the use of other fields not mentioned here (but using other capitalizations of these field names would cause confusion). Fields Note, Contact (for contacting the authors/developers) and MailingList are in common use. Some repositories (including CRAN and R-forge) add their own fields.

1.1.2 Licensing

Licensing for a package which might be distributed is an important but potentially complex subject.

It is very important that you include license information! Otherwise, it may not even be legally correct for others to distribute copies of the package, let alone use it.

The package management tools use the concept of 'free or open source software' (FOSS, e.g., https://en.wikipedia.org/wiki/FOSS) licenses: the idea being that some users of R and its packages want to restrict themselves to such software. Others need to ensure that there are no restrictions stopping them using a package, e.g. forbidding commercial or military use. It is a central tenet of FOSS software that there are no restrictions on users nor usage.

Do not use the 'License' field for information on copyright holders: if needed, use a 'Copyright' field.

The mandatory 'License' field in the DESCRIPTION file should specify the license of the package in a standardized form. Alternatives are indicated via vertical bars. Individual specifications must be one of

• One of the "standard" short specifications

```
GPL-2 GPL-3 LGPL-2 LGPL-2.1 LGPL-3 AGPL-3 Artistic-2.0
BSD_2_clause BSD_3_clause MIT
```

as made available *via* https://www.R-project.org/Licenses/ and contained in subdirectory share/licenses of the R source or home directory.

• The names or abbreviations of other licenses contained in the license data base in file share/licenses/license.db in the R source or home directory, possibly (for versioned licenses) followed by a version restriction of the form '(op v)' with 'op' one of the comparison operators '<', '<=', '>', '>=', '==', or '!=' and 'v' a numeric version specification (strings of non-negative integers separated by '.'), possibly combined via ',' (see below for an example). For versioned licenses, one can also specify the name followed by the version, or combine an existing abbreviation and the version with a '-'.

Abbreviations GPL and LGPL are ambiguous and usually taken to mean any version of the license: but it is better not to use them.

- One of the strings 'file LICENSE' or 'file LICENCE' referring to a file named LICENSE or LICENCE in the package (source and installation) top-level directory.
- The string 'Unlimited', meaning that there are no restrictions on distribution or use other than those imposed by relevant laws (including copyright laws).

If a package license restricts a base license (where permitted, e.g., using GPL-3 or AGPL-3 with an attribution clause), the additional terms should be placed in file LICENSE (or LICENCE), and the string '+ file LICENSE' (or '+ file LICENCE', respectively) should be appended to the corresponding individual license specification. Note that several commonly used licenses do not permit restrictions: this includes GPL-2 and hence any specification which includes it.

Examples of standardized specifications include

```
License: GPL-2
```

License: LGPL (>= 2.0, < 3) | Mozilla Public License

License: GPL-2 | file LICENCE

License: GPL (>= 2) | BSD_3_clause + file LICENSE

License: Artistic-2.0 | AGPL-3 + file LICENSE

Please note in particular that "Public domain" is not a valid license, since it is not recognized in some jurisdictions.

Please ensure that the license you choose also covers any dependencies (including system dependencies) of your package: it is particularly important that any restrictions on the use of such dependencies are evident to people reading your DESCRIPTION file.

Fields 'License_is_FOSS' and 'License_restricts_use' may be added by repositories where information cannot be computed from the name of the license. 'License_is_FOSS: yes' is used for licenses which are known to be FOSS, and 'License_restricts_use' can have values 'yes' or 'no' if the LICENSE file is known to restrict users or usage, or known not to. These are used by, e.g., the available.packages filters.

The optional file LICENSE/LICENCE contains a copy of the license of the package. To avoid any confusion only include such a file if it is referred to in the 'License' field of the DESCRIPTION file.

Whereas you should feel free to include a license file in your *source* distribution, please do not arrange to *install* yet another copy of the GNU COPYING or COPYING.LIB files but refer to the copies on https://www.R-project.org/Licenses/ and included in the R distribution (in directory share/licenses). Since files named LICENSE or LICENCE *will* be installed, do not use these names for standard license files. To include comments about the licensing rather than the body of a license, use a file named something like LICENSE.note.

A few "standard" licenses are rather license templates which need additional information to be completed via '+ file LICENSE'.

1.1.3 Package Dependencies

The 'Depends' field gives a comma-separated list of package names which this package depends on. Those packages will be attached before the current package when library or require is called. Each package name may be optionally followed by a comment in parentheses specifying a version requirement. The comment should contain a comparison operator, whitespace and a valid version number, e.g. 'MASS (>= 3.1-20)'.

The 'Depends' field can also specify a dependence on a certain version of R — e.g., if the package works only with R version 3.0.0 or later, include 'R (>= 3.0.0)' in the 'Depends' field. You can also require a certain SVN revision for R-devel or R-patched, e.g. 'R (>= 2.14.0), R (>= r56550)' requires a version later than R-devel of late July 2011 (including released versions of 2.14.0).

It makes no sense to declare a dependence on R without a version specification, nor on the package base: this is an R package and package base is always available.

A package or 'R' can appear more than once in the 'Depends' field, for example to give upper and lower bounds on acceptable versions.

Both library and the R package checking facilities use this field: hence it is an error to use improper syntax or misuse the 'Depends' field for comments on other software that might be needed. The R INSTALL facilities check if the version of R used is recent enough for the package being installed, and the list of packages which is specified will be attached (after checking version requirements) before the current package.

The 'Imports' field lists packages whose namespaces are imported from (as specified in the NAMESPACE file) but which do not need to be attached. Namespaces accessed by the '::' and ':::' operators must be listed here, or in 'Suggests' or 'Enhances' (see below). Ideally this field will include all the standard packages that are used, and it is important to include S4-using packages (as their class definitions can change and the DESCRIPTION file is used to decide which packages to re-install when this happens). Packages declared in the 'Depends' field should not

also be in the 'Imports' field. Version requirements can be specified and are checked when the namespace is loaded (since $R \ge 3.0.0$).

The 'Suggests' field uses the same syntax as 'Depends' and lists packages that are not necessarily needed. This includes packages used only in examples, tests or vignettes (see Section 1.4 [Writing package vignettes], page 35), and packages loaded in the body of functions. E.g., suppose an example from package foo uses a dataset from package bar. Then it is not necessary to have bar use foo unless one wants to execute all the examples/tests/vignettes: it is useful to have bar, but not necessary. Version requirements can be specified, and will be used by R CMD check.

Finally, the 'Enhances' field lists packages "enhanced" by the package at hand, e.g., by providing methods for classes from these packages, or ways to handle objects from these packages (so several packages have 'Enhances: chron' because they can handle datetime objects from chron (https://CRAN.R-project.org/package=chron) even though they prefer R's native datetime functions). Version requirements can be specified, but are currently not used. Such packages cannot be required to check the package: any tests which use them must be conditional on the presence of the package. (If your tests use e.g. a dataset from another package it should be in 'Suggests' and not 'Enhances'.)

The general rules are

- A package should be listed in only one of these fields.
- Packages whose namespace only is needed to load the package using library(pkgname) should be listed in the 'Imports' field and not in the 'Depends' field. Packages listed in imports or importFrom directives in the NAMESPACE file should almost always be in 'Imports' and not 'Depends'.
- Packages that need to be attached to successfully load the package using library(pkgname) must be listed in the 'Depends' field.
- All packages that are needed¹⁰ to successfully run R CMD check on the package must be listed in one of 'Depends' or 'Suggests' or 'Imports'. Packages used to run examples or tests conditionally (e.g. via if(require(pkgname))) should be listed in 'Suggests' or 'Enhances'. (This allows checkers to ensure that all the packages needed for a complete check are installed.)

In particular, packages providing "only" data for examples or vignettes should be listed in 'Suggests' rather than 'Depends' in order to make lean installations possible.

Version dependencies in the 'Depends' and 'Imports' fields are used by library when it loads the package, and install.packages checks versions for the 'Depends', 'Imports' and (for dependencies = TRUE) 'Suggests' fields.

It is increasingly important that the information in these fields is complete and accurate: it is for example used to compute which packages depend on an updated package and which packages can safely be installed in parallel.

This scheme was developed before all packages had namespaces (R 2.14.0 in October 2011), and good practice changed once that was in place.

Field 'Depends' should nowadays be used rarely, only for packages which are intended to be put on the search path to make their facilities available to the end user (and not to the

⁹ even one wrapped in \donttest.

This includes all packages directly called by library and require calls, as well as data obtained via data(theirdata, package = "somepkg") calls: R CMD check will warn about all of these. But there are subtler uses which it will not detect: e.g. if package A uses package B and makes use of functionality in package B which uses package C which package B suggests or enhances, then package C needs to be in the 'Suggests' list for package A. Nor will undeclared uses in included files be reported, nor unconditional uses of packages listed under 'Enhances'.

package itself): for example it makes sense that a user of package latticeExtra (https://CRAN.R-project.org/package=latticeExtra) would want the functions of package lattice (https://CRAN.R-project.org/package=lattice) made available.

Almost always packages mentioned in 'Depends' should also be imported from in the NAMESPACE file: this ensures that any needed parts of those packages are available when some other package imports the current package.

The 'Imports' field should not contain packages which are not imported from (via the NAMESPACE file or :: or ::: operators), as all the packages listed in that field need to be installed for the current package to be installed. (This is checked by R CMD check.)

R code in the package should call library or require only exceptionally. Such calls are never needed for packages listed in 'Depends' as they will already be on the search path. It used to be common practice to use require calls for packages listed in 'Suggests' in functions which used their functionality, but nowadays it is better to access such functionality via: calls.

A package that wishes to make use of header files in other packages needs to declare them as a comma-separated list in the field 'LinkingTo' in the DESCRIPTION file. For example

```
LinkingTo: link1, link2
```

As from R 3.0.2 the 'LinkingTo' field can have a version requirement which is checked at installation. (In earlier versions of R it would cause the specification to be ignored.)

Specifying a package in 'LinkingTo' suffices if these are C++ headers containing source code or static linking is done at installation: the packages do not need to be (and usually should not be) listed in the 'Depends' or 'Imports' fields. This includes CRAN package BH (https://CRAN.R-project.org/package=BH) and almost all users of RcppArmadillo (https://CRAN.R-project.org/package=RcppArmadillo) and RcppEigen (https://CRAN.R-project.org/package=RcppEigen).

For another use of 'LinkingTo' see Section 5.4.2 [Linking to native routines in other packages], page 105.

The 'Additional_repositories' field is a comma-separated list of repository URLs where the packages named in the other fields may be found. It is currently used by R CMD check to check that the packages can be found, at least as source packages (which can be installed on any platform).

1.1.3.1 Suggested packages

Note that someone wanting to run the examples/tests/vignettes may not have a suggested package available (and it may not even be possible to install it for that platform). The recommendation used to be to make their use conditional *via* if(require("pkgname"))): this is fine if that conditioning is done in examples/tests/vignettes.

However, using require for conditioning in package code is not good practice as it alters the search path for the rest of the session and relies on functions in that package not being masked by other require or library calls. It is better practice to use code like

```
if (requireNamespace("rgl", quietly = TRUE)) {
   rgl::plot3d(...)
} else {
   ## do something else not involving rgl.
}
```

Note the use of rgl:: as that object would not necessarily be visible (and if it is, it need not be the one from that namespace: plot3d occurs in several other packages). If the intention is to give an error if the suggested package is not available, simply use e.g. rgl::plot3d.

Note that the recommendation to use suggested packages conditionally in tests does also apply to packages used to manage test suites: a notorious example was **testthat** which in version 1.0.0 contained illegal C++ code and hence could not be installed on standards-compliant platforms.

As noted above, packages in 'Enhances' must be used conditionally and hence objects within them should always be accessed via ::.

1.1.4 The INDEX file

The optional file INDEX contains a line for each sufficiently interesting object in the package, giving its name and a description (functions such as print methods not usually called explicitly might not be included). Normally this file is missing and the corresponding information is automatically generated from the documentation sources (using tools::Rdindex()) when installing from source.

The file is part of the information given by library(help = pkgname).

Rather than editing this file, it is preferable to put customized information about the package into an overview help page (see Section 2.1.4 [Documenting packages], page 65) and/or a vignette (see Section 1.4 [Writing package vignettes], page 35).

1.1.5 Package subdirectories

The R subdirectory contains R code files, only. The code files to be installed must start with an ASCII (lower or upper case) letter or digit and have one of the extensions 11 .R, .S, .q, .r, or .s. We recommend using .R, as this extension seems to be not used by any other software. It should be possible to read in the files using source(), so R objects must be created by assignments. Note that there need be no connection between the name of the file and the R objects created by it. Ideally, the R code files should only directly assign R objects and definitely should not call functions with side effects such as require and options. If computations are required to create objects these can use code 'earlier' in the package (see the 'Collate' field) plus functions in the 'Depends' packages provided that the objects created do not depend on those packages except via namespace imports.

Two exceptions are allowed: if the R subdirectory contains a file sysdata.rda (a saved image of one or more R objects: please use suitable compression as suggested by tools::resaveRdaFiles, and see also the 'SysDataCompression' DESCRIPTION field.) this will be lazy-loaded into the namespace environment — this is intended for system datasets that are not intended to be user-accessible *via* data. Also, files ending in '.in' will be allowed in the R directory to allow a configure script to generate suitable files.

Only ASCII characters (and the control characters tab, formfeed, LF and CR) should be used in code files. Other characters are accepted in comments¹², but then the comments may not be readable in e.g. a UTF-8 locale. Non-ASCII characters in object names will normally¹³ fail when the package is installed. Any byte will be allowed in a quoted character string but \uxxxx escapes should be used for non-ASCII characters. However, non-ASCII character strings may not be usable in some locales and may display incorrectly in others.

Various R functions in a package can be used to initialize and clean up. See Section 1.5.3 [Load hooks], page 40.

The man subdirectory should contain (only) documentation files for the objects in the package in R documentation (Rd) format. The documentation filenames must start with an ASCII (lower or upper case) letter or digit and have the extension .Rd (the default) or .rd. Further, the names must be valid in 'file://' URLs, which means¹⁴ they must be entirely ASCII and not contain '%'. See Chapter 2 [Writing R documentation files], page 58, for more information. Note that all

Extensions .S and .s arise from code originally written for S(-PLUS), but are commonly used for assembler code. Extension .q was used for S, which at one time was tentatively called QPE.

 $^{^{\}rm 12}\,$ but they should be in the encoding declared in the <code>DESCRIPTION</code> file.

¹³ This is true for OSes which implement the 'C' locale: Windows' idea of the 'C' locale uses the WinAnsi charset.

More precisely, they can contain the English alphanumeric characters and the symbols '- - + !'(),; = &'.

user-level objects in a package should be documented; if a package pkg contains user-level objects which are for "internal" use only, it should provide a file pkg-internal.Rd which documents all such objects, and clearly states that these are not meant to be called by the user. See e.g. the sources for package grid in the R distribution. Note that packages which use internal objects extensively should not export those objects from their namespace, when they do not need to be documented (see Section 1.5 [Package namespaces], page 39).

Having a man directory containing no documentation files may give an installation error.

The man subdirectory may contain a subdirectory named macros; this will contain source for user-defined Rd macros. (See Section 2.13 [User-defined macros], page 73.) These use the Rd format, but may not contain anything but macro definitions, comments and whitespace.

The R and man subdirectories may contain OS-specific subdirectories named unix or windows.

The sources and headers for the compiled code are in src, plus optionally a file Makevars or Makefile. When a package is installed using R CMD INSTALL, make is used to control compilation and linking into a shared object for loading into R. There are default make variables and rules for this (determined when R is configured and recorded in R_HOME/etcR_ARCH/Makeconf), providing support for C, C++, FORTRAN 77, Fortran 9x¹⁵, Objective C and Objective C++¹⁶ with associated extensions .c, .cc or .cpp, .f, .f90 or .f95, .m, and .mm, respectively. We recommend using .h for headers, also for C++¹⁷ or Fortran 9x include files. (Use of extension .C for C++ is no longer supported.) Files in the src directory should not be hidden (start with a dot), and hidden files will under some versions of R be ignored.

It is not portable (and may not be possible at all) to mix all these languages in a single package, and we do not support using both C++ and Fortran 9x. Because R itself uses it, we know that C and FORTRAN 77 can be used together and mixing C and C++ seems to be widely successful.

If your code needs to depend on the platform there are certain defines which can used in C or C++. On all Windows builds (even 64-bit ones) '_WIN32' will be defined: on 64-bit Windows builds also '_WIN64', and on OS X '__APPLE__' is defined. 18

The default rules can be tweaked by setting macros¹⁹ in a file src/Makevars (see Section 1.2.1 [Using Makevars], page 20). Note that this mechanism should be general enough to eliminate the need for a package-specific src/Makefile. If such a file is to be distributed, considerable care is needed to make it general enough to work on all R platforms. If it has any targets at all, it should have an appropriate first target named 'all' and a (possibly empty) target 'clean' which removes all files generated by running make (to be used by 'R CMD INSTALL --clean' and 'R CMD INSTALL --preclean'). There are platform-specific file names on Windows: src/Makevars.win takes precedence over src/Makevars and src/Makefile.win must be used. Some make programs require makefiles to have a complete final line, including a newline.

A few packages use the src directory for purposes other than making a shared object (e.g. to create executables). Such packages should have files src/Makefile and src/Makefile.win (unless intended for only Unix-alikes or only Windows).

In very special cases packages may create binary files other than the shared objects/DLLs in the src directory. Such files will not be installed in a multi-architecture setting since R CMD

Note that Ratfor is not supported. If you have Ratfor source code, you need to convert it to FORTRAN. Only FORTRAN 77 (which we write in upper case) is supported on all platforms, but most also support Fortran-95 (for which we use title case). If you want to ship Ratfor source files, please do so in a subdirectory of src and not in the main subdirectory.

 $^{^{16}}$ either or both of which may not be supported on particular platforms

 $^{^{17}\,}$ Using .hpp is not guaranteed to be portable.

¹⁸ There is also '__APPLE_CC__', but that indicates a compiler with Apple-specific features, not the OS. It is used in Rinlinedfuns.h.

 $^{^{\}rm 19}\,$ the POSIX terminology, called 'make variables' by GNU make.

INSTALL --libs-only is used to merge multiple sub-architectures and it only copies shared objects/DLLs. If a package wants to install other binaries (for example executable programs), it should provide an R script src/install.libs.R which will be run as part of the installation in the src build directory instead of copying the shared objects/DLLs. The script is run in a separate R environment containing the following variables: R_PACKAGE_NAME (the name of the package), R_PACKAGE_SOURCE (the path to the source directory of the package), R_PACKAGE_DIR (the path of the target installation directory of the package), R_ARCH (the arch-dependent part of the path, often empty), SHLIB_EXT (the extension of shared objects) and WINDOWS (TRUE on Windows, FALSE elsewhere). Something close to the default behavior could be replicated with the following src/install.libs.R file:

```
files <- Sys.glob(paste0("*", SHLIB_EXT))
  dest <- file.path(R_PACKAGE_DIR, paste0('libs', R_ARCH))
  dir.create(dest, recursive = TRUE, showWarnings = FALSE)
  file.copy(files, dest, overwrite = TRUE)
  if(file.exists("symbols.rds"))
     file.copy("symbols.rds", dest, overwrite = TRUE)

On the other hand, executable programs could be installed along the lines of
  execs <- c("one", "two", "three")
  if(WINDOWS) execs <- paste0(execs, ".exe")
  if ( any(file.exists(execs)) ) {
    dest <- file.path(R_PACKAGE_DIR, paste0('bin', R_ARCH))
    dir.create(dest, recursive = TRUE, showWarnings = FALSE)
    file.copy(execs, dest, overwrite = TRUE)
}</pre>
```

Note the use of architecture-specific subdirectories of bin where needed.

The data subdirectory is for data files: See Section 1.1.6 [Data in packages], page 15.

The demo subdirectory is for R scripts (for running via demo()) that demonstrate some of the functionality of the package. Demos may be interactive and are not checked automatically, so if testing is desired use code in the tests directory to achieve this. The script files must start with a (lower or upper case) letter and have one of the extensions .R or .r. If present, the demo subdirectory should also have a OOIndex file with one line for each demo, giving its name and a description separated by a tab or at least three spaces. (This index file is not generated automatically.) Note that a demo does not have a specified encoding and so should be an ASCII file (see Section 1.6.3 [Encoding issues], page 50). As from R 3.0.0 demo() will use the package encoding if there is one, but this is mainly useful for non-ASCII comments.

The contents of the inst subdirectory will be copied recursively to the installation directory. Subdirectories of inst should not interfere with those used by R (currently, R, data, demo, exec, libs, man, help, html and Meta, and earlier versions used latex, R-ex). The copying of the inst happens after src is built so its Makefile can create files to be installed. To exclude files from being installed, one can specify a list of exclude patterns in file .Rinstignore in the top-level source directory. These patterns should be Perl-like regular expressions (see the help for regexp in R for the precise details), one per line, to be matched case-insensitively 20 against the file and directory paths, e.g. doc/.*[.]png\$ will exclude all PNG files in inst/doc based on the extension.

Note that with the exceptions of INDEX, LICENSE/LICENCE and NEWS, information files at the top level of the package will *not* be installed and so not be known to users of Windows and OS X compiled packages (and not seen by those who use R CMD INSTALL or install.packages on the tarball). So any information files you wish an end user to see should be included in inst.

 $^{^{20}\,}$ on all platforms from R 3.1.0

Note that if the named exceptions also occur in inst, the version in inst will be that seen in the installed package.

Things you might like to add to inst are a CITATION file for use by the citation function, and a NEWS.Rd file for use by the news function. See its help page for the specific format restrictions of the NEWS.Rd file.

Another file sometimes needed in inst is AUTHORS or COPYRIGHTS to specify the authors or copyright holders when this is too complex to put in the DESCRIPTION file.

Subdirectory tests is for additional package-specific test code, similar to the specific tests that come with the R distribution. Test code can either be provided directly in a .R file, or via a .R in file containing code which in turn creates the corresponding .R file (e.g., by collecting all function objects in the package and then calling them with the strangest arguments). The results of running a .R file are written to a .Rout file. If there is a corresponding²¹ .Rout.save file, these two are compared, with differences being reported but not causing an error. The directory tests is copied to the check area, and the tests are run with the copy as the working directory and with R_LIBS set to ensure that the copy of the package installed during testing will be found by library(pkg_name). Note that the package-specific tests are run in a vanilla R session without setting the random-number seed, so tests which use random numbers will need to set the seed to obtain reproducible results (and it can be helpful to do so in all cases, to avoid occasional failures when tests are run).

If directory tests has a subdirectory Examples containing a file pkg-Ex.Rout.save, this is compared to the output file for running the examples when the latter are checked. Reference output should be produced without having the --timings option set (and note that --as-cran sets it).

Subdirectory exec could contain additional executable scripts the package needs, typically scripts for interpreters such as the shell, Perl, or Tcl. NB: only files (and not directories) under exec are installed (and those with names starting with a dot are ignored), and they are all marked as executable (mode 755, moderated by 'umask') on POSIX platforms. Note too that this is not suitable for executable *programs* since some platforms (including Windows) support multiple architectures using the same installed package directory.

Subdirectory po is used for files related to *localization*: see Section 1.8 [Internationalization], page 54.

Subdirectory tools is the preferred place for auxiliary files needed during configuration, and also for sources need to re-create scripts (e.g. M4 files for autoconf).

1.1.6 Data in packages

The data subdirectory is for data files, either to be made available *via* lazy-loading or for loading using data(). (The choice is made by the 'LazyData' field in the DESCRIPTION file: the default is not to do so.) It should not be used for other data files needed by the package, and the convention has grown up to use directory inst/extdata for such files.

Data files can have one of three types as indicated by their extension: plain R code (.R or .r), tables (.tab, .txt, or .csv, see ?data for the file formats, and note that .csv is not the standard²² CSV format), or save() images (.RData or .rda). The files should not be hidden (have names starting with a dot). Note that R code should be "self-sufficient" and not make use of extra functionality provided by the package, so that the data file can also be used without having to load the package or its namespace.

The best way to generate such a file is to copy the .Rout from a successful run of R CMD check. If you want to generate it separately, do run R with options --vanilla --slave and with environment variable LANGUAGE=en set to get messages in English. Be careful not to use output with the option --timings (and note that --as-cran sets it).

 $^{^{22}}$ e.g. https://tools.ietf.org/html/rfc4180.

Images (extensions .RData²³ or .rda) can contain references to the namespaces of packages that were used to create them. Preferably there should be no such references in data files, and in any case they should only be to packages listed in the Depends and Imports fields, as otherwise it may be impossible to install the package. To check for such references, load all the images into a vanilla R session, and look at the output of loadedNamespaces().

If your data files are large and you are not using 'LazyData' you can speed up installation by providing a file datalist in the data subdirectory. This should have one line per topic that data() will find, in the format 'foo' if data(foo) provides 'foo', or 'foo: bar bah' if data(foo) provides 'bar' and 'bah'. R CMD build will automatically add a datalist file to data directories of over 1Mb, using the function tools::add_datalist.

Tables (.tab, .txt, or .csv files) can be compressed by gzip, bzip2 or xz, optionally with additional extension .gz, .bz2 or .xz.

If your package is to be distributed, do consider the resource implications of large datasets for your users: they can make packages very slow to download and use up unwelcome amounts of storage space, as well as taking many seconds to load. It is normally best to distribute large datasets as .rda images prepared by save(, compress = TRUE) (the default). Using bzip2 or xz compression will usually reduce the size of both the package tarball and the installed package, in some cases by a factor of two or more.

Package **tools** has a couple of functions to help with data images: **checkRdaFiles** reports on the way the image was saved, and **resaveRdaFiles** will re-save with a different type of compression, including choosing the best type for that particular image.

Some packages using 'LazyData' will benefit from using a form of compression other than gzip in the installed lazy-loading database. This can be selected by the --data-compress option to R CMD INSTALL or by using the 'LazyDataCompression' field in the DESCRIPTION file. Useful values are bzip2, xz and the default, gzip. The only way to discover which is best is to try them all and look at the size of the pkgname/data/Rdata.rdb file.

Lazy-loading is not supported for very large datasets (those which when serialized exceed 2GB, the limit for the format on 32-bit platforms and all platforms prior to R 3.0.0).

The analogue for sysdata.rda is field 'SysDataCompression': the default (since R 2.12.2) is xz for files bigger than 1MB otherwise gzip.

1.1.7 Non-R scripts in packages

Code which needs to be compiled (C, C++, FORTRAN, Fortran 95 . . .) is included in the src subdirectory and discussed elsewhere in this document.

Subdirectory exec could be used for scripts for interpreters such as the shell, BUGS, JavaScript, Matlab, Perl, php (amap (https://CRAN.R-project.org/package=amap)), Python or Tcl (Simile (https://CRAN.R-project.org/package=Simile)), or even R. However, it seems more common to use the inst directory, for example WriteXLS/inst/Perl, NMF/inst/m-files, RnavGraph/inst/tcl, RProtoBuf/inst/python and emdbook/inst/BUGS and gridSVG/inst/js.

Java code is a special case: except for very small programs, .java files should be byte-compiled (to a .class file) and distributed as part of a .jar file: the conventional location for the .jar file(s) is inst/java. It is desirable (and required under an Open Source license) to make the Java source files available: this is best done in a top-level java directory in the package—the source files should not be installed.

If your package requires one of these interpreters or an extension then this should be declared in the 'SystemRequirements' field of its DESCRIPTION file. (Users of Java most often do so via

People who have trouble with case are advised to use .rda as a common error is to refer to abc.RData as abc.Rdata!

rJava (https://CRAN.R-project.org/package=rJava), when depending on/importing that suffices.)

Windows and Mac users should be aware that the Tcl extensions 'BWidget' and 'Tktable' which are currently included with the R for Windows and in the OS X installers *are* extensions and do need to be declared for users of other platforms (and that 'Tktable' is less widely available than it used to be, including not in the main repositories for major Linux distributions).

'BWidget' needs to be installed by the user on other OSes. This is fairly easy to do: first find the Tcl/Tk search path:

```
library(tcltk)
strsplit(tclvalue('auto_path'), " ")[[1]]
```

then download the sources from http://sourceforge.net/projects/tcllib/files/BWidget/and at the command line run something like

```
tar xf bwidget-1.9.8.tar.gz
sudo mv bwidget-1.9.8 /usr/local/lib
```

substituting a location on the Tcl/Tk search path for /usr/local/lib if needed.

1.1.8 Specifying URLs

URLs in many places in the package documentation will be converted to clickable hyperlinks in at least some of their renderings. So care is needed that their forms are correct and portable.

The full URL should be given, including the scheme (often 'http://' or 'https://') and a final '/' for references to directories.

Spaces in URLs are not portable and how they are handled does vary by HTTP server and by client. There should be no space in the host part of an 'http://' URL, and spaces in the remainder should be encoded, with each space replaced by '%20'.

Other characters may benefit from being encoded: see the help on URLencode().

The canonical URL for a CRAN package is

```
https://cran.r-project.org/package=pkgname
```

and not a version starting 'http://cran.r-project.org/web/packages/pkgname'.

1.2 Configure and cleanup

Note that most of this section is specific to Unix-alikes: see the comments later on about the Windows port of R.

If your package needs some system-dependent configuration before installation you can include an executable (Bourne²⁴) shell script configure in your package which (if present) is executed by R CMD INSTALL before any other action is performed. This can be a script created by the Autoconf mechanism, but may also be a script written by yourself. Use this to detect if any nonstandard libraries are present such that corresponding code in the package can be disabled at install time rather than giving error messages when the package is compiled or used. To summarize, the full power of Autoconf is available for your extension package (including variable substitution, searching for libraries, etc.).

Under a Unix-alike only, an executable (Bourne shell) script cleanup is executed as the last thing by R CMD INSTALL if option --clean was given, and by R CMD build when preparing the package for building from its source.

The script should only assume a POSIX-compliant /bin/sh - see http://pubs.opengroup.org/onlinepubs/9699919799/utilities/V3_chap02.html. In particular bash extensions must not be used, and not all R platforms have a bash command, let alone one at /bin/bash. All known shells used with R support the use of backticks, but not all support '\$(cmd)'.

As an example consider we want to use functionality provided by a (C or FORTRAN) library foo. Using Autoconf, we can create a configure script which checks for the library, sets variable HAVE_FOO to TRUE if it was found and to FALSE otherwise, and then substitutes this value into output files (by replacing instances of '@HAVE_FOO@' in input files with the value of HAVE_FOO). For example, if a function named bar is to be made available by linking against library foo (i.e., using -lfoo), one could use

```
AC_CHECK_LIB(foo, fun, [HAVE_FOO=TRUE], [HAVE_FOO=FALSE])
AC_SUBST(HAVE_FOO)
.....
AC_CONFIG_FILES([foo.R])
AC_OUTPUT
in configure.ac (assuming Autoconf 2.50 or later).
The definition of the respective R function in foo.R.in could be
foo <- function(x) {
    if(!@HAVE_FOO@)
        stop("Sorry, library 'foo' is not available"))
    ...
From this file configure creates the actual R source file foo.R looking like
foo <- function(x) {
    if(!FALSE)
        stop("Sorry, library 'foo' is not available"))
```

if library foo was not found (with the desired functionality). In this case, the above R code effectively disables the function.

One could also use different file fragments for available and missing functionality, respectively.

You will very likely need to ensure that the same C compiler and compiler flags are used in the configure tests as when compiling R or your package. Under a Unix-alike, you can achieve this by including the following fragment early in configure.ac (before calling AC_PROG_CC)

```
: ${R_HOME='R RHOME'}
if test -z "${R_HOME}"; then
   echo "could not determine R_HOME"
   exit 1
fi
CC='"${R_HOME}/bin/R" CMD config CC'
CFLAGS='"${R_HOME}/bin/R" CMD config CFLAGS'
CPPFLAGS='"${R_HOME}/bin/R" CMD config CPPFLAGS'
```

(Using '\${R_HOME}/bin/R' rather than just 'R' is necessary in order to use the correct version of R when running the script as part of R CMD INSTALL, and the quotes since '\${R_HOME}' might contain spaces.)

If your code does load checks then you may also need

```
LDFLAGS='"${R_HOME}/bin/R" CMD config LDFLAGS'
```

and packages written with C++ need to pick up the details for the C++ compiler and switch the current language to C++ by something like

```
CXX='"${R_HOME}/bin/R" CMD config CXX'
CXXFLAGS='"${R_HOME}/bin/R" CMD config CXXFLAGS'
AC_LANG(C++)
```

The latter is important, as for example C headers may not be available to C++ programs or may not be written to avoid C++ name-mangling.

You can use R CMD config for getting the value of the basic configuration variables, and also the header and library flags necessary for linking a front-end executable program against R, see R CMD config --help for details.

To check for an external BLAS library using the ACX_BLAS macro from the official Autoconf Macro Archive, one can simply do

```
F77='"\R_HOME}/bin/R" CMD config F77'
AC_PROG_F77
FLIBS='"\R_HOME}/bin/R" CMD config FLIBS'
ACX_BLAS([], AC_MSG_ERROR([could not find your BLAS library], 1))
```

Note that FLIBS as determined by R must be used to ensure that FORTRAN 77 code works on all R platforms. Calls to the Autoconf macro AC_F77_LIBRARY_LDFLAGS, which would overwrite FLIBS, must not be used (and hence e.g. removed from ACX_BLAS). (Recent versions of Autoconf in fact allow an already set FLIBS to override the test for the FORTRAN linker flags.)

N.B.: If the configure script creates files, e.g. src/Makevars, you do need a cleanup script to remove them. Otherwise R CMD build may ship the files that are created. For example, package RODBC (https://CRAN.R-project.org/package=RODBC) has

```
#!/bin/sh
```

```
rm -f config.* src/Makevars src/config.h
```

As this example shows, configure often creates working files such as config.log.

If your configure script needs auxiliary files, it is recommended that you ship them in a tools directory (as R itself does).

You should bear in mind that the configure script will not be used on Windows systems. If your package is to be made publicly available, please give enough information for a user on a non-Unix-alike platform to configure it manually, or provide a configure.win script to be used on that platform. (Optionally, there can be a cleanup.win script. Both should be shell scripts to be executed by ash, which is a minimal version of Bourne-style sh.) When configure.win is run the environment variables R_HOME (which uses '/' as the file separator), R_ARCH and Use R_ARCH_BIN will be set. Use R_ARCH to decide if this is a 64-bit build (its value there is '/x64') and to install DLLs to the correct place (\${R_HOME}/libs\${R_ARCH}). Use R_ARCH_BIN to find the correct place under the bin directory, e.g. \${R_HOME}/bin\${R_ARCH_BIN}/Rscript.exe.

In some rare circumstances, the configuration and cleanup scripts need to know the location into which the package is being installed. An example of this is a package that uses C code and creates two shared object/DLLs. Usually, the object that is dynamically loaded by R is linked against the second, dependent, object. On some systems, we can add the location of this dependent object to the object that is dynamically loaded by R. This means that each user does not have to set the value of the LD_LIBRARY_PATH (or equivalent) environment variable, but that the secondary object is automatically resolved. Another example is when a package installs support files that are required at run time, and their location is substituted into an R data structure at installation time. The names of the top-level library directory (i.e., specifiable via the '-1' argument) and the directory of the package itself are made available to the installation scripts via the two shell/environment variables R_LIBRARY_DIR and R_PACKAGE_DIR. Additionally, the name of the package (e.g. 'survival' or 'MASS') being installed is available from the environment variable R_PACKAGE_NAME. (Currently the value of R_PACKAGE_DIR is always \${R_LIBRARY_DIR}/\${R_ PACKAGE_NAME}, but this used not to be the case when versioned installs were allowed. Its main use is in configure. win scripts for the installation path of external software's DLLs.) Note that the value of R_PACKAGE_DIR may contain spaces and other shell-unfriendly characters, and so should be quoted in makefiles and configure scripts.

One of the more tricky tasks can be to find the headers and libraries of external software. One tool which is increasingly available on Unix-alikes (but not by default on OS X) to do this is pkg-config. The configure script will need to test for the presence of the command itself (see for example package Cairo (https://CRAN.R-project.org/package=Cairo)), and if present it can be asked if the software is installed, of a suitable version and for compilation/linking flags by e.g.

```
$ pkg-config --exists 'QtCore >= 4.0.0' # check the status
$ pkg-config --modversion QtCore
4.7.1
$ pkg-config --cflags QtCore
-DQT_SHARED -I/usr/include/QtCore
$ pkg-config --libs QtCore
-lQtCore
```

Note that pkg-config --libs gives the information required to link against the default version of that library (usually the dynamic one), and pkg-config --static is needed if the static library is to be used.

Sometimes the name by which the software is known to pkg-config is not what one might expect (e.g. 'gtk+-2.0' even for 2.22). To get a complete list use

```
pkg-config --list-all | sort
```

1.2.1 Using Makevars

Sometimes writing your own configure script can be avoided by supplying a file Makevars: also one of the most common uses of a configure script is to make Makevars from Makevars.in.

A Makevars file is a makefile and is used as one of several makefiles by R CMD SHLIB (which is called by R CMD INSTALL to compile code in the src directory). It should be written if at all possible in a portable style, in particular (except for Makevars.win) without the use of GNU extensions.

The most common use of a Makevars file is to set additional preprocessor options (for example include paths) for C/C++ files *via* PKG_CPPFLAGS, and additional compiler flags by setting PKG_CFLAGS, PKG_CXXFLAGS, PKG_FFLAGS or PKG_FCFLAGS, for C, C++, FORTRAN or Fortran 9x respectively (see Section 5.5 [Creating shared objects], page 106).

N.B.: Include paths are preprocessor options, not compiler options, and **must** be set in PKG_CPPFLAGS as otherwise platform-specific paths (e.g. '-I/usr/local/include') will take precedence.

Makevars can also be used to set flags for the linker, for example '-L' and '-l' options, *via* PKG_LIBS.

When writing a Makevars file for a package you intend to distribute, take care to ensure that it is not specific to your compiler: flags such as -02 -Wall -pedantic (and all other -W flags: for the Solaris compiler these are used to pass arguments to compiler phases) are all specific to GCC.

Also, do not set variables such as CPPFLAGS, CFLAGS etc.: these should be settable by users (sites) through appropriate personal (site-wide) Makevars files. See Section "Customizing package compilation" in *R Installation and Administration*,

There are some macros²⁵ which are set whilst configuring the building of R itself and are stored in R_HOME/etcR_ARCH/Makeconf. That makefile is included as a Makefile after Makevars[.win], and the macros it defines can be used in macro assignments and make command lines in the latter. These include

FLIBS A macro containing the set of libraries need to link FORTRAN code. This may need to be included in PKG_LIBS: it will normally be included automatically if the package contains FORTRAN source files.

 $^{^{25}\,}$ in POSIX parlance: GNU make calls these 'make variables'.

BLAS_LIBS

A macro containing the BLAS libraries used when building R. This may need to be included in PKG_LIBS. Beware that if it is empty then the R executable will contain all the double-precision and double-complex BLAS routines, but no single-precision nor complex routines. If BLAS_LIBS is included, then FLIBS also needs to be²⁶ included following it, as most BLAS libraries are written at least partially in FORTRAN.

LAPACK_LIBS

A macro containing the LAPACK libraries (and paths where appropriate) used when building R. This may need to be included in PKG_LIBS. It may point to a dynamic library libRlapack which contains the main double-precision LAPACK routines as well as those double-complex LAPACK routines needed to build R, or it may point to an external LAPACK library, or may be empty if an external BLAS library also contains LAPACK.

[libRlapack includes all the double-precision LAPACK routines which were current in 2003: a list of which routines are included is in file src/modules/lapack/README. Note that an external LAPACK/BLAS library need not do so, as some were 'deprecated' (and not compiled by default) in LAPACK 3.6.0 in late 2015.]

For portability, the macros BLAS_LIBS and FLIBS should always be included *after* LAPACK_LIBS (and in that order).

SAFE_FFLAGS

A macro containing flags which are needed to circumvent over-optimization of FOR-TRAN code: it is typically '-g -02 -ffloat-store' on 'ix86' platforms using gfortran. Note that this is **not** an additional flag to be used as part of PKG_FFLAGS, but a replacement for FFLAGS, and that it is intended for the FORTRAN 77 compiler 'F77' and not necessarily for the Fortran 90/95 compiler 'FC'. See the example later in this section.

Setting certain macros in Makevars will prevent R CMD SHLIB setting them: in particular if Makevars sets 'OBJECTS' it will not be set on the make command line. This can be useful in conjunction with implicit rules to allow other types of source code to be compiled and included in the shared object. It can also be used to control the set of files which are compiled, either by excluding some files in src or including some files in subdirectories. For example

```
OBJECTS = 4dfp/endianio.o 4dfp/Getifh.o R4dfp-object.o
```

Note that Makevars should not normally contain targets, as it is included before the default makefile and make will call the first target, intended to be all in the default makefile. If you really need to circumvent that, use a suitable (phony) target all before any actual targets in Makevars.[win]: for example package fastICA (https://CRAN.R-project.org/package=fastICA) used to have

```
PKG_LIBS = @BLAS_LIBS@

SLAMC_FFLAGS=$(R_XTRA_FFLAGS) $(FPICFLAGS) $(SHLIB_FFLAGS) $(SAFE_FFLAGS)
all: $(SHLIB)

slamc.o: slamc.f
$(F77) $(SLAMC_FFLAGS) -c -o slamc.o slamc.f
```

²⁶ at least on Unix-alikes: the Windows build currently resolves such dependencies to a static FORTRAN library when Rblas.dll is built.

needed to ensure that the LAPACK routines find some constants without infinite looping. The Windows equivalent was

```
all: $(SHLIB)
slamc.o: slamc.f
$(F77) $(SAFE_FFLAGS) -c -o slamc.o slamc.f
```

(since the other macros are all empty on that platform, and R's internal BLAS was not used). Note that the first target in Makevars will be called, but for back-compatibility it is best named all.

If you want to create and then link to a library, say using code in a subdirectory, use something like

```
.PHONY: all mylibs
all: $(SHLIB)
$(SHLIB): mylibs
mylibs:
(cd subdir; $(MAKE))
```

Be careful to create all the necessary dependencies, as there is no guarantee that the dependencies of all will be run in a particular order (and some of the CRAN build machines use multiple CPUs and parallel makes). In particular,

```
all: mylibs does not suffice.
```

Note that on Windows it is required that Makevars[.win] does create a DLL: this is needed as it is the only reliable way to ensure that building a DLL succeeded. If you want to use the src directory for some purpose other than building a DLL, use a Makefile.win file.

It is sometimes useful to have a target 'clean' in Makevars or Makevars.win: this will be used by R CMD build to clean up (a copy of) the package sources. When it is run by build it will have fewer macros set, in particular not \$(SHLIB), nor \$(OBJECTS) unless set in the file itself. It would also be possible to add tasks to the target 'shlib-clean' which is run by R CMD INSTALL and R CMD SHLIB with options --clean and --preclean.

If you want to run R code in Makevars, e.g. to find configuration information, please do ensure that you use the correct copy of R or Rscript: there might not be one in the path at all, or it might be the wrong version or architecture. The correct way to do this is *via*

```
"$(R_HOME)/bin$(R_ARCH_BIN)/Rscript" filename
"$(R_HOME)/bin$(R_ARCH_BIN)/Rscript" -e 'R expression'
```

where (R_ARCH_BIN) is only needed currently on Windows.

Environment or make variables can be used to select different macros for 32- and 64-bit code, for example (GNU make syntax, allowed on Windows)

```
ifeq "$(WIN)" "64"
PKG_LIBS = value for 64-bit Windows
else
PKG_LIBS = value for 32-bit Windows
endif
```

On Windows there is normally a choice between linking to an import library or directly to a DLL. Where possible, the latter is much more reliable: import libraries are tied to a specific toolchain, and in particular on 64-bit Windows two different conventions have been commonly used. So for example instead of

```
PKG_LIBS = -L\$(XML_DIR)/lib -lxml2
```

```
one can use
    PKG_LIBS = -L$(XML_DIR)/bin -lxm12
since on Windows -lxxx will look in turn for
    libxxx.dll.a
    xxx.dll.a
    libxxx.a
    xxx.lib
    libxxx.dll
    xxx.dll
```

where the first and second are conventionally import libraries, the third and fourth often static libraries (with .lib intended for Visual C++), but might be import libraries. See for example https://sourceware.org/binutils/docs-2.20/ld/WIN32.html#WIN32.

The fly in the ointment is that the DLL might not be named libxxx.dll, and in fact on 32-bit Windows there is a libxml2.dll whereas on one build for 64-bit Windows the DLL is called libxml2-2.dll. Using import libraries can cover over these differences but can cause equal difficulties.

If static libraries are available they can save a lot of problems with run-time finding of DLLs, especially when binary packages are to be distributed and even more when these support both architectures. Where using DLLs is unavoidable we normally arrange (via configure.win) to ship them in the same directory as the package DLL.

1.2.1.1 OpenMP support

There is some support for packages which wish to use OpenMP²⁷. The make macros

```
SHLIB_OPENMP_CFLAGS
SHLIB_OPENMP_CXXFLAGS
SHLIB_OPENMP_FCFLAGS
SHLIB_OPENMP_FFLAGS
```

are available for use in src/Makevars or src/Makevars.win. Include the appropriate macro in PKG_CFLAGS, PKG_CPPFLAGS and so on, and also in PKG_LIBS. C/C++ code that needs to be conditioned on the use of OpenMP can be used inside #ifdef _OPENMP: note that some toolchains used for R (including many of those using clang²⁸) have no OpenMP support at all, not even omp.h.

For example, a package with C code written for OpenMP should have in ${\tt src/Makevars}$ the lines

```
PKG_CFLAGS = $(SHLIB_OPENMP_CFLAGS)
PKG_LIBS = $(SHLIB_OPENMP_CFLAGS)
```

Note that the macro SHLIB_OPENMP_CXXFLAGS applies to the C++98 compiler and not necessarily to the C++11 compiler: users of the latter should do their own configure checks.

Some care is needed when compilers are from different families which may use different OpenMP runtimes (e.g. clang vs GCC including gfortran, although it is currently possible to use the clang runtime with GCC but not vice versa). For a package with Fortran 77 code using OpenMP the appropriate lines are

```
PKG_FFLAGS = $(SHLIB_OPENMP_FFLAGS)
PKG_LIBS = $(SHLIB_OPENMP_CFLAGS)
```

http://www.openmp.org/, https://en.wikipedia.org/wiki/OpenMP, https://computing.llnl.gov/tutorials/openMP/

²⁸ Some builds of clang 3.7.0 have support for OpenMP 3.1 (this is built by default as from 3.8.0), but even if the compiler has, the libomp library may not be installed. At the time of writing Apple builds of clang for OS X had no support.

as the C compiler will be used to link the package code (and there is no guarantee that this will work everywhere).

There is nothing to say what version of OpenMP is supported: version 3.0 (May 2008) is supported by recent versions of the Linux, Windows and Solaris platforms, but portable packages cannot assume that end users have recent versions. OS X currently uses Apple builds of clang with no OpenMP support.

The performance of OpenMP varies substantially between platforms. Both the Windows and earlier Apple OS X implementations have substantial overheads and are only beneficial if quite substantial tasks are run in parallel. Also, on Windows new threads are started with the default²⁹ FPU control word, so computations done on OpenMP threads will not make use of extended-precision arithmetic which is the default for the main process.

Calling any of the R API from threaded code is 'for experts only': they will need to read the source code to determine if it is thread-safe. In particular, code which makes use of the stack-checking mechanism must not be called from threaded code.

Packages are not standard-alone programs, and an R process could contain more than one OpenMP-enabled package as well as other components (for example, an optimized BLAS) making use of OpenMP. So careful consideration needs to be given to resource usage. OpenMP works with parallel regions, and for most implementations the default is to use as many threads as 'CPUs' for such regions. Parallel regions can be nested, although it is common to use only a single thread below the first level. The correctness of the detected number of 'CPUs' and the assumption that the R process is entitled to use them all are both dubious assumptions. The best way to limit resources is to limit the overall number of threads available to OpenMP in the R process: this can be done via environment variable OMP_THREAD_LIMIT, where implemented. Alternatively, the number of threads per region can be limited by the environment variable OMP_NUM_THREADS or API call omp_set_num_threads, or, better, for the regions in your code as part of their specification. E.g. R uses

```
#pragma omp parallel for num_threads(nthreads) ...
```

That way you only control your own code and not that of other OpenMP users.

1.2.1.2 Using pthreads

There is no direct support for the POSIX threads (more commonly known as pthreads): by the time we considered adding it several packages were using it unconditionally so it seems that nowadays it is universally available on POSIX operating systems (hence not Windows).

For reasonably recent versions of gcc and clang the correct specification is

```
PKG_CPPFLAGS = -pthread
PKG_LIBS = -pthread
```

(and the plural version is also accepted on some systems/versions). For other platforms the specification is

```
PKG_CPPFLAGS = -D_REENTRANT
PKG_LIBS = -lpthread
```

(and note that the library name is singular). This is what **-pthread** does on all known current platforms (although earlier versions of OpenBSD used a different library name).

For a tutorial see https://computing.llnl.gov/tutorials/pthreads/.

POSIX threads are not normally used on Windows, which has its own native concepts of threads. However, there are two projects implementing pthreads on top of Windows, pthreads-w32 and winpthreads (part of the MinGW-w64 project).

 $^{^{29}\,}$ Windows default, not MinGW-w64 default.

 $^{^{30}}$ Which it was at the time of writing with GCC, Solaris Studio, Intel and Clang 3.7.x compilers.

Whether Windows toolchains implement pthreads is up to the toolchain provider. A make variable SHLIB_PTHREAD_FLAGS is available: this should be included in both PKG_CPPFLAGS (or the Fortran or F9x equivalents) and PKG_LIBS.

The presence of a working pthreads implementation cannot be unambiguously determined without testing for yourself: however, that '_REENTRANT' is defined³¹ in C/C++ code is a good indication.

See also the comments on thread-safety and performance under OpenMP: on all known R platforms OpenMP is implemented *via* pthreads and the known performance issues are in the latter.

1.2.1.3 Compiling in sub-directories

Package authors fairly often want to organize code in sub-directories of src, for example if they are including a separate piece of external software to which this is an R interface.

One simple way is simply to set OBJECTS to be all the objects that need to be compiled, including in sub-directories. For example, CRAN package RSiena (https://CRAN.R-project.org/package=RSiena) has

```
SOURCES = $(wildcard data/*.cpp network/*.cpp utils/*.cpp model/*.cpp model/*/*.cpp model/*/*.cpp)

OBJECTS = siena07utilities.o siena07internals.o siena07setup.o siena07models.o $(SOURCES:.cpp=.o)
```

One problem with that approach is that unless GNU make extensions are used, the source files need to be listed and kept up-to-date. As in the following from CRAN package lossDev (https://CRAN.R-project.org/package=lossDev):

```
OBJECTS.samplers = samplers/ExpandableArray.o samplers/Knots.o \
    samplers/RJumpSpline.o samplers/RJumpSplineFactory.o \
    samplers/RealSlicerOV.o samplers/SliceFactoryOV.o samplers/MNorm.o

OBJECTS.distributions = distributions/DSpline.o \
    distributions/DChisqrOV.o distributions/DTOV.o \
    distributions/DNormOV.o distributions/DUnifOV.o distributions/RScalarDist.o

OBJECTS.root = RJump.o

OBJECTS = $(OBJECTS.samplers) $(OBJECTS.distributions) $(OBJECTS.root)
```

Where the subdirectory is self-contained code with a suitable makefile, the best approach is something like

Note the quotes: the macros can contain spaces, e.g. CC = "gcc -m64 -std=gnu99". Several authors have forgotten about parallel makes: the static library in the subdirectory must be made before the shared object (\$(SHLIB)) and so the latter must depend on the former. Others forget the need for position-independent code.

We really do not recommend using src/Makefile instead of src/Makevars, and as the example above shows, it is not necessary.

1.2.2 Configure example

It may be helpful to give an extended example of using a configure script to create a src/Makevars file: this is based on that in the RODBC (https://CRAN.R-project.org/package=RODBC) package.

 $^{^{31}\,}$ some Windows toolchains had the typo '_REENTRANCE' instead.

The configure.ac file follows: configure is created from this by running autoconf in the top-level package directory (containing configure.ac).

```
AC_INIT([RODBC], 1.1.8) dnl package name, version
dnl A user-specifiable option
odbc_mgr=""
AC_ARG_WITH([odbc-manager],
    AC_HELP_STRING([--with-odbc-manager=MGR],
   [specify the ODBC manager, e.g. odbc or iodbc]),
    [odbc_mgr=$withval])
if test "$odbc_mgr" = "odbc" ; then
  AC_PATH_PROGS(ODBC_CONFIG, odbc_config)
dnl Select an optional include path, from a configure option
dnl or from an environment variable.
AC_ARG_WITH([odbc-include],
   AC_HELP_STRING([--with-odbc-include=INCLUDE_PATH],
   [the location of ODBC header files]),
    [odbc_include_path=$withval])
RODBC_CPPFLAGS="-I."
if test [ -n "$odbc_include_path" ] ; then
  RODBC_CPPFLAGS="-I. -I${odbc_include_path}"
  if test [ -n "${ODBC_INCLUDE}" ] ; then
     RODBC_CPPFLAGS="-I. -I${ODBC_INCLUDE}"
 fi
fi
dnl ditto for a library path
AC_ARG_WITH([odbc-lib],
   AC_HELP_STRING([--with-odbc-lib=LIB_PATH],
   [the location of ODBC libraries]),
    [odbc_lib_path=$withval])
if test [ -n "$odbc_lib_path" ] ; then
  LIBS="-L$odbc_lib_path ${LIBS}"
  if test [ -n "${ODBC_LIBS}" ] ; then
     LIBS="-L${ODBC_LIBS} ${LIBS}"
 else
    if test -n "${ODBC_CONFIG}"; then
      odbc_lib_path='odbc_config --libs | sed s/-lodbc//'
      LIBS="${odbc_lib_path} ${LIBS}"
   fi
 fi
fi
dnl Now find the compiler and compiler flags to use
: ${R_HOME='R RHOME'}
if test -z "R_HOME"; then
 echo "could not determine R_HOME"
 exit 1
fi
CC="\${R_HOME}/bin/R" CMD config CC"
CPP='"${R_HOME}/bin/R" CMD config CPP'
CFLAGS="${R_HOME}/bin/R" CMD config CFLAGS"
CPPFLAGS="${R_HOME}/bin/R" CMD config CPPFLAGS"
AC_PROG_CC
AC_PROG_CPP
if test -n "${ODBC_CONFIG}"; then
 RODBC_CPPFLAGS='odbc_config --cflags'
```

```
fi
            CPPFLAGS="${CPPFLAGS} ${RODBC_CPPFLAGS}"
            dnl Check the headers can be found
            AC_CHECK_HEADERS(sql.h sqlext.h)
            if test "${ac_cv_header_sql_h}" = no ||
               test "${ac_cv_header_sqlext_h}" = no; then
               AC_MSG_ERROR("ODBC headers sql.h and sqlext.h not found")
            fi
            dnl search for a library containing an ODBC function
            if test [ -n "${odbc_mgr}" ] ; then
              AC_SEARCH_LIBS(SQLTables, ${odbc_mgr},
                  AC_MSG_ERROR("ODBC driver manager ${odbc_mgr} not found"))
              AC_SEARCH_LIBS(SQLTables, odbc odbc32 iodbc,
                  AC_MSG_ERROR("no ODBC driver manager found"))
            fi
            dnl for 64-bit ODBC need SQL[U]LEN, and it is unclear where they are defined.
            AC_CHECK_TYPES([SQLLEN, SQLULEN], , , [# include <sql.h>])
            dnl for unixODBC header
            AC_CHECK_SIZEOF(long, 4)
            dnl substitute RODBC_CPPFLAGS and LIBS
            AC_SUBST(RODBC_CPPFLAGS)
            AC_SUBST(LIBS)
            AC_CONFIG_HEADERS([src/config.h])
            dnl and do substitution in the src/Makevars.in and src/config.h
            AC_CONFIG_FILES([src/Makevars])
            AC_OUTPUT
where src/Makevars.in would be simply
            PKG_CPPFLAGS = @RODBC_CPPFLAGS@
            PKG_LIBS = @LIBS@
```

A user can then be advised to specify the location of the ODBC driver manager files by options like (lines broken for easier reading)

```
R CMD INSTALL \
    --configure-args='--with-odbc-include=/opt/local/include \
    --with-odbc-lib=/opt/local/lib --with-odbc-manager=iodbc' \
RODBC
```

or by setting the environment variables ODBC_INCLUDE and ODBC_LIBS.

1.2.3 Using F95 code

R assumes that source files with extension .f are FORTRAN 77, and passes them to the compiler specified by 'F77'. On most but not all platforms that compiler will accept Fortran 90/95 code: some platforms have a separate Fortran 90/95 compiler and a few (by now quite rare³²) platforms have no Fortran 90/95 support.

This means that portable packages need to be written in correct FORTRAN 77, which will also be valid Fortran 95. See https://developer.R-project.org/Portability.html for reference resources. In particular, free source form F95 code is not portable.

On some systems an alternative F95 compiler is available: from the gcc family this might be gfortran or g95. Configuring R will try to find a compiler which (from its name) appears to be a Fortran 90/95 compiler, and set it in macro 'FC'. Note that it does not check that such a compiler is fully (or even partially) compliant with Fortran 90/95. Packages making use of

 $^{^{32}\,}$ Cygwin used g77 up to 2011, and some pre-built versions of R for Unix OSes still do.

Fortran 90/95 features should use file extension .f90 or .f95 for the source files: the variable PKG_FCFLAGS specifies any special flags to be used. There is no guarantee that compiled Fortran 90/95 code can be mixed with any other type of compiled code, nor that a build of R will have support for such packages.

Some (but not) all compilers specified by the 'FC' macro will accept Fortran 2003 or 2008 code: such code should still use file extension .f90 or .f95. For platforms using gfortran, you may need to include -std=f2003 or -std=f2008 in PKG_FCFLAGS: the default is 'GNU Fortran', Fortran 95 with non-standard extensions. The Solaris f95 compiler 'accepts some Fortran 2003 features'.

Modern versions of Fortran support modules, whereby compiling one source file creates a module file which is then included in others. (Module files typically have a .mod extension: they do depend on the compiler used and so should never be included in a package.) This creates a dependence which make will not know about and often causes installation with a parallel make to fail. Thus it is necessary to add explicit dependencies to src/Makevars to tell make the constraints on the order of compilation. For example, if file iface.f90 creates a module 'iface' used by files cmi.f90 and dmi.f90 then src/Makevars needs to contain something like

cmi.o dmi.o: iface.o

1.2.4 Using C++11 code

R can be built without a C++ compiler although one is available (but not necessarily installed) on all known R platforms. For full portability across platforms, all that can be assumed is approximate support for the C++98 standard (the widely used g++ deviates considerably from the standard). Some compilers have a concept of 'C++03' ('essentially a bug fix') or 'C++ Technical Report 1' (TR1), an optional addition to the 'C++03' revision which was published in 2007. A revised standard was published in 2011 and compilers with pretty much complete implementations are becoming available. C++11 added all of the C99 features which are not otherwise implemented in C++, and C++ compilers commonly accept C99 extensions to C++98. A minor update to C++11 (often known as C++14) was approved in August 2014.

What standard a C++ compiler aims to support can be hard to determine. As from version 6 (to be released in 2016), g++ will default to C++14: earlier versions aim to support C++03 with many extensions (including support for TR1). clang with its native³³ libcxx headers and library includes many C++11 features, and does not support TR1.

Since version 3.1.0, R has provided support for C++11 in packages in addition to C++98. This support is not uniform across platforms as it depends on the capabilities of the compiler (see below). When R is configured, it will determine whether the C++ compiler supports C++11 and which compiler flags, if any, are required to enable C++11 support. For example, recent versions of g++ or clang++ accept the compiler flag -std=c++11, and earlier versions support a flag -std=c++0x, but the latter only provides partial support for the C++11 standard.

In order to use C++11 code in a package, the package's Makevars file (or Makevars.win on Windows) should include the line

```
CXX\_STD = CXX11
```

Compilation and linking will then be done with the C++11 compiler. If any other value is given to the 'CXX_STD' macro it will be ignored. (Further options may become available in the future as the C++ standard evolves.)

Packages without a Makevars file may specify that they require C++11 by including 'C++11' in the 'SystemRequirements' field of the DESCRIPTION file, e.g.

SystemRequirements: C++11

³³ Some distributions, notably Debian, have supplied a build of clang with g++'s headers and library. Conversely, Apple's command named g++ is based on clang using libcxx.

If a package does have a Makevars[.win] file then setting the make variable 'CXX_STD' is preferred, as it allows R CMD SHLIB to work correctly in the package's src directory.

The C++11 compiler will be used systematically by R for all C++ code if the environment variable USE_CXX1X is defined (with any value). Hence this environment variable should be defined when invoking R CMD SHLIB in the absence of a Makevars file (or Makevars.win on Windows) if a C++11 compiler is required.

Further control over compilation of C++11 code can be obtained by specifying the macros 'CXX1X' and 'CXX1XSTD' when R is configured³⁴, or in a personal or site Makevars file. See Section "Customizing package compilation" in R Installation and Administration. If C++11 support is not available then these macros are both empty. Otherwise, 'CXX1X' defaults to the same value as the C++ compiler 'CXX' and the flag 'CXX1XSTD' defaults to -std=c++11 or -std=c++0x (the latter on Windows). It is possible to specify 'CXX1X' to be a distinct compiler just for C++11-using packages, e.g. g++ on Solaris. Note however that different C++ compilers (and even different versions of the same compiler) often differ in their ABI so their outputs can rarely be mixed. By setting 'CXX1XSTD' it is also possible to choose a different dialect of the standard, such as -std=gnu++11, or enable support for the 2014 revision using something like -std=c++14 or -std=c++1y.

As noted above, support for C++11 varies across platforms. The default compiler on Windows is GCC 4.6.x and supports the -std=c++0x flag and some C++11 features (see https://gcc.gnu.org/gcc-4.6/cxx0x_status.html). On some platforms, it may be possible or necessary to select a different compiler for C++11, via personal or site Makevars files.

There is no guarantee that C++11 can be used in a package in combination with any other compiled language (even C), as the C++11 compiler may be incompatible with the native compilers for the platform. (There are known problems mixing C++11 with Fortran.)

If a package using C++11 has a configure script it is essential that it selects the correct compiler, *via* something like

```
CXX1X='"${R_HOME}/bin/R" CMD config CXX1X'
CXX1XSTD='"${R_HOME}/bin/R" CMD config CXX1XSTD'
CXX="$(CXX1X) $(CXX1XSTD)"
CXXFLAGS='"${R_HOME}/bin/R" CMD config CXX1XFLAGS'
AC_LANG(C++)
```

(paying attention to all the quotes required).

1.3 Checking and building packages

Before using these tools, please check that your package can be installed (which checked it can be loaded). R CMD check will *inter alia* do this, but you may get more detailed error messages doing the install directly.

If your package specifies an encoding in its DESCRIPTION file, you should run these tools in a locale which makes use of that encoding: they may not work at all or may work incorrectly in other locales (although UTF-8 locales will most likely work).

Note: R CMD check and R CMD build run R processes with --vanilla in which none of the user's startup files are read. If you need R_LIBS set (to find packages in a non-standard library) you can set it in the environment: also you can use the check and build environment files (as specified by the environment variables R_CHECK_ENVIRON and R_BUILD_ENVIRON; if unset, files 35 ~/.R/check.Renviron

³⁴ For details of these and related macros, see file config.site in the R sources.

³⁵ On systems which use sub-architectures, architecture-specific versions such as ~/.R/check.Renviron.i386 take precedence.

and ~/.R/build.Renviron are used) to set environment variables when using these utilities.

Note to Windows users: R CMD build may make use of the Windows toolset (see the "R Installation and Administration" manual) if present and in your path, and it is required for packages which need it to install (including those with configure.win or cleanup.win scripts or a src directory) and e.g. need vignettes built.

You may need to set the environment variable TMPDIR to point to a suitable writable directory with a path not containing spaces — use forward slashes for the separators. Also, the directory needs to be on a case-honouring file system (some network-mounted file systems are not).

1.3.1 Checking packages

Using R CMD check, the R package checker, one can test whether *source* R packages work correctly. It can be run on one or more directories, or compressed package tar archives with extension .tar.gz, .tgz, .tar.bz2 or .tar.xz.

It is strongly recommended that the final checks are run on a tar archive prepared by R CMD build.

This runs a series of checks, including

- 1. The package is installed. This will warn about missing cross-references and duplicate aliases in help files.
- 2. The file names are checked to be valid across file systems and supported operating system platforms.
- 3. The files and directories are checked for sufficient permissions (Unix-alikes only).
- 4. The files are checked for binary executables, using a suitable version of file if available³⁶. (There may be rare false positives.)
- 5. The DESCRIPTION file is checked for completeness, and some of its entries for correctness. Unless installation tests are skipped, checking is aborted if the package dependencies cannot be resolved at run time. (You may need to set R_LIBS in the environment if dependent packages are in a separate library tree.) One check is that the package name is not that of a standard package, nor one of the defunct standard packages ('ctest', 'eda', 'lqs', 'mle', 'modreg', 'mva', 'nls', 'stepfun' and 'ts'). Another check is that all packages mentioned in library or requires or from which the NAMESPACE file imports or are called via:: or :: are listed (in 'Depends', 'Imports', 'Suggests'): this is not an exhaustive check of the actual imports.
- 6. Available index information (in particular, for demos and vignettes) is checked for completeness.
- 7. The package subdirectories are checked for suitable file names and for not being empty. The checks on file names are controlled by the option --check-subdirs=value. This defaults to 'default', which runs the checks only if checking a tarball: the default can be overridden by specifying the value as 'yes' or 'no'. Further, the check on the src directory is only run if the package does not contain a configure script (which corresponds to the value 'yes-maybe') and there is no src/Makefile or src/Makefile.in.

To allow a configure script to generate suitable files, files ending in '.in' will be allowed in the R directory.

A warning is given for directory names that look like R package check directories – many packages have been submitted to CRAN containing these.

A suitable file.exe is part of the Windows toolset: it checks for gfile if a suitable file is not found: the latter is available in the OpenCSW collection for Solaris at http://www.opencsw.org. The source repository is ftp://ftp.astron.com/pub/file/.

- 8. The R files are checked for syntax errors. Bytes which are non-ASCII are reported as warnings, but these should be regarded as errors unless it is known that the package will always be used in the same locale.
- 9. It is checked that the package can be loaded, first with the usual default packages and then only with package base already loaded. It is checked that the namespace this can be loaded in an empty session with only the base namespace loaded. (Namespaces and packages can be loaded very early in the session, before the default packages are available, so packages should work then.)
- 10. The R files are checked for correct calls to library.dynam. Package startup functions are checked for correct argument lists and (incorrect) calls to functions which modify the search path or inappropriately generate messages. The R code is checked for possible problems using codetools (https://CRAN.R-project.org/package=codetools). In addition, it is checked whether S3 methods have all arguments of the corresponding generic, and whether the final argument of replacement functions is called 'value'. All foreign function calls (.C, .Fortran, .Call and .External calls) are tested to see if they have a PACKAGE argument, and if not, whether the appropriate DLL might be deduced from the namespace of the package. Any other calls are reported. (The check is generous, and users may want to supplement this by examining the output of tools::checkFF("mypkg", verbose=TRUE), especially if the intention were to always use a PACKAGE argument)
- 11. The Rd files are checked for correct syntax and metadata, including the presence of the mandatory fields (\name, \alias, \title and \description). The Rd name and title are checked for being non-empty, and there is a check for missing cross-references (links).
- 12. A check is made for missing documentation entries, such as undocumented user-level objects in the package.
- 13. Documentation for functions, data sets, and S4 classes is checked for consistency with the corresponding code.
- 14. It is checked whether all function arguments given in \usage sections of Rd files are documented in the corresponding \arguments section.
- 15. The data directory is checked for non-ASCII characters and for the use of reasonable levels of compression.
- 16. C, C++ and FORTRAN source and header files³⁷ are tested for portable (LF-only) line endings. If there is a Makefile or Makefile.in or Makevars or Makevars.in file under the src directory, it is checked for portable line endings and the correct use of '\$(BLAS_LIBS)' and '\$(LAPACK_LIBS)'
 - Compiled code is checked for symbols corresponding to functions which might terminate R or write to stdout/stderr instead of the console. Note that the latter might give false positives in that the symbols might be pulled in with external libraries and could never be called. Windows³⁸ users should note that the Fortran and C++ runtime libraries are examples of such external libraries.
- 17. Some checks are made of the contents of the inst/doc directory. These always include checking for files that look like leftovers, and if suitable tools (such as qpdf) are available, checking that the PDF documentation is of minimal size.
- 18. The examples provided by the package's documentation are run. (see Chapter 2 [Writing R documentation files], page 58, for information on using \examples to create executable example code.) If there is a file tests/Examples/pkg-Ex.Rout.save, the output of running the examples is compared to that file.

³⁷ An exception is made for subdirectories with names starting 'win' or 'Win'.

³⁸ on most other platforms such runtime libraries are dynamic, but static libraries are currently used on Windows because the toolchain is not a standard part of the OS.

Of course, released packages should be able to run at least their own examples. Each example is run in a 'clean' environment (so earlier examples cannot be assumed to have been run), and with the variables T and F redefined to generate an error unless they are set in the example: See Section "Logical vectors" in An Introduction to R.

- 19. If the package sources contain a tests directory then the tests specified in that directory are run. (Typically they will consist of a set of .R source files and target output files .Rout.save.) Please note that the comparison will be done in the end user's locale, so the target output files should be ASCII if at all possible. (The command line option --test-dir=foo may be used to specify tests in a non-standard location. For example, unusually slow tests could be placed in inst/slowTests and then R CMD check --test-dir=inst/slowTests would be used to run them. Other names that have been suggested are, for example, inst/testWithOracle for tests that require Oracle to be installed, inst/randomTests for tests which use random values and may occasionally fail by chance, etc.)
- 20. The code in package vignettes (see Section 1.4 [Writing package vignettes], page 35) is executed, and the vignette PDFs re-made from their sources as a check of completeness of the sources (unless there is a 'BuildVignettes' field in the package's DESCRIPTION file with a false value). If there is a target output file .Rout.save in the vignette source directory, the output from running the code in that vignette is compared with the target output file and any differences are reported (but not recorded in the log file). (If the vignette sources are in the deprecated location inst/doc, do mark such target output files to not be installed in .Rinstignore.)

If there is an error³⁹ in executing the R code in vignette foo.ext, a log file foo.ext.log is created in the check directory. The vignette PDFs are re-made in a copy of the package sources in the vign_test subdirectory of the check directory, so for further information on errors look in directory pkgname/vign_test/vignettes. (It is only retained if there are errors or if environment variable _R_CHECK_CLEAN_VIGN_TEST_ is set to a false value.)

21. The PDF version of the package's manual is created (to check that the Rd files can be converted successfully). This needs LATEX and suitable fonts and LATEX packages to be installed. See Section "Making the manuals" in R Installation and Administration.

All these tests are run with collation set to the C locale, and for the examples and tests with environment variable LANGUAGE=en: this is to minimize differences between platforms.

Use *R CMD check* --help to obtain more information about the usage of the R package checker. A subset of the checking steps can be selected by adding command-line options. It also allows customization by setting environment variables _R_CHECK_*_:, as described in Section "Tools" in *R Internals*: a set of these customizations similar to those used by CRAN can be selected by the option --as-cran (which works best if Internet access is available). Some Windows users may need to set environment variable R_WIN_NO_JUNCTIONS to a non-empty value. The test of cyclic declarations⁴⁰in DESCRIPTION files needs repositories (including CRAN) set: do this in ~/.Rprofile, by e.g.

options(repos = c(CRAN="https://cran.r-project.org"))

One check customization which can be revealing is

_R_CHECK_CODETOOLS_PROFILE_="suppressLocalUnused=FALSE"

which reports unused local assignments. Not only does this point out computations which are unnecessary because their results are unused, it also can show errors. (Two such are to intend

³⁹ or if option --use-valgrind is used or environment variable _R_CHECK_ALWAYS_LOG_VIGNETTE_OUTPUT_ is set to a true value or if there are differences from a target output file

⁴⁰ For example, in early 2014 gdata (https://CRAN.R-project.org/package=gdata) declared 'Imports: gtools' and gtools (https://CRAN.R-project.org/package=gtools) declared 'Imports: gdata'.

to update an object by assigning a value but mistype its name or assign in the wrong scope, for example using <- where <<- was intended.) This can give false positives, most commonly because of non-standard evaluation for formulae and because the intention is to return objects in the environment of a function for later use.

Complete checking of a package which contains a file README.md needs pandoc installed: see http://johnmacfarlane.net/pandoc/installing.html. This should be reasonably current: at the time of writing CRAN used version 1.12.4.2 to process these files.

You do need to ensure that the package is checked in a suitable locale if it contains non-ASCII characters. Such packages are likely to fail some of the checks in a C locale, and R CMD check will warn if it spots the problem. You should be able to check any package in a UTF-8 locale (if one is available). Beware that although a C locale is rarely used at a console, it may be the default if logging in remotely or for batch jobs.

Multiple sub-architectures: On systems which support multiple sub-architectures (principally Windows), R CMD check will install and check a package which contains compiled code under all available sub-architectures. (Use option --force-multiarch to force this for packages without compiled code, which are otherwise only checked under the main sub-architecture.) This will run the loading tests, examples and tests directory under each installed sub-architecture in turn, and give an error if any fail. Where environment variables (including perhaps PATH) need to be set differently for each sub-architecture, these can be set in architecture-specific files such as R_HOME/etc/i386/Renviron.site.

An alternative approach is to use R CMD check --no-multiarch to check the primary sub-architecture, and then to use something like R --arch=x86_64 CMD check --extra-arch or (Windows) /path/to/R/bin/x64/Rcmd check --extra-arch to run for each additional sub-architecture just the checks⁴¹ which differ by sub-architecture. (This approach is required for packages which are installed by R CMD INSTALL --merge-multiarch.)

Where packages need additional commands to install all the sub-architectures these can be supplied by e.g. --install-args=--force-biarch.

1.3.2 Building package tarballs

Packages may be distributed in source form as "tarballs" (.tar.gz files) or in binary form. The source form can be installed on all platforms with suitable tools and is the usual form for Unix-like systems; the binary form is platform-specific, and is the more common distribution form for the Windows and OS X platforms.

Using R CMD build, the R package builder, one can build R package tarballs from their sources (for example, for subsequent release).

Prior to actually building the package in the standard gzipped tar file format, a few diagnostic checks and cleanups are performed. In particular, it is tested whether object indices exist and can be assumed to be up-to-date, and C, C++ and FORTRAN source files and relevant makefiles in a src directory are tested and converted to LF line-endings if necessary.

Run-time checks whether the package works correctly should be performed using R CMD check prior to invoking the final build procedure.

To exclude files from being put into the package, one can specify a list of exclude patterns in file .Rbuildignore in the top-level source directory. These patterns should be Perl-like regular expressions (see the help for regexp in R for the precise details), one per line, to be matched case-insensitively 42 against the file and directory names relative to the top-level package source

 $^{^{41}\,}$ loading, examples, tests, running vignette code

 $^{^{42}}$ on all platforms from R 3.1.0.

directory. In addition, directories from source control systems⁴³ or from eclipse⁴⁴, directories with names ending .Rcheck or Old or old and files GNUMakefile⁴⁵, Read-and-delete-me or with base names starting with '.#', or starting and ending with '#', or ending in '~', '.bak' or '.swp', are excluded by default. In addition, those files in the R, demo and man directories which are flagged by R CMD check as having invalid names will be excluded.

Use R CMD build --help to obtain more information about the usage of the R package builder.

Unless R CMD build is invoked with the --no-build-vignettes option (or the package's DESCRIPTION contains 'BuildVignettes: no' or similar), it will attempt to (re)build the vignettes (see Section 1.4 [Writing package vignettes], page 35) in the package. To do so it installs the current package into a temporary library tree, but any dependent packages need to be installed in an available library tree (see the Note: at the top of this section).

Similarly, if the .Rd documentation files contain any \Sexpr macros (see Section 2.12 [Dynamic pages], page 72), the package will be temporarily installed to execute them. Post-execution binary copies of those pages containing build-time macros will be saved in build/partial.rdb. If there are any install-time or render-time macros, a .pdf version of the package manual will be built and installed in the build subdirectory. (This allows CRAN or other repositories to display the manual even if they are unable to install the package.) This can be suppressed by the option --no-manual or if package's DESCRIPTION contains 'BuildManual: no' or similar.

One of the checks that R CMD build runs is for empty source directories. These are in most (but not all) cases unintentional, if they are intentional use the option --keep-empty-dirs (or set the environment variable _R_BUILD_KEEP_EMPTY_DIRS_ to 'TRUE', or have a 'BuildKeepEmpty' field with a true value in the DESCRIPTION file).

The --resave-data option allows saved images (.rda and .RData files) in the data directory to be optimized for size. It will also compress tabular files and convert .R files to saved images. It can take values no, gzip (the default if this option is not supplied, which can be changed by setting the environment variable _R_BUILD_RESAVE_DATA_) and best (equivalent to giving it without a value), which chooses the most effective compression. Using best adds a dependence on R (>= 2.10) to the DESCRIPTION file if bzip2 or xz compression is selected for any of the files. If this is thought undesirable, --resave-data=gzip (which is the default if that option is not supplied) will do what compression it can with gzip. A package can control how its data is resaved by supplying a 'BuildResaveData' field (with one of the values given earlier in this paragraph) in its DESCRIPTION file.

The --compact-vignettes option will run tools::compactPDF over the PDF files in inst/doc (and its subdirectories) to losslessly compress them. This is not enabled by default (it can be selected by environment variable _R_BUILD_COMPACT_VIGNETTES_) and needs qpdf (http://qpdf.sourceforge.net/) to be available.

It can be useful to run R CMD check --check-subdirs=yes on the built tarball as a final check on the contents.

Where a non-POSIX file system is in use which does not utilize execute permissions, some care is needed with permissions. This applies on Windows and to e.g. FAT-formatted drives and SMB-mounted file systems on other OSes. The 'mode' of the file recorded in the tarball will be whatever file.info() returns. On Windows this will record only directories as having execute permission and on other OSes it is likely that all files have reported 'mode' 0777. A particular issue is packages being built on Windows which are intended to contain executable scripts such as configure and cleanup: R CMD build ensures those two are recorded with execute permission.

 $^{^{43}}$ called CVS or .svn or .arch-ids or .bzr or .git (but not files called .git) or .hg.

 $^{^{44}}$ called .metadata.

 $^{^{45}}$ which is an error: GNU make uses GNUmakefile.

Directory build of the package sources is reserved for use by R CMD build: it contains information which may not easily be created when the package is installed, including index information on the vignettes and, rarely, information on the help pages and perhaps a copy of the PDF reference manual (see above).

1.3.3 Building binary packages

Binary packages are compressed copies of installed versions of packages. They contain compiled shared libraries rather than C, C++ or Fortran source code, and the R functions are included in their installed form. The format and filename are platform-specific; for example, a binary package for Windows is usually supplied as a .zip file, and for the OS X platform the default binary package file extension is .tgz.

The recommended method of building binary packages is to use

R CMD INSTALL --build pkg where pkg is either the name of a source tarball (in the usual .tar.gz format) or the location of the directory of the package source to be built. This operates by first installing the package and then packing the installed binaries into the appropriate binary package file for the particular platform.

By default, R CMD INSTALL --build will attempt to install the package into the default library tree for the local installation of R. This has two implications:

- If the installation is successful, it will overwrite any existing installation of the same package.
- The default library tree must have write permission; if not, the package will not install and the binary will not be created.

To prevent changes to the present working installation or to provide an install location with write access, create a suitably located directory with write access and use the -1 option to build the package in the chosen location. The usage is then

R CMD INSTALL -1 location --build pkg

where location is the chosen directory with write access. The package will be installed as a subdirectory of location, and the package binary will be created in the current directory.

Other options for R CMD INSTALL can be found using R CMD INSTALL --help, and platform-specific details for special cases are discussed in the platform-specific FAQs.

Finally, at least one web-based service is available for building binary packages from (checked) source code: WinBuilder (see http://win-builder.R-project.org/) is able to build Windows binaries. Note that this is intended for developers on other platforms who do not have access to Windows but wish to provide binaries for the Windows platform.

1.4 Writing package vignettes

In addition to the help files in Rd format, R packages allow the inclusion of documents in arbitrary other formats. The standard location for these is subdirectory <code>inst/doc</code> of a source package, the contents will be copied to subdirectory <code>doc</code> when the package is installed. Pointers from package help indices to the installed documents are automatically created. Documents in <code>inst/doc</code> can be in arbitrary format, however we strongly recommend providing them in PDF format, so users on almost all platforms can easily read them. To ensure that they can be accessed from a browser (as an HTML index is provided), the file names should start with an ASCII letter and be comprised entirely of ASCII letters or digits or hyphen or underscore.

A special case is *package vignettes*. Vignettes are documents in PDF or HTML format obtained from plain text literate source files from which R knows how to extract R code and create output (in PDF/HTML or intermediate (La)TEX). Vignette engines do this work, using "tangle" and "weave" functions respectively. Sweave, provided by the R distribution, is the default engine. Since R version 3.0.0, other vignette engines besides Sweave are supported; see Section 1.4.2 [Non-Sweave vignettes], page 38.

Package vignettes have their sources in subdirectory vignettes of the package sources. Note that the location of the vignette sources only affects R CMD build and R CMD check: the tarball built by R CMD build includes in inst/doc the components intended to be installed.

Sweave vignette sources are normally given the file extension .Rnw or .Rtex, but for historical reasons extensions 46 .Snw and .Stex are also recognized. Sweave allows the integration of LATEX documents: see the Sweave help page in R and the Sweave vignette in package utils for details on the source document format.

Package vignettes are tested by R CMD check by executing all R code chunks they contain (except those marked for non-evaluation, e.g., with option eval=FALSE for Sweave). The R working directory for all vignette tests in R CMD check is a copy of the vignette source directory. Make sure all files needed to run the R code in the vignette (data sets, ...) are accessible by either placing them in the inst/doc hierarchy of the source package or by using calls to system.file(). All other files needed to re-make the vignettes (such as LATEX style files, BibTEX input files and files for any figures not created by running the code in the vignette) must be in the vignette source directory.

R CMD build will automatically⁴⁷ create the (PDF or HTML versions of the) vignettes in inst/doc for distribution with the package sources. By including the vignette outputs in the package sources it is not necessary that these can be re-built at install time, i.e., the package author can use private R packages, screen snapshots and LATEX extensions which are only available on his machine.⁴⁸

By default R CMD build will run Sweave on all Sweave vignette source files in vignettes. If Makefile is found in the vignette source directory, then R CMD build will try to run make after the Sweave runs, otherwise texi2pdf is run on each .tex file produced.

The first target in the Makefile should take care of both creation of PDF/HTML files and cleaning up afterwards (including after Sweave), i.e., delete all files that shall not appear in the final package archive. Note that if the make step runs R it needs to be careful to respect the environment values of R_LIBS and R_HOME⁴⁹. Finally, if there is a Makefile and it has a 'clean:' target, make clean is run.

All the usual *caveats* about including a Makefile apply. It must be portable (no GNU extensions), use LF line endings and must work correctly with a parallel make: too many authors have written things like

```
## BAD EXAMPLE
all: pdf clean

pdf: ABC-intro.pdf ABC-details.pdf
%.pdf: %.tex
texi2dvi --pdf $*

clean:
rm *.tex ABC-details-*.pdf
```

which will start removing the source files whilst pdflatex is working.

Metadata lines can be placed in the source file, preferably in LATEX comments in the preamble. One such is a \V ignetteIndexEntry of the form

⁴⁶ and to avoid problems with case-insensitive file systems, lower-case versions of all these extensions.

 $^{^{\}rm 47}$ unless inhibited by using 'BuildVignettes: no' in the <code>DESCRIPTION</code> file.

⁴⁸ provided the conditions of the package's license are met: many, including CRAN, see the omission of source components as incompatible with an Open Source license.

⁴⁹ R_HOME/bin is prepended to the PATH so that references to R or Rscript in the Makefile do make use of the currently running version of R.

%\VignetteIndexEntry{Using Animal}

Others you may see are \VignettePackage (currently ignored), \VignetteDepends and \VignetteKeyword (which replaced \VignetteKeywords). These are processed at package installation time to create the saved data frame Meta/vignette.rds, but only the \VignetteIndexEntry and \VignetteKeyword statements are currently used. The \VignetteEngine statement is described in Section 1.4.2 [Non-Sweave vignettes], page 38.

At install time an HTML index for all vignettes in the package is automatically created from the \VignetteIndexEntry statements unless a file index.html exists in directory inst/doc. This index is linked from the HTML help index for the package. If you do supply a inst/doc/index.html file it should contain relative links only to files under the installed doc directory, or perhaps (not really an index) to HTML help files or to the DESCRIPTION file, and be valid HTML as confirmed via the W3C Markup Validation Service (https://validator.w3.org) or Validator.nu (https://validator.nu/).

Sweave/Stangle allows the document to specify the split=TRUE option to create a single R file for each code chunk: this will not work for vignettes where it is assumed that each vignette source generates a single file with the vignette extension replaced by .R.

Do watch that PDFs are not too large – one in a CRAN package was 72MB! This is usually caused by the inclusion of overly detailed figures, which will not render well in PDF viewers. Sometimes it is much better to generate fairly high resolution bitmap (PNG, JPEG) figures and include those in the PDF document.

When R CMD build builds the vignettes, it copies these and the vignette sources from directory vignettes to inst/doc. To install any other files from the vignettes directory, include a file vignettes/.install_extras which specifies these as Perl-like regular expressions on one or more lines. (See the description of the .Rinstignore file for full details.)

1.4.1 Encodings and vignettes

Vignettes will in general include descriptive text, R input, R output and figures, IATEX include files and bibliographic references. As any of these may contain non-ASCII characters, the handling of encodings can become very complicated.

The vignette source file should be written in ASCII or contain a declaration of the encoding (see below). This applies even to comments within the source file, since vignette engines process comments to look for options and metadata lines. When an engine's weave and tangle functions are called on the vignette source, it will be converted to the encoding of the current R session.

Stangle() will produce an R code file in the current locale's encoding: for a non-ASCII vignette what that is recorded in a comment at the top of the file.

Sweave() will produce a .tex file in the current encoding, or in UTF-8 if that is declared. Non-ASCII encodings need to be declared to LATEX via a line like

\usepackage[utf8]{inputenc}

(It is also possible to use the more recent 'inputenx' LATEX package.) For files where this line is not needed (e.g. chapters included within the body of a larger document, or non-Sweave vignettes), the encoding may be declared using a comment like

%!\VignetteEncoding{UTF-8}

If the encoding is UTF-8, this can also be declared using the declaration

%!\SweaveUTF8

If no declaration is given in the vignette, it will be assumed to be in the encoding declared for the package. If there is no encoding declared in either place, then it is an error to use non-ASCII characters in the vignette.

In any case, be aware that IATEX may require the 'usepackage' declaration.

Sweave() will also parse and evaluate the R code in each chunk. The R output will also be in the current locale (or UTF-8 if so declared), and should be covered by the 'inputenc' declaration. One thing people often forget is that the R output may not be ASCII even for ASCII R sources, for many possible reasons. One common one is the use of 'fancy' quotes: see the R help on sQuote: note carefully that it is not portable to declare UTF-8 or CP1252 to cover such quotes, as their encoding will depend on the locale used to run Sweave(): this can be circumvented by setting options(useFancyQuotes="UTF-8") in the vignette.

The final issue is the encoding of figures – this applies only to PDF figures and not PNG etc. The PDF figures will contain declarations for their encoding, but the Sweave option pdf.encoding may need to be set appropriately: see the help for the pdf() graphics device.

As a real example of the complexities, consider the **fortunes** (https://CRAN.R-project.org/package=fortunes) package version '1.4-0'. That package did not have a declared encoding, and its vignette was in ASCII. However, the data it displays are read from a UTF-8 CSV file and will be assumed to be in the current encoding, so fortunes.tex will be in UTF-8 in any locale. Had read.table been told the data were UTF-8, fortunes.tex would have been in the locale's encoding.

1.4.2 Non-Sweave vignettes

R 3.0.0 and later allow vignettes in formats other than Sweave by means of "vignette engines". For example knitr (https://CRAN.R-project.org/package=knitr) version 1.1 or later can create .tex files from a variation on Sweave format, and .html files from a variation on "markdown" format. These engines replace the Sweave() function with other functions to convert vignette source files into LATEX files for processing into .pdf, or directly into .pdf or .html files. The Stangle() function is replaced with a function that extracts the R source from a vignette.

R recognizes non-Sweave vignettes using filename extensions specified by the engine. For example, the **knitr** (https://CRAN.R-project.org/package=knitr) package supports the extension .Rmd (standing for "R markdown"). The user indicates the vignette engine within the vignette source using a \VignetteEngine line, for example

```
%\VignetteEngine{knitr::knitr}
```

This specifies the name of a package and an engine to use in place of Sweave in processing the vignette. As Sweave is the only engine supplied with the R distribution, the package providing any other engine must be specified in the 'VignetteBuilder' field of the package DESCRIPTION file, and also specified in the 'Suggests', 'Imports' or 'Depends' field (since its namespace must be available to build or check your package). If more than one package is specified as a builder, they will be searched in the order given there. The utils package is always implicitly appended to the list of builder packages, but may be included earlier to change the search order.

Note that a package with non-Sweave vignettes should always have a 'VignetteBuilder' field in the DESCRIPTION file, since this is how R CMD check recognizes that there are vignettes to be checked: packages listed there are required when the package is checked.

The vignette engine can produce .tex, .pdf, or .html files as output. If it produces .tex files, R will call texi2pdf to convert them to .pdf for display to the user (unless there is a Makefile in the vignettes directory).

Package writers who would like to supply vignette engines need to register those engines in the package .onLoad function. For example, that function could make the call

(The actual registration in **knitr** (https://CRAN.R-project.org/package=knitr) is more complicated, because it supports other input formats.) See the ?tools::vignetteEngine help topic for details on engine registration.

1.5 Package namespaces

R has a namespace management system for code in packages. This system allows the package writer to specify which variables in the package should be *exported* to make them available to package users, and which variables should be *imported* from other packages.

The namespace for a package is specified by the NAMESPACE file in the top level package directory. This file contains namespace directives describing the imports and exports of the namespace. Additional directives register any shared objects to be loaded and any S3-style methods that are provided. Note that although the file looks like R code (and often has R-style comments) it is not processed as R code. Only very simple conditional processing of if statements is implemented.

Packages are loaded and attached to the search path by calling library or require. Only the exported variables are placed in the attached frame. Loading a package that imports variables from other packages will cause these other packages to be loaded as well (unless they have already been loaded), but they will *not* be placed on the search path by these implicit loads. Thus code in the package can only depend on objects in its own namespace and its imports (including the base namespace) being visible⁵⁰.

Namespaces are *sealed* once they are loaded. Sealing means that imports and exports cannot be changed and that internal variable bindings cannot be changed. Sealing allows a simpler implementation strategy for the namespace mechanism. Sealing also allows code analysis and compilation tools to accurately identify the definition corresponding to a global variable reference in a function body.

The namespace controls the search strategy for variables used by functions in the package. If not found locally, R searches the package namespace first, then the imports, then the base namespace and then the normal search path.

Prior to R 2.14.0, namespaces were optional in packages: a default namespace was generated on installation in 2.14.x and 2.15.x. As from 3.0.0 a namespace is mandatory.

1.5.1 Specifying imports and exports

Exports are specified using the export directive in the NAMESPACE file. A directive of the form

specifies that the variables f and g are to be exported. (Note that variable names may be quoted, and reserved words and non-standard names such as [<-.fractions must be.)

For packages with many variables to export it may be more convenient to specify the names to export with a regular expression using exportPattern. The directive

```
exportPattern("^[^\\.]")
```

exports all variables that do not start with a period. However, such broad patterns are not recommended for production code: it is better to list all exports or use narrowly-defined groups. (This pattern applies to S4 classes.) Beware of patterns which include names starting with a period: some of these are internal-only variables and should never be exported, e.g. '.__S3MethodsTable__.' (and the code nowadays excludes known cases).

Packages implicitly import the base namespace. Variables exported from other packages with namespaces need to be imported explicitly using the directives import and importFrom. The import directive imports all exported variables from the specified package(s). Thus the directives

```
import(foo, bar)
```

⁵⁰ Note that lazy-loaded datasets are *not* in the package's namespace so need to be accessed *via* ::, e.g. survival::survexp.us.

specifies that all exported variables in the packages **foo** and **bar** are to be imported. If only some of the exported variables from a package are needed, then they can be imported using **importFrom**. The directive

```
importFrom(foo, f, g)
```

specifies that the exported variables f and g of the package foo are to be imported. Using importFrom selectively rather than import is good practice and recommended notably when importing from packages with more than a dozen exports.

To import every symbol from a package but for a few exceptions, pass the except argument to import. The directive

```
import(foo, except=c(bar, baz))
```

imports every symbol from **foo** except **bar** and **baz**. The value of **except** should evaluate to something coercible to a character vector, after substituting each symbol for its corresponding string.

It is possible to export variables from a namespace which it has imported from other namespaces: this has to be done explicitly and not *via* exportPattern.

If a package only needs a few objects from another package it can use a fully qualified variable reference in the code instead of a formal import. A fully qualified reference to the function f in package foo is of the form foo::f. This is slightly less efficient than a formal import and also loses the advantage of recording all dependencies in the NAMESPACE file (but they still need to be recorded in the DESCRIPTION file). Evaluating foo::f will cause package foo to be loaded, but not attached, if it was not loaded already—this can be an advantage in delaying the loading of a rarely used package.

Using foo:::f instead of foo::f allows access to unexported objects. This is generally not recommended, as the semantics of unexported objects may be changed by the package author in routine maintenance.

1.5.2 Registering S3 methods

The standard method for S3-style UseMethod dispatching might fail to locate methods defined in a package that is imported but not attached to the search path. To ensure that these methods are available the packages defining the methods should ensure that the generics are imported and register the methods using S3method directives. If a package defines a function print.foo intended to be used as a print method for class foo, then the directive

```
S3method(print, foo)
```

ensures that the method is registered and available for UseMethod dispatch, and the function print.foo does not need to be exported. Since the generic print is defined in base it does not need to be imported explicitly.

(Note that function and class names may be quoted, and reserved words and non-standard names such as [<- and function must be.)

It is possible to specify a third argument to S3method, the function to be used as the method, for example

```
S3method(print, check_so_symbols, .print.via.format)
```

when print.check_so_symbols is not needed.

There used to be a limit on the number of S3method directives: it was 500 prior to R 3.0.2.

1.5.3 Load hooks

There are a number of hooks called as packages are loaded, attached, detached, and unloaded. See help(".onLoad") for more details.

Since loading and attaching are distinct operations, separate hooks are provided for each. These hook functions are called .onLoad and .onAttach. They both take arguments⁵¹ libname and pkgname; they should be defined in the namespace but not exported.

Packages can use a .onDetach (as from R 3.0.0) or .Last.lib function (provided the latter is exported from the namespace) when detach is called on the package. It is called with a single argument, the full path to the installed package. There is also a hook .onUnload which is called when the namespace is unloaded (via a call to unloadNamespace, perhaps called by detach(unload = TRUE)) with argument the full path to the installed package's directory. .onUnload and .onDetach should be defined in the namespace and not exported, but .Last.lib does need to be exported.

Packages are not likely to need .onAttach (except perhaps for a start-up banner); code to set options and load shared objects should be placed in a .onLoad function, or use made of the useDynLib directive described next.

User-level hooks are also available: see the help on function setHook.

These hooks are often used incorrectly. People forget to export .Last.lib. Compiled code should be loaded in .onLoad (or *via* a useDynLb directive: see below) and unloaded in .onUnload. Do remember that a package's namespace can be loaded without the namespace being attached (e.g. by pkgname::fun) and that a package can be detached and re-attached whilst its namespace remains loaded.

1.5.4 useDynLib

A NAMESPACE file can contain one or more useDynLib directives which allows shared objects that need to be loaded.⁵² The directive

```
useDynLib(foo)
```

registers the shared object foo⁵³ for loading with library.dynam. Loading of registered object(s) occurs after the package code has been loaded and before running the load hook function. Packages that would only need a load hook function to load a shared object can use the useDynLib directive instead.

The useDynLib directive also accepts the names of the native routines that are to be used in R via the .C, .Call, .Fortran and .External interface functions. These are given as additional arguments to the directive, for example,

```
useDynLib(foo, myRoutine, myOtherRoutine)
```

By specifying these names in the useDynLib directive, the native symbols are resolved when the package is loaded and R variables identifying these symbols are added to the package's namespace with these names. These can be used in the .C, .Call, .Fortran and .External calls in place of the name of the routine and the PACKAGE argument. For instance, we can call the routine myRoutine from R with the code

```
.Call(myRoutine, x, y)
```

rather than

```
.Call("myRoutine", x, y, PACKAGE = "foo")
```

There are at least two benefits to this approach. Firstly, the symbol lookup is done just once for each symbol rather than each time the routine is invoked. Secondly, this removes any ambiguity in resolving symbols that might be present in several compiled DLLs.

 $^{^{51}}$ they will be called with two unnamed arguments, in that order.

 $^{^{52}}$ NB: this will only be read in all versions of R if the package contains R code in a R directory.

⁵³ Note that this is the basename of the shared object, and the appropriate extension (.so or .dll) will be added.

In some circumstances, there will already be an R variable in the package with the same name as a native symbol. For example, we may have an R function in the package named myRoutine. In this case, it is necessary to map the native symbol to a different R variable name. This can be done in the useDynLib directive by using named arguments. For instance, to map the native symbol name myRoutine to the R variable myRoutine_sym, we would use

```
useDynLib(foo, myRoutine_sym = myRoutine, myOtherRoutine)
```

We could then call that routine from R using the command

```
.Call(myRoutine_sym, x, y)
```

Symbols without explicit names are assigned to the R variable with that name.

In some cases, it may be preferable not to create R variables in the package's namespace that identify the native routines. It may be too costly to compute these for many routines when the package is loaded if many of these routines are not likely to be used. In this case, one can still perform the symbol resolution correctly using the DLL, but do this each time the routine is called. Given a reference to the DLL as an R variable, say dll, we can call the routine myRoutine using the expression

```
.Call(dll$myRoutine, x, y)
```

The \$ operator resolves the routine with the given name in the DLL using a call to getNativeSymbol. This is the same computation as above where we resolve the symbol when the package is loaded. The only difference is that this is done each time in the case of dll\$myRoutine.

In order to use this dynamic approach (e.g., dll\$myRoutine), one needs the reference to the DLL as an R variable in the package. The DLL can be assigned to a variable by using the variable = dllName format used above for mapping symbols to R variables. For example, if we wanted to assign the DLL reference for the DLL foo in the example above to the variable myDLL, we would use the following directive in the NAMESPACE file:

```
myDLL = useDynLib(foo, myRoutine_sym = myRoutine, myOtherRoutine)
```

Then, the R variable myDLL is in the package's namespace and available for calls such as myDLL\$dynRoutine to access routines that are not explicitly resolved at load time.

If the package has registration information (see Section 5.4 [Registering native routines], page 102), then we can use that directly rather than specifying the list of symbols again in the useDynLib directive in the NAMESPACE file. Each routine in the registration information is specified by giving a name by which the routine is to be specified along with the address of the routine and any information about the number and type of the parameters. Using the .registration argument of useDynLib, we can instruct the namespace mechanism to create R variables for these symbols. For example, suppose we have the following registration information for a DLL named myDLL:

```
static R_NativePrimitiveArgType foo_t[] = {
    REALSXP, INTSXP, STRSXP, LGLSXP
};

static R_CMethodDef cMethods[] = {
    {"foo", (DL_FUNC) &foo, 4, foo_t},
    {"bar_sym", (DL_FUNC) &bar, 0},
    {NULL, NULL, 0}
};

static R_CallMethodDef callMethods[] = {
    {"R_call_sym", (DL_FUNC) &R_call, 4},
    {"R_version_sym", (DL_FUNC) &R_version, 0},
    {NULL, NULL, 0}
```

};

Then, the directive in the NAMESPACE file

```
useDynLib(myDLL, .registration = TRUE)
```

causes the DLL to be loaded and also for the R variables foo, bar_sym, R_call_sym and R_version_sym to be defined in the package's namespace.

Note that the names for the R variables are taken from the entry in the registration information and do not need to be the same as the name of the native routine. This allows the creator of the registration information to map the native symbols to non-conflicting variable names in R, e.g. R_version to R_version_sym for use in an R function such as

```
R_version <- function()
{
    .Call(R_version_sym)
}</pre>
```

Using argument .fixes allows an automatic prefix to be added to the registered symbols, which can be useful when working with an existing package. For example, package **KernSmooth** (https://CRAN.R-project.org/package=KernSmooth) has

```
useDynLib(KernSmooth, .registration = TRUE, .fixes = "F_")
```

which makes the R variables corresponding to the FORTRAN symbols F_bkde and so on, and so avoid clashes with R code in the namespace.

1.5.5 An example

As an example consider two packages named **foo** and **bar**. The R code for package **foo** in file **foo**.R is

```
x <- 1
f <- function(y) c(x,y)
foo <- function(x) .Call("foo", x, PACKAGE="foo")
print.foo <- function(x, ...) cat("<a foo>\n")
```

Some C code defines a C function compiled into DLL foo (with an appropriate extension). The NAMESPACE file for this package is

```
useDynLib(foo)
export(f, foo)
S3method(print, foo)
```

The second package bar has code file bar.R

```
c <- function(...) sum(...)
g <- function(y) f(c(y, 7))
h <- function(y) y+9</pre>
```

and NAMESPACE file

```
import(foo)
export(g, h)
```

Calling library(bar) loads bar and attaches its exports to the search path. Package foo is also loaded but not attached to the search path. A call to g produces

```
> g(6)
[1] 1 13
```

This is consistent with the definitions of c in the two settings: in **bar** the function c is defined to be equivalent to **sum**, but in **foo** the variable c refers to the standard function c in **base**.

1.5.6 Namespaces with S4 classes and methods

Some additional steps are needed for packages which make use of formal (S4-style) classes and methods (unless these are purely used internally). The package should have Depends: methods in its DESCRIPTION file⁵⁴ and import(methods) or importFrom(methods, ...) plus any classes and methods which are to be exported need to be declared in the NAMESPACE file. For example, the stats4 package has

All S4 classes to be used outside the package need to be listed in an exportClasses directive. Alternatively, they can be specified using exportClassPattern⁵⁵ in the same style as for exportPattern. To export methods for generics from other packages an exportMethods directive can be used.

Note that exporting methods on a generic in the namespace will also export the generic, and exporting a generic in the namespace will also export its methods. If the generic function is not local to this package, either because it was imported as a generic function or because the non-generic version has been made generic solely to add S4 methods to it (as for functions such as plot in the example above), it can be declared via either or both of export or exportMethods, but the latter is clearer (and is used in the stats4 example above). In particular, for primitive functions there is no generic function, so export would export the primitive, which makes no sense. On the other hand, if the generic is local to this package, it is more natural to export the function itself using export(), and this must be done if an implicit generic is created without setting any methods for it (as is the case for AIC in stats4).

A non-local generic function is only exported to ensure that calls to the function will dispatch the methods from this package (and that is not done or required when the methods are for primitive functions). For this reason, you do not need to document such implicitly created generic functions, and undoc in package tools will not report them.

If a package uses S4 classes and methods exported from another package, but does not import the entire namespace of the other package⁵⁶, it needs to import the classes and methods explicitly, with directives

```
importClassesFrom(package, ...)
```

 $^{^{54}}$ This was necessary at least prior to R 3.0.2 as the **methods** package looked for its own R code on the search path.

This defaults to the same pattern as exportPattern: use something like exportClassPattern("^\$") to override this

⁵⁶ if it does, there will be opaque warnings about replacing imports if the classes/methods are also imported.

```
importMethodsFrom(package, ...)
```

listing the classes and functions with methods respectively. Suppose we had two small packages $\bf A$ and $\bf B$ with $\bf B$ using $\bf A$. Then they could have NAMESPACE files

```
export(f1, ng1)
exportMethods("[")
exportClasses(c1)
```

and

```
importFrom(A, ng1)
importClassesFrom(A, c1)
importMethodsFrom(A, f1)
export(f4, f5)
exportMethods(f6, "[")
exportClasses(c1, c2)
```

respectively.

Note that importMethodsFrom will also import any generics defined in the namespace on those methods.

It is important if you export S4 methods that the corresponding generics are available. You may for example need to import plot from **graphics** to make visible a function to be converted into its implicit generic. But it is better practice to make use of the generics exported by **stats4** as this enables multiple packages to unambiguously set methods on those generics.

1.6 Writing portable packages

This section contains advice on writing packages to be used on multiple platforms or for distribution (for example to be submitted to a package repository such as CRAN).

Portable packages should have simple file names: use only alphanumeric ASCII characters and period (.), and avoid those names not allowed under Windows which are mentioned above.

Many of the graphics devices are platform-specific: even X11() (aka x11()) which although emulated on Windows may not be available on a Unix-alike (and is not the preferred screen device on OS X). It is rarely necessary for package code or examples to open a new device, but if essential, 57 use dev.new().

Use R CMD build to make the release .tar.gz file.

R CMD check provides a basic set of checks, but often further problems emerge when people try to install and use packages submitted to CRAN – many of these involve compiled code. Here are some further checks that you can do to make your package more portable.

- If your package has a configure script, provide a configure.win script to be used on Windows (an empty file if no actions are needed).
- If your package has a Makevars or Makefile file, make sure that you use only portable make features. Such files should be LF-terminated⁵⁸ (including the final line of the file) and not make use of GNU extensions. (The POSIX specification is available at http://pubs.opengroup.org/onlinepubs/9699919799/utilities/make.html; anything not documented there should be regarded as an extension to be avoided.) Commonly misused GNU

⁵⁷ People use dev.new() to open a device at a particular size: that is not portable but using dev.new(noRStudioGD = TRUE) helps.

⁵⁸ Solaris make does not accept CRLF-terminated Makefiles; Solaris warns about and some other makes ignore incomplete final lines.

extensions are conditional inclusions (ifeq and the like), \${shell ...} and \${wildcard ...}, and the use of +=⁵⁹ and :=. Also, the use of \$< other than in implicit rules is a GNU extension, as is the \$^ macro. Unfortunately makefiles which use GNU extensions often run on other platforms but do not have the intended results.

The use of \${shell ...} can be avoided by using backticks, e.g.

```
PKG_CPPFLAGS = 'gsl-config --cflags'
```

which works in all versions of make known⁶⁰ to be used with R.

If you really must require GNU make, declare it in the DESCRIPTION file by

SystemRequirements: GNU make

and ensure that you use the value of environment variable MAKE (and not just make) in your scripts. (On some platforms GNU make is available under a name such as gmake, and there SystemRequirements is used to set MAKE.)

If you only need GNU make for parts of the package which are rarely needed (for example to create bibliography files under vignettes), use a file called GNUmakefile rather than Makefile as GNU make (only) will use the former.

Since the only viable make for Windows is GNU make, it is permissible to use GNU extensions in files Makevars.win or Makefile.win.

- Bash extensions also need to be avoided in shell scripts, including expressions in Makefiles (which are passed to the shell for processing). Some R platforms use strict⁶¹ Bourne shells: the R toolset on Windows and some Unix-alike OSes use ash (https://en.wikipedia.org/wiki/Almquist_shell), a rather minimal shell with few builtins. Beware of assuming that all the POSIX command-line utilities are available, especially on Windows where only a minimal set is provided for use with R. (See Section "The command line tools" in R Installation and Administration.) One particular issue is the use of echo, for which two behaviours are allowed (http://pubs.opengroup.org/onlinepubs/9699919799/utilities/echo. html) and both occur as defaults on R platforms: portable applications should not use -n (as the first argument) nor escape sequences.
- Make use of the abilities of your compilers to check the standards-conformance of your code. For example, gcc can be used with options -Wall -pedantic to alert you to potential problems. This is particularly important for C++, where g++ -Wall -pedantic will alert you to the use of GNU extensions which fail to compile on most other C++ compilers. If R was not configured accordingly, one can achieve this via personal Makevars files. See Section "Customizing package compilation" in R Installation and Administration,

Although there is a 2011 version of the C++ standard, even partial implementations are not universally available. Portable C++ code needs to follow the 1998 standard (and not use features from C99). See also Section 1.2.4 [Using C++11 code], page 28, to specify a C++11 compiler.

If you use FORTRAN 77, ftnchek (http://www.dsm.fordham.edu/~ftnchek/) provides thorough testing of conformance to the standard.

Not all common R platforms conform to the expected standards, e.g. C99 for C code. One common area of problems is the *printf functions where Windows does not support %lld, %Lf and similar formats (and has its own formats such as %I64d for 64-bit integers). It is very rare to need to output such types, and 64-bit integers can usually be converted to doubles for output.

 $^{^{59}}$ This was apparently introduced in SunOS 4, and is available elsewhere *provided* it is surrounded by spaces.

⁶⁰ GNU make, BSD make formerly in FreeBSD and OS X, AT&T make as implemented on Solaris, pmake in FreeBSD, 'Distributed Make' (dmake), part of Solaris Studio and available in other versions.

For example, test options -a and -e are not portable, and not supported in the AT&T Bourne shell used on Solaris 10/11, even though they are in the 2008 POSIX standard. Nor does Solaris support '\$(cmd)'.

• Do be very careful with passing arguments between R, C and FORTRAN code. In particular, long in C will be 32-bit on some R platforms (including 64-bit Windows), but 64-bit on most modern Unix and Linux platforms. It is rather unlikely that the use of long in C code has been thought through: if you need a longer type than int you should use a configure test for a C99 type such as int_fast64_t (and failing that, long long⁶²) and typedef your own type to be long or long long, or use another suitable type (such as size_t).

It is not safe to assume that long and pointer types are the same size, and they are not on 64-bit Windows. If you need to convert pointers to and from integers use the C99 integer types intptr_t and uintptr_t (which are defined in the header <stdint.h> and are not required to be implemented by the C99 standard).

Note that integer in FORTRAN corresponds to int in C on all R platforms.

• Under no circumstances should your compiled code ever call abort or exit⁶³: these terminate the user's R process, quite possibly including all his unsaved work. One usage that could call abort is the assert macro in C or C++ functions, which should never be active in production code. The normal way to ensure that is to define the macro NDEBUG, and R CMD INSTALL does so as part of the compilation flags. If you wish to use assert during development. you can include -UNDEBUG in PKG_CPPFLAGS. Note that your own src/Makefile or makefiles in sub-directories may also need to define NDEBUG.

This applies not only to your own code but to any external software you compile in or link to.

- Compiled code should not write to stdout or stderr and C++ and Fortran I/O should not be used. As with the previous item such calls may come from external software and may never be called, but package authors are often mistaken about that.
- Compiled code should not call the system random number generators such as rand, drand48 and random⁶⁴, but rather use the interfaces to R's RNGs described in Section 6.3 [Random numbers], page 136. In particular, if more than one package initializes the system RNG (e.g. via srand), they will interfere with each other.

Nor should the C++11 random number library be used.

- Errors in memory allocation and reading/writing outside arrays are very common causes of crashes (e.g., segfaults) on some machines. See Section 4.3 [Checking memory access], page 89, for tools which can be used to look for this.
- Many platforms will allow unsatisfied entry points in compiled code, but will crash the application (here R) if they are ever used. Some (notably Windows) will not. Looking at the output of

and checking if any of the symbols marked U is unexpected is a good way to avoid this.

• Linkers have a lot of freedom in how to resolve entry points in dynamically-loaded code, so the results may differ by platform. One area that has caused grief is packages including copies of standard system software such as libz (especially those already linked into R). In the case in point, entry point gzgets was sometimes resolved against the old version compiled into the package, sometimes against the copy compiled into R and sometimes against the system dynamic library. The only safe solution is to rename the entry points in the copy in the package. We have even seen problems with entry point name myprintf, which is a system entry point⁶⁵ on some Linux systems.

⁶² but note that long long is not a standard C++ type, and C++ compilers set up for strict checking will reject it.

 $^{^{63}}$ or where supported the variants $\tt _Exit$ and $\tt _exit.$

 $^{^{64}}$ This and srandom are in any case not portable. They are in POSIX but not in the C99 standard, and not available on Windows.

 $^{^{65}}$ in libselinux.

- Conflicts between symbols in DLLs are handled in very platform-specific ways. Good ways
 to avoid trouble are to make as many symbols as possible static (check with nm -pg), and
 to use names which are clearly tied to your package (which also helps users if anything does
 go wrong). Note that symbol names starting with R_ are regarded as part of R's namespace
 and should not be used in packages.
- It is not portable to call compiled code in R or other packages *via* .Internal, .C, .Fortran, .Call or .External, since such interfaces are subject to change without notice and will probably result in your code terminating the R process.
- Do not use (hard or symbolic) file links in your package sources. Where possible R CMD build will replace them by copies.
- If you do not yourself have a Windows system, consider submitting your source package to WinBuilder (http://win-builder.r-project.org/) before distribution.
- It is bad practice for package code to alter the search path using library, require or attach and this often does not work as intended. For alternatives, see Section 1.1.3.1 [Suggested packages], page 11, and with.
- Examples can be run interactively *via* example as well as in batch mode when checking. So they should behave appropriately in both scenarios, conditioning by interactive() the parts which need an operator or observer. For instance, progress bars⁶⁶ are only appropriate in interactive use, as is displaying help pages or calling View() (see below).
- \bullet Be careful with the order of entries in macros such as PKG_LIBS. Some linkers will re-order the entries, and behaviour can differ between dynamic and static libraries. Generally <code>-L</code> options should precede 67 the libraries (typically specified by <code>-l</code> options) to be found from those directories, and libraries are searched once in the order they are specified. Not all linkers allow a space after <code>-L</code> .
- The ar utility is often used in makefiles to make static libraries. Its modifier u is defined by POSIX but is disabled in GNU ar on some recent Linux distributions which use 'deterministic mode'. The safest way to make a static library is to first remove any existing file of that name then use ar -cr and then ranlib if needed (which is system-dependent: on most systems⁶⁸ ar always maintains a symbol table). The POSIX standard says options should be preceded by a hyphen (as in -cr), although most OSes accept them without. Note that on some systems ar -cr must have at least one file specified.
- Some people have a need to set a locale. Locale names are not portable, and e.g. 'fr_FR.utf8' is commonly used on Linux but not accepted on either Solaris or OS X. 'fr_FR.UTF-8' is more portable, being accepted on recent Linux, AIX, FreeBSD, OS X and Solaris (at least). However, some Linux distributions micro-package, so locales defined by glibc (including these examples) may not be installed.

Do be careful in what your tests (and examples) actually test. Bad practice seen in distributed packages include:

- It is not reasonable to test the time taken by a command: you cannot know how fast or how heavily loaded an R platform might be. At best you can test a ratio of times, and even that is fraught with difficulties.
- Do not test the exact format of R error messages: They change, and they can be translated. Worse, packages have tested the exact format of system error messages, which are platform-dependent and perhaps locale-dependent.

⁶⁶ except perhaps the simplest kind as used by download.file() in non-interactive use.

 $^{^{67}}$ Whereas the GNU linker reorders so -L options are processed first, the Solaris one does not.

⁶⁸ some versions of OS X did not.

• If you use functions such as View, remember that in testing there is no one to look at the output. It is better to use something like one of

```
if(interactive()) View(obj) else print(head(obj))
if(interactive()) View(obj) else str(obj)
```

• Only test the accuracy of results if you have done a formal error analysis. Things such as checking that probabilities numerically sum to one are silly: numerical tests should always have a tolerance. That the tests on your platform achieve a particular tolerance says little about other platforms. R is configured by default to make use of long doubles where available, but they may not be available or be too slow for routine use. Most R platforms use 'ix86' or 'x86_64' CPUs: these use extended precision registers on some but not all of their FPU instructions. Thus the achieved precision can depend on the compiler version and optimization flags—our experience is that 32-bit builds tend to be less precise than 64-bit ones. But not all platforms use those CPUs, and not all⁶⁹ which use them configure them to allow the use of extended precision. In particular, ARM CPUs do not (currently) have extended precision nor long doubles, and long double was 64-bit on HP/PA Linux.

If you must try to establish a tolerance empirically, configure and build R with --disable-long-double and use appropriate compiler flags (such as -ffloat-store and -fexcess-precision=standard for gcc, depending on the CPU type⁷⁰) to mitigate the effects of extended-precision calculations.

Tests which involve random inputs or non-deterministic algorithms should normally set a seed or be tested for many seeds.

1.6.1 PDF size

There are a several tools available to reduce the size of PDF files: often the size can be reduced substantially with no or minimal loss in quality. Not only do large files take up space: they can stress the PDF viewer and take many minutes to print (if they can be printed at all).

qpdf (http://qpdf.sourceforge.net/) can compress losslessly. It is fairly readily available (e.g. it has binaries for Windows and packages in Debian/Ubuntu/Fedora, and is installed as part of the CRAN OS X distribution of R). R CMD build has an option to run qpdf over PDF files under inst/doc and replace them if at least 10Kb and 10% is saved. The full path to the qpdf command can be supplied as environment variable R_QPDF (and is on the CRAN binary of R for OS X). It seems MiKTeX does not use PDF object compression and so qpdf can reduce considerably the files it outputs: MiKTeX can be overridden by code in the preamble of an Sweave or IATEX file — see how this is done for the R reference manual at https://svn.r-project.org/R/trunk/doc/manual/refman.top.

Other tools can reduce the size of PDFs containing bitmap images at excessively high resolution. These are often best re-generated (for example Sweave defaults to 300 ppi, and 100–150 is more appropriate for a package manual). These tools include Adobe Acrobat (not Reader), Apple's Preview⁷¹ and Ghostscript (which converts PDF to PDF by

```
ps2pdf options -dAutoRotatePages=/None in.pdf out.pdf
```

and suitable options might be

```
-dPDFSETTINGS=/ebook
```

-dPDFSETTINGS=/screen

 $^{^{69}}$ Not doing so is the default on Windows, overridden for the R executables. It is also the default on some Solaris compilers.

 $^{^{70}}$ These are not needed for the default compiler settings on 'x86_64' but are likely to be needed on 'ix86'.

⁷¹ Select 'Save as', and select 'Reduce file size' from the 'Quartz filter' menu': this can be accessed in other ways, for example by Automator.

; see http://www.ghostscript.com/doc/current/Ps2pdf.htm for more such and consider all the options for image downsampling). There have been examples in CRAN packages for which Ghostscript 9.06 and later produced much better reductions than 9.05 or earlier.

We come across occasionally large PDF files containing excessively complicated figures using PDF vector graphics: such figures are often best redesigned or failing that, output as PNG files.

Option --compact-vignettes to R CMD build defaults to value 'qpdf': use 'both' to try harder to reduce the size, provided you have Ghostscript available (see the help for tools::compactPDF).

1.6.2 Check timing

There are several ways to find out where time is being spent in the check process. Start by setting the environment variable <code>_R_CHECK_TIMINGS_</code> to '0'. This will report the total CPU times (not Windows) and elapsed times for installation and running examples, tests and vignettes, under each sub-architecture if appropriate. For tests and vignettes, it reports the time for each as well as the total.

Setting <code>_R_CHECK_TIMINGS_</code> to a positive value sets a threshold (in seconds elapsed time) for reporting timings.

If you need to look in more detail at the timings for examples, use option --timings to R CMD check (this is set by --as-cran). This adds a summary to the check output for all the examples with CPU or elapsed time of more than 5 seconds. It produces a file mypkg.Rcheck/mypkg-Ex.timings containing timings for each help file: it is a tab-delimited file which can be read into R for further analysis.

Timings for the tests and vignette runs are given at the bottom of the corresponding log file: note that log files for successful vignette runs are only retained if environment variable <code>_R_CHECK_ALWAYS_LOG_VIGNETTE_OUTPUT_</code> is set to a true value.

1.6.3 Encoding issues

Care is needed if your package contains non-ASCII text, and in particular if it is intended to be used in more than one locale. It is possible to mark the encoding used in the DESCRIPTION file and in .Rd files, as discussed elsewhere in this manual.

First, consider carefully if you really need non-ASCII text. Many users of R will only be able to view correctly text in their native language group (e.g. Western European, Eastern European, Simplified Chinese) and ASCII.⁷². Other characters may not be rendered at all, rendered incorrectly, or cause your R code to give an error. For .Rd documentation, marking the encoding and including ASCII transliterations is likely to do a reasonable job. The set of characters which is commonly supported is wider than it used to be around 2000, but non-Latin alphabets (Greek, Russian, Georgian, . . .) are still often problematic and those with double-width characters (Chinese, Japanese, Korean) often need specialist fonts to render correctly.

Several CRAN packages have messages in their R code in French (and a few in German). A better way to tackle this is to use the internationalization facilities discussed elsewhere in this manual.

Function showNonASCIIfile in package tools can help in finding non-ASCII bytes in files.

There is a portable way to have arbitrary text in character strings (only) in your R code, which is to supply them in Unicode as \uxxxx escapes. If there are any characters not in the current encoding the parser will encode the character string as UTF-8 and mark it as such. This applies also to character strings in datasets: they can be prepared using \uxxxx escapes or

⁷² except perhaps some special characters such as backslash and hash which may be taken over for currency symbols.

encoded in UTF-8 in a UTF-8 locale, or even converted to UTF-8 via 'iconv()'. If you do this, make sure you have 'R (>= 2.10)' (or later) in the 'Depends' field of the DESCRIPTION file.

R sessions running in non-UTF-8 locales will if possible re-encode such strings for display (and this is done by RGui on Windows, for example). Suitable fonts will need to be selected or made available⁷³ both for the console/terminal and graphics devices such as 'X11()' and 'windows()'. Using 'postscript' or 'pdf' will choose a default 8-bit encoding depending on the language of the UTF-8 locale, and your users would need to be told how to select the 'encoding' argument.

If you want to run R CMD check on a Unix-alike over a package that sets a package encoding in its DESCRIPTION file and do not use a UTF-8 locale you may need to specify a suitable locale via environment variable R_ENCODING_LOCALES. The default is equivalent to the value

"latin1=en_US:latin2=pl_PL:UTF-8=en_US.UTF-8:latin9=fr_FR.iso885915@euro" (which is appropriate for a system based on glibc: OS X requires latin9=fr_FR.ISO8859-15) except that if the current locale is UTF-8 then the package code is translated to UTF-8 for syntax checking, so it is strongly recommended to check in a UTF-8 locale.

1.6.4 Portable C and C++ code

Writing portable C and C++ code is mainly a matter of observing the standards (C99, C++98 or where declared C++11) and testing that extensions (such as POSIX functions) are supported. However, some common errors are worth pointing out here. It can be helpful to look up functions at http://www.cplusplus.com/reference/ or http://en.cppreference.com/w/ and compare what is defined in the various standards.

Both the compiler and OS (via system header files, which differ by architecture even for nominally the same OS) affect the compilability of C/C++ code. Compilers from the GCC, clang, Intel and Solaris Studio suites are routinely used with R, and both clang and Solaris have more than one implementation of C++ headers and library. The range of possibilities makes comprehensive empirical checking impossible, and regrettably compilers are patchy at best on warning about non-standard code.

• Mathematical functions such as sqrt are defined in C++ for floating-point arguments. It is legitimate in C++ to overload these with versions for types float, double, long double and possibly more. This means that calling sqrt on an integer type may have 'overloading ambiguity' as it could be promoted to any of the supported floating-point types: this is commonly seen on Solaris, but for pow also seen on OS X. (C++98 has an overload for std::pow(<double>, <int>), but this may not be visible from the main namespace. C++11 requires additional overloads for integer types, and ambiguous overloads are more common in C++11 (and later) compiler modes.)

A not-uncommonly-seen problem is to mistakenly call floor(x/y) or ceil(x/y) for int arguments x and y. Since x/y does integer division, the result is an int and 'overloading ambiguity' may be reported. Some people have (pointlessly) called floor and ceil on integer arguments, which may have an 'overloading ambiguity'.

A surprising common misuse is things like pow(10, -3): this should be the constant 1e-3.

• Function fabs is defined only for floating-point types, except in C++11 which has overloads for std::fabs in <cmath> for integer types. Function abs is defined in C99's <stdlib.h> for int and in C++98's <cstdlib> for integer types, overloaded in <cmath> for floating-point types. C++11 has additional overloads for std::abs in <cmath> for integer types. The effect of calling abs with a floating-point type is implementation-specific: it may truncate to an integer.

⁷³ Typically on a Unix-alike this is done by telling fontconfig where to find suitable fonts to select glyphs from.

• Functions/macros such as isnan, isinf and isfinite are not required by C++98: where compilers support them they may be only in the std namespace or only in the main namespace. There is no way to make use of these functions which works with all C++ compilers currently in use on R platforms: use R's versions such as ISNAN and R_FINITE instead.

If you must use them in C++11, beware that some compilers⁷⁴ provide both std::isnan and ::isnan, so using

using namespace std;

may cause 'overloading ambiguity' and you must use std::isnan etc explicitly.

It is an error (and make little sense, although has been seen) to call these functions for integer arguments.

• The GNU compilers support a large number of non-portable extensions. For example, INFINITY (which is in C99 but not C++98), for which R provides the portable R_PosInf (and R_NegInf for -INFINITY). And NAN is just one NaN value: in R code NA_REAL is usually what is intended, but R_NaN is also available.

Some (but not all) extensions are listed at https://gcc.gnu.org/onlinedocs/gcc/C-Extensions.html and https://gcc.gnu.org/onlinedocs/gcc/C_002b_002b-Extensions.html.

- Including C-style headers in C++ code is not portable. Including the legacy header⁷⁵ math.h in C++ code may conflict with cmath which may be included by other headers. This is particularly problematic with C++11 compilers, as functions like sqrt and isnan are defined for double arguments in math.h and for a range of types including double in cmath. Similar issues have been seen for stdlib.h and cstdlib. Historically, including the C++ version first was a sufficient workaround but for some 2016 compilers only one can be included.
- Variable-length arrays are C99, not supported by C++98 nor by the C++ compilers in use with R on some platforms.
- The restrict qualifier is C99/C11 but not part of C++11 and not supported by some C++ compilers used with R.
- Be careful to include the headers which define the functions you use. Some compilers/OSes include other system headers in their headers which are not required by the standards, and so code may compile on such systems and not on others. (A prominent example is the C++11 header <random> which is indirectly included by <algorithm> by g++. Another frequent issue is the C header <time.h> which is included by other headers on Linux and Windows but not OS X nor Solaris.) Another common issue is that malloc, calloc, realloc and free are defined by C99 in the header stdlib.h and (in the std:: namespace) by C++ header cstdlib. Some earlier implementations used a header malloc.h, but that is not portable and does not exist on OS X.
- For C++ code, be careful to specify name spaces where needed. Many functions are defined by the standards to be in the std name space, but g++ puts many such also in the C++ main name space. One way to do so is to use declarations such as

using std::floor;

• Macros defined by the compiler/OS can cause problems. Identifiers starting with an underscore followed by an upper-case letter or another underscore are reserved for system macros and should not be used in portable code (including not as guards in C/C++ headers). Other macros, typically upper-case, may be defined by the compiler or system headers and can cause problems. The most common issue involves the names of the Intel CPU registers

⁷⁴ E.g. gcc 5.3 in C++11 mode.

⁷⁵ which often is the same as the header included by the C compiler, but some compilers have wrappers for some of the C headers.

such as CS, DS, ES, FS, GS and SS (and more with longer abbreviations) defined on i586/x64 Solaris in <sys/regset.h> and often included indirectly by <stdlib.h> and other core headers. Further examples are ERR, zero and I (which is defined in Solaris' <complex.h> as a compiler intrinsic for the imaginary unit).

• typedefs in OS headers can conflict with those in the package: an example is index_t defined in <sys/types.h> on Solaris.

Some additional information for C++ is available at http://journal.r-project.org/archive/2011-2/RJournal_2011-2_Plummer.pdf by Martyn Plummer.

1.6.5 Binary distribution

If you want to distribute a binary version of a package on Windows or OS X, there are further checks you need to do to check it is portable: it is all too easy to depend on external software on your own machine that other users will not have.

For Windows, check what other DLLs your package's DLL depends on ('imports' from in the DLL tools' parlance). A convenient GUI-based tool to do so is 'Dependency Walker' (http://www.dependencywalker.com/) for both 32-bit and 64-bit DLLs – note that this will report as missing links to R's own DLLs such as R.dll and Rblas.dll. For 32-bit DLLs only, the command-line tool pedump.exe -i (in Rtools*.exe) can be used, and for the brave, the objdump tool in the appropriate toolchain will also reveal what DLLs are imported from. If you use a toolchain other than one provided by the R developers or use your own makefiles, watch out in particular for dependencies on the toolchain's runtime DLLs such as libgfortran, libstdc++ and libgcc_s.

For OS X, using R CMD otool -L on the package's shared object(s) in the libs directory will show what they depend on: watch for any dependencies in /usr/local/lib, notably libgfortran.2.dylib, libgfortran.3.dylib or libquadmath.0.dylib.

Many people (including the CRAN package repository) will not accept source packages containing binary files as the latter are a security risk. If you want to distribute a source package which needs external software on Windows or OS X, options include

- To arrange for installation of the package to download the additional software from a URL, as e.g. package Cairo (https://CRAN.R-project.org/package=Cairo) does.
- (For CRAN.) To negotiate with Uwe Ligges to host the additional components on Win-Builder, and write a configure.win file to install them. There used to be many examples, e.g. package rgdal (https://CRAN.R-project.org/package=rgdal) (however nowadays CRAN prefers to use a uniform cross-compilation approach for software such as GDAL).

Be aware that license requirements will need to be met so you may need to supply the sources for the additional components (and will if your package has a GPL-like license).

1.7 Diagnostic messages

Diagnostic messages can be made available for translation, so it is important to write them in a consistent style. Using the tools described in the next section to extract all the messages can give a useful overview of your consistency (or lack of it). Some guidelines follow.

- Messages are sentence fragments, and not viewed in isolation. So it is conventional not to capitalize the first word and not to end with a period (or other punctuation).
- Try not to split up messages into small pieces. In C error messages use a single format string containing all English words in the messages.
 - In R error messages do not construct a message with paste (such messages will not be translated) but *via* multiple arguments to stop or warning, or *via* gettextf.
- Do not use colloquialisms such as "can't" and "don't".

Conventionally single quotation marks are used for quotations such as

'ord' must be a positive integer, at most the number of knots and double quotation marks when referring to an R character string or a class, such as

```
'format' must be "normal" or "short" - using "normal"
```

Since ASCII does not contain directional quotation marks, it is best to use '' and let the translator (including automatic translation) use directional quotations where available. The range of quotation styles is immense: unfortunately we cannot reproduce them in a portable texinfo document. But as a taster, some languages use 'up' and 'down' (comma) quotes rather than left or right quotes, and some use guillemets (and some use what Adobe calls 'guillemotleft' to start and others use it to end).

In R messages it is also possible to use sQuote or dQuote as in

• Occasionally messages need to be singular or plural (and in other languages there may be no such concept or several plural forms – Slovenian has four). So avoid constructions such as was once used in library

```
if((length(nopkgs) > 0) && !missing(lib.loc)) {
         if(length(nopkgs) > 1)
     warning("libraries ",
     paste(sQuote(nopkgs), collapse = ", "),
     " contain no packages")
         else
     warning("library ", paste(sQuote(nopkgs)),
     " contains no package")
     }
and was replaced by
     if((length(nopkgs) > 0) && !missing(lib.loc)) {
         pkglist <- paste(sQuote(nopkgs), collapse = ", ")</pre>
         msg <- sprintf(ngettext(length(nopkgs),</pre>
         "library %s contains no packages",
         "libraries %s contain no packages",
         domain = "R-base"),
        pkglist)
         warning(msg, domain=NA)
```

Note that it is much better to have complete clauses as here, since in another language one might need to say 'There is no package in library %s' or 'There are no packages in libraries %s'.

1.8 Internationalization

There are mechanisms to translate the R- and C-level error and warning messages. There are only available if R is compiled with NLS support (which is requested by configure option --enable-nls, the default).

The procedures make use of msgfmt and xgettext which are part of GNU gettext and this will need to be installed: Windows users can find pre-compiled binaries at https://www.stats.ox.ac.uk/pub/Rtools/goodies/gettext-tools.zip.

1.8.1 C-level messages

The process of enabling translations is

• In a header file that will be included in all the C (or C++ or Objective C/C++) files containing messages that should be translated, declare

```
#include <R.h> /* to include Rconfig.h */
       #ifdef ENABLE_NLS
       #include <libintl.h>
       #define _(String) dgettext ("pkg", String)
       /* replace pkg as appropriate */
       #else
       #define _(String) (String)
       #endif
• For each message that should be translated, wrap it in _(...), for example
       error(_("'ord' must be a positive integer"));
  If you want to use different messages for singular and plural forms, you need to add
       #ifndef ENABLE NLS
       #define dngettext(pkg, String, StringP, N) (N > 1 ? StringP : String)
       #endif
  and mark strings by
       dngettext(("pkg", <singular string>, <plural string>, n)
• In the package's src directory run
       xgettext --keyword=_ -o pkg.pot *.c
```

The file src/pkg.pot is the template file, and conventionally this is shipped as po/pkg.pot.

1.8.2 R messages

Mechanisms are also available to support the automatic translation of R stop, warning and message messages. They make use of message catalogs in the same way as C-level messages, but using domain R-pkg rather than pkg. Translation of character strings inside stop, warning and message calls is automatically enabled, as well as other messages enclosed in calls to gettext or gettextf. (To suppress this, use argument domain=NA.)

Tools to prepare the R-pkg.pot file are provided in package tools: xgettext2pot will prepare a file from all strings occurring inside gettext/gettextf, stop, warning and message calls. Some of these are likely to be spurious and so the file is likely to need manual editing. xgettext extracts the actual calls and so is more useful when tidying up error messages.

The R function ngettext provides an interface to the C function of the same name: see example in the previous section. It is safest to use domain="R-pkg" explicitly in calls to ngettext, and necessary for earlier versions of R unless they are calls directly from a function in the package.

1.8.3 Preparing translations

Once the template files have been created, translations can be made. Conventional translations have file extension .po and are placed in the po subdirectory of the package with a name that is either '11.po' or 'R-11.po' for translations of the C and R messages respectively to language with code '11'.

See Section "Localization of messages" in R Installation and Administration, for details of language codes.

There is an R function, update_pkg_po in package tools, to automate much of the maintenance of message translations. See its help for what it does in detail.

If this is called on a package with no existing translations, it creates the directory <code>pkgdir/po</code>, creates a template file of R messages, <code>pkgdir/po/R-pkg.pot</code>, within it, creates the 'en@quot' translation and installs that. (The 'en@quot' pseudo-language interprets quotes in their directional forms in suitable (e.g. UTF-8) locales.)

If the package has C source files in its src directory that are marked for translation, use

```
touch pkgdir/po/pkg.pot
```

to create a dummy template file, then call update_pkg_po again (this can also be done before it is called for the first time).

When translations to new languages are added in the <code>pkgdir/po</code> directory, running the same command will check and then install the translations.

If the package sources are updated, the same command will update the template files, merge the changes into the translation .po files and then installed the updated translations. You will often see that merging marks translations as 'fuzzy' and this is reported in the coverage statistics. As fuzzy translations are *not* used, this is an indication that the translation files need human attention.

The merged translations are run through tools::checkPofile to check that C-style formats are used correctly: if not the mismatches are reported and the broken translations are not installed.

This function needs the GNU gettext-tools installed and on the path: see its help page.

1.9 CITATION files

An installed file named CITATION will be used by the citation() function. (It should be in the inst subdirectory of the package sources.)

The CITATION file is parsed as R code (in the package's declared encoding, or in ASCII if none is declared). If no such file is present, citation auto-generates citation information from the package DESCRIPTION metadata, and an example of what that would look like as a CITATION file can be seen in recommended package nlme (https://CRAN.R-project.org/package=nlme) (see below): recommended packages boot (https://CRAN.R-project.org/package=boot), cluster (https://CRAN.R-project.org/package=cluster) and mgcv (https://CRAN.R-project.org/package=mgcv) have further examples.

A CITATION file will contain calls to function bibentry.

Note the way that information that may need to be updated is picked up from object meta, a parsed version of the DESCRIPTION file⁷⁶ – it is tempting to hardcode such information, but it normally then gets outdated. See ?bibentry for further details of the information which can be provided.

In case a bibentry contains LATEX markup (e.g., for accented characters or mathematical symbols), it may be necessary to provide a text representation to be used for printing via the textVersion argument to bibentry. E.g., earlier versions of nlme (https://CRAN.R-project.org/package=nlme) additionally used

The CITATION file should itself produce no output when source-d.

It is desirable (and essential for CRAN) that the CITATION file does not contain calls to functions such as packageDescription which assume the package is installed in a library tree on the package search path.

1.10 Package types

The DESCRIPTION file has an optional field Type which if missing is assumed to be 'Package', the sort of extension discussed so far in this chapter. Currently one other type is recognized; there used also to be a 'Translation' type.

1.10.1 Frontend

This is a rather general mechanism, designed for adding new front-ends such as the former **gnomeGUI** package (see the Archive area on CRAN). If a configure file is found in the top-level directory of the package it is executed, and then if a Makefile is found (often generated by configure), make is called. If R CMD INSTALL --clean is used make clean is called. No other action is taken.

R CMD build can package up this type of extension, but R CMD check will check the type and skip it.

Many packages of this type need write permission for the R installation directory.

1.11 Services

Several members of the R project have set up services to assist those writing R packages, particularly those intended for public distribution.

win-builder.r-project.org (http://win-builder.r-project.org) offers the automated preparation of (32/64-bit) Windows binaries from well-tested source packages.

R-Forge (R-Forge.r-project.org (https://R-Forge.r-project.org)) and RForge (www.rforge.net (https://www.rforge.net)) are similar services with similar names. Both provide source-code management through SVN, daily building and checking, mailing lists and a repository that can be accessed *via* install.packages (they can be selected by setRepositories and the GUI menus that use it). Package developers have the opportunity to present their work on the basis of project websites or news announcements. Mailing lists, forums or wikis provide useRs with convenient instruments for discussions and for exchanging information between developers and/or interested useRs.

⁷⁶ this object is available since R 2.8.0, so the 'Depends' field in the DESCRIPTION file should contain something at least as restrictive as 'R (>= 2.8'.

2 Writing R documentation files

2.1 Rd format

R objects are documented in files written in "R documentation" (Rd) format, a simple markup language much of which closely resembles (La)TeX, which can be processed into a variety of formats, including IATeX, HTML and plain text. The translation is carried out by functions in the **tools** package called by the script Rdconv in R_HOME/bin and by the installation scripts for packages.

The R distribution contains more than 1300 such files which can be found in the src/library/pkg/man directories of the R source tree, where pkg stands for one of the standard packages which are included in the R distribution.

As an example, let us look at a simplified version of src/library/base/man/load.Rd which documents the R function load.

```
% File src/library/base/man/load.Rd
\name{load}
\alias{load}
\title{Reload Saved Datasets}
\description{
 Reload the datasets written to a file with the function
  \code{save}.
\usage{
load(file, envir = parent.frame())
\arguments{
  \item{file}{a connection or a character string giving the
   name of the file to load.}
  \item{envir}{the environment where the data should be
    loaded.}
\seealso{
 \code{\link{save}}.
\examples{
## save all data
save(list = ls(), file= "all.RData")
## restore the saved values to the current environment
load("all.RData")
## restore the saved values to the workspace
load("all.RData", .GlobalEnv)
\keyword{file}
```

An Rd file consists of three parts. The header gives basic information about the name of the file, the topics documented, a title, a short textual description and R usage information for the objects documented. The body gives further information (for example, on the function's arguments and return value, as in the above example). Finally, there is an optional footer with keyword information. The header is mandatory.

Information is given within a series of *sections* with standard names (and user-defined sections are also allowed). Unless otherwise specified¹ these should occur only once in an Rd file (in any

¹ e.g. \alias, \keyword and \note sections.

order), and the processing software will retain only the first occurrence of a standard section in the file, with a warning.

See "Guidelines for Rd files" (https://developer.r-project.org/Rds.html) for guidelines for writing documentation in Rd format which should be useful for package writers. The R generic function prompt is used to construct a bare-bones Rd file ready for manual editing. Methods are defined for documenting functions (which fill in the proper function and argument names) and data frames. There are also functions promptData, promptPackage, promptClass, and promptMethods for other types of Rd file.

The general syntax of Rd files is summarized below. For a detailed technical discussion of current Rd syntax, see "Parsing Rd files" (https://developer.r-project.org/parseRd.pdf).

Rd files consist of four types of text input. The most common is LATEX-like, with the backslash used as a prefix on markup (e.g. \alias), and braces used to indicate arguments (e.g. {load}). The least common type of text is 'verbatim' text, where no markup other than the comment marker (%) is processed. There is also a rare variant of 'verbatim' text (used in \eqn, \deqn, \figure, and \newcommand) where comment markers need not be escaped. The final type is R-like, intended for R code, but allowing some embedded macros. Quoted strings within R-like text are handled specially: regular character escapes such as \n may be entered as-is. Only markup starting with \1 (e.g. \link) or \v (e.g. \var) will be recognized within quoted strings. The rarely used vertical tab \v must be entered as \\v.

Each macro defines the input type for its argument. For example, the file initially uses IATEX-like syntax, and this is also used in the \description section, but the \usage section uses R-like syntax, and the \alias macro uses 'verbatim' syntax. Comments run from a percent symbol % to the end of the line in all types of text except the rare 'verbatim' variant (as on the first line of the load example).

Because backslashes, braces and percent symbols have special meaning, to enter them into text sometimes requires escapes using a backslash. In general balanced braces do not need to be escaped, but percent symbols always do, except in the 'verbatim' variant. For the complete list of macros and rules for escapes, see "Parsing Rd files" (https://developer.r-project.org/parseRd.pdf).

2.1.1 Documenting functions

The basic markup commands used for documenting R objects (in particular, functions) are given in this subsection.

\name{name}

name typically² is the basename of the Rd file containing the documentation. It is the "name" of the Rd object represented by the file and has to be unique in a package. To avoid problems with indexing the package manual, it may not contain '!' 'l' nor '@', and to avoid possible problems with the HTML help system it should not contain '/' nor a space. (IATEX special characters are allowed, but may not be collated correctly in the index.) There can only be one \name entry in a file, and it must not contain any markup. Entries in the package manual will be in alphabetic³ order of the \name entries.

\alias{topic}

The \alias sections specify all "topics" the file documents. This information is collected into index data bases for lookup by the on-line (plain text and HTML)

² There can be exceptions: for example Rd files are not allowed to start with a dot, and have to be uniquely named on a case-insensitive file system.

³ in the current locale, and with special treatment for IAT_EX special characters and with any 'pkgname-package' topic moved to the top of the list.

help systems. The *topic* can contain spaces, but (for historical reasons) leading and trailing spaces will be stripped. Percent and left brace need to be escaped by a backslash.

There may be several \alias entries. Quite often it is convenient to document several R objects in one file. For example, file Normal.Rd documents the density, distribution function, quantile function and generation of random variates for the normal distribution, and hence starts with

```
\name{Normal}
\alias{Normal}
\alias{dnorm}
\alias{pnorm}
\alias{qnorm}
\alias{rnorm}
```

Also, it is often convenient to have several different ways to refer to an R object, and an **\alias** does not need to be the name of an object.

Note that the \name is not necessarily a topic documented, and if so desired it needs to have an explicit \alias entry (as in this example).

\title{Title}

Title information for the Rd file. This should be capitalized and not end in a period; try to limit its length to at most 65 characters for widest compatibility.

Markup is supported in the text, but use of characters other than English text and punctuation (e.g., '<') may limit portability.

There must be one (and only one) \title section in a help file.

\description{...}

A short description of what the function(s) do(es) (one paragraph, a few lines only). (If a description is too long and cannot easily be shortened, the file probably tries to document too much at once.) This is mandatory except for package-overview files.

\usage{fun(arg1, arg2, ...)}

One or more lines showing the synopsis of the function(s) and variables documented in the file. These are set in typewriter font. This is an R-like command.

The usage information specified should match the function definition *exactly* (such that automatic checking for consistency between code and documentation is possible).

It is no longer advisable to use \synopsis for the actual synopsis and show modified synopses in the \usage. Support for \synopsis will be removed in \R 3.1.0. To indicate that a function can be used in several different ways, depending on the named arguments specified, use section \details. E.g., abline.Rd contains

```
\details{
   Typical usages are
\preformatted{abline(a, b, untf = FALSE, \dots)
.....
}
```

Use \method{generic}{class} to indicate the name of an S3 method for the generic function generic for objects inheriting from class "class". In the printed versions, this will come out as generic (reflecting the understanding that methods should not be invoked directly but via method dispatch), but codoc() and other QC tools always have access to the full name.

For example, print.ts.Rd contains

```
\usage{
  \method{print}{ts}(x, calendar, \dots)
}
which will print as
  Usage:
  ## S3 method for class 'ts':
    print(x, calendar, ...)
```

Usage for replacement functions should be given in the style of $\dim(x)$ <- value rather than explicitly indicating the name of the replacement function ("dim<-" in the above). Similarly, one can use \method{generic}{class}(arglist) <- value to indicate the usage of an S3 replacement method for the generic replacement function "generic<-" for objects inheriting from class "class".

Usage for S3 methods for extracting or replacing parts of an object, S3 methods for members of the Ops group, and S3 methods for user-defined (binary) infix operators ('%xxx%') follows the above rules, using the appropriate function names. E.g., Extract.factor.Rd contains

```
\usage{
  \method{[]{factor}(x, \dots, drop = FALSE)
  \method{[[]{factor}(x, \dots)
  \method{[]{factor}(x, \dots) <- value
  }
which will print as
  Usage:

  ## S3 method for class 'factor':
   x[..., drop = FALSE]
  ## S3 method for class 'factor':
   x[[...]]
  ## S3 replacement method for class 'factor':
   x[...] <- value</pre>
```

\S3method is accepted as an alternative to \method.

\arguments{...}

Description of the function's arguments, using an entry of the form

```
\item{arg_i}{Description of arg_i.}
```

for each element of the argument list. (Note that there is no whitespace between the three parts of the entry.) There may be optional text outside the *\item* entries, for example to give general information about groups of parameters.

\details{...}

A detailed if possible precise description of the functionality provided, extending the basic information in the \description slot.

\value{...}

Description of the function's return value.

If a list with multiple values is returned, you can use entries of the form

```
\item{comp_i}{Description of comp_i.}
```

for each component of the list returned. Optional text may precede⁴ this list (see for example the help for rle). Note that \value is implicitly a \describe environ-

⁴ Text between or after list items is discouraged.

ment, so that environment should not be used for listing components, just individual \item{}{} entries.

\references{...}

A section with references to the literature. Use \url{} or \href{}{} for web pointers.

\note{...}

Use this for a special note you want to have pointed out. Multiple \note sections are allowed, but might be confusing to the end users.

For example, pie.Rd contains

```
\note{
   Pie charts are a very bad way of displaying information.
   The eye is good at judging linear measures and bad at judging relative areas.
   ......
}
```

\author{...}

Information about the author(s) of the Rd file. Use \email{} without extra delimiters (such as '()' or '< >') to specify email addresses, or \url{} or \href{}{} for web pointers.

\seealso{...}

Pointers to related R objects, using \code{\link{...}} to refer to them (\code is the correct markup for R object names, and \link produces hyperlinks in output formats which support this. See Section 2.3 [Marking text], page 66, and Section 2.5 [Cross-references], page 68).

\examples{...}

Examples of how to use the function. Code in this section is set in typewriter font without reformatting and is run by example() unless marked otherwise (see below).

Examples are not only useful for documentation purposes, but also provide test code used for diagnostic checking of R code. By default, text inside \examples{} will be displayed in the output of the help page and run by example() and by R CMD check. You can use \dontrun{} for text that should only be shown, but not run, and \dontshow{} for extra commands for testing that should not be shown to users, but will be run by example(). (Previously this was called \testonly, and that is still accepted.)

Text inside \dontrun{} is 'verbatim', but the other parts of the \examples section are R-like text.

For example,

```
x <- runif(10)  # Shown and run.
\dontrun{plot(x)}  # Only shown.
\dontshow{log(x)}  # Only run.</pre>
```

Thus, example code not included in \dontrun must be executable! In addition, it should not use any system-specific features or require special facilities (such as Internet access or write permission to specific directories). Text included in \dontrun is indicated by comments in the processed help files: it need not be valid R code but the escapes must still be used for \%, \ and unpaired braces as in other 'verbatim' text.

Example code must be capable of being run by example, which uses source. This means that it should not access stdin, e.g. to scan() data from the example file.

Data needed for making the examples executable can be obtained by random number generation (for example, $x \leftarrow rnorm(100)$), or by using standard data sets listed by data() (see ?data for more info).

Finally, there is \donttest, used (at the beginning of a separate line) to mark code that should be run by example() but not by R CMD check (by default: the option --run-donttest can be used). This should be needed only occasionally but can be used for code which might fail in circumstances that are hard to test for, for example in some locales. (Use e.g. capabilities() or nzchar(Sys.which("someprogram")) to test for features needed in the examples wherever possible, and you can also use try() or tryCatch(). Use interactive() to condition examples which need someone to interact with.) Note that code included in \donttest must be correct R code, and any packages used should be declared in the DESCRIPTION file. It is good practice to include a comment in the \donttest section explaining why it is needed.

\keyword{key}

There can be zero or more \keyword sections per file. Each \keyword section should specify a single keyword, preferably one of the standard keywords as listed in file KEYWORDS in the R documentation directory (default R_HOME/doc). Use e.g. RShowDoc("KEYWORDS") to inspect the standard keywords from within R. There can be more than one \keyword entry if the R object being documented falls into more than one category, or none.

Do strongly consider using \concept (see Section 2.9 [Indices], page 71) instead of \keyword if you are about to use more than very few non-standard keywords.

The special keyword 'internal' marks a page of internal objects that are not part of the package's API. If the help page for object foo has keyword 'internal', then help(foo) gives this help page, but foo is excluded from several object indices, including the alphabetical list of objects in the HTML help system.

help.search() can search by keyword, including user-defined values: however the 'Search Engine & Keywords' HTML page accessed *via* help.start() provides single-click access only to a pre-defined list of keywords.

2.1.2 Documenting data sets

The structure of Rd files which document R data sets is slightly different. Sections such as \arguments and \value are not needed but the format and source of the data should be explained.

As an example, let us look at src/library/datasets/man/rivers.Rd which documents the standard R data set rivers.

```
\name{rivers}
\docType{data}
\alias{rivers}
\title{Lengths of Major North American Rivers}
\description{
   This data set gives the lengths (in miles) of 141 \dQuote{major}
    rivers in North America, as compiled by the US Geological
    Survey.
}
\usage{rivers}
\format{A vector containing 141 observations.}
\source{World Almanac and Book of Facts, 1975, page 406.}
\references{
   McNeil, D. R. (1977) \emph{Interactive Data Analysis}.
   New York: Wiley.
}
\keyword{datasets}
```

This uses the following additional markup commands.

\docType{...}

Indicates the "type" of the documentation object. Always 'data' for data sets, and 'package' for pkg-package.Rd overview files. Documentation for S4 methods and classes uses 'methods' (from promptMethods()) and 'class' (from promptClass()).

\format{...}

A description of the format of the data set (as a vector, matrix, data frame, time series, ...). For matrices and data frames this should give a description of each column, preferably as a list or table. See Section 2.4 [Lists and tables], page 68, for more information.

\source{...}

Details of the original source (a reference or URL, see Section 1.1.8 [Specifying URLs], page 17). In addition, section \references could give secondary sources and usages.

Note also that when documenting data set bar,

- The \usage entry is always bar or (for packages which do not use lazy-loading of data) data(bar). (In particular, only document a single data object per Rd file.)
- The \keyword entry should always be 'datasets'.

If bar is a data frame, documenting it as a data set can be initiated via prompt(bar). Otherwise, the promptData function may be used.

2.1.3 Documenting S4 classes and methods

There are special ways to use the '?' operator, namely 'class?topic' and 'methods?topic', to access documentation for S4 classes and methods, respectively. This mechanism depends on conventions for the topic names used in \alias entries. The topic names for S4 classes and methods respectively are of the form

```
class-class
generic,signature_list-method
```

where signature_list contains the names of the classes in the signature of the method (without quotes) separated by ',' (without whitespace), with 'ANY' used for arguments without an explicit specification. E.g., 'genericFunction-class' is the topic name for documentation for the S4 class "genericFunction", and 'coerce, ANY, NULL-method' is the topic name for documentation for the S4 method for coerce for signature c("ANY", "NULL").

Skeletons of documentation for S4 classes and methods can be generated by using the functions promptClass() and promptMethods() from package methods. If it is necessary or desired

to provide an explicit function declaration (in a \usage section) for an S4 method (e.g., if it has "surprising arguments" to be mentioned explicitly), one can use the special markup

```
\label{list} $$ \end{generic} signature_list (argument_list) $$ (e.g., 'S4method{coerce}{ANY,NULL}(from, to)'). $$
```

To make full use of the potential of the on-line documentation system, all user-visible S4 classes and methods in a package should at least have a suitable \alias entry in one of the package's Rd files. If a package has methods for a function defined originally somewhere else, and does not change the underlying default method for the function, the package is responsible for documenting the methods it creates, but not for the function itself or the default method.

An S4 replacement method is documented in the same way as an S3 one: see the description of \method in Section 2.1.1 [Documenting functions], page 59.

See help("Documentation", package = "methods") for more information on using and creating on-line documentation for S4 classes and methods.

2.1.4 Documenting packages

Packages may have an overview help page with an \alias pkgname-package, e.g. 'utils-package' for the utils package, when package?pkgname will open that help page. If a topic named pkgname does not exist in another Rd file, it is helpful to use this as an additional \alias.

Skeletons of documentation for a package can be generated using the function promptPackage(). If the final = LIBS argument is used, then the Rd file will be generated in final form, containing the information that would be produced up to library(help = pkgname). Otherwise (the default) comments will be inserted giving suggestions for content.

Apart from the mandatory \name and \title and the pkgname-package alias, the only requirement for the package overview page is that it include a \docType{package} statement. All other content is optional. We suggest that it should be a short overview, to give a reader unfamiliar with the package enough information to get started. More extensive documentation is better placed into a package vignette (see Section 1.4 [Writing package vignettes], page 35) and referenced from this page, or into individual man pages for the functions, datasets, or classes.

2.2 Sectioning

To begin a new paragraph or leave a blank line in an example, just insert an empty line (as in (La)TEX). To break a line, use \cr.

In addition to the predefined sections (such as \description{}, \value{}, etc.), you can "define" arbitrary ones by \section{section_title}{...}. For example

```
\section{Warning}{
   You must not call this function unless ...
}
```

For consistency with the pre-assigned sections, the section name (the first argument to \section) should be capitalized (but not all upper case). Whitespace between the first and second braced expressions is not allowed. Markup (e.g. \code) within the section title may cause problems with the latex conversion (depending on the version of macro packages such as 'hyperref') and so should be avoided.

The \subsection macro takes arguments in the same format as \section, but is used within a section, so it may be used to nest subsections within sections or other subsections. There is no predefined limit on the nesting level, but formatting is not designed for more than 3 levels (i.e. subsections within subsections within sections).

Note that additional named sections are always inserted at a fixed position in the output (before \note, \seealso and the examples), no matter where they appear in the input (but in the same order amongst themselves as in the input).

2.3 Marking text

The following logical markup commands are available for emphasizing or quoting text.

\emph{text}

\strong{text}

Emphasize text using italic and **bold** font if possible; \strong is regarded as stronger (more emphatic).

\bold{text}

Set text in **bold** font where possible.

\sQuote{text}

\dQuote{text}

Portably single or double quote text (without hard-wiring the characters used for quotation marks).

Each of the above commands takes \LaTeX Larger Larger

The following logical markup commands are available for indicating specific kinds of text. Except as noted, these take 'verbatim' text input, and so other macros may not be used within them. Some characters will need to be escaped (see Section 2.8 [Insertions], page 70).

\code{text}

Indicate text that is a literal example of a piece of an R program, e.g., a fragment of R code or the name of an R object. Text is entered in R-like syntax, and displayed using typewriter font where possible. Macros \var and \link are interpreted within text.

\preformatted{text}

Indicate text that is a literal example of a piece of a program. Text is displayed using typewriter font where possible. Formatting, e.g. line breaks, is preserved. (Note that this includes a line break after the initial $\{$, so typically text should start on the same line as the command.)

Due to limitations in IATEX as of this writing, this macro may not be nested within other markup macros other than \dQuote and \sQuote, as errors or bad formatting may result.

\kbd{keyboard-characters}

Indicate keyboard input, using *slanted typewriter* font if possible, so users can distinguish the characters they are supposed to type from computer output. Text is entered 'verbatim'.

\samp{text}

Indicate text that is a literal example of a sequence of characters, entered 'verbatim'. No wrapping or reformatting will occur. Displayed using typewriter font where possible.

\verb{text}

Indicate text that is a literal example of a sequence of characters, with no interpretation of e.g. \var, but which will be included within word-wrapped text. Displayed using typewriter font if possible.

\pkg{package_name}

Indicate the name of an R package. LATEX-like.

\file{file_name}

Indicate the name of a file. Text is LATEX-like, so backslash needs to be escaped. Displayed using a distinct font where possible.

\email{email_address}

Indicate an electronic mail address. LAT_EX-like, will be rendered as a hyperlink in HTML and PDF conversion. Displayed using typewriter font where possible.

\url{uniform_resource_locator}

Indicate a uniform resource locator (URL) for the World Wide Web. The argument is handled as 'verbatim' text (with percent and braces escaped by backslash), and rendered as a hyperlink in HTML and PDF conversion. Linefeeds are removed, and as from R 3.2.0 leading and trailing whitespace⁵ is removed. See Section 1.1.8 [Specifying URLs], page 17.

Displayed using typewriter font where possible.

\href{uniform_resource_locator}{text}

Indicate a hyperlink to the World Wide Web. The first argument is handled as 'verbatim' text (with percent and braces escaped by backslash) and is used as the URL in the hyperlink, with the second argument of LATEX-like text displayed to the user. Linefeeds are removed from the first argument, and as from R 3.2.0 leading and trailing whitespace is removed.

Note that RFC3986-encoded URLs (e.g. using '\%28VS.85\\%29' in place of '(VS.85)') may not work correctly in versions of R before 3.1.3 and are best avoided—use URLdecode() to decode them.

\var{metasyntactic_variable}

Indicate a metasyntactic variable. In some cases this will be rendered distinctly, e.g. in italic, but not in all⁶. LAT_FX-like.

\env{environment_variable}

Indicate an environment variable. 'Verbatim'. Displayed using typewriter font where possible

\option{option}

Indicate a command-line option. 'Verbatim'. Displayed using typewriter font where possible.

\command{command_name}

Indicate the name of a command. IATEX-like, so \var is interpreted. Displayed using typewriter font where possible.

\dfn{term}

Indicate the introductory or defining use of a term. IATEX-like.

\cite{reference}

Indicate a reference without a direct cross-reference *via* \link (see Section 2.5 [Cross-references], page 68), such as the name of a book. LATEX-like.

\acronym{acronym}

Indicate an acronym (an abbreviation written in all capital letters), such as GNU. \LaTeX

⁵ as defined by the R function trimws.

 $^{^6}$ Currently it is rendered differently only in HTML conversions, and LATEX conversion outside '\usage' and '\examples' environments.

2.4 Lists and tables

The \itemize and \enumerate commands take a single argument, within which there may be one or more \item commands. The text following each \item is formatted as one or more paragraphs, suitably indented and with the first paragraph marked with a bullet point (\itemize) or a number (\enumerate).

Note that unlike argument lists, \item in these formats is followed by a space and the text (not enclosed in braces). For example

```
\enumerate{
  \item A database consists of one or more records, each with one or
  more named fields.
  \item Regular lines start with a non-whitespace character.
  \item Records are separated by one or more empty lines.
}
```

\itemize and \enumerate commands may be nested.

The \describe command is similar to \itemize but allows initial labels to be specified. Each \item takes two arguments, the label and the body of the item, in exactly the same way as an argument or value \item. \describe commands are mapped to <DL> lists in HTML and \description lists in LATEX.

The \tabular command takes two arguments. The first gives for each of the columns the required alignment ('1' for left-justification, 'r' for right-justification or 'c' for centring.) The second argument consists of an arbitrary number of lines separated by \cr, and with fields separated by \tab. For example:

```
\tabular{rll1}{
  [,1] \tab Ozone \tab numeric \tab Ozone (ppb)\cr
  [,2] \tab Solar.R \tab numeric \tab Solar R (lang)\cr
  [,3] \tab Wind \tab numeric \tab Wind (mph)\cr
  [,4] \tab Temp \tab numeric \tab Temperature (degrees F)\cr
  [,5] \tab Month \tab numeric \tab Month (1--12)\cr
  [,6] \tab Day \tab numeric \tab Day of month (1--31)
}
```

There must be the same number of fields on each line as there are alignments in the first argument, and they must be non-empty (but can contain only spaces). (There is no whitespace between **\tabular** and the first argument, nor between the two arguments.)

2.5 Cross-references

The markup \link{foo} (usually in the combination \code{\link{foo}}) produces a hyperlink to the help for foo. Here foo is a topic, that is the argument of \alias markup in another Rd file (possibly in another package). Hyperlinks are supported in some of the formats to which Rd files are converted, for example HTML and PDF, but ignored in others, e.g. the text format.

One main usage of \link is in the \seealso section of the help page, see Section 2.1 [Rd format], page 58.

Note that whereas leading and trailing spaces are stripped when extracting a topic from a \alias, they are not stripped when looking up the topic of a \link.

You can specify a link to a different topic than its name by \link[=dest]{name} which links to topic dest with name name. This can be used to refer to the documentation for S3/4 classes, for example \code{"\link[=abc-class]{abc}"} would be a way to refer to the documentation of an S4 class "abc" defined in your package, and \code{"\link[=terms.object]{terms}"} to the S3 "terms" class (in package stats). To make these easy to read in the source file, \code{"\linkS4class{abc}"} expands to the form given above.

There are two other forms of optional argument specified as \link[pkg]{foo} and \link[pkg:bar]{foo} to link to the package pkg, to files foo.html and bar.html respectively. These are rarely needed, perhaps to refer to not-yet-installed packages (but there the HTML help system will resolve the link at run time) or in the normally undesirable event that more than one package offers help on a topic⁷ (in which case the present package has precedence so this is only needed to refer to other packages). They are currently only used in HTML help (and ignored for hyperlinks in LATEX conversions of help pages), and link to the file rather than the topic (since there is no way to know which topics are in which files in an uninstalled package). The only reason to use these forms for base and recommended packages is to force a reference to a package that might be further down the search path. Because they have been frequently misused, the HTML help system looks for topic foo in package pkg if it does not find file foo.html.

2.6 Mathematics

Mathematical formulae should be set beautifully for printed documentation yet we still want something useful for text and HTML online help. To this end, the two commands \eqn{latex}{ascii} and \deqn{latex}{ascii} are used. Whereas \eqn is used for "inline" formulae (corresponding to TEX's \$...\$), \deqn gives "displayed equations" (as in LATEX's displaymath environment, or TEX's \$\$...\$). Both arguments are treated as 'verbatim' text.

Both commands can also be used as \eqn{latexascii} (only one argument) which then is used for both latex and ascii. No whitespace is allowed between command and the first argument, nor between the first and second arguments.

The following example is from Poisson.Rd:

```
\label{eq:lambda} $$ \left( \sum_{x \in \mathbb{N}} x e^{-\lambda x} e^{-\lambda x} \right) = \lambda x \exp(-\lambda x) / x! $$ for \leq 0, 1, 2, \lambda .
```

For the LATEX manual, this becomes

$$p(x) = \lambda^x \frac{e^{-\lambda}}{x!}$$
 for $x = 0, 1, 2, \dots$

For text on-line help we get

```
p(x) = lambda^x exp(-lambda)/x!
for x = 0, 1, 2, ....
```

Greek letters (both cases) will be rendered in HTML if preceded by a backslash, \dots and \ldots will be rendered as ellipses and \sqrt, \ge and \le as mathematical symbols.

Note that only basic IATEX can be used, there being no provision to specify IATEX style files such as the AMS extensions.

2.7 Figures

To include figures in help pages, use the \figure markup. There are three forms.

⁷ a common example in CRAN packages is \link[mgcv]{gam}.

The two commonly used simple forms are \figure{filename} and \figure{filename} {alternate text}. This will include a copy of the figure in either HTML or LATEX output. In text output, the alternate text will be displayed instead. (When the second argument is omitted, the filename will be used.) Both the filename and the alternate text will be parsed verbatim, and should not include special characters that are significant in HTML or LATEX.

The expert form is \figure{filename}{options: string}. (The word 'options:' must be typed exactly as shown and followed by at least one space.) In this form, the string is copied into the HTML img tag as attributes following the src attribute, or into the second argument of the \Figure macro in LATEX, which by default is used as options to an \includegraphics call. As it is unlikely that any single string would suffice for both display modes, the expert form would normally be wrapped in conditionals. It is up to the author to make sure that legal HTML/LATEX is used. For example, to include a logo in both HTML (using the simple form) and LATEX (using the expert form), the following could be used:

```
\if{\tml}{\figure{Rlogo.svg}{options: width=100 alt="R logo"}} \if{\latex}{\figure{Rlogo.pdf}{options: width=0.5in}}
```

The files containing the figures should be stored in the directory man/figures. Files with extensions .jpg, .jpeg, .pdf, .png and .svg from that directory will be copied to the help/figures directory at install time. (Figures in PDF format will not display in most HTML browsers, but might be the best choice in reference manuals.) Specify the filename relative to man/figures in the \figure directive.

2.8 Insertions

Use \R for the R system itself. Use \dots for the dots in function argument lists '...', and \dots for ellipsis dots in ordinary text.⁸ These can be followed by $\{\}$, and should be unless followed by whitespace.

After an unescaped '%', you can put your own comments regarding the help text. The rest of the line (but not the newline at the end) will be completely disregarded. Therefore, you can also use it to make part of the "help" invisible.

You can produce a backslash ('\') by escaping it by another backslash. (Note that \cr is used for generating line breaks.)

The "comment" character "%" and unpaired braces almost always need to be escaped by '\', and '\\' can be used for backslash and needs to be when there two or more adjacent backslashes). In R-like code quoted strings are handled slightly differently; see "Parsing Rd files" (https://developer.r-project.org/parseRd.pdf) for details — in particular braces should not be escaped in quoted strings.

All of '% { } \' should be escaped in LATEX-like text.

Text which might need to be represented differently in different encodings should be marked by \enc, e.g. \enc{Jöreskog}{Joreskog} (with no whitespace between the braces) where the first argument will be used where encodings are allowed and the second should be ASCII (and is used for e.g. the text conversion in locales that cannot represent the encoded form). (This is intended to be used for individual words, not whole sentences or paragraphs.)

⁸ There is only a fine distinction between \dots and \ldots. It is technically incorrect to use \ldots in code blocks and tools::checkRd will warn about this—on the other hand the current converters treat them the same way in code blocks, and elsewhere apart from the small distinction between the two in IATEX.

 $^{^{9}}$ See the examples section in the file Paren.Rd for an example.

2.9 Indices

The \alias command (see Section 2.1.1 [Documenting functions], page 59) is used to specify the "topics" documented, which should include all R objects in a package such as functions and variables, data sets, and S4 classes and methods (see Section 2.1.3 [Documenting S4 classes and methods], page 64). The on-line help system searches the index data base consisting of all alias topics.

In addition, it is possible to provide "concept index entries" using \concept, which can be used for help.search() lookups. E.g., file cor.test.Rd in the standard package stats contains

```
\concept{Kendall correlation coefficient}
\concept{Pearson correlation coefficient}
\concept{Spearman correlation coefficient}
```

so that e.g. ??Spearman will succeed in finding the help page for the test for association between paired samples using Spearman's ρ .

(Note that help.search() only uses "sections" of documentation objects with no additional markup.)

If you want to cross reference such items from other help files *via* \link, you need to use \alias and not \concept.

2.10 Platform-specific documentation

Sometimes the documentation needs to differ by platform. Currently two OS-specific options are available, 'unix' and 'windows', and lines in the help source file can be enclosed in

```
#ifdef OS
...
#endif
or
#ifndef OS
...
#endif
```

for OS-specific inclusion or exclusion. Such blocks should not be nested, and should be entirely within a block (that, is between the opening and closing brace of a section or item), or at top-level contain one or more complete sections.

If the differences between platforms are extensive or the R objects documented are only relevant to one platform, platform-specific Rd files can be put in a unix or windows subdirectory.

2.11 Conditional text

Occasionally the best content for one output format is different from the best content for another. For this situation, the \if{format}{text} or \ifelse{format}{text}{alternate} markup is used. Here format is a comma separated list of formats in which the text should be rendered. The alternate will be rendered if the format does not match. Both text and alternate may be any sequence of text and markup.

Currently the following formats are recognized: example, html, latex and text. These select output for the corresponding targets. (Note that example refers to extracted example code rather than the displayed example in some other format.) Also accepted are TRUE (matching all formats) and FALSE (matching no formats). These could be the output of the \Sexpr macro (see Section 2.12 [Dynamic pages], page 72).

The \out{literal} macro would usually be used within the text part of \if{format}{text}. It causes the renderer to output the literal text exactly, with

no attempt to escape special characters. For example, use the following to output the markup necessary to display the Greek letter in LATEX or HTML, and the text string alpha in other formats:

\if{latex}{\out{\alpha}}\ifelse{html}{\out{α}}{alpha}

2.12 Dynamic pages

Two macros supporting dynamically generated man pages are \Sexpr and \RdOpts. These are modelled after Sweave, and are intended to contain executable R expressions in the Rd file.

The main argument to \Sexpr must be valid R code that can be executed. It may also take options in square brackets before the main argument. Depending on the options, the code may be executed at package build time, package install time, or man page rendering time.

The options follow the same format as in Sweave, but different options are supported. Currently the allowed options and their defaults are:

- eval=TRUE Whether the R code should be evaluated.
- echo=FALSE Whether the R code should be echoed. If TRUE, a display will be given in a preformatted block. For example, \Sexpr[echo=TRUE] { x <- 1 } will be displayed as

- keep.source=TRUE Whether to keep the author's formatting when displaying the code, or throw it away and use a departed version.
- results=text How should the results be displayed? The possibilities are:
 - results=text Apply as.character() to the result of the code, and insert it as a text element.
 - results=verbatim Print the results of the code just as if it was executed at the console,
 and include the printed results verbatim. (Invisible results will not print.)
 - results=rd The result is assumed to be a character vector containing markup to be passed to parse_Rd(), with the result inserted in place. This could be used to insert computed aliases, for instance. parse_Rd() is called first with fragment = FALSE to allow a single Rd section macro to be inserted. If that fails, it is called again with fragment = TRUE, the older behavior.
 - results=hide Insert no output.
- strip.white=TRUE Remove leading and trailing white space from each line of output if strip.white=TRUE. With strip.white=all, also remove blank lines.
- stage=install Control when this macro is run. Possible values are
 - stage=build The macro is run when building a source tarball.
 - stage=install The macro is run when installing from source.
 - stage=render The macro is run when displaying the help page.

Conditionals such as **#ifdef** (see Section 2.10 [Platform-specific sections], page 71) are applied after the **build** macros but before the **install** macros. In some situations (e.g. installing directly from a source directory without a tarball, or building a binary package) the above description is not literally accurate, but authors can rely on the sequence being **build**, **#ifdef**, **install**, **render**, with all stages executed.

Code is only run once in each stage, so a \Sexpr[results=rd] macro can output an \Sexpr macro designed for a later stage, but not for the current one or any earlier stage.

• width, height, fig These options are currently allowed but ignored.

The \RdOpts macro is used to set new defaults for options to apply to following uses of \Sexpr .

For more details, see the online document "Parsing Rd files" (https://developer.r-project.org/parseRd.pdf).

2.13 User-defined macros

The \newcommand and \renewcommand macros allow new macros to be defined within an Rd file. These are similar but not identical to the same-named LATEX macros.

They each take two arguments which are parsed verbatim. The first is the name of the new macro including the initial backslash, and the second is the macro definition. As in LATEX, \newcommand requires that the new macro not have been previously defined, whereas \renewcommand allows existing macros (including all built-in ones) to be replaced. (As from version 3.2.0, this test is disabled by default, but may be enabled by setting the environment variable _WARN_DUPLICATE_RD_MACROS_ to a true value.)

Also as in IATEX, the new macro may be defined to take arguments, and numeric placeholders such as #1 are used in the macro definition. However, unlike IATEX, the number of arguments is determined automatically from the highest placeholder number seen in the macro definition. For example, a macro definition containing #1 and #3 (but no other placeholders) will define a three argument macro (whose second argument will be ignored). As in IATEX, at most 9 arguments may be defined. If the # character is followed by a non-digit it will have no special significance. All arguments to user-defined macros will be parsed as verbatim text, and simple text-substitution will be used to replace the place-holders, after which the replacement text will be parsed.

As of R version 3.2.0, a number of macros are defined in the file share/Rd/macros/system.Rd of the R source or home directory, and these will normally be available in all .Rd files. For example, that file contains the definition

```
\newcommand{\PR}{\Sexpr[results=rd]{tools:::Rd_expr_PR(#1)}}
```

which defines \PR to be a single argument macro; then code (typically used in the NEWS.Rd file) like

\PR{1234}

will expand to

\Sexpr[results=rd] {tools:::Rd_expr_PR(1234)}

when parsed.

Some macros that might be of general use are:

\CRANpkg{pkg}

A package on CRAN

\sspace A single space (used after a period that does not end a sentence).

\doi{numbers}

A digital object identifier (DOI).

See the system.Rd file in share/Rd/macros for more details and macro definitions, including macros \packageTitle, \packageDescription, \packageAuthor, \packageMaintainer, \packageDESCRIPTION and \packageIndices.

Packages may also define their own common macros; these would be stored in an .Rd file in man/macros in the package source and will be installed into help/macros when the package is installed. A package may also use the macros from a different package by listing the other package in the 'RdMacros' field in the DESCRIPTION file.

2.14 Encoding

Rd files are text files and so it is impossible to deduce the encoding they are written in unless ASCII: files with 8-bit characters could be UTF-8, Latin-1, Latin-9, KOI8-R, EUC-JP, etc. So an \encoding{} section must be used to specify the encoding if it is not ASCII. (The \encoding{} section must be on a line by itself, and in particular one containing no non-ASCII characters. The encoding declared in the DESCRIPTION file will be used if none is declared in the file.) The Rd files are converted to UTF-8 before parsing and so the preferred encoding for the files themselves is now UTF-8.

For convenience, encoding names 'latin1' and 'latin2' are always recognized: these and 'UTF-8' are likely to work fairly widely. However, this does not mean that all characters in UTF-8 will be recognized, and the coverage of non-Latin characters 10 is fairly low. Using LATEX inputenx (see ?Rd2pdf in R) will give greater coverage of UTF-8.

The \enc command (see Section 2.8 [Insertions], page 70) can be used to provide transliterations which will be used in conversions that do not support the declared encoding.

The LATEX conversion converts the file to UTF-8 from the declared encoding, and includes a

\inputencoding{utf8}

command, and this needs to be matched by a suitable invocation of the \usepackage{inputenc} command. The R utility R CMD Rd2pdf looks at the converted code and includes the encodings used: it might for example use

\usepackage[utf8]{inputenc}

(Use of utf8 as an encoding requires IATEX dated 2003/12/01 or later. Also, the use of Cyrillic characters in 'UTF-8' appears to also need '\usepackage[T2A]{fontenc}', and R CMD Rd2pdf includes this conditionally on the file t2aenc.def being present and environment variable _R_CYRILLIC_TEX_ being set.)

Note that this mechanism works best with Latin letters: the coverage of UTF-8 in \LaTeX is quite low.

2.15 Processing documentation files

There are several commands to process Rd files from the system command line.

Using R CMD Rdconv one can convert R documentation format to other formats, or extract the executable examples for run-time testing. The currently supported conversions are to plain text, HTML and LATEX as well as extraction of the examples.

R CMD Rd2pdf generates PDF output from documentation in Rd files, which can be specified either explicitly or by the path to a directory with the sources of a package. In the latter case, a reference manual for all documented objects in the package is created, including the information in the DESCRIPTION files.

R CMD Sweave and R CMD Stangle process vignette-like documentation files (e.g. Sweave vignettes with extension '.Snw' or '.Rnw', or other non-Sweave vignettes). R CMD Stangle is used to extract the R code fragments.

 $^{^{10}}$ R 2.9.0 added support for UTF-8 Cyrillic characters in IATEX, but on some OSes this will need Cyrillic support added to IATEX, so environment variable <code>_R_CYRILLIC_TEX_</code> may need to be set to a non-empty value to enable this.

The exact usage and a detailed list of available options for all of these commands can be obtained by running R CMD command --help, e.g., R CMD Rdconv --help. All available commands can be listed using R --help (or Rcmd --help under Windows).

All of these work under Windows. You may need to have installed the the tools to build packages from source as described in the "R Installation and Administration" manual, although typically all that is needed is a LATEX installation.

2.16 Editing Rd files

It can be very helpful to prepare .Rd files using a editor which knows about their syntax and will highlight commands, indent to show the structure and detect mis-matched braces, and so on.

The system most commonly used for this is some version of Emacs (including XEmacs) with the ESS package (http://ess.r-project.org/: it is often is installed with Emacs but may need to be loaded, or even installed, separately).

Another is the Eclipse IDE with the Stat-ET plugin (http://www.walware.de/goto/statet), and (on Windows only) Tinn-R (http://sourceforge.net/projects/tinn-r/).

People have also used LATEX mode in a editor, as .Rd files are rather similar to LATEX files.

Some R front-ends provide editing support for .Rd files, for example RStudio (https://rstudio.org/).

3 Tidying and profiling R code

R code which is worth preserving in a package and perhaps making available for others to use is worth documenting, tidying up and perhaps optimizing. The last two of these activities are the subject of this chapter.

3.1 Tidying R code

R treats function code loaded from packages and code entered by users differently. By default code entered by users has the source code stored internally, and when the function is listed, the original source is reproduced. Loading code from a package (by default) discards the source code, and the function listing is re-created from the parse tree of the function.

Normally keeping the source code is a good idea, and in particular it avoids comments being removed from the source. However, we can make use of the ability to re-create a function listing from its parse tree to produce a tidy version of the function, for example with consistent indentation and spaces around operators. If the original source does not follow the standard format this tidied version can be much easier to read.

We can subvert the keeping of source in two ways.

- 1. The option keep.source can be set to FALSE before the code is loaded into R.
- 2. The stored source code can be removed by calling the removeSource() function, for example by

```
myfun <- removeSource(myfun)</pre>
```

In each case if we then list the function we will get the standard layout.

Suppose we have a file of functions myfuns.R that we want to tidy up. Create a file tidy.R containing

```
source("myfuns.R", keep.source = FALSE)
dump(ls(all = TRUE), file = "new.myfuns.R")
```

and run R with this as the source file, for example by R --vanilla < tidy.R or by pasting into an R session. Then the file new.myfuns.R will contain the functions in alphabetical order in the standard layout. Warning: comments in your functions will be lost.

The standard format provides a good starting point for further tidying. Although the deparsing cannot do so, we recommend the consistent use of the preferred assignment operator '<-' (rather than '=') for assignment. Many package authors use a version of Emacs (on a Unix-alike or Windows) to edit R code, using the ESS[S] mode of the ESS Emacs package. See Section "R coding standards" in R Internals for style options within the ESS[S] mode recommended for the source code of R itself.

3.2 Profiling R code for speed

It is possible to profile R code on Windows and most¹ Unix-alike versions of R.

The command Rprof is used to control profiling, and its help page can be consulted for full details. Profiling works by recording at fixed intervals² (by default every 20 msecs) which line in which R function is being used, and recording the results in a file (default Rprof.out in the working directory). Then the function summaryRprof or the command-line utility R CMD Rprof Rprof.out can be used to summarize the activity.

 $^{^{1}}$ R has to be built to enable this, but the option --enable-R-profiling is the default.

² For Unix-alikes these are intervals of CPU time, and for Windows of elapsed time.

```
As an example, consider the following code (from Venables & Ripley, 2002, pp. 225–6).
      library(MASS); library(boot)
      storm.fm <- nls(Time ~ b*Viscosity/(Wt - c), stormer,</pre>
      start = c(b=30.401, c=2.2183))
      st <- cbind(stormer, fit=fitted(storm.fm))</pre>
      storm.bf <- function(rs, i) {</pre>
          st$Time <- st$fit + rs[i]
          tmp <- nls(Time ~ (b * Viscosity)/(Wt - c), st,</pre>
             start = coef(storm.fm))
          tmp$m$getAllPars()
      }
      rs <- scale(resid(storm.fm), scale = FALSE) # remove the mean
      Rprof("boot.out")
      storm.boot <- boot(rs, storm.bf, R = 4999) # slow enough to profile
      Rprof(NULL)
Having run this we can summarize the results by
      R CMD Rprof boot.out
      Each sample represents 0.02 seconds.
      Total run time: 22.52 seconds.
      Total seconds: time spent in function and callees.
      Self seconds: time spent in function alone.
                total
                                      self
       total
                seconds
                           self seconds
                25.22
                                     0.04
       100.0
                            0.2
                                               "boot."
        99.8
                 25.18
                            0.6
                                               "statistic"
                                      0.16
        96.3
                 24.30
                             4.0
                                      1.02
                                               "nls"
                                               "<Anonymous>"
        33.9
                  8.56
                             2.2
                                      0.56
        32.4
                 8.18
                                               "eval"
                             1.4
                                      0.36
        31.8
                 8.02
                             1.4
                                      0.34
                                               ".Call"
        28.6
                 7.22
                             0.0
                                      0.00
                                               "eval.parent"
        28.5
                 7.18
                             0.3
                                      0.08
                                               "model.frame"
        28.1
                             3.5
                                      0.88
                 7.10
                                               "model.frame.default"
        17.4
                  4.38
                             0.7
                                      0.18
                                               "sapply"
        15.0
                  3.78
                             3.2
                                      0.80
                                               "nlsModel"
                  3.16
                                      0.46
                                               "lapply"
        12.5
                             1.8
                  3.10
                             2.7
                                      0.68
                                               "assign"
        12.3
                             %
         %
                  self
                                     total
        self
                seconds
                            total
                                    seconds
                                               name
         5.7
                  1.44
                             7.5
                                      1.88
                                               "inherits"
                                               "nls"
         4.0
                  1.02
                            96.3
                                     24.30
                                               "$"
         3.6
                  0.92
                            3.6
                                     0.92
         3.5
                  0.88
                            28.1
                                      7.10
                                               "model.frame.default"
         3.2
                  0.80
                           15.0
                                      3.78
                                               "nlsModel"
         2.8
                  0.70
                            9.8
                                      2.46
                                               "qr.coef"
         2.7
                  0.68
                            12.3
                                               "assign"
                                      3.10
                  0.64
         2.5
                             2.5
                                      0.64
                                               ".Fortran"
         2.5
                  0.62
                             7.1
                                      1.80
                                               "qr.default"
         2.2
                  0.56
                            33.9
                                      8.56
                                               "<Anonymous>"
         2.1
                  0.54
                                      1.48
                                               "unlist"
                            5.9
                  0.52
                             7.9
                                      2.00
                                               "FUN"
         2.1
```

This often produces surprising results and can be used to identify bottlenecks or pieces of R code that could benefit from being replaced by compiled code.

Two warnings: profiling does impose a small performance penalty, and the output files can be very large if long runs are profiled at the default sampling interval.

Profiling short runs can sometimes give misleading results. R from time to time performs quarbage collection to reclaim unused memory, and this takes an appreciable amount of time

which profiling will charge to whichever function happens to provoke it. It may be useful to compare profiling code immediately after a call to gc() with a profiling run without a preceding call to gc.

More detailed analysis of the output can be achieved by the tools in the CRAN packages proftools (https://CRAN.R-project.org/package=proftools) and profr (https://CRAN.R-project.org/package=profr): in particular these allow call graphs to be studied.

3.3 Profiling R code for memory use

Measuring memory use in R code is useful either when the code takes more memory than is conveniently available or when memory allocation and copying of objects is responsible for slow code. There are three ways to profile memory use over time in R code. All three require R to have been compiled with --enable-memory-profiling, which is not the default, but is currently used for the OS X and Windows binary distributions. All can be misleading, for different reasons.

In understanding the memory profiles it is useful to know a little more about R's memory allocation. Looking at the results of gc() shows a division of memory into Vcells used to store the contents of vectors and Ncells used to store everything else, including all the administrative overhead for vectors such as type and length information. In fact the vector contents are divided into two pools. Memory for small vectors (by default 128 bytes or less) is obtained in large chunks and then parcelled out by R; memory for larger vectors is obtained directly from the operating system.

Some memory allocation is obvious in interpreted code, for example,

```
y < -x + 1
```

allocates memory for a new vector y. Other memory allocation is less obvious and occurs because R is forced to make good on its promise of 'call-by-value' argument passing. When an argument is passed to a function it is not immediately copied. Copying occurs (if necessary) only when the argument is modified. This can lead to surprising memory use. For example, in the 'survey' package we have

```
print.svycoxph <- function (x, ...)
{
    print(x$survey.design, varnames = FALSE, design.summaries = FALSE,
    ...)
         x$call <- x$printcall
        NextMethod()
}</pre>
```

It may not be obvious that the assignment to x call will cause the entire object x to be copied. This copying to preserve the call-by-value illusion is usually done by the internal C function duplicate.

The main reason that memory-use profiling is difficult is garbage collection. Memory is allocated at well-defined times in an R program, but is freed whenever the garbage collector happens to run.

3.3.1 Memory statistics from Rprof

The sampling profiler Rprof described in the previous section can be given the option memory.profiling=TRUE. It then writes out the total R memory allocation in small vectors, large vectors, and cons cells or nodes at each sampling interval. It also writes out the number of calls to the internal function duplicate, which is called to copy R objects. summaryRprof provides summaries of this information. The main reason that this can be misleading is that the memory use is attributed to the function running at the end of the sampling interval. A second reason is that garbage collection can make the amount of memory in use decrease, so a function appears to use little memory. Running under gctorture helps with both problems: it

slows down the code to effectively increase the sampling frequency and it makes each garbage collection release a smaller amount of memory. Changing the memory limits with mem.limits() may also be useful, to see how the code would run under different memory conditions.

3.3.2 Tracking memory allocations

The second method of memory profiling uses a memory-allocation profiler, Rprofmem(), which writes out a stack trace to an output file every time a large vector is allocated (with a user-specified threshold for 'large') or a new page of memory is allocated for the R heap. Summary functions for this output are still being designed.

Running the example from the previous section with

```
> Rprofmem("boot.memprof",threshold=1000)
> storm.boot <- boot(rs, storm.bf, R = 4999)
> Rprofmem(NULL)
```

shows that apart from some initial and final work in **boot** there are no vector allocations over 1000 bytes.

3.3.3 Tracing copies of an object

The third method of memory profiling involves tracing copies made of a specific (presumably large) R object. Calling tracemem on an object marks it so that a message is printed to standard output when the object is copied *via* duplicate or coercion to another type, or when a new object of the same size is created in arithmetic operations. The main reason that this can be misleading is that copying of subsets or components of an object is not tracked. It may be helpful to use tracemem on these components.

In the example above we can run tracemem on the data frame st

```
> tracemem(st)
[1] "<0x9abd5e0>"
> storm.boot <- boot(rs, storm.bf, R = 4)
memtrace[0x9abd5e0->0x92a6d08]: statistic boot
memtrace[0x92a6d08->0x92a6d80]: $<-.data.frame $<- statistic boot
memtrace[0x92a6d80->0x92a6df8]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x9271318]: statistic boot
memtrace[0x9271318->0x9271390]: $<-.data.frame $<- statistic boot
memtrace[0x9271390->0x9271408]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x914f558]: statistic boot
memtrace[0x914f558->0x914f5f8]: $<-.data.frame $<- statistic boot
memtrace[0x914f5f8->0x914f670]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x972cbf0]: statistic boot
memtrace[0x972cbf0->0x972cc68]: $<-.data.frame $<- statistic boot
memtrace[0x972cc68->0x972cd08]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x98ead98]: statistic boot
memtrace[0x98ead98->0x98eae10]: $<-.data.frame $<- statistic boot</pre>
memtrace[0x98eae10->0x98eae88]: $<-.data.frame $<- statistic boot
```

The object is duplicated fifteen times, three times for each of the R+1 calls to storm.bf. This is surprising, since none of the duplications happen inside nls. Stepping through storm.bf in the debugger shows that all three happen in the line

```
st$Time <- st$fit + rs[i]</pre>
```

Data frames are slower than matrices and this is an example of why. Using tracemem(st\$Viscosity) does not reveal any additional copying.

3.4 Profiling compiled code

Profiling compiled code is highly system-specific, but this section contains some hints gleaned from various R users. Some methods need to be different for a compiled executable and for dynamic/shared libraries/objects as used by R packages. We know of no good way to profile DLLs on Windows.

3.4.1 Linux

Options include using sprof for a shared object, and oprofile (see http://oprofile.sourceforge.net/) and perf (see https://perf.wiki.kernel.org/index.php/Tutorial) for any executable or shared object.

3.4.1.1 sprof

You can select shared objects to be profiled with **sprof** by setting the environment variable LD_PROFILE. For example

```
% setenv LD_PROFILE /path/to/R_HOME/library/stats/libs/stats.so
R
... run the boot example
% sprof /path/to/R_HOME/library/stats/libs/stats.so \
   /var/tmp/path/to/R_HOME/library/stats/libs/stats.so.profile
```

Flat profile:

Each sample counts as 0.01 seconds.

% (cumulative	self		self	total	
time	seconds	seconds	calls	us/call	us/call	name
76.19	0.32	0.32	0	0.00		numeric_deriv
16.67	0.39	0.07	0	0.00		nls_iter
7.14	0.42	0.03	0	0.00		getListElement

```
rm /var/tmp/path/to/R_HOME/library/stats/libs/stats.so.profile
... to clean up ...
```

It is possible that root access is needed to create the directories used for the profile data.

3.4.1.2 oprofile and operf

"sprintf"

The oprofile project has two modes of operation. In what is now called 'legacy' mode, it is uses a daemon to collect information on a process (see below). Since version 0.9.8 (August 2012), the preferred mode is to use operf, so we discuss that first. The modes differ in how the profiling data is collected: it is analysed by tools such as opreport and oppannote in both.

Here is an example on x86_64 Linux using R 3.0.2. File pvec.R contains the part of the examples from pvec in package parallel:

```
library(parallel)
  N <- 1e6
  dates <- sprintf('%04d-%02d-%02d', as.integer(2000+rnorm(N)),</pre>
                     as.integer(runif(N, 1, 12)), as.integer(runif(N, 1, 28)))
  system.time(a <- as.POSIXct(dates, format = "%Y-%m-%d"))</pre>
with timings from the final step
            system elapsed
     user
    0.371
             0.237
                      0.612
R-level profiling by Rprof shows
                         self.time self.pct total.time total.pct
  "strptime"
                              1.70
                                       41.06
                                                    1.70
                                                             41.06
  "as.POSIXct.POSIX1t"
                                       33.82
                                                    1.42
                                                             34.30
                              1.40
```

17.87

0.98

23.67

so the conversion from character to POSIX1t takes most of the time.

0.74

```
This can be run under operf and analysed by
  operf R -f pvec.R
  opreport
  opreport -1 /path/to/R_HOME/bin/exec/R
  opannotate --source /path/to/R_HOME/bin/exec/R
  ## And for the system time
  opreport -1 /lib64/libc.so.6
The first report shows where (which library etc) the time was spent:
  CPU_CLK_UNHALT...
    samples|
  _____
     166761 99.9161 Rdev
  CPU_CLK_UNHALT...
    samples|
      70586 42.3276 no-vmlinux
      56963 34.1585 libc-2.16.so
      36922 22.1407 R
       1584 0.9499 stats.so
        624 0.3742 libm-2.16.so
```

The rest of the output is voluminous, and only extracts are shown below.

Most of the time within R is spent in

samples	%	image nam	ne symbol name
10397	28.5123	R	$R_gc_internal$
5683	15.5848	R	do_sprintf
3036	8.3258	R	do_asPOSIXct
2427	6.6557	R	do_strptime
2421	6.6392	R	Rf_mkCharLenCE
1480	4.0587	R	$w_strptime_internal$
1202	3.2963	R	Rf_qnorm5
1165	3.1948	R	unif_rand
675	1.8511	R	mktime0
617	1.6920	R	makelt
617	1.6920	R	validate_tm
584	1.6015	R	${\tt day_of_the_week}$

opannotate shows that 31% of the time in R is spent in memory.c, 21% in datetime.c and 7% in Rstrptime.h. The analysis for libc showed that calls to wcsftime dominated, so those calls were cached for R 3.0.3: the time spent in no-vmlinux (the kernel) was reduced dramatically.

On platforms which support it, call graphs can be produced by opcontrol --callgraph if collected via operf --callgraph.

The profiling data is by default stored in sub-directory oprofile_data of the current directory, which can be removed at the end of the session.

Another example, from sm (https://CRAN.R-project.org/package=sm) version 2.2-5.4. The example for sm.variogram took a long time:

```
system.time(example(sm.variogram))
...
```

```
user system elapsed 5.543 3.202 8.785
```

including a lot of system time. Profiling just the slow part, the second plot, showed

so the system time was almost all in the Linux kernel. It is possible to dig deeper if you have a matching uncompressed kernel with debug symbols to specify *via* --vmlinux: we did not.

In 'legacy' mode oprofile works by running a daemon which collects information. The daemon must be started as root, e.g.

```
% su
  % opcontrol --no-vmlinux
  % (optional, some platforms) opcontrol --callgraph=5
  % opcontrol --start
  % exit
Then as a user
  % R.
  ... run the boot example
  % opcontrol --dump
  % opreport -1 /path/to/R_HOME/library/stats/libs/stats.so
  samples %
                    symbol name
  1623
           75.5939 anonymous symbol from section .plt
           16.2552 numeric_deriv
  349
           5.2632 nls_iter
  113
            2.8878 getListElement
  % opreport -1 /path/to/R_HOME/bin/exec/R
  samples %
                    symbol name
  76052
           11.9912 Rf_eval
           8.6198 Rf_findVarInFrame3
  54670
            5.9622 Rf_allocVector
  37814
            4.9649 Rf_duplicate
  31489
            4.4496 Rf_protect
  28221
  26485
            4.1759 Rf_cons
            3.7289 Rf_matchArgs
  23650
            3.3250 Rf_findFun
  21088
            3.1526 findVarLocInFrame
  19995
  14871
            2.3447 Rf_evalList
            2.1749 R_Newhashpjw
  13794
  13522
            2.1320 R_gc_internal
```

Shutting down the profiler and clearing the records needs to be done as root.

3.4.2 Solaris

On 64-bit (only) Solaris, the standard profiling tool gprof collects information from shared objects compiled with -pg.

3.4.3 OS X

Developers have recommended sample (or Sampler.app, which is a GUI version), Shark (in version of Xcode up to those for Snow Leopard), and Instruments (part of Xcode, see https://developer.apple.com/library/mac/#documentation/DeveloperTools/Conceptual/InstrumentsUserGuide/Introduction/Introduction.html).

4 Debugging

This chapter covers the debugging of R extensions, starting with the ways to get useful error information and moving on to how to deal with errors that crash R. For those who prefer other styles there are contributed packages such as **debug** (https://CRAN.R-project.org/package=debug) on CRAN (described in an article in R-News 3/3 (https://CRAN.R-project.org/doc/Rnews/Rnews_2003-3.pdf)). (There are notes from 2002 provided by Roger Peng at http://www.biostat.jhsph.edu/~rpeng/docs/R-debug-tools.pdf which provide complementary examples to those given here.)

4.1 Browsing

Most of the R-level debugging facilities are based around the built-in browser. This can be used directly by inserting a call to browser() into the code of a function (for example, using fix(my_function)). When code execution reaches that point in the function, control returns to the R console with a special prompt. For example

```
> fix(summary.data.frame) ## insert browser() call after for() loop
> summary(women)
Called from: summary.data.frame(women)
Browse[1] > ls()
 [1] "digits" "i"
                        "lbs"
                                           "maxsum" "nm"
                                                              "nr"
                                                                       "nv"
 [9] "object" "sms"
                        "z"
Browse[1] > maxsum
[1] 7
Browse[1]>
     height
                     weight
Min.
        :58.0
                Min.
                       :115.0
 1st Qu.:61.5
                1st Qu.:124.5
Median:65.0
                Median :135.0
Mean
        :65.0
                Mean
                        :136.7
 3rd Qu.:68.5
                3rd Qu.:148.0
        :72.0
Max.
                Max.
                        :164.0
> rm(summary.data.frame)
```

At the browser prompt one can enter any R expression, so for example 1s() lists the objects in the current frame, and entering the name of an object will¹ print it. The following commands are also accepted

• n

Enter 'step-through' mode. In this mode, hitting return executes the next line of code (more precisely one line and any continuation lines). Typing c will continue to the end of the current context, e.g. to the end of the current loop or function.

C

In normal mode, this quits the browser and continues execution, and just return works in the same way. cont is a synonym.

where

This prints the call stack. For example

```
> summary(women)
Called from: summary.data.frame(women)
```

¹ With the exceptions of the commands listed below: an object of such a name can be printed *via* an explicit call to print.

```
Browse[1]> where
where 1: summary.data.frame(women)
where 2: summary(women)
Browse[1]>
```

• Q

Quit both the browser and the current expression, and return to the top-level prompt.

Errors in code executed at the browser prompt will normally return control to the browser prompt. Objects can be altered by assignment, and will keep their changed values when the browser is exited. If really necessary, objects can be assigned to the workspace from the browser prompt (by using <<- if the name is not already in scope).

4.2 Debugging R code

Suppose your R program gives an error message. The first thing to find out is what R was doing at the time of the error, and the most useful tool is traceback(). We suggest that this is run whenever the cause of the error is not immediately obvious. Daily, errors are reported to the R mailing lists as being in some package when traceback() would show that the error was being reported by some other package or base R. Here is an example from the regression suite.

The calls to the active frames are given in reverse order (starting with the innermost). So we see the error message comes from an explicit check in glm.fit. (traceback() shows you all the lines of the function calls, which can be limited by setting option "deparse.max.lines".)

Sometimes the traceback will indicate that the error was detected inside compiled code, for example (from ?nls)

```
Error in nls(y ~ a + b * x, start = list(a = 0.12345, b = 0.54321), trace = TRUE) : step factor 0.000488281 reduced below 'minFactor' of 0.000976563 > traceback()  
2: .Call(R_nls_iter, m, ctrl, trace)  
1: nls(y ~ a + b * x, start = list(a = 0.12345, b = 0.54321), trace = TRUE)
```

This will be the case if the innermost call is to .C, .Fortran, .Call, .External or .Internal, but as it is also possible for such code to evaluate R expressions, this need not be the innermost call, as in

```
assign("resid", .swts * (lhs - assign("rhs", getRHS(), envir = thisEnv)),
envir = thisEnv)
    assign("dev", sum(resid^2), envir = thisEnv)
    assign("QR", qr(.swts * attr(rhs, "gradient")), envir = thisEnv)
    return(QR$rank < min(dim(QR$qr)))
}(c(-0.00760232418963883, 1.00119632515036))
2: .Call(R_nls_iter, m, ctrl, trace)
1: nls(yeps ~ gm(a, b, x), start = list(a = 0.12345, b = 0.54321))</pre>
```

Occasionally traceback() does not help, and this can be the case if S4 method dispatch is involved. Consider the following example

```
> xyd <- new("xyloc", x=runif(20), y=runif(20))
Error in as.environment(pkg) : no item called "package:S4nswv"
on the search list
Error in initialize(value, ...) : S language method selection got
an error when called from internal dispatch for function 'initialize'
> traceback()
2: initialize(value, ...)
1: new("xyloc", x = runif(20), y = runif(20))
```

which does not help much, as there is no call to as.environment in initialize (and the note "called from internal dispatch" tells us so). In this case we searched the R sources for the quoted call, which occurred in only one place, methods:::asEnvironmentPackage. So now we knew where the error was occurring. (This was an unusually opaque example.)

The error message

evaluation nested too deeply: infinite recursion / options(expressions=)? can be hard to handle with the default value (5000). Unless you know that there actually is deep recursion going on, it can help to set something like

```
options(expressions=500)
```

and re-run the example showing the error.

Sometimes there is warning that clearly is the precursor to some later error, but it is not obvious where it is coming from. Setting options(warn = 2) (which turns warnings into errors) can help here.

Once we have located the error, we have some choices. One way to proceed is to find out more about what was happening at the time of the crash by looking a *post-mortem* dump. To do so, set options(error=dump.frames) and run the code again. Then invoke debugger() and explore the dump. Continuing our example:

```
> options(error = dump.frames)
> glm(resp ~ 0 + predictor, family = binomial(link ="log"))
Error: no valid set of coefficients has been found: please supply starting values
```

which is the same as before, but an object called last.dump has appeared in the workspace. (Such objects can be large, so remove it when it is no longer needed.) We can examine this at a later time by calling the function debugger.

```
> debugger()
Message: Error: no valid set of coefficients has been found: please supply starting values
Available environments had calls:
1: glm(resp ~ 0 + predictor, family = binomial(link = "log"))
2: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart, mus
3: stop("no valid set of coefficients has been found: please supply starting values
Enter an environment number, or 0 to exit Selection:
```

which gives the same sequence of calls as traceback, but in outer-first order and with only the first line of the call, truncated to the current width. However, we can now examine in more detail what was happening at the time of the error. Selecting an environment opens the browser

in that frame. So we select the function call which spawned the error message, and explore some of the variables (and execute two function calls).

```
Enter an environment number, or 0 to exit Selection: 2
Browsing in the environment with call:
   glm.fit(x = X, y = Y, weights = weights, start = start, etas
Called from: debugger.look(ind)
Browse[1] > ls()
 [1] "aic"
                  "boundary"
                               "coefold"
                                            "control"
                                                          "conv"
[6] "dev"
                  "dev.resids" "devold"
                                            "EMPTY"
                                                          "eta"
[11] "etastart"
                  "family"
                               "fit"
                                            "good"
                                                          "intercept"
[16] "iter"
                  "linkinv"
                                            "mu.eta"
                               "mu"
                                                          "mu.eta.val"
                  "n"
                               "ngoodobs"
[21] "mustart"
                                            "nobs"
                                                          "nvars"
[26] "offset"
                  "start"
                               "valideta"
                                            "validmu"
                                                          "variance"
[31] "varmu"
                  "w"
                               "weights"
                                            "x"
                                                          "xnames"
[36] "y"
                  "ynames"
                               "z"
Browse[1]> eta
                                3
   1
  0.000000e + 00 - 2.235357e - 06 - 1.117679e - 05 - 5.588393e - 05 - 2.794197e - 04 \\
                  7
                               8
                                              9
-1.397098e-03 -6.985492e-03 -3.492746e-02 -1.746373e-01
Browse[1]> valideta(eta)
[1] TRUE
Browse[1]> mu
        2
1.0000000 0.9999978 0.9999888 0.9999441 0.9997206 0.9986039 0.9930389 0.9656755
0.8397616
Browse[1] > validmu(mu)
[1] FALSE
Browse[1]> c
Available environments had calls:
1: glm(resp ~ 0 + predictor, family = binomial(link = "log"))
2: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart
3: stop("no valid set of coefficients has been found: please supply starting v
Enter an environment number, or 0 to exit Selection: 0
> rm(last.dump)
```

Because last.dump can be looked at later or even in another R session, post-mortem debugging is possible even for batch usage of R. We do need to arrange for the dump to be saved: this can be done either using the command-line flag --save to save the workspace at the end of the run, or via a setting such as

```
> options(error = quote({dump.frames(to.file=TRUE); q()}))
```

See the help on dump.frames for further options and a worked example.

An alternative error action is to use the function recover():

```
> options(error = recover)
> glm(resp ~ 0 + predictor, family = binomial(link = "log"))
Error: no valid set of coefficients has been found: please supply starting values
Enter a frame number, or 0 to exit

1: glm(resp ~ 0 + predictor, family = binomial(link = "log"))
2: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart
Selection:
```

which is very similar to dump.frames. However, we can examine the state of the program directly, without dumping and re-loading the dump. As its help page says, recover can be routinely used as the error action in place of dump.calls and dump.frames, since it behaves like dump.frames in non-interactive use.

Post-mortem debugging is good for finding out exactly what went wrong, but not necessarily why. An alternative approach is to take a closer look at what was happening just before the error, and a good way to do that is to use debug. This inserts a call to the browser at the beginning of the function, starting in step-through mode. So in our example we could use

```
> debug(glm.fit)
> glm(resp ~ 0 + predictor, family = binomial(link ="log"))
debugging in: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart,
    mustart = mustart, offset = offset, family = family, control = control,
    intercept = attr(mt, "intercept") > 0)
debug: {
## lists the whole function
Browse[1]>
debug: x <- as.matrix(x)</pre>
Browse[1]> start
[1] -2.235357e-06
debug: eta <- drop(x %*% start)</pre>
Browse[1]> eta
                                3
0.000000e+00 -2.235357e-06 -1.117679e-05 -5.588393e-05 -2.794197e-04
                7
                             8
                                             9
-1.397098e-03 -6.985492e-03 -3.492746e-02 -1.746373e-01
Browse[1]>
debug: mu <- linkinv(eta <- eta + offset)</pre>
Browse[1]> mu
                                        5
                                                  6
1.0000000 0.9999978 0.9999888 0.9999441 0.9997206 0.9986039 0.9930389 0.9656755
```

(The prompt Browse[1] > indicates that this is the first level of browsing: it is possible to step into another function that is itself being debugged or contains a call to browser().)

debug can be used for hidden functions and S3 methods by e.g. debug(stats:::predict.Arima). (It cannot be used for S4 methods, but an alternative is given on the help page for debug.) Sometimes you want to debug a function defined inside another function, e.g. the function arimafn defined inside arima. To do so, set debug on the outer function (here arima) and step through it until the inner function has been defined. Then call debug on the inner function (and use c to get out of step-through mode in the outer function).

To remove debugging of a function, call undebug with the argument previously given to debug; debugging otherwise lasts for the rest of the R session (or until the function is edited or otherwise replaced).

trace can be used to temporarily insert debugging code into a function, for example to insert a call to browser() just before the point of the error. To return to our running example

```
## first get a numbered listing of the expressions of the function
> page(as.list(body(glm.fit)), method="print")
> trace(glm.fit, browser, at=22)
Tracing function "glm.fit" in package "stats"
[1] "glm.fit"
> glm(resp ~ 0 + predictor, family = binomial(link ="log"))
Tracing glm.fit(x = X, y = Y, weights = weights, start = start,
    etastart = etastart, .... step 22
Called from: eval(expr, envir, enclos)
Browse[1] > n
## and single-step from here.
> untrace(glm.fit)
```

For your own functions, it may be as easy to use fix to insert temporary code, but trace can help with functions in a namespace (as can fixInNamespace). Alternatively, use trace(,edit=TRUE) to insert code visually.

4.3 Checking memory access

Errors in memory allocation and reading/writing outside arrays are very common causes of crashes (e.g., segfaults) on some machines. Often the crash appears long after the invalid memory access: in particular damage to the structures which R itself has allocated may only become apparent at the next garbage collection (or even at later garbage collections after objects have been deleted).

Note that memory access errors may be seen with LAPACK, BLAS, OpenMP and Javausing packages: some at least of these seem to be intentional, and some are related to passing characters to Fortran.

Some of these tools can detect mismatched allocation and deallocation. C++ programmers should note that memory allocated by new [] must be freed by delete [], other uses of new by delete, and memory allocated by malloc, calloc and realloc by free. Some platforms will tolerate mismatches (perhaps with memory leaks) but others will segfault.

4.3.1 Using gctorture

We can help to detect memory problems in R objects earlier by running garbage collection as often as possible. This is achieved by gctorture(TRUE), which as described on its help page

Provokes garbage collection on (nearly) every memory allocation. Intended to ferret out memory protection bugs. Also makes R run *very* slowly, unfortunately.

The reference to 'memory protection' is to missing C-level calls to PROTECT/UNPROTECT (see Section 5.9.1 [Garbage Collection], page 112) which if missing allow R objects to be garbage-collected when they are still in use. But it can also help with other memory-related errors.

Normally running under gctorture (TRUE) will just produce a crash earlier in the R program, hopefully close to the actual cause. See the next section for how to decipher such crashes.

It is possible to run all the examples, tests and vignettes covered by R CMD check under gctorture(TRUE) by using the option --use-gct.

The function gctorture2 provides more refined control over the GC torture process. Its arguments step, wait and inhibit_release are documented on its help page. Environment variables can also be used at the start of the R session to turn on GC torture: R_GCTORTURE corresponds to the step argument to gctorture2, R_GCTORTURE_WAIT to wait, and R_GCTORTURE_INHIBIT_RELEASE to inhibit_release.

If R is configured with --enable-strict-barrier then a variety of tests for the integrity of the write barrier are enabled. In addition tests to help detect protect issues are enabled:

- All GCs are full GCs.
- New nodes in small node pages are marked as NEWSXP on creation.
- After a GC all free nodes that are not of type NEWSXP are marked as type FREESXP and their previous type is recorded.
- Most calls to accessor functions check their SEXP inputs and SEXP outputs and signal an error if a FREESXP is found. The address of the node and the old type are included in the error message.

R CMD check --use-gct can be set to use gctorture2(n) rather than gctorture(TRUE) by setting environment variable _R_CHECK_GCT_N_ to a positive integer value to be used as n.

Used with a debugger and with gctorture or gctorture2 this mechanism can be helpful in isolating memory protect problems.

4.3.2 Using valgrind

If you have access to Linux on a common CPU type or supported versions of OS X^2 you can use valgrind (http://www.valgrind.org/, pronounced to rhyme with 'tinned') to check for possible problems. To run some examples under valgrind use something like

```
R -d valgrind --vanilla < mypkg-Ex.R
R -d "valgrind --tool=memcheck --leak-check=full" --vanilla < mypkg-Ex.R
```

where mypkg-Ex.R is a set of examples, e.g. the file created in mypkg.Rcheck by R CMD check. Occasionally this reports memory reads of 'uninitialised values' that are the result of compiler optimization, so can be worth checking under an unoptimized compile: for maximal information use a build with debugging symbols. We know there will be some small memory leaks from readline and R itself — these are memory areas that are in use right up to the end of the R session. Expect this to run around 20x slower than without valgrind, and in some cases much slower than that. Several versions of valgrind were not happy with some optimized BLASes that use CPU-specific instructions so you may need to build a version of R specifically to use with valgrind.

On platforms where valgrind is installed you can build a version of R with extra instrumentation to help valgrind detect errors in the use of memory allocated from the R heap. The configure option is --with-valgrind-instrumentation=level, where level is 0, 1 or 2. Level 0 is the default and does not add any anything. Level 1 will detect some uses³ of uninitialised memory and has little impact on speed (compared to level 0). Level 2 will detect many other memory-use bugs⁴ but make R much slower when running under valgrind. Using this in conjunction with gctorture can be even more effective (and even slower).

An example of valgrind output is

```
==12539== Invalid read of size 4
==12539==
             at 0x1CDF6CBE: csc_compTr (Mutils.c:273)
==12539==
             by 0x1CE07E1E: tsc_transpose (dtCMatrix.c:25)
==12539==
             by 0x80A67A7: do_dotcall (dotcode.c:858)
==12539==
             by 0x80CACE2: Rf_eval (eval.c:400)
==12539==
             by 0x80CB5AF: R_execClosure (eval.c:658)
==12539==
            by 0x80CB98E: R_execMethod (eval.c:760)
==12539==
            by Ox1B93DEFA: R_standardGeneric (methods_list_dispatch.c:624)
==12539==
            by 0x810262E: do_standardGeneric (objects.c:1012)
==12539==
            by 0x80CAD23: Rf_eval (eval.c:403)
==12539==
             by 0x80CB2F0: Rf_applyClosure (eval.c:573)
==12539==
             by 0x80CADCC: Rf_eval (eval.c:414)
            by 0x80CAA03: Rf_eval (eval.c:362)
==12539==
==12539== Address 0x1COD2EA8 is 280 bytes inside a block of size 1996 alloc'd
==12539==
             at 0x1B9008D1: malloc (vg_replace_malloc.c:149)
==12539==
             by 0x80F1B34: GetNewPage (memory.c:610)
==12539==
             by 0x80F7515: Rf_allocVector (memory.c:1915)
```

This example is from an instrumented version of R, while tracking down a bug in the Matrix (https://CRAN.R-project.org/package=Matrix) package in 2006. The first line indicates that R has tried to read 4 bytes from a memory address that it does not have access to. This is followed by a C stack trace showing where the error occurred. Next is a description of the memory that was accessed. It is inside a block allocated by malloc, called from GetNewPage, that is, in the internal R heap. Since this memory all belongs to R, valgrind would not (and did not) detect the problem in an uninstrumented build of R. In this example the stack trace was enough to isolate and fix the bug, which was in tsc_transpose, and in this example running under gctorture() did not provide any additional information. When the stack trace is not

 $^{^{2}}$ at the time of writing mainly for 10.9 with some support for 10.8, none for the current 10.10.

 $^{^3}$ Those in some numeric, logical, integer, raw, complex vectors and in memory allocated by R_alloc.

⁴ including using the data sections of R vectors after they are freed.

sufficiently informative the option --db-attach=yes to valgrind may be helpful. This starts a post-mortem debugger (by default gdb) so that variables in the C code can be inspected (see Section 4.4.2 [Inspecting R objects], page 97).

valgrind is good at spotting the use of uninitialized values: use option --track-origins=yes to show where these originated from. What it cannot detect is the misuse of arrays allocated on the stack: this includes C automatic variables and some⁵ Fortran arrays.

It is possible to run all the examples, tests and vignettes covered by R CMD check under valgrind by using the option --use-valgrind. If you do this you will need to select the valgrind options some other way, for example by having a ~/.valgrindrc file containing

- --leak-check=full
- --track-origins=yes

or setting the environment variable VALGRIND_OPTS.

On OS X you may need to ensure that debugging symbols are made available (so valgrind reports line numbers in files). This can usually be done with the valgrind option --dsymutil=yes to ask for the symbols to be dumped when the .so file is loaded. This will not work where packages are installed into a system area (such as the R.framework) and can be slow. Installing packages with R CMD INSTALL --dsym installs the dumped symbols. (This can also be done by setting environment variable PKG_MAKE_DSYM to a non-empty value before the INSTALL.)

This section has described the use of memtest, the default (and most useful) of valgrind's tools. There are others described in its documentation: helgrind can be useful for threaded programs.

4.3.3 Using the Address Sanitizer

AddressSanitizer ('ASan') is a tool with similar aims to the memory checker in valgrind. It is available with suitable builds⁶ of gcc and clang on common Linux and OS X platforms. See http://clang.llvm.org/docs/UsersManual.html#controlling-code-generation, http://clang.llvm.org/docs/AddressSanitizer.html and https://code.google.com/p/address-sanitizer/.

More thorough checks of C++ code are done if the C++ library has been 'annotated': at the time of writing this applied to std::vector in libc++ for use with clang and gives rise to 'container-overflow' reports.

It requires code to have been compiled and linked with -fsanitize=address and compiling with -fno-omit-frame-pointer will give more legible reports. It has a runtime penalty of 2–3x, extended compilation times and uses substantially more memory, often 1–2GB, at run time. On 64-bit platforms it reserves (but does not allocate) 16–20TB of virtual memory: restrictive shell settings can cause problems.

By comparison with valgrind, ASan can detect misuse of stack and global variables but not the use of uninitialized memory.

Recent versions return symbolic addresses for the location of the error provided llvm-symbolizer⁷ is on the path: if it is available but not on the path or has been renamed⁸, one can use an environment variable, e.g.

ASAN_SYMBOLIZER_PATH=/path/to/llvm-symbolizer

 $^{^{5}\,}$ small fixed-size arrays by default in ${\tt gfortran},$ for example.

⁶ currently only on 'ix86'/'x86_64' Linux and OS X (including the builds in Xcode 7 but not earlier Apple releases). On some platforms the runtime library, **libasan**, needs to be installed separately, and for checking C++ you may also need **libubsan**.

⁷ part of the LLVM project and in distributed in llvm RPMs and .debs on Linux. It is not currently shipped by Apple.

 $^{^{8}}$ as Ubuntu does.

An alternative is to pipe the output through asan_symbolize.py⁹ and perhaps then (for compiled C++ code) c++filt. (On OS X, you may need to run dsymutil to get line-number reports.)

The simplest way to make use of this is to build a version of R with something like

```
CC="gcc -std=gnu99 -fsanitize=address"

CFLAGS="-fno-omit-frame-pointer -g -02 -Wall -pedantic -mtune=native"
```

which will ensure that the libasan run-time library is compiled into the R executable. However this check can be enabled on a per-package basis by using a ~/.R/Makevars file like

```
CC = gcc -std=gnu99 -fsanitize=address -fno-omit-frame-pointer
CXX = g++ -fsanitize=address -fno-omit-frame-pointer
F77 = gfortran -fsanitize=address
FC = gfortran -fsanitize=address
```

(Note that -fsanitize=address has to be part of the compiler specification to ensure it is used for linking. These settings will not be honoured by packages which ignore ~/.R/Makevars.) It will be necessary to build R with

```
MAIN_LDFLAGS = -fsanitize=address
```

to link the runtime libraries into the R executable if it was not specified as part of 'CC' when R was built.

For options available *via* the environment variable ASAN_OPTIONS see https://code.google.com/p/address-sanitizer/wiki/Flags#Run-time_flags. With gcc additional control is available *via* the --params flag: see its man page.

For more detailed information on an error, R can be run under a debugger with a breakpoint set before the address sanitizer report is produced: for gdb or lldb you could use

```
break __asan_report_error
(See https://code.google.com/p/address-sanitizer/wiki/AddressSanitizer#gdb.)
```

4.3.3.1 Using the Leak Sanitizer

For x86_64 Linux there is a leak sanitizer, 'LSan': see https://code.google.com/p/address-sanitizer/wiki/LeakSanitizer. This is available on recent versions of gcc and clang, and where available is compiled in as part of ASan.

One way to invoke this from an ASan-enabled build is by the environment variable

```
ASAN_OPTIONS='detect_leaks=1'
```

However, this was made the default for clang 3.5 and gcc 5.1.0.

When LSan is enabled, leaks give the process a failure error status (by default 23). For an R package this means the R process, and as the parser retains some memory to the end of the process, if R itself was built against ASan, all runs will have a failure error status (which may include running R as part of building R itself).

To disable both this and some strict checking use

```
setenv ASAN_OPTIONS 'alloc_dealloc_mismatch=0:detect_leaks=0:detect_odr_violation=0'
```

LSan also has a 'stand-alone' mode where it is compiled in using -fsanitize=leak and avoids the run-time overhead of ASan.

⁹ installed on some Linux systems as asan_symbolize, and obtainable from https://llvm.org/svn/llvm-project/compiler-rt/trunk/lib/asan/scripts/asan_symbolize.py: it makes use of llvm-symbolizer if available.

4.3.4 Using the Undefined Behaviour Sanitizer

'Undefined behaviour' is where the language standard does not require particular behaviour from the compiler. Examples include division by zero (where for doubles R requires the ISO/IEC 60559 behaviour but C/C++ do not), use of zero-length arrays, shifts too far for signed types (e.g. int x, y; y = x << 31;), out-of-range coercion, invalid C++ casts and mis-alignment. Not uncommon examples of out-of-range coercion in R packages are attempts to coerce a NaN or infinity to type int or NA_INTEGER to an unsigned type such as size_t. Also common is y[x-1] forgetting that x might be NA_INTEGER.

'UBSanitizer' is a tool for C/C++ source code selected by -fsanitize=undefined in suitable builds¹⁰ of clang and GCC. Its (main) runtime library is linked into each package's DLL, so it is less often needed to be included in MAIN_LDFLAGS.

Some versions have greatly increased compilation times on a few files¹¹.

This sanitizer can be combined with the Address Sanitizer by -fsanitize=undefined,address (where both are supported).

Finer control of what is checked can be achieved by other options: for clang see http://clang.llvm.org/docs/UsersManual.html#controlling-code-generation.¹² The current set for clang is (on a single line):

```
-fsanitize=alignment,bool,bounds,enum,float-cast-overflow,
float-divide-by-zero,function,integer-divide-by-zero,non-null-attribute,
null,object-size,return,returns-nonnull-attribute,shift,
signed-integer-overflow,unreachable,vla-bound,vptr
```

a subset of which could be combined with address, or use something like

-fsanitize=undefined -fno-sanitize=float-divide-by-zero (function, return and vptr apply only to C++). In addition,

-fsanitize=unsigned-integer-overflow

is available as a separate option in some versions of clang (not enabled by -fsanitize=undefined).

clang 3.5 and later may need

```
-fsanitize=undefined -fno-sanitize=float-divide-by-zero,vptr
```

for C++ code (in CXX and CXX1X) as the run-time library for vptr needs to be linked into the main R executable (and that would need to be linked by clang++, not clang: you could try building R with something like

```
MAIN_LD="clang++ -fsanitize=undefined" and perhaps for clang 3.7.x

R_OPENMP_CFLAGS="-fopenmp=libomp" or add -lclang_rt.asan_cxx-x86_64<sup>13</sup> or similar to LD_FLAGS).
```

See https://gcc.gnu.org/onlinedocs/gcc/Debugging-Options.html (or the manual for your version of GCC, installed or via https://gcc.gnu.org/onlinedocs/) for the options supported by GCC: 5.3 supports

```
-fsanitize=alignment,bool,bounds,enum,float-cast-overflow,
integer-divide-by-zero,non-null-attribute,null,object-size,
return,returns-nonnull-attribute,shift,signed-integer-overflow,
```

 $^{^{10}}$ On some platforms the runtime library, **libubsan**, needs to be installed separately.

e.g. src/main/dotcode.c and parts of the Matrix sources with clang 3.7.0 amd later.

or the user manual for your version of clang, e.g. http://llvm.org/releases/3.6.2/tools/docs/ UsersManual.html.

 $^{^{13}}$ This includes the C++ UBSAN handlers, despite its name.

unreachable, vla-bound, vptr

with

```
-fsanitize=float-divide-by-zero
```

as a separate option not enabled by -fsanitize=undefined (and not desirable for R uses). At the time of writing the object-size and vptr checks produced many warnings on GCC's own C++ headers, so should be disabled.

GCC 6 will add

-fsanitize=bounds-strict

an extension of bounds to 'flexible array member-like arrays'.

Other useful flags include

```
-no-fsanitize-recover
```

which causes the first report to be fatal (it always is for the unreachable and return suboptions). For more detailed information on where the runtime error occurs, R can be run under a debugger with a breakpoint set before the sanitizer report is produced: for gdb or lldb you could use

```
break __ubsan_handle_float_cast_overflow
break __ubsan_handle_float_cast_overflow_abort
```

or similar (there are handlers for each type of undefined behaviour).

There are also the compiler flags -fcatch-undefined-behavior and -ftrapv, said to be more reliable in clang than gcc.

For more details on the topic see http://blog.regehr.org/archives/213 and http://blog.llvm.org/2011/05/what-every-c-programmer-should-know.html (which has 3 parts).

4.3.5 Other analyses with 'clang'

Recent versions of clang on 'x86_64' Linux have 'ThreadSanitizer' (https://code.google.com/p/thread-sanitizer/), a 'data race detector for C/C++ programs', and 'MemorySanitizer' (http://clang.llvm.org/docs/MemorySanitizer.html, https://code.google.com/p/memory-sanitizer/wiki/MemorySanitizer) for the detection of uninitialized memory. Both are based on and provide similar functionality to tools in valgrind.

clang has a 'Static Analyser' which can be run on the source files during compilation: see http://clang-analyzer.llvm.org/.

4.3.6 Using 'Dr. Memory'

'Dr. Memory' from http://www.drmemory.org/ is a memory checker for (currently) 32-bit Windows, Linux and OS X with similar aims to valgrind. It works with unmodified executables 14 and detects memory access errors, uninitialized reads and memory leaks.

4.3.7 Fortran array bounds checking

Most of the Fortran compilers used with R allow code to be compiled with checking of array bounds: for example gfortran has option -fbounds-check and Solaris Studio has -C. This will give an error when the upper or lower bound is exceeded, e.g.

```
At line 97 of file .../src/appl/dqrdc2.f
```

Fortran runtime error: Index '1' of dimension 1 of array 'x' above upper bound of 0

One does need to be aware that lazy programmers often specify Fortran dimensions as 1 rather than * or a real bound and these will be reported.

 $^{^{14}}$ but works better if inlining and frame pointer optimizations are disabled.

It is easy to arrange to use this check on just the code in your package: add to ~/.R/Makevars something like (for gfortran)

This may report incorrectly errors with the way that Fortran character variables are passed, particularly when Fortran subroutines are called from C code. This may include the use of BLAS and LAPACK subroutines in R, so it is not advisable to build R itself with bounds checking (and may not even be possible as these subroutines are called during the R build).

4.4 Debugging compiled code

Sooner or later programmers will be faced with the need to debug compiled code loaded into R. This section is geared to platforms using gdb with code compiled by gcc, but similar things are possible with other debuggers such as lldb (http://lldb.llvm.org/, used on OS X) and Sun's dbx: some debuggers have graphical front-ends available.

Consider first 'crashes', that is when R terminated unexpectedly with an illegal memory access (a 'segfault' or 'bus error'), illegal instruction or similar. Unix-alike versions of R use a signal handler which aims to give some basic information. For example

```
*** caught segfault ***
address 0x20000028, cause 'memory not mapped'

Traceback:
1: .identC(class1[[1]], class2)
2: possibleExtends(class(sloti), classi, ClassDef2 = getClassDef(classi, where = where))
3: validObject(t(cu))
4: stopifnot(validObject(cu <- as(tu, "dtCMatrix")), validObject(t(cu)), validObject(t(tu)))

Possible actions:
1: abort (with core dump)
2: normal R exit
3: exit R without saving workspace
4: exit R saving workspace
Selection: 3</pre>
```

Since the R process may be damaged, the only really safe options are the first or third. (Note that a core dump is only produced where enabled: a common default in a shell is to limit its size to 0, thereby disabling it.)

A fairly common cause of such crashes is a package which uses .C or .Fortran and writes beyond (at either end) one of the arguments it is passed. As from R 3.0.0 there is a good way to detect this: using options(CBoundsCheck = TRUE) (which can be selected *via* the environment variable R_C_BOUNDS_CHECK=yes) changes the way .C and .Fortran work to check if the compiled code writes in the 64 bytes at either end of an argument.

Another cause of a 'crash' is to overrun the C stack. R tries to track that in its own code, but it may happen in third-party compiled code. For modern POSIX-compliant OSes R can safely catch that and return to the top-level prompt, so one gets something like

```
> .C("aaa")
Error: segfault from C stack overflow
>
```

However, C stack overflows are fatal under Windows and normally defeat attempts at debugging on that platform. Further, the size of the stack is set when R is compiled, whereas on POSIX OSes it can be set in the shell from which R is launched.

If you have a crash which gives a core dump you can use something like

```
gdb /path/to/R/bin/exec/R core.12345
```

to examine the core dump. If core dumps are disabled or to catch errors that do not generate a dump one can run R directly under a debugger by for example

```
$ R -d gdb --vanilla
...
gdb> run
```

at which point R will run normally, and hopefully the debugger will catch the error and return to its prompt. This can also be used to catch infinite loops or interrupt very long-running code. For a simple example

```
> for(i in 1:1e7) x <- rnorm(100)
[hit Ctrl-C]
Program received signal SIGINT, Interrupt.
0x00397682 in _int_free () from /lib/tls/libc.so.6
(gdb) where
   0x00397682 in _int_free () from /lib/tls/libc.so.6
#0
   0x00397eba in free () from /lib/tls/libc.so.6
#2  0xb7cf2551 in R_gc_internal (size_needed=313)
    at /users/ripley/R/svn/R-devel/src/main/memory.c:743
   Oxb7cf3617 in Rf_allocVector (type=13, length=626)
    at /users/ripley/R/svn/R-devel/src/main/memory.c:1906
#4 0xb7c3f6d3 in PutRNGstate ()
    at /users/ripley/R/svn/R-devel/src/main/RNG.c:351
#5 0xb7d6c0a5 in do_random2 (call=0x94bf7d4, op=0x92580e8, args=0x9698f98,
    rho=0x9698f28) at /users/ripley/R/svn/R-devel/src/main/random.c:183
```

In many cases it is possible to attach a debugger to a running process: this is helpful if an alternative front-end is in use or to investigate a task that seems to be taking far too long. This is done by something like

```
gdb -p pid
```

where *pid* is the id of the R executable or front-end. This stops the process so its state can be examined: use **continue** to resume execution.

Some "tricks" worth knowing follow:

4.4.1 Finding entry points in dynamically loaded code

Under most compilation environments, compiled code dynamically loaded into R cannot have breakpoints set within it until it is loaded. To use a symbolic debugger on such dynamically loaded code under Unix-alikes use

- Call the debugger on the R executable, for example by R -d gdb.
- Start R.
- At the R prompt, use dyn.load or library to load your shared object.
- Send an interrupt signal. This will put you back to the debugger prompt.
- Set the breakpoints in your code.
- Continue execution of R by typing signal ORET.

Under Windows signals may not be able to be used, and if so the procedure is more complicated. See the rw-FAQ and www.stats.uwo.ca/faculty/murdoch/software/debuggingR/gdb.shtml (http://www.stats.uwo.ca/faculty/murdoch/software/debuggingR/gdb.shtml).

4.4.2 Inspecting R objects when debugging

The key to inspecting R objects from compiled code is the function PrintValue(SEXP s) which uses the normal R printing mechanisms to print the R object pointed to by s, or the safer version R_PV(SEXP s) which will only print 'objects'.

One way to make use of PrintValue is to insert suitable calls into the code to be debugged.

Another way is to call R_PV from the symbolic debugger. (PrintValue is hidden as Rf_PrintValue.) For example, from gdb we can use

```
(gdb) p R_PV(ab)
```

using the object **ab** from the convolution example, if we have placed a suitable breakpoint in the convolution C code.

To examine an arbitrary R object we need to work a little harder. For example, let

```
R> DF <- data.frame(a = 1:3, b = 4:6)
```

By setting a breakpoint at do_get and typing get("DF") at the R prompt, one can find out the address in memory of DF, for example

```
Value returned is $1 = (SEXPREC *) 0x40583e1c
(gdb) p *$1
$2 = {
  sxpinfo = \{type = 19, obj = 1, named = 1, gp = 0,
    mark = 0, debug = 0, trace = 0, = 0},
  attrib = 0x40583e80,
  u = {
    vecsxp = {
      length = 2,
      type = \{c = 0x40634700 "0>X@D>X@0>X@", i = 0x40634700,
f = 0x40634700, z = 0x40634700, s = 0x40634700},
      truelength = 1075851272,
    },
    primsxp = {offset = 2},
    symsxp = \{pname = 0x2, value = 0x40634700, internal = 0x40203008\},
    listsxp = {carval = 0x2, cdrval = 0x40634700, tagval = 0x40203008},
    envsxp = \{frame = 0x2, enclos = 0x40634700\},
    closxp = \{formals = 0x2, body = 0x40634700, env = 0x40203008\},\
    promsxp = {value = 0x2, expr = 0x40634700, env = 0x40203008}
  }
}
```

(Debugger output reformatted for better legibility).

Using R_PV() one can "inspect" the values of the various elements of the SEXP, for example,

```
(gdb) p R_PV($1->attrib)
     $names
     [1] "a" "b"
     $row.names
     [1] "1" "2" "3"
     $class
     [1] "data.frame"
     $3 = void
  To find out where exactly the corresponding information is stored, one needs to go "deeper":
     (gdb) set $a = $1->attrib
     (gdb) p $a->u.listsxp.tagval->u.symsxp.pname->u.vecsxp.type.c
     $4 = 0x405d40e8 "names"
     (gdb) p $a->u.listsxp.carval->u.vecsxp.type.s[1]->u.vecsxp.type.c
     $5 = 0x40634378 "b"
     (gdb) p $1->u.vecsxp.type.s[0]->u.vecsxp.type.i[0]
     $6 = 1
     (gdb) p $1->u.vecsxp.type.s[1]->u.vecsxp.type.i[1]
     $7 = 5
   Another alternative is the R_inspect function which shows the low-level structure of the ob-
jects recursively (addresses differ from the above as this example is created on another machine):
     (gdb) p R_inspect($1)
     @100954d18 19 VECSXP g0c2 [OBJ,NAM(2),ATT] (len=2, tl=0)
       @100954d50 13 INTSXP gOc2 [NAM(2)] (len=3, tl=0) 1,2,3
       @100954d88 13 INTSXP gOc2 [NAM(2)] (len=3, tl=0) 4,5,6
     ATTRIB:
       @102a70140 02 LISTSXP g0c0 []
         TAG: @10083c478 01 SYMSXP g0c0 [MARK,NAM(2),gp=0x4000] "names"
         @100954dc0 16 STRSXP g0c2 [NAM(2)] (len=2, t1=0)
           @10099df28 09 CHARSXP g0c1 [MARK,gp=0x21] "a"
           @10095e518 09 CHARSXP gOc1 [MARK,gp=0x21] "b"
         TAG: @100859e60 01 SYMSXP g0c0 [MARK,NAM(2),gp=0x4000] "row.names"
         @102a6f868 13 INTSXP g0c1 [NAM(2)] (len=2, t1=1) -2147483648,-3
         TAG: @10083c948 01 SYMSXP g0c0 [MARK,gp=0x4000] "class"
         @102a6f838 16 STRSXP g0c1 [NAM(2)] (len=1, tl=1)
           @1008c6d48 09 CHARSXP g0c2 [MARK,gp=0x21,ATT] "data.frame"
```

For a more fine-grained control over the depth of the recursion and the output of vectors R_inspect3 takes additional two character() parameters: maximum depth and the maximal number of elements that will be printed for scalar vectors. The defaults in R_inspect are currently -1 (no limit) and 5 respectively.

In general the representation of each object follows the format: @<address> <type-nr> <type-name> <gc-info> [<flags>] ...

5 System and foreign language interfaces

5.1 Operating system access

Access to operating system functions is via the R functions system and system2. The details will differ by platform (see the on-line help), and about all that can safely be assumed is that the first argument will be a string command that will be passed for execution (not necessarily by a shell) and the second argument to system will be internal which if true will collect the output of the command into an R character vector.

On POSIX-compliant OSes these commands pass a command-line to a shell: Windows is not POSIX-compliant and there is a separate function shell to do so.

The function system.time is available for timing. Timing on child processes is only available on Unix-alikes, and may not be reliable there.

5.2 Interface functions .C and .Fortran

These two functions provide an interface to compiled code that has been linked into R, either at build time or *via* dyn.load (see Section 5.3 [dyn.load and dyn.unload], page 101). They are primarily intended for compiled C and FORTRAN 77 code respectively, but the .C function can be used with other languages which can generate C interfaces, for example C++ (see Section 5.6 [Interfacing C++ code], page 107).

The first argument to each function is a character string specifying the symbol name as known¹ to C or FORTRAN, that is the function or subroutine name. (That the symbol is loaded can be tested by, for example, is.loaded("cg"). Use the name you pass to .C or .Fortran rather than the translated symbol name.)

There can be up to 65 further arguments giving R objects to be passed to compiled code. Normally these are copied before being passed in, and copied again to an R list object when the compiled code returns. If the arguments are given names, these are used as names for the components in the returned list object (but not passed to the compiled code).

The following table gives the mapping between the modes of R atomic vectors and the types of arguments to a C function or FORTRAN subroutine.

R storage mode	C type	FORTRAN type
logical	int *	INTEGER
integer	int *	INTEGER
double	double *	DOUBLE PRECISION
complex	Rcomplex *	DOUBLE COMPLEX
character	char **	CHARACTER*255
raw	unsigned char *	none

Do please note the first two. On the 64-bit Unix/Linux/OS X platforms, long is 64-bit whereas int and INTEGER are 32-bit. Code ported from S-PLUS (which uses long * for logical and integer) will not work on all 64-bit platforms (although it may appear to work on some, including Windows). Note also that if your compiled code is a mixture of C functions and FORTRAN subprograms the argument types must match as given in the table above.

C type Rcomplex is a structure with double members r and i defined in the header file R_ext/Complex.h included by R.h. (On most platforms this is stored in a way compatible with the C99 double complex type: however, it may not be possible to pass Rcomplex to a C99 function

¹ possibly after some platform-specific translation, e.g. adding leading or trailing underscores.

expecting a double complex argument. Nor need it be compatible with a C++ complex type. Moreover, the compatibility can depends on the optimization level set for the compiler.)

Only a single character string can be passed to or from FORTRAN, and the success of this is compiler-dependent. Other R objects can be passed to .C, but it is much better to use one of the other interfaces.

It is possible to pass numeric vectors of storage mode double to C as float * or to FORTRAN as REAL by setting the attribute Csingle, most conveniently by using the R functions as.single, single or mode. This is intended only to be used to aid interfacing existing C or FORTRAN code.

Logical values are sent as 0 (FALSE), 1 (TRUE) or INT_MIN = -2147483648 (NA, but only if NAOK is true), and the compiled code should return one of these three values. (Non-zero values other than INT_MIN are mapped to TRUE.)

Unless formal argument NAOK is true, all the other arguments are checked for missing values NA and for the IEEE special values NaN, Inf and -Inf, and the presence of any of these generates an error. If it is true, these values are passed unchecked.

Argument PACKAGE confines the search for the symbol name to a specific shared object (or use "base" for code compiled into R). Its use is highly desirable, as there is no way to avoid two package writers using the same symbol name, and such name clashes are normally sufficient to cause R to crash. (If it is not present and the call is from the body of a function defined in a package namespace, the shared object loaded by the first (if any) useDynLib directive will be used. However, prior to R 2.15.2 the detection of the correct namespace is unreliable and you are strongly recommended to use the PACKAGE argument for packages to be used with earlier versions of R.

Note that the compiled code should not return anything except through its arguments: C functions should be of type void and FORTRAN subprograms should be subroutines.

To fix ideas, let us consider a very simple example which convolves two finite sequences. (This is hard to do fast in interpreted R code, but easy in C code.) We could do this using .C by

```
void convolve(double *a, int *na, double *b, int *nb, double *ab)
     {
         int nab = *na + *nb - 1;
         for(int i = 0; i < nab; i++)
     ab[i] = 0.0;
         for(int i = 0; i < *na; i++)
     for(int j = 0; j < *nb; j++)
         ab[i + j] += a[i] * b[j];
called from R by
     conv <- function(a, b)</pre>
         .C("convolve",
            as.double(a),
            as.integer(length(a)),
            as.double(b),
            as.integer(length(b)),
            ab = double(length(a) + length(b) - 1))$ab
```

Note that we take care to coerce all the arguments to the correct R storage mode before calling .C; mistakes in matching the types can lead to wrong results or hard-to-catch errors.

Special care is needed in handling character vector arguments in C (or C++). On entry the contents of the elements are duplicated and assigned to the elements of a char ** array, and on

exit the elements of the C array are copied to create new elements of a character vector. This means that the contents of the character strings of the char ** array can be changed, including to \0 to shorten the string, but the strings cannot be lengthened. It is possible to allocate a new string via R_alloc and replace an entry in the char ** array by the new string. However, when character vectors are used other than in a read-only way, the .Call interface is much to be preferred.

Passing character strings to FORTRAN code needs even more care, and should be avoided where possible. Only the first element of the character vector is passed in, as a fixed-length (255) character array. Up to 255 characters are passed back to a length-one character vector. How well this works (or even if it works at all) depends on the C and FORTRAN compilers on each platform (including on their options). Often what is being passed to FORTRAN is one of a small set of possible values (a factor in R terms) which could alternatively be passed as an integer code: similarly FORTRAN code that wants to generate diagnostic messages can pass an integer code to a C or R wrapper which will convert it to a character string.

It is possible to pass some R objects other than atomic vectors via .C, but this is only supported for historical compatibility: use the .Call or .External interfaces for such objects. Any C/C++ code that includes Rinternals.h should be called via .Call or .External.

5.3 dyn.load and dyn.unload

Compiled code to be used with R is loaded as a shared object (Unix-alikes including OS X, see Section 5.5 [Creating shared objects], page 106, for more information) or DLL (Windows).

The shared object/DLL is loaded by dyn.load and unloaded by dyn.unload. Unloading is not normally necessary, but it is needed to allow the DLL to be re-built on some platforms, including Windows.

The first argument to both functions is a character string giving the path to the object. Programmers should not assume a specific file extension for the object/DLL (such as .so) but use a construction like

```
file.path(path1, path2, paste0("mylib", .Platform$dynlib.ext))
```

for platform independence. On Unix-alike systems the path supplied to dyn.load can be an absolute path, one relative to the current directory or, if it starts with '~', relative to the user's home directory.

Loading is most often done automatically based on the useDynLib() declaration in the NAMESPACE file, but may be done explicitly via a call to library.dynam. This has the form

```
library.dynam("libname", package, lib.loc)
```

where libname is the object/DLL name with the extension omitted. Note that the first argument, chname, should **not** be package since this will not work if the package is installed under another name.

Under some Unix-alike systems there is a choice of how the symbols are resolved when the object is loaded, governed by the arguments local and now. Only use these if really necessary: in particular using now=FALSE and then calling an unresolved symbol will terminate R unceremoniously.

R provides a way of executing some code automatically when a object/DLL is either loaded or unloaded. This can be used, for example, to register native routines with R's dynamic symbol mechanism, initialize some data in the native code, or initialize a third party library. On loading a DLL, R will look for a routine within that DLL named R_init_lib where lib is the name of the DLL file with the extension removed. For example, in the command

```
library.dynam("mylib", package, lib.loc)
```

 $^{^2}$ Note that this is then not checked for over-runs by option CBoundsCheck = TRUE.

R looks for the symbol named R_init_mylib. Similarly, when unloading the object, R looks for a routine named R_unload_lib, e.g., R_unload_mylib. In either case, if the routine is present, R will invoke it and pass it a single argument describing the DLL. This is a value of type DllInfo which is defined in the Rdynload.h file in the R_ext directory.

Note that there are some implicit restrictions on this mechanism as the basename of the DLL needs to be both a valid file name and valid as part of a C entry point (e.g. it cannot contain '.'): for portable code it is best to confine DLL names to be ASCII alphanumeric plus underscore. If entry point R_init_lib is not found it is also looked for with '.' replaced by '_'.

The following example shows templates for the initialization and unload routines for the mylib DLL.

```
#include <R.h>
#include <Rinternals.h>
#include <R_ext/Rdynload.h>

void
R_init_mylib(DllInfo *info)
{
    /* Register routines,
        allocate resources. */
}

void
R_unload_mylib(DllInfo *info)
{
    /* Release resources. */
}
```

If a shared object/DLL is loaded more than once the most recent version is used. More generally, if the same symbol name appears in several shared objects, the most recently loaded occurrence is used. The PACKAGE argument and registration (see the next section) provide good ways to avoid any ambiguity in which occurrence is meant.

On Unix-alikes the paths used to resolve dynamically linked dependent libraries are fixed (for security reasons) when the process is launched, so dyn.load will only look for such libraries in the locations set by the R shell script (via etc/ldpaths) and in the OS-specific defaults.

Windows allows more control (and less security) over where dependent DLLs are looked for. On all versions this includes the PATH environment variable, but with lowest priority: note that it does not include the directory from which the DLL was loaded. It is possible to add a single path with quite high priority *via* the DLLpath argument to dyn.load. This is (by default) used by library.dynam to include the package's libs/i386 or libs/x64 directory in the DLL search path.

5.4 Registering native routines

By 'native' routine, we mean an entry point in compiled code.

In calls to .C, .Call, .Fortran and .External, R must locate the specified native routine by looking in the appropriate shared object/DLL. By default, R uses the operating system-specific dynamic loader to lookup the symbol in all loaded DLLs and elsewhere. Alternatively, the author of the DLL can explicitly register routines with R and use a single, platform-independent mechanism for finding the routines in the DLL. One can use this registration mechanism to provide additional information about a routine, including the number and type of the arguments,

and also make it available to R programmers under a different name. In the future, registration may be used to implement a form of "secure" or limited native access.

To register routines with R, one calls the C routine R_registerRoutines. This is typically done when the DLL is first loaded within the initialization routine R_init_dll name described in Section 5.3 [dyn.load and dyn.unload], page 101. R_registerRoutines takes 5 arguments. The first is the DllInfo object passed by R to the initialization routine. This is where R stores the information about the methods. The remaining 4 arguments are arrays describing the routines for each of the 4 different interfaces: .C, .Call, .Fortran and .External. Each argument is a FIND-terminated array of the element types given in the following table:

```
.C R_CMethodDef
.Call R_CallMethodDef
.Fortran R_FortranMethodDef
.External R_ExternalMethodDef
```

Currently, the R_ExternalMethodDef is the same as R_CallMethodDef type and contains fields for the name of the routine by which it can be accessed in R, a pointer to the actual native symbol (i.e., the routine itself), and the number of arguments the routine expects to be passed from R. For example, if we had a routine named myCall defined as

along with any other routines for the .Call interface. For routines with a variable number of arguments invoked *via* the .External interface, one specifies -1 for the number of arguments which tells R not to check the actual number passed. Note that the number of arguments passed to .External were not checked prior to R 3.0.0.

Routines for use with the .C and .Fortran interfaces are described with similar data structures, but which have two additional fields for describing the type and "style" of each argument. Each of these can be omitted. However, if specified, each should be an array with the same number of elements as the number of parameters for the routine. The types array should contain the SEXP types describing the expected type of the argument. (Technically, the elements of the types array are of type R_NativePrimitiveArgType which is just an unsigned integer.) The R types and corresponding type identifiers are provided in the following table:

```
REALSXP
  numeric
  integer
               INTSXP
  logical
               LGLSXP
  single
               SINGLESXP
  character
               STRSXP
               VECSXP
  list
Consider a C routine, myC, declared as
  void myC(double *x, int *n, char **names, int *status);
We would register it as
  static R_NativePrimitiveArgType myC_t[] = {
      REALSXP, INTSXP, STRSXP, LGLSXP
  };
  static R_CMethodDef cMethods[] = {
     {"myC", (DL_FUNC) &myC, 4, myC_t}
```

```
{NULL, NULL, O};
```

One can also specify whether each argument is used simply as input, or as output, or as both input and output. The style field in the description of a method is used for this. The purpose is to allow³ R to transfer values more efficiently across the R-C/FORTRAN interface by avoiding copying values when it is not necessary. Typically, one omits this information in the registration data.

Having created the arrays describing each routine, the last step is to actually register them with R. We do this by calling R_registerRoutines. For example, if we have the descriptions above for the routines accessed by the .C and .Call we would use the following code:

```
void
R_init_myLib(DllInfo *info)
{
    R_registerRoutines(info, cMethods, callMethods, NULL, NULL);
}
```

This routine will be invoked when R loads the shared object/DLL named myLib. The last two arguments in the call to R_registerRoutines are for the routines accessed by .Fortran and .External interfaces. In our example, these are given as NULL since we have no routines of these types.

When R unloads a shared object/DLL, its registrations are automatically removed. There is no other facility for unregistering a symbol.

Examples of registering routines can be found in the different packages in the R source tree (e.g., stats). Also, there is a brief, high-level introduction in R News (volume 1/3, September 2001, pages 20–23, https://www.r-project.org/doc/Rnews/Rnews_2001-3.pdf).

Once routines are registered, they can be referred to as R objects if they this is arranged in the useDynLib call in the package's NAMESPACE file (see Section 1.5.4 [useDynLib], page 41). This avoids the overhead of looking up an entry point each time it is used, and ensure that the entry point in the package is the one used (without a PACKAGE = "pkg" argument). So for example the stats package has

5.4.1 Speed considerations

Sometimes registering native routines or using a PACKAGE argument can make a large difference. The results can depend quite markedly on the OS (and even if it is 32- or 64-bit), on the version of R and what else is loaded into R at the time.

To fix ideas, first consider $\tt x84_64$ OS 10.7 and R 2.15.2. A simple .Call function might be

```
foo <- function(x) .Call("foo", x)
with C code
   SEXP foo(SEXP x)
{
     return x;
}</pre>
```

 $^{^{3}}$ but this is not currently done.

If we compile with by R CMD SHLIB foo.c, load the code by dyn.load("foo.so") and run foo(pi) it took around 22 microseconds (us). Specifying the DLL by

```
foo2 <- function(x) .Call("foo", x, PACKAGE = "foo")
reduced the time to 1.7 us.</pre>
```

Now consider making these functions part of a package whose NAMESPACE file uses useDynlib(foo). This immediately reduces the running time as "foo" will be preferentially looked for foo.dll. Without specifying PACKAGE it took about 5 us (it needs to fathom out the appropriate DLL each time it is invoked but it does not need to search all DLLs), and with the PACKAGE argument it is again about 1.7 us.

Next suppose the package has registered the native routine foo. Then foo() still has to find the appropriate DLL but can get to the entry point in the DLL faster, in about 4.2 us. And foo2() now takes about 1 us. If we register the symbols in the NAMESPACE file and use

```
foo3 <- function(x) .Call(C_foo, x)</pre>
```

then the address for the native routine is looked up just once when the package is loaded, and foo3(pi) takes about 0.8 us.

Versions using .C() rather than .Call() take about 0.2 us longer.

These are all quite small differences, but C routines are not uncommonly invoked millions of times for run times of a few microseconds, and those doing such things may wish to be aware of the differences.

On Linux and Solaris there is a much smaller overhead in looking up symbols so foo(pi) takes around 5 times as long as foo3(pi).

Symbol lookup on Windows used to be far slower, so R maintains a small cache. If the cache is currently empty enough that the symbol can be stored in the cache then the performance is similar to Linux and Solaris: if not it may be slower. R's own code always uses registered symbols and so these never contribute to the cache: however many other packages do rely on symbol lookup.

5.4.2 Linking to native routines in other packages

In addition to registering C routines to be called by R, it can at times be useful for one package to make some of its C routines available to be called by C code in another package. The interface consists of two routines declared in header R_ext/Rdynload.h as

```
void R_RegisterCCallable(const char *package, const char *name,
DL_FUNC fptr);
DL_FUNC R_GetCCallable(const char *package, const char *name);
```

A package **packA** that wants to make a C routine myCfun available to C code in other packages would include the call

```
R_RegisterCCallable("packA", "myCfun", myCfun);
```

in its initialization function R_init_packA. A package packB that wants to use this routine would retrieve the function pointer with a call of the form

```
p_myCfun = R_GetCCallable("packA", "myCfun");
```

The author of **packB** is responsible for ensuring that **p_myCfun** has an appropriate declaration. In the future R may provide some automated tools to simplify exporting larger numbers of routines.

A package that wishes to make use of header files in other packages needs to declare them as a comma-separated list in the field 'LinkingTo' in the DESCRIPTION file. This then arranges that the include directories in the installed linked-to packages are added to the include paths for C and C++ code.

It must specify⁴ 'Imports' or 'Depends' of those packages, for they have to be loaded⁵ prior to this one (so the path to their compiled code has been registered).

A CRAN example of the use of this mechanism is package lme4 (https://CRAN.R-project.org/package=lme4), which links to Matrix (https://CRAN.R-project.org/package=Matrix).

5.5 Creating shared objects

Shared objects for loading into R can be created using R CMD SHLIB. This accepts as arguments a list of files which must be object files (with extension .o) or sources for C, C++, FORTRAN 77, Fortran 9x, Objective C or Objective C++ (with extensions .c, .cc or .cpp, .f, .f90 or .f95, .m, and .mm or .M, respectively), or commands to be passed to the linker. See R CMD SHLIB --help (or the R help for SHLIB) for usage information.

If compiling the source files does not work "out of the box", you can specify additional flags by setting some of the variables PKG_CPPFLAGS (for the C preprocessor, typically '-I' flags), PKG_CFLAGS, PKG_CXXFLAGS, PKG_FFLAGS, PKG_DBJCFLAGS, and PKG_DBJCXXFLAGS (for the C, C++, FORTRAN 77, Fortran 9x, Objective C, and Objective C++ compilers, respectively) in the file Makevars in the compilation directory (or, of course, create the object files directly from the command line). Similarly, variable PKG_LIBS in Makevars can be used for additional '-1' and '-L' flags to be passed to the linker when building the shared object. (Supplying linker commands as arguments to R CMD SHLIB will take precedence over PKG_LIBS in Makevars.)

It is possible to arrange to include compiled code from other languages by setting the macro 'OBJECTS' in file Makevars, together with suitable rules to make the objects.

Flags which are already set (for example in file etcR_ARCH/Makeconf) can be overridden by the environment variable MAKEFLAGS (at least for systems using a POSIX-compliant make), as in (Bourne shell syntax)

```
MAKEFLAGS="CFLAGS=-03" R CMD SHLIB *.c
```

It is also possible to set such variables in personal Makevars files, which are read after the local Makevars and the system makefiles or in a site-wide Makevars.site file. See Section "Customizing package compilation" in R Installation and Administration,

Note that as R CMD SHLIB uses Make, it will not remake a shared object just because the flags have changed, and if test.c and test.f both exist in the current directory

```
R CMD SHLIB test.f
```

will compile test.c!

If the src subdirectory of an add-on package contains source code with one of the extensions listed above or a file Makevars but not a file Makefile, R CMD INSTALL creates a shared object (for loading into R through useDynlib in the NAMESPACE, or in the .onLoad function of the package) using the R CMD SHLIB mechanism. If file Makevars exists it is read first, then the system makefile and then any personal Makevars files.

If the src subdirectory of package contains a file Makefile, this is used by R CMD INSTALL in place of the R CMD SHLIB mechanism. make is called with makefiles R_HOME/etcR_ARCH/Makeconf, src/Makefile and any personal Makevars files (in that order). The first target found in src/Makefile is used.

It is better to make use of a Makevars file rather than a Makefile: the latter should be needed only exceptionally.

Under Windows the same commands work, but Makevars.win will be used in preference to Makevars, and only src/Makefile.win will be used by R CMD INSTALL with src/Makefile

⁴ whether or not 'LinkingTo' is used.

 $^{^{5}}$ so there needs to be a corresponding import or importFrom entry in the NAMESPACE file.

being ignored. For past experiences of building DLLs with a variety of compilers, see file 'README.packages' and http://www.stats.uwo.ca/faculty/murdoch/software/compilingDLLs/. Under Windows you can supply an exports definitions file called dllname-win.def: otherwise all entry points in objects (but not libraries) supplied to R CMD SHLIB will be exported from the DLL. An example is stats-win.def for the stats package: a CRAN example in package fastICA (https://CRAN.R-project.org/package=fastICA).

If you feel tempted to read the source code and subvert these mechanisms, please resist. Far too much developer time has been wasted in chasing down errors caused by failures to follow this documentation, and even more by package authors demanding explanations as to why their packages no longer work. In particular, undocumented environment or make variables are not for use by package writers and are subject to change without notice.

5.6 Interfacing C++ code

Suppose we have the following hypothetical C++ library, consisting of the two files X.h and X.cpp, and implementing the two classes X and Y which we want to use in R.

```
// X.h

class X {
  public: X (); ~X ();
  };

class Y {
  public: Y (); ~Y ();
  };
```

```
// X.cpp

#include <R.h>
#include "X.h"

static Y y;

X::X() { REprintf("constructor X\n"); }
X::~X() { REprintf("destructor X\n"); }
Y::Y() { REprintf("constructor Y\n"); }
Y::~Y() { REprintf("destructor Y\n"); }
```

To use with R, the only thing we have to do is writing a wrapper function and ensuring that the function is enclosed in

```
extern "C" {
}
```

For example,

```
// X_main.cpp:
#include "X.h"

extern "C" {

void X_main () {
    X x;
}

} // extern "C"
```

Compiling and linking should be done with the C++ compiler-linker (rather than the C compiler-linker or the linker itself); otherwise, the C++ initialization code (and hence the constructor of the static variable Y) are not called. On a properly configured system, one can simply use

```
R CMD SHLIB X.cpp X_main.cpp
```

to create the shared object, typically X.so (the file name extension may be different on your platform). Now starting R yields

```
R version 2.14.1 Patched (2012-01-16 r58124)
Copyright (C) 2012 The R Foundation for Statistical Computing
...
Type    "q()" to quit R.

R> dyn.load(paste("X", .Platform$dynlib.ext, sep = ""))
constructor Y
R> .C("X_main")
constructor X
destructor X
list()
R> q()
Save workspace image? [y/n/c]: y
destructor Y
```

The R for Windows FAQ (rw-FAQ) contains details of how to compile this example under Windows.

Earlier version of this example used C++ iostreams: this is best avoided. There is no guarantee that the output will appear in the R console, and indeed it will not on the R for Windows console. Use R code or the C entry points (see Section 6.5 [Printing], page 137) for all I/O if at all possible. Examples have been seen where merely loading a DLL that contained calls to C++ I/O upset R's own C I/O (for example by resetting buffers on open files).

Most R header files can be included within C++ programs but they should **not** be included within an extern "C" block (as they include system headers⁶). It may not be possible to include some R headers as they in turn include system header files that may cause conflicts—if this happens, try defining 'NO_C_HEADERS' before including the R headers, and include C++ versions (such as 'cmath' and 'cstdlib') of the appropriate headers yourself before the R headers. (Headers R.h and Rmath.h support 'NO_C_HEADERS': the legacy header S.h does not. Header Rinternals.h does as from R 3.3.0.)

By default header Rmath.h includes math.h or cmath. Header R.h includes

```
limits.h math.h stddef.h stdio.h stdlib.h string.h
```

⁶ Even including C system headers in such a block has caused compilation errors.

or their C++ equivalents directly or indirectly, and either stddef.h or cstddef needs to be included before R.h if 'NO_C_HEADERS' is defined.

5.7 Fortran I/O

We have already warned against the use of C++ iostreams not least because output is not guaranteed to appear on the R console, and this warning applies equally to Fortran (77 or 9x) output to units * and 6. See Section 6.5.1 [Printing from FORTRAN], page 137, which describes workarounds.

In the past most Fortran compilers implemented I/O on top of the C I/O system and so the two interworked successfully. This was true of g77, but it is less true of gfortran as used in gcc 4.y.z. In particular, any package that makes use of Fortran I/O will when compiled on Windows interfere with C I/O: when the Fortran I/O is initialized (typically when the package is loaded) the C stdout and stderr are switched to LF line endings. (Function init in file src/modules/lapack/init_win.c shows how to mitigate this.)

5.8 Linking to other packages

It is not in general possible to link a DLL in package **packA** to a DLL provided by package **packB** (for the security reasons mentioned in Section 5.3 [dyn.load and dyn.unload], page 101, and also because some platforms distinguish between shared objects and dynamic libraries), but it is on Windows.

Note that there can be tricky versioning issues here, as package \mathbf{packB} could be re-installed after package \mathbf{packA} — it is desirable that the API provided by package \mathbf{packB} remains backwards-compatible.

Shipping a static library in package **packB** for other packages to link to avoids most of the difficulties.

5.8.1 Unix-alikes

It is possible to link a shared object in package **packA** to a library provided by package **packB** under limited circumstances on a Unix-alike OS. There are severe portability issues, so this is not recommended for a distributed package.

This is easiest if **packB** provides a static library **packB/lib/libpackB.a.** (Note using directory **lib** rather than **libs** is conventional, and architecture-specific sub-directories may be needed and are assumed in the sample code below. The code in the static library will need to be compiled with **PIC** flags on platforms where it matters.) Then as the code from package **packB** is incorporated when package **packA** is installed, we only need to find the static library at install time for package **packA**. The only issue is to find package **packB**, and for that we can ask R by something like (long lines broken for display here)

For a dynamic library packB/lib/libpackB.so (packB/lib/libpackB.dylib on OS X: note that you cannot link to a shared object, .so, on that platform) we could use

This will work for installation, but very likely not when package packB is loaded, as the path to package packB's lib directory is not in the ld.so⁷ search path. You can arrange to put it there before R is launched by setting (on some platforms) LD_RUN_PATH or LD_LIBRARY_PATH or adding to the ld.so cache (see man ldconfig). On platforms that support it, the path to the directory containing the dynamic library can be hardcoded at install time (which assumes that the location of package packB will not be changed nor the package updated to a changed API). On systems with the gcc or clang and the GNU linker (e.g. Linux) and some others this can be done by e.g.

It may be possible to figure out what is required semi-automatically from the result of R CMD libtool --config (look for 'hardcode').

Making headers provided by package **packB** available to the code to be compiled in package **packA** can be done by the LinkingTo mechanism (see Section 5.4 [Registering native routines], page 102).

5.8.2 Windows

Suppose package **packA** wants to make use of compiled code provided by **packB** in DLL packB/libs/exB.dll, possibly the package's DLL packB/libs/packB.dll. (This can be extended to linking to more than one package in a similar way.) There are three issues to be addressed:

 Making headers provided by package packB available to the code to be compiled in package packA.

This is done by the LinkingTo mechanism (see Section 5.4 [Registering native routines], page 102).

• preparing packA.dll to link to packB/libs/exB.dll.

This needs an entry in Makevars.win of the form

```
PKG_LIBS= -L<something> -lexB
```

and one possibility is that <something> is the path to the installed pkgB/libs directory. To find that we need to ask R where it is by something like

```
PKGB_PATH='echo 'library(packB);
  cat(system.file("libs", package="packB", mustWork=TRUE))' \
  | rterm --vanilla --slave'
PKG_LIBS= -L"$(PKGB_PATH)$(R_ARCH)" -lexB
```

Another possibility is to use an import library, shipping with package **packA** an exports file exB.def. Then Makevars.win could contain

```
PKG_LIBS= -L. -lexB

all: $(SHLIB) before

before: libexB.dll.a
libexB.dll.a: exB.def
```

and then installing package **packA** will make and use the import library for **exB.dll**. (One way to prepare the exports file is to use **pexports.exe**.)

⁷ dyld on OS X, and DYLD_LIBRARY_PATHS below.

• loading packA.dll which depends on exB.dll.

If exB.dll was used by package packB (because it is in fact packB.dll or packB.dll depends on it) and packB has been loaded before packA, then nothing more needs to be done as exB.dll will already be loaded into the R executable. (This is the most common scenario.)

More generally, we can use the DLLpath argument to library.dynam to ensure that exB.dll is found, for example by setting

Note that DLLpath can only set one path, and so for linking to two or more packages you would need to resort to setting environment variable PATH.

5.9 Handling R objects in C

Using C code to speed up the execution of an R function is often very fruitful. Traditionally this has been done via the .C function in R. However, if a user wants to write C code using internal R data structures, then that can be done using the .Call and .External functions. The syntax for the calling function in R in each case is similar to that of .C, but the two functions have different C interfaces. Generally the .Call interface is simpler to use, but .External is a little more general.

```
A call to .Call is very similar to .C, for example .Call("convolve2", a, b)
```

The first argument should be a character string giving a C symbol name of code that has already been loaded into R. Up to 65 R objects can passed as arguments. The C side of the interface is

```
#include <R.h>
#include <Rinternals.h>

SEXP convolve2(SEXP a, SEXP b)
...
A call to .External is almost identical
.External("convolveE", a, b)
but the C side of the interface is different, having only one argument
#include <R.h>
#include <Rinternals.h>

SEXP convolveE(SEXP args)
...
```

Here args is a LISTSXP, a Lisp-style pairlist from which the arguments can be extracted.

In each case the R objects are available for manipulation via a set of functions and macros defined in the header file Rinternals.h or some S-compatibility macros⁸ defined in Rdefines.h. See Section 5.10 [Interface functions .Call and .External], page 120, for details on .Call and .External.

Before you decide to use .Call or .External, you should look at other alternatives. First, consider working in interpreted R code; if this is fast enough, this is normally the best option. You should also see if using .C is enough. If the task to be performed in C is simple enough involving only atomic vectors and requiring no call to R, .C suffices. A great deal of useful

⁸ That is, similar to those defined in S version 4 from the 1990s: these are not kept up to date and are not recommended for new projects. Prior to R 3.3.0 it was not compatible with defining R_NO_REMAP.

code was written using just .C before .Call and .External were available. These interfaces allow much more control, but they also impose much greater responsibilities so need to be used with care. Neither .Call nor .External copy their arguments: you should treat arguments you receive through these interfaces as read-only.

To handle R objects from within C code we use the macros and functions that have been used to implement the core parts of R. A public⁹ subset of these is defined in the header file Rinternals.h in the directory R_INCLUDE_DIR (default R_HOME/include) that should be available on any R installation.

A substantial amount of R, including the standard packages, is implemented using the functions and macros described here, so the R source code provides a rich source of examples and "how to do it": do make use of the source code for inspirational examples.

It is necessary to know something about how R objects are handled in C code. All the R objects you will deal with will be handled with the type $SEXP^{10}$, which is a pointer to a structure with typedef SEXPREC. Think of this structure as a variant type that can handle all the usual types of R objects, that is vectors of various modes, functions, environments, language objects and so on. The details are given later in this section and in Section "R Internal Structures" in R Internals, but for most purposes the programmer does not need to know them. Think rather of a model such as that used by Visual Basic, in which R objects are handed around in C code (as they are in interpreted R code) as the variant type, and the appropriate part is extracted for, for example, numerical calculations, only when it is needed. As in interpreted R code, much use is made of coercion to force the variant object to the right type.

5.9.1 Handling the effects of garbage collection

We need to know a little about the way R handles memory allocation. The memory allocated for R objects is not freed by the user; instead, the memory is from time to time garbage collected. That is, some or all of the allocated memory not being used is freed or marked as re-usable.

The R object types are represented by a C structure defined by a typedef SEXPREC in Rinternals.h. It contains several things among which are pointers to data blocks and to other SEXPRECs. A SEXP is simply a pointer to a SEXPREC.

If you create an R object in your C code, you must tell R that you are using the object by using the PROTECT macro on a pointer to the object. This tells R that the object is in use so it is not destroyed during garbage collection. Notice that it is the object which is protected, not the pointer variable. It is a common mistake to believe that if you invoked PROTECT(p) at some point then p is protected from then on, but that is not true once a new object is assigned to p.

Protecting an R object automatically protects all the R objects pointed to in the corresponding SEXPREC, for example all elements of a protected list are automatically protected.

The programmer is solely responsible for housekeeping the calls to PROTECT. There is a corresponding macro UNPROTECT that takes as argument an int giving the number of objects to unprotect when they are no longer needed. The protection mechanism is stack-based, so UNPROTECT(n) unprotects the last n objects which were protected. The calls to PROTECT and UNPROTECT must balance when the user's code returns. R will warn about "stack imbalance in .Call" (or .External) if the housekeeping is wrong.

Here is a small example of creating an R numeric vector in C code:

⁹ see Chapter 6 [The R API], page 134: note that these are not all part of the API.

 $^{^{10}\,}$ SEXP is an acronym for Simple EXP ression, common in LISP-like language syntaxes.

```
#include <R.h>
#include <Rinternals.h>

SEXP ab;
....
ab = PROTECT(allocVector(REALSXP, 2));
REAL(ab)[0] = 123.45;
REAL(ab)[1] = 67.89;
UNPROTECT(1);
```

Now, the reader may ask how the R object could possibly get removed during those manipulations, as it is just our C code that is running. As it happens, we can do without the protection in this example, but in general we do not know (nor want to know) what is hiding behind the R macros and functions we use, and any of them might cause memory to be allocated, hence garbage collection and hence our object ab to be removed. It is usually wise to err on the side of caution and assume that any of the R macros and functions might remove the object.

In some cases it is necessary to keep better track of whether protection is really needed. Be particularly aware of situations where a large number of objects are generated. The pointer protection stack has a fixed size (default 10,000) and can become full. It is not a good idea then to just PROTECT everything in sight and UNPROTECT several thousand objects at the end. It will almost invariably be possible to either assign the objects as part of another object (which automatically protects them) or unprotect them immediately after use.

Protection is not needed for objects which R already knows are in use. In particular, this applies to function arguments.

There is a less-used macro UNPROTECT_PTR(s) that unprotects the object pointed to by the SEXP s, even if it is not the top item on the pointer protection stack. This is rarely needed outside the parser (the R sources currently have three examples, one in src/main/plot3d.c).

Sometimes an object is changed (for example duplicated, coerced or grown) yet the current value needs to be protected. For these cases PROTECT_WITH_INDEX saves an index of the protection location that can be used to replace the protected value using REPROTECT. For example (from the internal code for optim)

```
PROTECT_INDEX ipx;
....
s = PROTECT_WITH_INDEX(eval(OS->R_fcall, OS->R_env), &ipx);
s = REPROTECT(coerceVector(s, REALSXP), ipx);
```

Note that it is dangerous to mix UNPROTECT_PTR with PROTECT_WITH_INDEX, as the former changes the protection locations of objects that were protected after the one being unprotected.

There is another way to avoid the affects of garbage collection: a call to R_PreserveObject adds an object to an internal list of objects not to be collects, and a subsequent call to R_ReleaseObject removes it from that list. This provides a way for objects which are not returned as part of R objects to be protected across calls to compiled code: on the other hand it becomes the user's responsibility to release them when they are no longer needed (and this often requires the use of a finalizer). It is less efficient that the normal protection mechanism, and should be used sparingly.

5.9.2 Allocating storage

For many purposes it is sufficient to allocate R objects and manipulate those. There are quite a few allocXxx functions defined in Rinternals.h—you may want to explore them.

One that is commonly used is allocVector, the C-level equivalent of R-level vector() and its wrappers such as integer() and character(). One distinction is that whereas the R functions

always initialize the elements of the vector, allocVector only does so for lists, expressions and character vectors (the cases where the elements are themselves R objects).

If storage is required for C objects during the calculations this is best allocating by calling R_alloc; see Section 6.1 [Memory allocation], page 134. All of these memory allocation routines do their own error-checking, so the programmer may assume that they will raise an error and not return if the memory cannot be allocated.

5.9.3 Details of R types

Users of the Rinternals.h macros will need to know how the R types are known internally. The different R data types are represented in C by SEXPTYPE. Some of these are familiar from R and some are internal data types. The usual R object modes are given in the table.

SEXPTYPE	R equivalent
REALSXP	numeric with storage mode double
INTSXP	integer
CPLXSXP	complex
LGLSXP	logical
STRSXP	character
VECSXP	list (generic vector)
LISTSXP	pairlist
DOTSXP	a '' object
NILSXP	NULL
SYMSXP	name/symbol
CLOSXP	function or function closure
ENVSXP	environment

Among the important internal SEXPTYPEs are LANGSXP, CHARSXP, PROMSXP, etc. (N.B.: although it is possible to return objects of internal types, it is unsafe to do so as assumptions are made about how they are handled which may be violated at user-level evaluation.) More details are given in Section "R Internal Structures" in R Internals.

Unless you are very sure about the type of the arguments, the code should check the data types. Sometimes it may also be necessary to check data types of objects created by evaluating an R expression in the C code. You can use functions like isReal, isInteger and isString to do type checking. See the header file Rinternals.h for definitions of other such functions. All of these take a SEXP as argument and return 1 or 0 to indicate TRUE or FALSE.

What happens if the SEXP is not of the correct type? Sometimes you have no other option except to generate an error. You can use the function error for this. It is usually better to coerce the object to the correct type. For example, if you find that an SEXP is of the type INTEGER, but you need a REAL object, you can change the type by using

```
newSexp = PROTECT(coerceVector(oldSexp, REALSXP));
```

Protection is needed as a new *object* is created; the object formerly pointed to by the SEXP is still protected but now unused.¹¹

All the coercion functions do their own error-checking, and generate NAs with a warning or stop with an error as appropriate.

Note that these coercion functions are *not* the same as calling as.numeric (and so on) in R code, as they do not dispatch on the class of the object. Thus it is normally preferable to do the coercion in the calling R code.

So far we have only seen how to create and coerce R objects from C code, and how to extract the numeric data from numeric R vectors. These can suffice to take us a long way in interfacing

¹¹ If no coercion was required, coerceVector would have passed the old object through unchanged.

R objects to numerical algorithms, but we may need to know a little more to create useful return objects.

5.9.4 Attributes

Many R objects have attributes: some of the most useful are classes and the dim and dimnames that mark objects as matrices or arrays. It can also be helpful to work with the names attribute of vectors.

To illustrate this, let us write code to take the outer product of two vectors (which outer and %o% already do). As usual the R code is simple

```
out <- function(x, y)
{
    storage.mode(x) <- storage.mode(y) <- "double"
    .Call("out", x, y)
}</pre>
```

where we expect x and y to be numeric vectors (possibly integer), possibly with names. This time we do the coercion in the calling R code.

C code to do the computations is

```
#include <R.h>
#include <Rinternals.h>

SEXP out(SEXP x, SEXP y)
{
    int nx = length(x), ny = length(y);
    SEXP ans = PROTECT(allocMatrix(REALSXP, nx, ny));
    double *rx = REAL(x), *ry = REAL(y), *rans = REAL(ans);
    for(int i = 0; i < nx; i++) {
    double tmp = rx[i];
    for(int j = 0; j < ny; j++)
        rans[i + nx*j] = tmp * ry[j];
    }
    UNPROTECT(1);
    return ans;
}</pre>
```

Note the way REAL is used: as it is a function call it can be considerably faster to store the result and index that.

However, we would like to set the dimnames of the result. We can use

```
#include <R.h>
#include <Rinternals.h>
```

```
SEXP out(SEXP x, SEXP y)
{
   int nx = length(x), ny = length(y);
   SEXP ans = PROTECT(allocMatrix(REALSXP, nx, ny));
   double *rx = REAL(x), *ry = REAL(y), *rans = REAL(ans);

   for(int i = 0; i < nx; i++) {
      double tmp = rx[i];
      for(int j = 0; j < ny; j++)
   rans[i + nx*j] = tmp * ry[j];
   }

   SEXP dimnames = PROTECT(allocVector(VECSXP, 2));
   SET_VECTOR_ELT(dimnames, 0, getAttrib(x, R_NamesSymbol));
   SET_VECTOR_ELT(dimnames, 1, getAttrib(y, R_NamesSymbol));
   setAttrib(ans, R_DimNamesSymbol, dimnames);

   UNPROTECT(3);
   return ans;
}</pre>
```

This example introduces several new features. The getAttrib and setAttrib functions get and set individual attributes. Their second argument is a SEXP defining the name in the symbol table of the attribute we want; these and many such symbols are defined in the header file Rinternals.h.

There are shortcuts here too: the functions namesgets, dimgets and dimnamesgets are the internal versions of the default methods of names<-, dim<- and dimnames<- (for vectors and arrays), and there are functions such as GetMatrixDimnames and GetArrayDimnames.

What happens if we want to add an attribute that is not pre-defined? We need to add a symbol for it *via* a call to **install**. Suppose for illustration we wanted to add an attribute "version" with value 3.0. We could use

```
SEXP version;
version = PROTECT(allocVector(REALSXP, 1));
REAL(version)[0] = 3.0;
setAttrib(ans, install("version"), version);
UNPROTECT(1);
```

Using install when it is not needed is harmless and provides a simple way to retrieve the symbol from the symbol table if it is already installed. However, the lookup takes a non-trivial amount of time, so consider code such as

```
static SEXP VerSymbol = NULL;
...
if (VerSymbol == NULL) VerSymbol = install("version");
if it is to be done frequently.
```

This example can be simplified by another convenience function:

```
SEXP version = PROTECT(ScalarReal(3.0));
setAttrib(ans, install("version"), version);
UNPROTECT(1);
```

5.9.5 Classes

In R the class is just the attribute named "class" so it can be handled as such, but there is a shortcut classgets. Suppose we want to give the return value in our example the class "mat". We can use

As the value is a character vector, we have to know how to create that from a C character array, which we do using the function mkChar.

5.9.6 Handling lists

Some care is needed with lists, as R moved early on from using LISP-like lists (now called "pairlists") to S-like generic vectors. As a result, the appropriate test for an object of mode list is isNewList, and we need allocVector(VECSXP, n) and not allocList(n).

List elements can be retrieved or set by direct access to the elements of the generic vector. Suppose we have a list object

This can rapidly become tedious, and the following function (based on one in package stats) is very useful:

```
/* get the list element named str, or return NULL */

SEXP getListElement(SEXP list, const char *str)
{
    SEXP elmt = R_NilValue, names = getAttrib(list, R_NamesSymbol);

    for (int i = 0; i < length(list); i++)
    if(strcmp(CHAR(STRING_ELT(names, i)), str) == 0) {
        elmt = VECTOR_ELT(list, i);
        break;
    }
        return elmt;
}

and enables us to say
    double g;
    g = REAL(getListElement(a, "g"))[0];</pre>
```

5.9.7 Handling character data

R character vectors are stored as STRSXPs, a vector type like VECSXP where every element is of type CHARSXP. The CHARSXP elements of STRSXPs are accessed using STRING_ELT and SET_STRING_ELT.

CHARSXPs are read-only objects and must never be modified. In particular, the C-style string contained in a CHARSXP should be treated as read-only and for this reason the CHAR function used to access the character data of a CHARSXP returns (const char *) (this also allows compilers to issue warnings about improper use). Since CHARSXPs are immutable, the same CHARSXP can be shared by any STRSXP needing an element representing the same string. R maintains a global cache of CHARSXPs so that there is only ever one CHARSXP representing a given string in memory.

You can obtain a CHARSXP by calling mkChar and providing a nul-terminated C-style string. This function will return a pre-existing CHARSXP if one with a matching string already exists, otherwise it will create a new one and add it to the cache before returning it to you. The variant mkCharLen can be used to create a CHARSXP from part of a buffer and will ensure null-termination.

Note that R character strings are restricted to 2³¹ - 1 bytes, and hence so should the input to mkChar be (C allows longer strings on 64-bit platforms).

5.9.8 Finding and setting variables

It will be usual that all the R objects needed in our C computations are passed as arguments to .Call or .External, but it is possible to find the values of R objects from within the C given their names. The following code is the equivalent of get(name, envir = rho).

```
SEXP getvar(SEXP name, SEXP rho)
{
    SEXP ans;

    if(!isString(name) || length(name) != 1)
error("name is not a single string");
    if(!isEnvironment(rho))
error("rho should be an environment");
    ans = findVar(installChar(STRING_ELT(name, 0)), rho);
    Rprintf("first value is %f\n", REAL(ans)[0]);
    return R_NilValue;
}
```

The main work is done by findVar, but to use it we need to install name as a name in the symbol table. As we wanted the value for internal use, we return NULL.

Similar functions with syntax

```
void defineVar(SEXP symbol, SEXP value, SEXP rho)
void setVar(SEXP symbol, SEXP value, SEXP rho)
```

can be used to assign values to R variables. defineVar creates a new binding or changes the value of an existing binding in the specified environment frame; it is the analogue of assign(symbol, value, envir = rho, inherits = FALSE), but unlike assign, defineVar does not make a copy of the object value. setVar searches for an existing binding for symbol in rho or its enclosing environments. If a binding is found, its value is changed to value. Otherwise, a new binding with the specified value is created in the global environment. This corresponds to assign(symbol, value, envir = rho, inherits = TRUE).

¹² You can assign a *copy* of the object in the environment frame rho using defineVar(symbol, duplicate(value), rho)).

5.9.9 Some convenience functions

Some operations are done so frequently that there are convenience functions to handle them. (All these are provided via the header file Rinternals.h.)

Suppose we wanted to pass a single logical argument ignore_quotes: we could use

```
int ign = asLogical(ignore_quotes);
if(ign == NA_LOGICAL) error("'ignore_quotes' must be TRUE or FALSE");
```

which will do any coercion needed (at least from a vector argument), and return NA_LOGICAL if the value passed was NA or coercion failed. There are also asInteger, asReal and asComplex. The function asChar returns a CHARSXP. All of these functions ignore any elements of an input vector after the first.

To return a length-one real vector we can use

```
double x;
...
return ScalarReal(x);
```

and there are versions of this for all the atomic vector types (those for a length-one character vector being ScalarString with argument a CHARSXP and mkString with argument const char *).

Some of the isXXXX functions differ from their apparent R-level counterparts: for example isVector is true for any atomic vector type (isVectorAtomic) and for lists and expressions (isVectorList) (with no check on attributes). isMatrix is a test of a length-2 "dim" attribute.

There are a series of small macros/functions to help construct pairlists and language objects (whose internal structures just differ by SEXPTYPE). Function CONS(u, v) is the basic building block: it constructs a pairlist from u followed by v (which is a pairlist or R_NilValue). LCONS is a variant that constructs a language object. Functions list1 to list5 construct a pairlist from one to five items, and lang1 to lang6 do the same for a language object (a function to call plus zero to five arguments). Functions elt and lastElt find the *i*th element and the last element of a pairlist, and nthcdr returns a pointer to the *n*th position in the pairlist (whose CAR is the *n*th item).

Functions str2type and type2str map R length-one character strings to and from SEXPTYPE numbers, and type2char maps numbers to C character strings.

5.9.9.1 Semi-internal convenience functions

There is quite a collection of functions that may be used in your C code *if* you are willing to adapt to rare "API" changes. These typically contain "workhorses" of their R counterparts.

Functions any_duplicated and any_duplicated3 are fast versions of R's any(duplicated(.)).

Function R_compute_identical corresponds to R's identical function.

5.9.10 Named objects and copying

When assignments are done in R such as

```
x <- 1:10
y <- x
```

the named object is not necessarily copied, so after those two assignments y and x are bound to the same SEXPREC (the structure a SEXP points to). This means that any code which alters one of them has to make a copy before modifying the copy if the usual R semantics are to apply. Note that whereas .C and .Fortran do copy their arguments (unless the dangerous dup = FALSE is used), .Call and .External do not. So duplicate is commonly called on arguments to .Call before modifying them.

However, at least some of this copying is unneeded. In the first assignment shown, x <-1:10, R first creates an object with value 1:10 and then assigns it to x but if x is modified no copy is necessary as the temporary object with value 1:10 cannot be referred to again. R distinguishes between named and unnamed objects via a field in a SEXPREC that can be accessed via the macros NAMED and SET_NAMED. This can take values

- O The object is not bound to any symbol
- 1 The object has been bound to exactly one symbol
- The object has potentially been bound to two or more symbols, and one should act as if another variable is currently bound to this value.

Note the past tenses: R does not do full reference counting and there may currently be fewer bindings.

It is safe to modify the value of any SEXP for which NAMED(foo) is zero, and if NAMED(foo) is two, the value should be duplicated (*via* a call to duplicate) before any modification. Note that it is the responsibility of the author of the code making the modification to do the duplication, even if it is x whose value is being modified after y <- x.

The case NAMED(foo) == 1 allows some optimization, but it can be ignored (and duplication done whenever NAMED(foo) > 0). (This optimization is not currently usable in user code.) It is intended for use within replacement functions. Suppose we used

$$x <- 1:10$$
 foo(x) <- 3

which is computed as

Then inside "foo<-" the object pointing to the current value of x will have NAMED(foo) as one, and it would be safe to modify it as the only symbol bound to it is x and that will be rebound immediately. (Provided the remaining code in "foo<-" make no reference to x, and no one is going to attempt a direct call such as y <- "foo<-"(x).)

This mechanism is likely to be replaced in future versions of R.

5.10 Interface functions .Call and .External

In this section we consider the details of the R/C interfaces.

These two interfaces have almost the same functionality. .Call is based on the interface of the same name in S version 4, and .External is based on R's .Internal. .External is more complex but allows a variable number of arguments.

5.10.1 Calling .Call

Let us convert our finite convolution example to use .Call. The calling function in R is

```
conv <- function(a, b) .Call("convolve2", a, b)</pre>
```

which could hardly be simpler, but as we shall see all the type coercion is transferred to the C code, which is

```
#include <R.h>
#include <Rinternals.h>
SEXP convolve2(SEXP a, SEXP b)
{
    int na, nb, nab;
    double *xa, *xb, *xab;
    SEXP ab;
    a = PROTECT(coerceVector(a, REALSXP));
    b = PROTECT(coerceVector(b, REALSXP));
    na = length(a); nb = length(b); nab = na + nb - 1;
    ab = PROTECT(allocVector(REALSXP, nab));
    xa = REAL(a); xb = REAL(b); xab = REAL(ab);
    for(int i = 0; i < nab; i++) xab[i] = 0.0;
    for(int i = 0; i < na; i++)
for(int j = 0; j < nb; j++) xab[i + j] += xa[i] * xb[j];
    UNPROTECT(3);
    return ab;
}
```

5.10.2 Calling .External

We can use the same example to illustrate .External. The R code changes only by replacing .Call by .External

```
conv <- function(a, b) .External("convolveE", a, b)</pre>
```

but the main change is how the arguments are passed to the C code, this time as a single SEXP. The only change to the C code is how we handle the arguments.

```
#include <R.h>
#include <Rinternals.h>

SEXP convolveE(SEXP args)
{
   int i, j, na, nb, nab;
   double *xa, *xb, *xab;
   SEXP a, b, ab;

   a = PROTECT(coerceVector(CADR(args), REALSXP));
   b = PROTECT(coerceVector(CADDR(args), REALSXP));
   ...
}
```

Once again we do not need to protect the arguments, as in the R side of the interface they are objects that are already in use. The macros

```
first = CADR(args);
second = CADDR(args);
third = CADDDR(args);
fourth = CAD4R(args);
```

provide convenient ways to access the first four arguments. More generally we can use the \mathtt{CDR} and \mathtt{CAR} macros as in

```
args = CDR(args); a = CAR(args);
args = CDR(args); b = CAR(args);
```

which clearly allows us to extract an unlimited number of arguments (whereas .Call has a limit, albeit at 65 not a small one).

More usefully, the .External interface provides an easy way to handle calls with a variable number of arguments, as length(args) will give the number of arguments supplied (of which the first is ignored). We may need to know the names ('tags') given to the actual arguments, which we can by using the TAG macro and using something like the following example, that prints the names and the first value of its arguments if they are vector types.

```
SEXP showArgs(SEXP args)
  {
      args = CDR(args); /* skip 'name' */
      for(int i = 0; args != R_NilValue; i++, args = CDR(args)) {
  const char *name =
      isNull(TAG(args)) ? "" : CHAR(PRINTNAME(TAG(args)));
  SEXP el = CAR(args);
  if (length(el) == 0) {
      Rprintf("[%d] '%s' R type, length 0\n", i+1, name);
     continue;
  switch(TYPEOF(el)) {
  case REALSXP:
      Rprintf("[%d] '%s' %f\n", i+1, name, REAL(el)[0]);
      break:
  case LGLSXP:
  case INTSXP:
      Rprintf("[%d] '%s' %d\n", i+1, name, INTEGER(el)[0]);
      break:
  case CPLXSXP:
      Rcomplex cpl = COMPLEX(el)[0];
      Rprintf("[%d] '%s' %f + %fi\n", i+1, name, cpl.r, cpl.i);
  }
      break;
  case STRSXP:
      Rprintf("[%d] '%s' %s\n", i+1, name,
     CHAR(STRING_ELT(el, 0)));
     break;
  default:
      Rprintf("[%d] '%s' R type\n", i+1, name);
         }
      }
      return R_NilValue;
  }
This can be called by the wrapper function
```

```
showArgs <- function(...) invisible(.External("showArgs", ...))</pre>
```

Note that this style of programming is convenient but not necessary, as an alternative style is

```
showArgs1 <- function(...) invisible(.Call("showArgs1", list(...)))</pre>
```

The (very similar) C code is in the scripts.

5.10.3 Missing and special values

One piece of error-checking the .C call does (unless NAOK is true) is to check for missing (NA) and IEEE special values (Inf, -Inf and NaN) and give an error if any are found. With the .Call interface these will be passed to our code. In this example the special values are no problem, as IEC60559 arithmetic will handle them correctly. In the current implementation this is also true of NA as it is a type of NaN, but it is unwise to rely on such details. Thus we will re-write the code to handle NAs using macros defined in R_ext/Arith.h included by R.h.

The code changes are the same in any of the versions of convolve2 or convolveE:

```
for(int i = 0; i < na; i++)
  for(int j = 0; j < nb; j++)
if(ISNA(xa[i]) || ISNA(xb[j]) || ISNA(xab[i + j]))
  xab[i + j] = NA_REAL;
else
  xab[i + j] += xa[i] * xb[j];</pre>
```

Note that the ISNA macro, and the similar macros ISNAN (which checks for NaN or NA) and R_FINITE (which is false for NA and all the special values), only apply to numeric values of type double. Missingness of integers, logicals and character strings can be tested by equality to the constants NA_INTEGER, NA_LOGICAL and NA_STRING. These and NA_REAL can be used to set elements of R vectors to NA.

The constants R_NaN, R_PosInf and R_NegInf can be used to set doubles to the special values.

5.11 Evaluating R expressions from C

The main function we will use is

```
SEXP eval(SEXP expr, SEXP rho);
```

the equivalent of the interpreted R code eval(expr, envir = rho) (so rho must be an environment), although we can also make use of findVar, defineVar and findFun (which restricts the search to functions).

To see how this might be applied, here is a simplified internal version of lapply for expressions, used as

```
a <- list(a = 1:5, b = rnorm(10), test = runif(100))
.Call("lapply", a, quote(sum(x)), new.env())</pre>
```

with C code

```
SEXP lapply(SEXP list, SEXP expr, SEXP rho)
{
   int n = length(list);
   SEXP ans;

   if(!isNewList(list)) error("'list' must be a list");
   if(!isEnvironment(rho)) error("'rho' should be an environment");
   ans = PROTECT(allocVector(VECSXP, n));
   for(int i = 0; i < n; i++) {
   defineVar(install("x"), VECTOR_ELT(list, i), rho);
   SET_VECTOR_ELT(ans, i, eval(expr, rho));
   }
   setAttrib(ans, R_NamesSymbol, getAttrib(list, R_NamesSymbol));
   UNPROTECT(1);
   return ans;
}</pre>
```

It would be closer to lapply if we could pass in a function rather than an expression. One way to do this is *via* interpreted R code as in the next example, but it is possible (if somewhat obscure) to do this in C code. The following is based on the code in src/main/optimize.c.

```
SEXP lapply2(SEXP list, SEXP fn, SEXP rho)
     {
         int n = length(list);
         SEXP R_fcall, ans;
         if(!isNewList(list)) error("'list' must be a list");
         if(!isFunction(fn)) error("'fn' must be a function");
         if(!isEnvironment(rho)) error("'rho' should be an environment");
         R_fcall = PROTECT(lang2(fn, R_NilValue));
         ans = PROTECT(allocVector(VECSXP, n));
         for(int i = 0; i < n; i++) {
     SETCADR(R_fcall, VECTOR_ELT(list, i));
     SET_VECTOR_ELT(ans, i, eval(R_fcall, rho));
         setAttrib(ans, R_NamesSymbol, getAttrib(list, R_NamesSymbol));
         UNPROTECT(2);
         return ans;
     }
used by
     .Call("lapply2", a, sum, new.env())
```

Function lang2 creates an executable pairlist of two elements, but this will only be clear to those with a knowledge of a LISP-like language.

As a more comprehensive example of constructing an R call in C code and evaluating, consider the following fragment of printAttributes in src/main/print.c.

```
/* Need to construct a call to
   print(CAR(a), digits=digits)
   based on the R_print structure, then eval(call, env).
   See do_docall for the template for this sort of thing.
*/
SEXP s, t;
t = s = PROTECT(allocList(3));
```

```
SET_TYPEOF(s, LANGSXP);
SETCAR(t, install("print")); t = CDR(t);
SETCAR(t, CAR(a)); t = CDR(t);
SETCAR(t, ScalarInteger(digits));
SET_TAG(t, install("digits"));
eval(s, env);
UNPROTECT(1);
```

At this point CAR(a) is the R object to be printed, the current attribute. There are three steps: the call is constructed as a pairlist of length 3, the list is filled in, and the expression represented by the pairlist is evaluated.

A pairlist is quite distinct from a generic vector list, the only user-visible form of list in R. A pairlist is a linked list (with CDR(t) computing the next entry), with items (accessed by CAR(t)) and names or tags (set by SET_TAG). In this call there are to be three items, a symbol (pointing to the function to be called) and two argument values, the first unnamed and the second named. Setting the type to LANGSXP makes this a call which can be evaluated.

5.11.1 Zero-finding

In this section we re-work the example of Becker, Chambers & Wilks (1988, pp.~205–10) on finding a zero of a univariate function. The R code and an example are

```
zero <- function(f, guesses, tol = 1e-7) {
         f.check <- function(x) {</pre>
     x \leftarrow f(x)
     if(!is.numeric(x)) stop("Need a numeric result")
     as.double(x)
         }
          .Call("zero", body(f.check), as.double(guesses), as.double(tol),
       new.env())
     }
     cube1 <- function(x) (x^2 + 1) * (x - 1.5)
     zero(cube1, c(0, 5))
where this time we do the coercion and error-checking in the R code. The C code is
     SEXP mkans(double x)
     {
         // no need for PROTECT() here, as REAL(.) does not allocate:
         SEXP ans = allocVector(REALSXP, 1);
         REAL(ans)[0] = x;
         return ans;
     }
     double feval(double x, SEXP f, SEXP rho)
         // a version with (too) much PROTECT()ion .. "better safe than sorry"
         SEXP symbol, value;
         PROTECT(symbol = install("x"));
         PROTECT(value = mkans(x));
         defineVar(symbol, value, rho);
         UNPROTECT(2);
         return(REAL(eval(f, rho))[0]);
     }
```

```
SEXP zero(SEXP f, SEXP guesses, SEXP stol, SEXP rho)
    double x0 = REAL(guesses)[0], x1 = REAL(guesses)[1],
   tol = REAL(stol)[0];
    double f0, f1, fc, xc;
    if(tol <= 0.0) error("non-positive tol value");</pre>
    f0 = feval(x0, f, rho); f1 = feval(x1, f, rho);
    if(f0 == 0.0) return mkans(x0);
    if(f1 == 0.0) return mkans(x1);
    if(f0*f1 > 0.0) error("x[0] and x[1] have the same sign");
    for(;;) {
xc = 0.5*(x0+x1);
if(fabs(x0-x1) < tol) return mkans(xc);</pre>
fc = feval(xc, f, rho);
if(fc == 0) return mkans(xc);
if(f0*fc > 0.0) {
    x0 = xc; f0 = fc;
} else {
    x1 = xc; f1 = fc;
}
    }
}
```

5.11.2 Calculating numerical derivatives

We will use a longer example (by Saikat DebRoy) to illustrate the use of evaluation and .External. This calculates numerical derivatives, something that could be done as effectively in interpreted R code but may be needed as part of a larger C calculation.

An interpreted R version and an example are

```
numeric.deriv <- function(expr, theta, rho=sys.frame(sys.parent()))</pre>
    eps <- sqrt(.Machine$double.eps)</pre>
    ans <- eval(substitute(expr), rho)</pre>
    grad <- matrix(, length(ans), length(theta),</pre>
   dimnames=list(NULL, theta))
    for (i in seq_along(theta)) {
old <- get(theta[i], envir=rho)</pre>
delta <- eps * max(1, abs(old))</pre>
assign(theta[i], old+delta, envir=rho)
ans1 <- eval(substitute(expr), rho)</pre>
assign(theta[i], old, envir=rho)
grad[, i] <- (ans1 - ans)/delta</pre>
    attr(ans, "gradient") <- grad
    ans
}
omega <- 1:5; x <- 1; y <- 2
numeric.deriv(sin(omega*x*y), c("x", "y"))
```

where expr is an expression, theta a character vector of variable names and rho the environment to be used.

```
For the compiled version the call from R will be
     .External("numeric_deriv", expr, theta, rho)
with example usage
     .External("numeric_deriv", quote(sin(omega*x*y)),
       c("x", "v"), .GlobalEnv)
Note the need to quote the expression to stop it being evaluated in the caller.
```

Here is the complete C code which we will explain section by section.

```
#include <R.h> /* for DOUBLE_EPS */
#include <Rinternals.h>
SEXP numeric_deriv(SEXP args)
    SEXP theta, expr, rho, ans, ans1, gradient, par, dimnames;
    double tt, xx, delta, eps = sqrt(DOUBLE_EPS), *rgr, *rans;
    int i, start;
    expr = CADR(args);
    if(!isString(theta = CADDR(args)))
error("theta should be of type character");
    if(!isEnvironment(rho = CADDDR(args)))
error("rho should be an environment");
    ans = PROTECT(coerceVector(eval(expr, rho), REALSXP));
    gradient = PROTECT(allocMatrix(REALSXP, LENGTH(ans), LENGTH(theta)));
    rgr = REAL(gradient); rans = REAL(ans);
    for(i = 0, start = 0; i < LENGTH(theta); i++, start += LENGTH(ans)) {</pre>
par = PROTECT(findVar(installChar(STRING_ELT(theta, i)), rho));
tt = REAL(par)[0];
xx = fabs(tt);
delta = (xx < 1) ? eps : xx*eps;
REAL(par)[0] += delta;
ans1 = PROTECT(coerceVector(eval(expr, rho), REALSXP));
for(int j = 0; j < LENGTH(ans); j++)</pre>
    rgr[j + start] = (REAL(ans1)[j] - rans[j])/delta;
REAL(par)[0] = tt;
UNPROTECT(2); /* par, ans1 */
    }
    dimnames = PROTECT(allocVector(VECSXP, 2));
    SET_VECTOR_ELT(dimnames, 1, theta);
    dimnamesgets(gradient, dimnames);
    setAttrib(ans, install("gradient"), gradient);
    UNPROTECT(3); /* ans gradient dimnames */
    return ans;
}
```

The code to handle the arguments is

```
expr = CADR(args);
  if(!isString(theta = CADDR(args)))
error("theta should be of type character");
  if(!isEnvironment(rho = CADDDR(args)))
error("rho should be an environment");
```

Note that we check for correct types of theta and rho but do not check the type of expr. That is because eval can handle many types of R objects other than EXPRSXP. There is no useful coercion we can do, so we stop with an error message if the arguments are not of the correct mode.

The first step in the code is to evaluate the expression in the environment rho, by

```
ans = PROTECT(coerceVector(eval(expr, rho), REALSXP));
```

We then allocate space for the calculated derivative by

```
gradient = PROTECT(allocMatrix(REALSXP, LENGTH(ans), LENGTH(theta)));
```

The first argument to allocMatrix gives the SEXPTYPE of the matrix: here we want it to be REALSXP. The other two arguments are the numbers of rows and columns. (Note that LENGTH is intended to be used for vectors: length is more generally applicable.)

```
for(i = 0, start = 0; i < LENGTH(theta); i++, start += LENGTH(ans)) {
par = PROTECT(findVar(installChar(STRING_ELT(theta, i)), rho));</pre>
```

Here, we are entering a for loop. We loop through each of the variables. In the for loop, we first create a symbol corresponding to the i'th element of the STRSXP theta. Here, STRING_ELT(theta, i) accesses the i'th element of the STRSXP theta. Macro CHAR() extracts the actual character representation of it: it returns a pointer. We then install the name and use findVar to find its value.

```
tt = REAL(par)[0];
xx = fabs(tt);
delta = (xx < 1) ? eps : xx*eps;
REAL(par)[0] += delta;
ans1 = PROTECT(coerceVector(eval(expr, rho), REALSXP));</pre>
```

We first extract the real value of the parameter, then calculate delta, the increment to be used for approximating the numerical derivative. Then we change the value stored in par (in environment rho) by delta and evaluate expr in environment rho again. Because we are directly dealing with original R memory locations here, R does the evaluation for the changed parameter value.

```
for(int j = 0; j < LENGTH(ans); j++)
    rgr[j + start] = (REAL(ans1)[j] - rans[j])/delta;
REAL(par)[0] = tt;
UNPROTECT(2);
}</pre>
```

Now, we compute the i'th column of the gradient matrix. Note how it is accessed: R stores matrices by column (like FORTRAN).

```
dimnames = PROTECT(allocVector(VECSXP, 2));
SET_VECTOR_ELT(dimnames, 1, theta);
dimnamesgets(gradient, dimnames);
setAttrib(ans, install("gradient"), gradient);
UNPROTECT(3);
return ans;
```

}

 $^{^{13}}$ see Section 5.15 [Character encoding issues], page 132, for why this might not be what is required.

First we add column names to the gradient matrix. This is done by allocating a list (a VECSXP) whose first element, the row names, is NULL (the default) and the second element, the column names, is set as theta. This list is then assigned as the attribute having the symbol R_DimNamesSymbol. Finally we set the gradient matrix as the gradient attribute of ans, unprotect the remaining protected locations and return the answer ans.

5.12 Parsing R code from C

Suppose an R extension want to accept an R expression from the user and evaluate it. The previous section covered evaluation, but the expression will be entered as text and needs to be parsed first. A small part of R's parse interface is declared in header file R_ext/Parse.h¹⁴.

An example of the usage can be found in the (example) Windows package **windlgs** included in the R source tree. The essential part is

```
#include <R.h>
#include <Rinternals.h>
#include <R_ext/Parse.h>
SEXP menu_ttest3()
{
    char cmd[256];
    SEXP cmdSexp, cmdexpr, ans = R_NilValue;
    ParseStatus status;
    if(done == 1) {
cmdSexp = PROTECT(allocVector(STRSXP, 1));
SET_STRING_ELT(cmdSexp, 0, mkChar(cmd));
cmdexpr = PROTECT(R_ParseVector(cmdSexp, -1, &status, R_NilValue));
if (status != PARSE_OK) {
    UNPROTECT(2);
    error("invalid call %s", cmd);
/* Loop is needed here as EXPSEXP will be of length > 1 */
for(int i = 0; i < length(cmdexpr); i++)</pre>
    ans = eval(VECTOR_ELT(cmdexpr, i), R_GlobalEnv);
UNPROTECT(2);
    }
    return ans;
}
```

Note that a single line of text may give rise to more than one R expression.

R_ParseVector is essentially the code used to implement parse(text=) at R level. The first argument is a character vector (corresponding to text) and the second the maximal number of expressions to parse (corresponding to n). The third argument is a pointer to a variable of an enumeration type, and it is normal (as parse does) to regard all values other than PARSE_OK as an error. Other values which might be returned are PARSE_INCOMPLETE (an incomplete expression was found) and PARSE_ERROR (a syntax error), in both cases the value returned being R_NilValue. The fourth argument is a length one character vector to be used as a filename in error messages, a srcfile object or the R NULL object (as in the example above). If a srcfile object was used, a srcref attribute would be attached to the result, containing a list of srcref objects of the same length as the expression, to allow it to be echoed with its original formatting.

¹⁴ This is only guaranteed to show the current interface: it is liable to change.

5.12.1 Accessing source references

The source references added by the parser are recorded by R's evaluator as it evaluates code. Two functions make these available to debuggers running C code:

```
SEXP R_GetCurrentSrcref(int skip);
```

This function checks R_Srcref and the current evaluation stack for entries that contain source reference information. The skip argument tells how many source references to skip before returning the SEXP of the srcref object, counting from the top of the stack. If skip < 0, abs(skip) locations are counted up from the bottom of the stack. If too few or no source references are found, NULL is returned.

```
SEXP R_GetSrcFilename(SEXP srcref);
```

This function extracts the filename from the source reference for display, returning a length 1 character vector containing the filename. If no name is found, "" is returned.

5.13 External pointers and weak references

The SEXPTYPES EXTPTRSXP and WEAKREFSXP can be encountered at R level, but are created in C code.

External pointer SEXPs are intended to handle references to C structures such as 'handles', and are used for this purpose in package RODBC (https://CRAN.R-project.org/package=RODBC) for example. They are unusual in their copying semantics in that when an R object is copied, the external pointer object is not duplicated. (For this reason external pointers should only be used as part of an object with normal semantics, for example an attribute or an element of a list.)

An external pointer is created by

```
SEXP R_MakeExternalPtr(void *p, SEXP tag, SEXP prot);
```

where p is the pointer (and hence this cannot portably be a function pointer), and tag and prot are references to ordinary R objects which will remain in existence (be protected from garbage collection) for the lifetime of the external pointer object. A useful convention is to use the tag field for some form of type identification and the prot field for protecting the memory that the external pointer represents, if that memory is allocated from the R heap. Both tag and prot can be R_NilValue, and often are.

The elements of an external pointer can be accessed and set via

```
void *R_ExternalPtrAddr(SEXP s);
SEXP R_ExternalPtrTag(SEXP s);
SEXP R_ExternalPtrProtected(SEXP s);
void R_ClearExternalPtr(SEXP s);
void R_SetExternalPtrAddr(SEXP s, void *p);
void R_SetExternalPtrTag(SEXP s, SEXP tag);
void R_SetExternalPtrProtected(SEXP s, SEXP p);
```

Clearing a pointer sets its value to the C NULL pointer.

An external pointer object can have a *finalizer*, a piece of code to be run when the object is garbage collected. This can be R code or C code, and the various interfaces are, respectively.

```
void R_RegisterFinalizerEx(SEXP s, SEXP fun, Rboolean onexit);
```

```
typedef void (*R_CFinalizer_t)(SEXP);
void R_RegisterCFinalizerEx(SEXP s, R_CFinalizer_t fun, Rboolean onexit);
```

The R function indicated by fun should be a function of a single argument, the object to be finalized. R does not perform a garbage collection when shutting down, and the onexit argument

of the extended forms can be used to ask that the finalizer be run during a normal shutdown of the R session. It is suggested that it is good practice to clear the pointer on finalization.

The only R level function for interacting with external pointers is reg.finalizer which can be used to set a finalizer.

It is probably not a good idea to allow an external pointer to be saved and then reloaded, but if this happens the pointer will be set to the C NULL pointer.

Finalizers can be run at many places in the code base and much of it, including the R interpreter, is not re-entrant. So great care is needed in choosing the code to be run in a finalizer. As from R 3.0.3 finalizers are marked to be run at garbage collection but only run at a somewhat safe point thereafter.

Weak references are used to allow the programmer to maintain information on entities without preventing the garbage collection of the entities once they become unreachable.

A weak reference contains a key and a value. The value is reachable is if it either reachable directly or *via* weak references with reachable keys. Once a value is determined to be unreachable during garbage collection, the key and value are set to R_NilValue and the finalizer will be run later in the garbage collection.

Weak reference objects are created by one of

```
SEXP R_MakeWeakRef(SEXP key, SEXP val, SEXP fin, Rboolean onexit);
SEXP R_MakeWeakRefC(SEXP key, SEXP val, R_CFinalizer_t fin,
    Rboolean onexit);
```

where the R or C finalizer are specified in exactly the same way as for an external pointer object (whose finalization interface is implemented *via* weak references).

The parts can be accessed via

```
SEXP R_WeakRefKey(SEXP w);
SEXP R_WeakRefValue(SEXP w);
void R_RunWeakRefFinalizer(SEXP w);
```

A toy example of the use of weak references can be found at homepage.stat.uiowa.edu/~luke/R/references/weakfinex.html (http://homepage.stat.uiowa.edu/~luke/R/references/weakfinex.html), but that is used to add finalizers to external pointers which can now be done more directly. At the time of writing no CRAN or Bioconductor package uses weak references.

5.13.1 An example

Package RODBC (https://CRAN.R-project.org/package=RODBC) uses external pointers to maintain its *channels*, connections to databases. There can be several connections open at once, and the status information for each is stored in a C structure (pointed to by this_handle) in the code extract below) that is returned *via* an external pointer as part of the RODBC 'channel' (as the "handle_ptr" attribute). The external pointer is created by

```
SEXP ans, ptr;
ans = PROTECT(allocVector(INTSXP, 1));
ptr = R_MakeExternalPtr(thisHandle, install("RODBC_channel"), R_NilValue);
PROTECT(ptr);
R_RegisterCFinalizerEx(ptr, chanFinalizer, TRUE);
...
/* return the channel no */
INTEGER(ans)[0] = nChannels;
/* and the connection string as an attribute */
setAttrib(ans, install("connection.string"), constr);
setAttrib(ans, install("handle_ptr"), ptr);
```

```
UNPROTECT(3);
return ans;
```

Note the symbol given to identify the usage of the external pointer, and the use of the finalizer. Since the final argument when registering the finalizer is TRUE, the finalizer will be run at the of the R session (unless it crashes). This is used to close and clean up the connection to the database. The finalizer code is simply

```
static void chanFinalizer(SEXP ptr)
{
   if(!R_ExternalPtrAddr(ptr)) return;
   inRODBCClose(R_ExternalPtrAddr(ptr));
   R_ClearExternalPtr(ptr); /* not really needed */
}
```

Clearing the pointer and checking for a NULL pointer avoids any possibility of attempting to close an already-closed channel.

R's connections provide another example of using external pointers, in that case purely to be able to use a finalizer to close and destroy the connection if it is no longer is use.

5.14 Vector accessor functions

The vector accessors like REAL and INTEGER and VECTOR_ELT are *functions* when used in R extensions. (For efficiency they are macros when used in the R source code, apart from SET_STRING_ELT and SET_VECTOR_ELT which are always functions.)

The accessor functions check that they are being used on an appropriate type of SEXP.

If efficiency is essential, the macro versions of the accessors can be obtained by defining 'USE_RINTERNALS' before including Rinternals.h. If you find it necessary to do so, please do test that your code compiles without 'USE_RINTERNALS' defined, as this provides a stricter test that the accessors have been used correctly. Note too that the use of 'USE_RINTERNALS' when the header is included in C++ code is not supported: doing so may use C99 features which are not necessarily supported by the C++ compiler. Nor is use with Rdefines.h supported.

5.15 Character encoding issues

CHARSXPs can be marked as coming from a known encoding (Latin-1 or UTF-8). This is mainly intended for human-readable output, and most packages can just treat such CHARSXPs as a whole. However, if they need to be interpreted as characters or output at C level then it would normally be correct to ensure that they are converted to the encoding of the current locale: this can be done by accessing the data in the CHARSXP by translateChar rather than by CHAR. If re-encoding is needed this allocates memory with R_alloc which thus persists to the end of the .Call/.External call unless vmaxset is used (see Section 6.1.1 [Transient storage allocation], page 134).

There is a similar function translateCharUTF8 which converts to UTF-8: this has the advantage that a faithful translation is almost always possible (whereas only a few languages can be represented in the encoding of the current locale unless that is UTF-8).

There is a public interface to the encoding marked on CHARXSXPs via

```
typedef enum {CE_NATIVE, CE_UTF8, CE_LATIN1, CE_SYMBOL, CE_ANY} cetype_t;
cetype_t getCharCE(SEXP);
SEXP mkCharCE(const char *, cetype_t);
```

Only CE_UTF8 and CE_LATIN1 are marked on CHARSXPs (and so Rf_getCharCE will only return one of the first three), and these should only be used on non-ASCII strings. Value CE_SYMBOL is used internally to indicate Adobe Symbol encoding. Value CE_ANY is used to indicate a character

string that will not need re-encoding – this is used for character strings known to be in ASCII, and can also be used as an input parameter where the intention is that the string is treated as a series of bytes. (See the comments under mkChar about the length of input allowed.)

Function

```
const char *reEnc(const char *x, cetype_t ce_in, cetype_t ce_out,
  int subst);
```

can be used to re-encode character strings: like translateChar it returns a string allocated by R_alloc. This can translate from CE_SYMBOL to CE_UTF8, but not conversely. Argument subst controls what to do with untranslatable characters or invalid input: this is done byte-by-byte with 1 indicates to output hex of the form <a0>, and 2 to replace by ., with any other value causing the byte to produce no output.

There is also

SEXP mkCharLenCE(const char *, size_t, cetype_t); to create marked character strings of a given length.

6 The R API: entry points for C code

There are a large number of entry points in the R executable/DLL that can be called from C code (and some that can be called from FORTRAN code). Only those documented here are stable enough that they will only be changed with considerable notice.

The recommended procedure to use these is to include the header file R.h in your C code by #include <R.h>

This will include several other header files from the directory $R_INCLUDE_DIR/R_ext$, and there are other header files there that can be included too, but many of the features they contain should be regarded as undocumented and unstable.

Most of these header files, including all those included by R.h, can be used from C++ code.

Note: Because R re-maps many of its external names to avoid clashes with user code, it is *essential* to include the appropriate header files when using these entry points.

This remapping can cause problems¹, and can be eliminated by defining R_NO_REMAP and prepending 'Rf_' to all the function names used from Rinternals.h and R_ext/Error.h. These problems can usually be avoided by including other headers (such as system headers and those for external software used by the package) before R.h.

We can classify the entry points as

API Entry points which are documented in this manual and declared in an installed header file. These can be used in distributed packages and will only be changed after deprecation.

public Entry points declared in an installed header file that are exported on all R platforms but are not documented and subject to change without notice.

private Entry points that are used when building R and exported on all R platforms but are not declared in the installed header files. Do not use these in distributed code.

hidden Entry points that are where possible (Windows and some modern Unix-alike compilers/loaders when using R as a shared library) not exported.

6.1 Memory allocation

There are two types of memory allocation available to the C programmer, one in which R manages the clean-up and the other in which user has full control (and responsibility).

6.1.1 Transient storage allocation

Here R will reclaim the memory at the end of the call to .C, .Call or .External. Use char $*R_alloc(size_t n, int size)$

which allocates n units of size bytes each. A typical usage (from package stats) is

```
x = (int *) R_alloc(nrows(merge)+2, sizeof(int));
```

(size_t is defined in stddef.h which the header defining R_alloc includes.)

There is a similar call, S_alloc (for compatibility with older versions of S) which zeroes the memory allocated,

```
char *S_alloc(long n, int size)
and
char *S_realloc(char *p, long new, long old, int size)
```

¹ Known problems are redefining LENGTH, error, length, vector and warning

which changes the allocation size from old to new units, and zeroes the additional units.

For compatibility with current versions of S, header S.h (only) defines wrapper macros equivalent to

```
type* Salloc(long n, int type)
type* Srealloc(char *p, long new, long old, int type)
```

This memory is taken from the heap, and released at the end of the .C, .Call or .External call. Users can also manage it, by noting the current position with a call to vmaxget and subsequently clearing memory allocated by a call to vmaxset. An example might be

```
void *vmax = vmaxget()
// a loop involving the use of R_alloc at each iteration
vmaxset(vmax)
```

This is only recommended for experts.

Note that this memory will be freed on error or user interrupt (if allowed: see Section 6.12 [Allowing interrupts], page 147).

Note that although n is $size_t$, there may be limits imposed by R's internal allocation mechanism. These will only come into play on 64-bit systems, where the limit for n prior to R 3.0.0 was just under 16Gb.

The memory returned is only guaranteed to be aligned as required for double pointers: take precautions if casting to a pointer which needs more. As from R 3.2.0 there is also

```
long double *R_allocLD(size_t n)
```

which is guaranteed to have the 16-byte alignment needed for long double pointers on some platforms.

These functions should only be used in code called by .C etc, never from front-ends. They are not thread-safe.

6.1.2 User-controlled memory

The other form of memory allocation is an interface to malloc, the interface providing R error handling. This memory lasts until freed by the user and is additional to the memory allocated for the R workspace.

The interface functions are

```
type* Calloc(size_t n, type)
type* Realloc(any *p, size_t n, type)
void Free(any *p)
```

providing analogues of calloc, realloc and free. If there is an error during allocation it is handled by R, so if these routines return the memory has been successfully allocated or freed. Free will set the pointer p to NULL. (Some but not all versions of S do so.)

Users should arrange to Free this memory when no longer needed, including on error or user interrupt. This can often be done most conveniently from an on.exit action in the calling R function – see pwilcox for an example.

Do not assume that memory allocated by Calloc/Realloc comes from the same pool as used by malloc: in particular do not use free or strdup with it.

Memory obtained by these functions should be aligned in the same way as malloc, that is 'suitably aligned for any kind of variable'.

These entry points need to be prefixed by R_ if STRICT_R_HEADERS has been defined.

6.2 Error handling

The basic error handling routines are the equivalents of stop and warning in R code, and use the same interface.

```
void error(const char * format, ...);
void warning(const char * format, ...);
```

These have the same call sequences as calls to printf, but in the simplest case can be called with a single character string argument giving the error message. (Don't do this if the string contains '%' or might otherwise be interpreted as a format.)

If $\mathtt{STRICT_R_HEADERS}$ is not defined there is also an S-compatibility interface which uses calls of the form

```
PROBLEM ..... ERROR
MESSAGE ..... WARN
PROBLEM ..... RECOVER(NULL_ENTRY)
MESSAGE ..... WARNING(NULL_ENTRY)
```

the last two being the forms available in all S versions. Here '.....' is a set of arguments to printf, so can be a string or a format string followed by arguments separated by commas.

6.2.1 Error handling from FORTRAN

There are two interface function provided to call error and warning from FORTRAN code, in each case with a simple character string argument. They are defined as

```
subroutine rexit(message)
subroutine rwarn(message)
```

Messages of more than 255 characters are truncated, with a warning.

6.3 Random number generation

The interface to R's internal random number generation routines is

```
double unif_rand();
double norm_rand();
double exp_rand();
```

giving one uniform, normal or exponential pseudo-random variate. However, before these are used, the user must call

```
GetRNGstate();
```

and after all the required variates have been generated, call

```
PutRNGstate();
```

These essentially read in (or create) .Random.seed and write it out after use.

File S.h defines seed_in and seed_out for S-compatibility rather than GetRNGstate and PutRNGstate. These take a long * argument which is ignored.

The random number generator is private to R; there is no way to select the kind of RNG or set the seed except by evaluating calls to the R functions.

The C code behind R's rxxx functions can be accessed by including the header file Rmath.h; See Section 6.7.1 [Distribution functions], page 139. Those calls generate a single variate and should also be enclosed in calls to GetRNGstate and PutRNGstate.

6.4 Missing and IEEE special values

A set of functions is provided to test for NA, Inf, -Inf and NaN. These functions are accessed via macros:

ISNA(x) True for R's NA only

ISNAN(x) True for R's NA and IEEE NaN R_FINITE(x) False for Inf, -Inf, NA, NaN

and via function R_IsNaN which is true for NaN but not NA.

Do use R_FINITE rather than isfinite or finite; the latter is often mendacious and isfinite is only available on a some platforms, on which R_FINITE is a macro expanding to isfinite.

Currently in C code ISNAN is a macro calling isnan. (Since this gives problems on some C++ systems, if the R headers is called from C++ code a function call is used.)

You can check for Inf or -Inf by testing equality to R_PosInf or R_NegInf, and set (but not test) an NA as NA_REAL.

All of the above apply to *double* variables only. For integer variables there is a variable accessed by the macro NA_INTEGER which can used to set or test for missingness.

6.5 Printing

The most useful function for printing from a C routine compiled into R is Rprintf. This is used in exactly the same way as printf, but is guaranteed to write to R's output (which might be a GUI console rather than a file, and can be re-directed by sink). It is wise to write complete lines (including the "\n") before returning to R. It is defined in R_ext/Print.h.

The function REprintf is similar but writes on the error stream (stderr) which may or may not be different from the standard output stream.

Functions Rvprintf and REvprintf are analogues using the vprintf interface. Because that is a C99 interface, they are only defined by R_ext/Print.h in C++ code if the macro R_USE_C99_IN_CXX is defined when it is included.

Another circumstance when it may be important to use these functions is when using parallel computation on a cluster of computational nodes, as their output will be re-directed/logged appropriately.

6.5.1 Printing from FORTRAN

On many systems FORTRAN write and print statements can be used, but the output may not interleave well with that of C, and will be invisible on GUI interfaces. They are not portable and best avoided.

Three subroutines are provided to ease the output of information from FORTRAN code.

```
subroutine dblepr(label, nchar, data, ndata) subroutine realpr(label, nchar, data, ndata) subroutine intpr (label, nchar, data, ndata)
```

Here *label* is a character label of up to 255 characters, *nchar* is its length (which can be -1 if the whole label is to be used), and *data* is an array of length at least *ndata* of the appropriate type (double precision, real and integer respectively). These routines print the label on one line and then print *data* as if it were an R vector on subsequent line(s). They work with zero *ndata*, and so can be used to print a label alone.

6.6 Calling C from FORTRAN and vice versa

Naming conventions for symbols generated by FORTRAN differ by platform: it is not safe to assume that FORTRAN names appear to C with a trailing underscore. To help cover up the platform-specific differences there is a set of macros that should be used.

```
F77_SUB(name)
to define a function in C to be called from FORTRAN

F77_NAME(name)
to declare a FORTRAN routine in C before use

F77_CALL(name)
to call a FORTRAN routine from C

F77_COMDECL(name)
to declare a FORTRAN common block in C

F77_COM(name)
to access a FORTRAN common block from C
```

On most current platforms these are all the same, but it is unwise to rely on this. Note that names with underscores are not legal in FORTRAN 77, and are not portably handled by the above macros. (Also, all FORTRAN names for use by R are lower case, but this is not enforced by the macros.)

For example, suppose we want to call R's normal random numbers from FORTRAN. We need a C wrapper along the lines of

```
#include <R.h>

void F77_SUB(rndstart)(void) { GetRNGstate(); }

void F77_SUB(rndend)(void) { PutRNGstate(); }

double F77_SUB(normrnd)(void) { return norm_rand(); }

to be called from FORTRAN as in

subroutine testit()

double precision normrnd, x

call rndstart()

x = normrnd()

call dblepr("X was", 5, x, 1)

call rndend()

end
```

Note that this is not guaranteed to be portable, for the return conventions might not be compatible between the C and FORTRAN compilers used. (Passing values via arguments is safer.)

The standard packages, for example **stats**, are a rich source of further examples.

Passing character strings from C to FORTRAN 77 or *vice versa* is not portable (and to Fortran 90 or later is even less so). We have found that it helps to ensure that a C string to be passed is followed by several nuls (and not just the one needed as a C terminator). But for maximal portability character strings in FORTRAN should be avoided.

6.7 Numerical analysis subroutines

R contains a large number of mathematical functions for its own use, for example numerical linear algebra computations and special functions.

The header files R_ext/BLAS.h, R_ext/Lapack.h and R_ext/Linpack.h contains declarations of the BLAS, LAPACK and LINPACK linear algebra functions included in R. These are

expressed as calls to FORTRAN subroutines, and they will also be usable from users' FORTRAN code. Although not part of the official API, this set of subroutines is unlikely to change (but might be supplemented).

The header file Rmath.h lists many other functions that are available and documented in the following subsections. Many of these are C interfaces to the code behind R functions, so the R function documentation may give further details.

6.7.1 Distribution functions

The routines used to calculate densities, cumulative distribution functions and quantile functions for the standard statistical distributions are available as entry points.

The arguments for the entry points follow the pattern of those for the normal distribution:

```
double dnorm(double x, double mu, double sigma, int give_log);
double pnorm(double x, double mu, double sigma, int lower_tail,
    int give_log);
double qnorm(double p, double mu, double sigma, int lower_tail,
    int log_p);
double rnorm(double mu, double sigma);
```

That is, the first argument gives the position for the density and CDF and probability for the quantile function, followed by the distribution's parameters. Argument *lower_tail* should be TRUE (or 1) for normal use, but can be FALSE (or 0) if the probability of the upper tail is desired or specified.

Finally, $give_log$ should be non-zero if the result is required on log scale, and log_p should be non-zero if p has been specified on log scale.

Note that you directly get the cumulative (or "integrated") hazard function, $H(t) = -\log(1 - F(t))$, by using

```
- pdist(t, ..., /*lower_tail = */ FALSE, /* give_log = */ TRUE) or shorter (and more cryptic) - pdist(t, ..., 0, 1).
```

The random-variate generation routine rnorm returns one normal variate. See Section 6.3 [Random numbers], page 136, for the protocol in using the random-variate routines.

Note that these argument sequences are (apart from the names and that **rnorm** has no n) mainly the same as the corresponding R functions of the same name, so the documentation of the R functions can be used. Note that the exponential and gamma distributions are parametrized by scale rather than rate.

For reference, the following table gives the basic name (to be prefixed by 'd', 'p', 'q' or 'r' apart from the exceptions noted) and distribution-specific arguments for the complete set of distributions.

beta	beta	a, b
non-central beta	nbeta	a, b, ncp
binomial	binom	n, p
Cauchy	cauchy	location, scale
chi-squared	chisq	df
non-central chi-squared	nchisq	df, ncp
exponential	exp	scale (and not rate)
\mathbf{F}	f	n1, n2
non-central F	nf	n1, n2, ncp
gamma	gamma	shape, scale
geometric	geom	p
hypergeometric	hyper	NR, NB, n
logistic	logis	location, scale

[Function]

lognormal	lnorm	logmean, logsd
negative binomial	nbinom	$\mathtt{size},\mathtt{prob}$
normal	norm	mu, sigma
Poisson	pois	lambda
Student's t	t	n
non-central t	nt	${ t df},{ t delta}$
Studentized range	tukey $(*)$	$\mathtt{rr},\mathtt{cc},\mathtt{df}$
uniform	unif	a, b
Weibull	weibull	${ t shape}, { t scale}$
Wilcoxon rank sum	wilcox	$\mathtt{m},\ \mathtt{n}$
Wilcoxon signed rank	signrank	n

Entries marked with an asterisk only have 'p' and 'q' functions available, and none of the non-central distributions have 'r' functions. After a call to dwilcox, pwilcox or qwilcox the function wilcox_free() should be called, and similarly for the signed rank functions.

(If remapping is suppressed, the Normal distribution names are Rf_dnorm4, Rf_pnorm5 and Rf_qnorm5.)

For the negative binomial distribution ('nbinom'), in addition to the (size, prob) parametrization, the alternative (size, mu) parametrization is provided as well by functions '[dpqr]nbinom_mu()', see ?NegBinomial in R.

Functions $dpois_raw(x, *)$ and $dbinom_raw(x, *)$ are versions of the Poisson and binomial probability mass functions which work continuously in x, whereas dbinom(x, *) and dpois(x, *) only return non zero values for integer x.

```
double dbinom_raw(double x, double n, double p, double q, int give_log)
double dpois_raw (double x, double lambda, int give_log)
```

Note that $dbinom_raw()$ gets both p and q=1-p which may be advantageous when one of them is close to 1.

6.7.2 Mathematical functions

double 1choose (double n, double k)

The Gamma function, the natural logarithm of its absolute value and first four derivatives and the n-th derivative of Psi, the digamma function, which is the derivative of lgammafn. In other words, digamma(x) is the same as (psigamma(x,0), trigamma(x) == psigamma(x,1), etc.

```
double beta (double a, double b) [Function]
double lbeta (double a, double b) [Function]
The (complete) Beta function and its natural logarithm.
```

```
double choose (double n, double k) [Function]
```

The number of combinations of k items chosen from n and the natural logarithm of its absolute value, generalized to arbitrary real n. k is rounded to the nearest integer (with a warning if needed).

```
double bessel_i (double x, double nu, double expo)[Function]double bessel_j (double x, double nu)[Function]double bessel_k (double x, double nu, double expo)[Function]double bessel_y (double x, double nu)[Function]
```

Bessel functions of types I, J, K and Y with index nu. For bessel_i and bessel_k there is the option to return $\exp(-x)$ I(x; nu) or $\exp(x)$ K(x; nu) if $\exp o$ is 2. (Use $\exp o == 1$ for unscaled values.)

6.7.3 Numerical Utilities

There are a few other numerical utility functions available as entry points.

```
double R_pow (double x, double y) [Function]
double R_pow_di (double x, int i) [Function]
```

 $R_{pow}(x, y)$ and $R_{pow_di}(x, i)$ compute x^y and x^i , respectively using $R_{pow_di}(x, i)$ and returning the proper result (the same as R) for the cases where x, y or i are 0 or missing or infinite or NaN.

double log1p (double x)

[Function]

Computes log(1 + x) (log 1 plus x), accurately even for small x, i.e., $|x| \ll 1$.

This should be provided by your platform, in which case it is not included in Rmath.h, but is (probably) in math.h which Rmath.h includes.

double log1pmx (double x)

[Function]

Computes log(1 + x) - x (log 1 plus x minus x), accurately even for small x, i.e., $|x| \ll 1$.

double log1pexp (double x)

[Function]

Computes log(1 + exp(x)) (log 1 plus exp), accurately, notably for large x, e.g., x > 720.

double expm1 (double x)

[Function]

Computes $\exp(x) - 1$ (exp x minus 1), accurately even for small x, i.e., $|x| \ll 1$.

This should be provided by your platform, in which case it is not included in Rmath.h, but is (probably) in math.h which Rmath.h includes.

double lgamma1p (double x)

[Function]

Computes log(gamma(x + 1)) (log(gamma(1 plus x))), accurately even for small x, i.e., 0 < x < 0.5.

double cospi (double x)

[Function]

Computes cos(pi * x) (where pi is 3.14159...), accurately, notably for half integer x.

This might be provided by your platform², in which case it is not included in Rmath.h, but is in math.h which Rmath.h includes. (Ensure that neither math.h nor cmath is included before Rmath.h or define

```
#define __STDC_WANT_IEC_60559_FUNCS_EXT__ 1
```

before the first inclusion.)

double sinpi (double x)

[Function]

Computes sin(pi * x) accurately, notably for (half) integer x.

This might be provided by your platform, in which case it is not included in Rmath.h, but is in math.h which Rmath.h includes (but see the comments for cospi).

 $^{^2}$ It is an optional C11 extension.

double tanpi (double x)

[Function]

Computes tan(pi * x) accurately, notably for (half) integer x.

This might be provided by your platform, in which case it is not included in Rmath.h, but is in math.h which Rmath.h includes (but see the comments for cospi).

```
double logspace_add (double logx, double logy)[Function]double logspace_sub (double logx, double logy)[Function]double logspace_sum (const double* logx, int n)[Function]
```

Compute the log of a sum or difference from logs of terms, i.e., "x + y" as log (exp(logx) + exp(logy)) and "x - y" as log (exp(logx) - exp(logy)), and " $sum_i x[i]$ " as log (sum[i = 1:n exp(logx[i])]) without causing unnecessary overflows or throwing away too much accuracy.

```
int imax2 (int x, int y) [Function]
int imin2 (int x, int y) [Function]
double fmax2 (double x, double y) [Function]
double fmin2 (double x, double y)
```

Return the larger (max) or smaller (min) of two integer or double numbers, respectively. Note that fmax2 and fmin2 differ from C99's fmax and fmin when one of the arguments is a NaN: these versions return NaN.

double sign (double x)

[Function]

Compute the *signum* function, where sign(x) is 1, 0, or -1, when x is positive, 0, or negative, respectively, and NaN if x is a NaN.

double fsign (double x, double y)

[Function]

Performs "transfer of sign" and is defined as |x| * sign(y).

double fprec (double x, double digits)

[Function]

Returns the value of x rounded to digits decimal digits (after the decimal point).

This is the function used by R's signif().

double fround (double x, double digits)

[Function]

Returns the value of x rounded to digits significant decimal digits.

This is the function used by R's round().

double ftrunc (double x)

[Function]

Returns the value of x truncated (to an integer value) towards zero.

Note that this is no longer needed in C code, as C99 provide a trunc function. It is needed for portable C++98 code.

6.7.4 Mathematical constants

R has a set of commonly used mathematical constants encompassing constants usually found math.h and contains further ones that are used in statistical computations. All these are defined to (at least) 30 digits accuracy in Rmath.h. The following definitions use ln(x) for the natural logarithm (log(x) in R).

Name	Definition (ln = log)	round(value, 7)
M_E	e	2.7182818
M_LOG2E	$\log 2(e)$	1.4426950
M_LOG10E	$\log 10(e)$	0.4342945
M_LN2	$\ln(2)$	0.6931472
M_LN10	ln(10)	2.3025851

π	3.1415927
π/Z	1.5707963
$\pi/4$	0.7853982
$1/\pi$	0.3183099
$2/\pi$	0.6366198
$2/\operatorname{sqrt}(\pi)$	1.1283792
$\operatorname{sqrt}(2)$	1.4142136
$1/\operatorname{sqrt}(2)$	0.7071068
$\operatorname{sqrt}(3)$	1.7320508
$\operatorname{sqrt}(32)$	5.6568542
$\log 10(2)$	0.3010300
2π	6.2831853
$\operatorname{sqrt}(\pi)$	1.7724539
$1/\operatorname{sqrt}(2\pi)$	0.3989423
$\operatorname{sqrt}(2/\pi)$	0.7978846
$\ln(\operatorname{sqrt}(\pi))$	0.5723649
$\ln(\operatorname{sqrt}(2\pi))$	0.9189385
$\ln(\operatorname{sqrt}(\pi/2))$	0.2257914
	$1/\pi$ $2/\pi$ $2/sqrt(\pi)$ $sqrt(2)$ $1/sqrt(2)$ $sqrt(3)$ $sqrt(32)$ $log10(2)$ 2π $sqrt(\pi)$ $1/sqrt(2\pi)$ $sqrt(2/\pi)$ $ln(sqrt(\pi))$ $ln(sqrt(2\pi))$

There are a set of constants (PI, DOUBLE_EPS) (and so on) defined (unless STRICT_R_HEADERS is defined) in the included header R_ext/Constants.h, mainly for compatibility with S.

Further, the included header R_ext/Boolean.h has enumeration constants TRUE and FALSE of type Rboolean in order to provide a way of using "logical" variables in C consistently. This can conflict with other software: for example it conflicts with the headers in IJG's jpeg-9 (but not earlier versions).

6.8 Optimization

The C code underlying optim can be accessed directly. The user needs to supply a function to compute the function to be minimized, of the type

```
typedef double optimfn(int n, double *par, void *ex);
```

where the first argument is the number of parameters in the second argument. The third argument is a pointer passed down from the calling routine, normally used to carry auxiliary information.

Some of the methods also require a gradient function

```
typedef void optimgr(int n, double *par, double *gr, void *ex);
```

which passes back the gradient in the gr argument. No function is provided for finite-differencing, nor for approximating the Hessian at the result.

The interfaces (defined in header R_ext/Applic.h) are

• Nelder Mead:

```
void nmmin(int n, double *xin, double *x, double *Fmin, optimfn fn,
  int *fail, double abstol, double intol, void *ex,
  double alpha, double beta, double gamma, int trace,
  int *fncount, int maxit);
```

• BFGS:

```
void vmmin(int n, double *x, double *Fmin,
  optimfn fn, optimgr gr, int maxit, int trace,
  int *mask, double abstol, double reltol, int nREPORT,
  void *ex, int *fncount, int *grcount, int *fail);
```

• Conjugate gradients:

```
void cgmin(int n, double *xin, double *x, double *Fmin,
   optimfn fn, optimgr gr, int *fail, double abstol,
   double intol, void *ex, int type, int trace,
   int *fncount, int *grcount, int maxit);
```

• Limited-memory BFGS with bounds:

```
void lbfgsb(int n, int lmm, double *x, double *lower,
    double *upper, int *nbd, double *Fmin, optimfn fn,
    optimgr gr, int *fail, void *ex, double factr,
    double pgtol, int *fncount, int *grcount,
    int maxit, char *msg, int trace, int nREPORT);
```

• Simulated annealing:

```
void samin(int n, double *x, double *Fmin, optimfn fn, int maxit,
  int tmax, double temp, int trace, void *ex);
```

Many of the arguments are common to the various methods. n is the number of parameters, x or xin is the starting parameters on entry and x the final parameters on exit, with final value returned in Fmin. Most of the other parameters can be found from the help page for optim: see the source code src/appl/lbfgsb.c for the values of nbd, which specifies which bounds are to be used.

6.9 Integration

The C code underlying integrate can be accessed directly. The user needs to supply a *vector-izing* C function to compute the function to be integrated, of the type

```
typedef void integr_fn(double *x, int n, void *ex);
```

where x[] is both input and output and has length n, i.e., a C function, say fn, of type integr_fn must basically do for(i in 1:n) x[i] := f(x[i], ex). The vectorization requirement can be used to speed up the integrand instead of calling it n times. Note that in the current implementation built on QUADPACK, n will be either 15 or 21. The ex argument is a pointer passed down from the calling routine, normally used to carry auxiliary information.

There are interfaces (defined in header R_ext/Applic.h) for integrals over finite and infinite intervals (or "ranges" or "integration boundaries").

• Finite:

```
void Rdqags(integr_fn f, void *ex, double *a, double *b,
    double *epsabs, double *epsrel,
    double *result, double *abserr, int *neval, int *ier,
    int *limit, int *lenw, int *last,
    int *iwork, double *work);
```

• Infinite:

```
void Rdqagi(integr_fn f, void *ex, double *bound, int *inf,
   double *epsabs, double *epsrel,
   double *result, double *abserr, int *neval, int *ier,
   int *limit, int *lenw, int *last,
   int *iwork, double *work);
```

Only the 3rd and 4th argument differ for the two integrators; for the finite range integral using Rdqags, a and b are the integration interval bounds, whereas for an infinite range integral using Rdqagi, bound is the finite bound of the integration (if the integral is not doubly-infinite) and inf is a code indicating the kind of integration range,

```
inf = 1 corresponds to (bound, +Inf),
```

```
inf = -1 corresponds to (-Inf, bound),
inf = 2 corresponds to (-Inf, +Inf),
```

f and ex define the integrand function, see above; epsabs and epsrel specify the absolute and relative accuracy requested, result, abserr and last are the output components value, abs.err and subdivisions of the R function integrate, where neval gives the number of integrand function evaluations, and the error code ier is translated to R's integrate() \$ message, look at that function definition. limit corresponds to integrate(..., subdivisions = *). It seems you should always define the two work arrays and the length of the second one as

```
lenw = 4 * limit;
iwork = (int *) R_alloc(limit, sizeof(int));
work = (double *) R_alloc(lenw, sizeof(double));
```

The comments in the source code in src/appl/integrate.c give more details, particularly about reasons for failure (ier >= 1).

6.10 Utility functions

R has a fairly comprehensive set of sort routines which are made available to users' C code. The following is declared in header file Rinternals.h.

Rboolean decreasing)

R_orderVector() corresponds to R's order(..., na.last, decreasing). More specifically, indx <- order(x, y, na.last, decreasing) corresponds to R_orderVector(indx, n, Rf_lang2(x, y), nalast, decreasing) and for three vectors, Rf_lang3(x,y,z) is used as arglist.

Both R_orderVector and R_orderVector1 assume the vector indx to be allocated to length $\geq n$. On return, indx[] contains a permutation of 0: (n-1), i.e., 0-based C indices (and not 1-based R indices, as R's order()).

When ordering only one vector, R_orderVector1 is faster and corresponds (but is 0-based) to R's indx <- order(x, na.last, decreasing). It was added in R 3.3.0.

All other sort routines are declared in header file R_ext/Utils.h (included by R.h) and include the following.

```
void R_isort (int* x, int n)
void R_rsort (double* x, int n)
void R_csort (Rcomplex* x, int n)
void rsort_with_index (double* x, int* index, int n)

[Function]
[Function]
```

The first three sort integer, real (double) and complex data respectively. (Complex numbers are sorted by the real part first then the imaginary part.) NAs are sorted last.

rsort_with_index sorts on x, and applies the same permutation to index. NAs are sorted last.

```
void iPsort (int* x, int n, int k)
void rPsort (double* x, int n, int k) [Function]
```

void cPsort (Rcomplex* x, int n, int k)

[Function]

These all provide (very) partial sorting: they permute x so that x[k] is in the correct place with smaller values to the left, larger ones to the right.

```
void R_qsort (double *v, size_t i, size_t j)
void R_qsort_I (double *v, int *I, int i, int j)
void R_qsort_int (int *iv, size_t i, size_t j)
void R_qsort_int_I (int *iv, int *I, int i, int j)
[Function]
[Function]
```

These routines sort v[i:j] or iv[i:j] (using 1-indexing, i.e., v[1] is the first element) calling the quicksort algorithm as used by R's sort(v, method = "quick") and documented on the help page for the R function sort. The ..._I() versions also return the sort.index() vector in I. Note that the ordering is not stable, so tied values may be permuted.

Note that NAs are not handled (explicitly) and you should use different sorting functions if NAs can be present.

subroutine qsort3 (double precision v, integer ii, integer jj) [Function]

The FORTRAN interface routines for sorting double precision vectors are qsort3 and qsort4, equivalent to R_qsort and R_qsort_I, respectively.

Given the *nr* by *nc* matrix matrix in column-major ("FORTRAN") order, R_max_col() returns in maxes[i-1] the column number of the maximal element in the *i*-th row (the same as R's max.col() function). In the case of ties (multiple maxima), *ties_meth is an integer code in 1:3 determining the method: 1 = "random", 2 = "first" and 3 = "last". See R's help page ?max.col.

Given the ordered vector xt of length n, return the interval or index of x in xt[], typically $\max(i; 1 \le i \le n \& xt[i] \le x)$ where we use 1-indexing as in R and FORTRAN (but not C). If $rightmost_closed$ is true, also returns n-1 if x equals xt[n]. If all_inside is not 0, the result is coerced to lie in 1: (n-1) even when x is outside the xt[] range. On return, *mflag equals -1 if x < xt[1], +1 if x >= xt[n], and 0 otherwise.

The algorithm is particularly fast when ilo is set to the last result of findInterval() and x is a value of a sequence which is increasing or decreasing for subsequent calls.

findInterval2() is a generalization of findInterval(), with an extra Rboolean argument $left_open$. Setting left_open = TRUE basically replaces all left-closed right-open intervals [s by left-open ones (s, see the help page of R function findInterval for details.

There is also an F77_CALL(interv)() version of findInterval() with the same arguments, but all pointers.

A system-independent interface to produce the name of a temporary file is provided as

Return a pathname for a temporary file with name beginning with *prefix* and ending with *fileext* in directory *tmpdir*. A NULL prefix or extension is replaced by "". Note that the

return value is malloced and should be freed when no longer needed (unlike the system call tmpnam).

There is also the internal function used to expand file names in several R functions, and called directly by path.expand.

const char * R_ExpandFileName (const char *fn)

[Function]

Expand a path name fn by replacing a leading tilde by the user's home directory (if defined). The precise meaning is platform-specific; it will usually be taken from the environment variable HOME if this is defined.

For historical reasons there are FORTRAN interfaces to functions D1MACH and I1MACH. These can be called from C code as e.g. F77_CALL(d1mach)(4). Note that these are emulations of the original functions by Fox, Hall and Schryer on NetLib at http://www.netlib.org/slatec/src/ for IEC 60559 arithmetic (required by R).

6.11 Re-encoding

R has its own C-level interface to the encoding conversion capabilities provided by iconv because there are incompatibilities between the declarations in different implementations of iconv.

These are declared in header file R_ext/Riconv.h.

```
void * Riconv_open (const char *to, const char *from)
```

[Function]

Set up a pointer to an encoding object to be used to convert between two encodings: "" indicates the current locale.

Convert as much as possible of inbuf to outbuf. Initially the int variables indicate the number of bytes available in the buffers, and they are updated (and the char pointers are updated to point to the next free byte in the buffer). The return value is the number of characters converted, or (size_t)-1 (beware: size_t is usually an unsigned type). It should be safe to assume that an error condition sets errno to one of E2BIG (the output buffer is full), EILSEQ (the input cannot be converted, and might be invalid in the encoding specified) or EINVAL (the input does not end with a complete multi-byte character).

```
int Riconv_close (void * cd)
```

[Function]

Free the resources of an encoding object.

6.12 Allowing interrupts

No port of R can be interrupted whilst running long computations in compiled code, so programmers should make provision for the code to be interrupted at suitable points by calling from C

```
#include <R_ext/Utils.h>
```

void R_CheckUserInterrupt(void);

and from FORTRAN

```
subroutine rchkusr()
```

These check if the user has requested an interrupt, and if so branch to R's error handling functions.

Note that it is possible that the code behind one of the entry points defined here if called from your C or FORTRAN code could be interruptible or generate an error and so not return to your code.

6.13 Platform and version information

The header files define USING_R, which can be used to test if the code is indeed being used with R.

Header file Rconfig.h (included by R.h) is used to define platform-specific macros that are mainly for use in other header files. The macro WORDS_BIGENDIAN is defined on big-endian³ systems (e.g. most OSes on Sparc and PowerPC hardware) and not on little-endian systems (nowadays all the commoner R platforms). It can be useful when manipulating binary files.

Header file Rversion.h (not included by R.h) defines a macro R_VERSION giving the version number encoded as an integer, plus a macro R_Version to do the encoding. This can be used to test if the version of R is late enough, or to include back-compatibility features. For protection against very old versions of R which did not have this macro, use a construction such as

```
#if defined(R_VERSION) && R_VERSION >= R_Version(3, 1, 0)
...
#endif
```

More detailed information is available in the macros R_MAJOR, R_MINOR, R_YEAR, R_MONTH and R_DAY: see the header file Rversion.h for their format. Note that the minor version includes the patchlevel (as in '2.2').

Packages which use alloca need to ensure it is defined: as it is part of neither C nor POSIX there is no standard way to do so. As from R 3.2.2 one can use

```
#include <Rconfig.h> // for HAVE_ALLOCA_H
#ifdef __GNUC__
// this covers gcc, clang, icc
# undef alloca
# define alloca(x) __builtin_alloca((x))
#elif defined(HAVE_ALLOCA_H)
// needed for native compilers on Solaris and AIX
# include <alloca.h>
#endif
```

(and this should be included before standard C headers such as **stdlib.h**, since on some platforms these include **malloc.h** which may have a conflicting definition), which suffices for known R platforms.

6.14 Inlining C functions

The C99 keyword inline should be recognized by all compilers nowadays used to build R. Portable code which might be used with earlier versions of R can be written using the macro R_INLINE (defined in file Rconfig.h included by R.h), as for example from package cluster (https://CRAN.R-project.org/package=cluster)

```
#include <R.h>
static R_INLINE int ind_2(int 1, int j)
{
...
}
```

Be aware that using inlining with functions in more than one compilation unit is almost impossible to do portably, see http://www.greenend.org.uk/rjk/2003/03/inline.html, so this usage is for static functions as in the example. All the R configure code has checked is that R_INLINE can be used in a single C file with the compiler used to build R. We recommend that packages making extensive use of inlining include their own configure code.

³ https://en.wikipedia.org/wiki/Endianness.

6.15 Controlling visibility

Header R_ext/Visibility has some definitions for controlling the visibility of entry points. These are only effective when 'HAVE_VISIBILITY_ATTRIBUTE' is defined – this is checked when R is configured and recorded in header Rconfig.h (included by R_ext/Visibility.h). It is generally defined on modern Unix-alikes with a recent compiler, but not supported on OS X nor Windows. Minimizing the visibility of symbols in a shared library will both speed up its loading (unlikely to be significant) and reduce the possibility of linking to other entry points of the same name.

C/C++ entry points prefixed by attribute_hidden will not be visible in the shared object. There is no comparable mechanism for FORTRAN entry points, but there is a more comprehensive scheme used by, for example package stats. Most compilers which allow control of visibility will allow control of visibility for all symbols *via* a flag, and where known the flag is encapsulated in the macros 'C_VISIBILITY' and F77_VISIBILITY for C and FORTRAN compilers. These are defined in etc/Makeconf and so available for normal compilation of package code. For example, src/Makevars could include

```
PKG_CFLAGS=$(C_VISIBILITY)
PKG_FFLAGS=$(F77_VISIBILITY)
```

This would end up with **no** visible entry points, which would be pointless. However, the effect of the flags can be overridden by using the attribute_visible prefix. A shared object which registers its entry points needs only for have one visible entry point, its initializer, so for example package **stats** has

```
void attribute_visible R_init_stats(DllInfo *dll)
{
    R_registerRoutines(dll, CEntries, CallEntries, FortEntries, NULL);
    R_useDynamicSymbols(dll, FALSE);
...
}
```

The visibility mechanism is not available on Windows, but there is an equally effective way to control which entry points are visible, by supplying a definitions file <code>pkgnme/src/pkgname-win.def</code>: only entry points listed in that file will be visible. Again using stats as an example, it has

```
LIBRARY stats.dll
EXPORTS
R_init_stats
```

6.16 Using these functions in your own C code

It is possible to build Mathlib, the R set of mathematical functions documented in Rmath.h, as a standalone library libRmath under both Unix-alikes and Windows. (This includes the functions documented in Section 6.7 [Numerical analysis subroutines], page 138, as from that header file.)

The library is not built automatically when R is installed, but can be built in the directory src/nmath/standalone in the R sources: see the file README there. To use the code in your own C program include

```
#define MATHLIB_STANDALONE
#include <Rmath.h>
and link against '-lRmath' (and perhaps '-lm'). There is an example file test.c.
```

A little care is needed to use the random-number routines. You will need to supply the uniform random number generator

```
double unif_rand(void)
```

or use the one supplied (and with a dynamic library or DLL you will have to use the one supplied, which is the Marsaglia-multicarry with an entry points

```
set_seed(unsigned int, unsigned int)
to set its seeds and
   get_seed(unsigned int *, unsigned int *)
to read the seeds).
```

6.17 Organization of header files

The header files which R installs are in directory $R_INCLUDE_DIR$ (default $R_HOME/include$). This currently includes

R.h includes many other files S.h different version for code ported from S definitions for using R's internal structures Rinternals.h macros for an S-like interface to the above (no Rdefines.h longer maintained) standalone math library Rmath.h R version information Rversion.h Rinterface.h for add-on front-ends (Unix-alikes only) Rembedded.h for add-on front-ends optimization and integration R_ext/Applic.h R_ext/BLAS.h C definitions for BLAS routines C (and R function) top-level task handlers R_ext/Callbacks.h R_ext/GetX11Image.h X11Image interface used by package trkplot C definitions for some LAPACK routines R_ext/Lapack.h R_ext/Linpack.h C definitions for some LINPACK routines, not all of which are included in R a small part of R's parse interface: not part of the R_ext/Parse.h stable API. for add-on front-ends R_ext/RStartup.h needed to register compiled code in packages R_ext/Rdynload.h R_ext/R-ftp-http.h interface to internal method of download.file R_ext/Riconv.h interface to iconv

R_ext/eventloop.h for add-on front-ends and for packages that need to share in the R event loops (on all platforms)

definitions controlling visibility

The following headers are included by R.h:

R_ext/Visibility.h

Rconfig.h configuration info that is made available R_ext/Arith.h handling for NAs, NaNs, Inf/-Inf TRUE/FALSE type R_ext/Boolean.h R_ext/Complex.h C typedefs for R's complex R_ext/Constants.h constants R_ext/Error.h error handling memory allocation R_ext/Memory.h R_ext/Print.h Rprintf and variations. definitions common to R.h and S.h, including R_ext/RS.h F77_CALL etc. R_ext/Random.h random number generation R_ext/Utils.h sorting and other utilities

R_ext/libextern.h definitions for exports from R.dll on Windows.

The graphics systems are exposed in headers R_ext/GraphicsEngine.h, R_ext/GraphicsDevice.h (which it includes) and R_ext/QuartzDevice.h. Facilities for defining custom connection implementations are provided in R_ext/Connections.h, but make sure you consult the file before use.

Let us re-iterate the advice to include system headers before the R header files, especially Rinternals.h (included by Rdefines.h) and Rmath.h, which redefine names which may be used in system headers (fewer if 'R_NO_REMAP' is defined, or 'R_NO_REMAP_RMATH' for Rmath.h, as from R 3.1.0).

7 Generic functions and methods

R programmers will often want to add methods for existing generic functions, and may want to add new generic functions or make existing functions generic. In this chapter we give guidelines for doing so, with examples of the problems caused by not adhering to them.

This chapter only covers the 'informal' class system copied from S3, and not with the S4 (formal) methods of package **methods**.

First, a *caveat*: a function named *gen.cl* will be invoked by the generic *gen* for class *cl*, so do not name functions in this style unless they are intended to be methods.

The key function for methods is NextMethod, which dispatches the next method. It is quite typical for a method function to make a few changes to its arguments, dispatch to the next method, receive the results and modify them a little. An example is

```
t.data.frame <- function(x)
{
    x <- as.matrix(x)
    NextMethod("t")
}</pre>
```

Note that the example above works because there is a *next* method, the default method, not that a new method is selected when the class is changed.

Any method a programmer writes may be invoked from another method by NextMethod, with the arguments appropriate to the previous method. Further, the programmer cannot predict which method NextMethod will pick (it might be one not yet dreamt of), and the end user calling the generic needs to be able to pass arguments to the next method. For this to work

A method must have all the arguments of the generic, including . . . if the generic does

It is a grave misunderstanding to think that a method needs only to accept the arguments it needs. The original S version of predict.lm did not have a ... argument, although predict did. It soon became clear that predict.glm needed an argument dispersion to handle over-dispersion. As predict.lm had neither a dispersion nor a ... argument, NextMethod could no longer be used. (The legacy, two direct calls to predict.lm, lives on in predict.glm in R, which is based on the workaround for S3 written by Venables & Ripley.)

Further, the user is entitled to use positional matching when calling the generic, and the arguments to a method called by UseMethod are those of the call to the generic. Thus

A method must have arguments in exactly the same order as the generic.

To see the scale of this problem, consider the generic function scale, defined as

```
scale <- function (x, center = TRUE, scale = TRUE)
    UseMethod("scale")</pre>
```

Suppose an unthinking package writer created methods such as

```
scale.foo <- function(x, scale = FALSE, ...) { }
Then for x of class "foo" the calls
    scale(x, , TRUE)
    scale(x, scale = TRUE)</pre>
```

would do most likely do different things, to the justifiable consternation of the end user.

To add a further twist, which default is used when a user calls scale(x) in our example? What if

```
scale.bar <- function(x, center, scale = TRUE) NextMethod("scale")</pre>
```

and x has class c("bar", "foo")? It is the default specified in the method that is used, but the default specified in the generic may be the one the user sees. This leads to the recommendation:

If the generic specifies defaults, all methods should use the same defaults.

An easy way to follow these recommendations is to always keep generics simple, e.g.

```
scale <- function(x, ...) UseMethod("scale")</pre>
```

Only add parameters and defaults to the generic if they make sense in all possible methods implementing it.

7.1 Adding new generics

When creating a new generic function, bear in mind that its argument list will be the maximal set of arguments for methods, including those written elsewhere years later. So choosing a good set of arguments may well be an important design issue, and there need to be good arguments not to include a . . . argument.

If a ... argument is supplied, some thought should be given to its position in the argument sequence. Arguments which follow ... must be named in calls to the function, and they must be named in full (partial matching is suppressed after ...). Formal arguments before ... can be partially matched, and so may 'swallow' actual arguments intended for Although it is commonplace to make the ... argument the last one, that is not always the right choice.

Sometimes package writers want to make generic a function in the base package, and request a change in R. This may be justifiable, but making a function generic with the old definition as the default method does have a small performance cost. It is never necessary, as a package can take over a function in the base package and make it generic by something like

```
foo <- function(object, ...) UseMethod("foo")
foo.default <- function(object, ...) base::foo(object)</pre>
```

Earlier versions of this manual suggested assigning foo.default <- base::foo. This is not a good idea, as it captures the base function at the time of installation and it might be changed as R is patched or updated.

The same idea can be applied for functions in other packages with namespaces.

8 Linking GUIs and other front-ends to R

There are a number of ways to build front-ends to R: we take this to mean a GUI or other application that has the ability to submit commands to R and perhaps to receive results back (not necessarily in a text format). There are other routes besides those described here, for example the package Rserve (https://CRAN.R-project.org/package=Rserve) (from CRAN, see also https://www.rforge.net/Rserve/) and connections to Java in 'JRI' (part of the rJava (https://CRAN.R-project.org/package=rJava) package on CRAN) and the Omegahat/Bioconductor package 'SJava'.

Note that the APIs described in this chapter are only intended to be used in an alternative front-end: they are not part of the API made available for R packages and can be dangerous to use in a conventional package (although packages may contain alternative front-ends). Conversely some of the functions from the API (such as R_alloc) should not be used in front-ends.

8.1 Embedding R under Unix-alikes

R can be built as a shared library¹ if configured with --enable-R-shlib. This shared library can be used to run R from alternative front-end programs. We will assume this has been done for the rest of this section. Also, it can be built as a static library if configured with --enable-R-static-lib, and that can be used in a very similar way (at least on Linux: on other platforms one needs to ensure that all the symbols exported by libR.a are linked into the front-end).

The command-line R front-end, R_HOME/bin/exec/R, is one such example, and the former GNOME (see package gnomeGUI on CRAN's 'Archive' area) and OS X consoles are others. The source for R_HOME/bin/exec/R is in file src/main/Rmain.c and is very simple

```
int Rf_initialize_R(int ac, char **av); /* in ../unix/system.c */
void Rf_mainloop(); /* in main.c */

extern int R_running_as_main_program; /* in ../unix/system.c */
int main(int ac, char **av)
{
    R_running_as_main_program = 1;
    Rf_initialize_R(ac, av);
    Rf_mainloop(); /* does not return */
    return 0;
}
```

indeed, misleadingly simple. Remember that $R_HOME/bin/exec/R$ is run from a shell script $R_HOME/bin/R$ which sets up the environment for the executable, and this is used for

- Setting R_HOME and checking it is valid, as well as the path R_SHARE_DIR and R_DOC_DIR to the installed share and doc directory trees. Also setting R_ARCH if needed.
- Setting LD_LIBRARY_PATH to include the directories used in linking R. This is recorded as the default setting of R_LD_LIBRARY_PATH in the shell script R_HOME/etcR_ARCH/ldpaths.
- Processing some of the arguments, for example to run R under a debugger and to launch alternative front-ends to provide GUIs.

The first two of these can be achieved for your front-end by running it via R CMD. So, for example

```
R CMD /usr/local/lib/R/bin/exec/R
R CMD exec/R
```

¹ In the parlance of OS X this is a *dynamic* library, and is the normal way to build R on that platform.

will both work in a standard R installation. (R CMD looks first for executables in R_HOME/bin. These command-lines need modification if a sub-architecture is in use.) If you do not want to run your front-end in this way, you need to ensure that R_HOME is set and LD_LIBRARY_PATH is suitable. (The latter might well be, but modern Unix/Linux systems do not normally include /usr/local/lib(/usr/local/lib64 on some architectures), and R does look there for system components.)

The other senses in which this example is too simple are that all the internal defaults are used and that control is handed over to the R main loop. There are a number of small examples in the tests/Embedding directory. These make use of Rf_initEmbeddedR in src/main/Rembedded.c, and essentially use

```
#include <Rembedded.h>
int main(int ac, char **av)
    /* do some setup */
    Rf_initEmbeddedR(argc, argv);
    /* do some more setup */
    /* submit some code to R, which is done interactively via
run_Rmainloop();
A possible substitute for a pseudo-console is
R_ReplDLLinit();
while(R_ReplDLLdo1() > 0) {
/* add user actions here if desired */
     */
    Rf_endEmbeddedR(0);
    /* final tidying up after R is shutdown */
    return 0;
}
```

If you do not want to pass R arguments, you can fake an argy array, for example by

```
char *argv[]= {"REmbeddedPostgres", "--silent"};
Rf_initEmbeddedR(sizeof(argv)/sizeof(argv[0]), argv);
```

However, to make a GUI we usually do want to run run_Rmainloop after setting up various parts of R to talk to our GUI, and arranging for our GUI callbacks to be called during the R mainloop.

One issue to watch is that on some platforms Rf_initEmbeddedR and Rf_endEmbeddedR change the settings of the FPU (e.g. to allow errors to be trapped and to make use of extended precision registers).

The standard code sets up a session temporary directory in the usual way, unless R_TempDir is set to a non-NULL value before Rf_initEmbeddedR is called. In that case the value is assumed to contain an existing writable directory (no check is done), and it is not cleaned up when R is shut down.

 $Rf_initEmbeddedR$ sets R to be in interactive mode: you can set $R_Interactive$ (defined in Rinterface.h) subsequently to change this.

² but these are not part of the automated test procedures and so little tested.

Note that R expects to be run with the locale category 'LC_NUMERIC' set to its default value of C, and so should not be embedded into an application which changes that.

It is the user's responsibility to attempt to initialize only once. To protect the R interpreter, Rf_initialize_R will exit the process if re-initialization is attempted.

8.1.1 Compiling against the R library

Suitable flags to compile and link against the R (shared or static) library can be found by

```
R CMD config --cppflags
R CMD config --ldflags
```

(These apply only to an uninstalled copy or a standard install.)

If R is installed, pkg-config is available and neither sub-architectures nor an OS X framework have been used, alternatives for a shared R library are

```
pkg-config --cflags libR
pkg-config --libs libR
and for a static R library
pkg-config --cflags libR
pkg-config --libs --static libR
```

(This may work for an installed OS framework if pkg-config is taught where to look for libR.pc: it is installed inside the framework.)

However, a more comprehensive way is to set up a Makefile to compile the front-end. Suppose file myfe.c is to be compiled to myfe. A suitable Makefile might be

R CMD make

R CMD myfe

Additional flags which \$(MAIN_LINK) includes are, amongst others, those to select OpenMP and --export-dynamic for the GNU linker on some platforms. In principle \$(LIBS) is not needed when using a shared R library as libR is linked against those libraries, but some platforms need the executable also linked against them.

8.1.2 Setting R callbacks

For Unix-alikes there is a public header file Rinterface.h that makes it possible to change the standard callbacks used by R in a documented way. This defines pointers (if R_INTERFACE_PTRS is defined)

```
extern void (*ptr_R_Suicide)(const char *);
extern void (*ptr_R_ShowMessage)(const char *);
extern int (*ptr_R_ReadConsole)(const char *, unsigned char *, int, int);
extern void (*ptr_R_WriteConsole)(const char *, int);
extern void (*ptr_R_WriteConsoleEx)(const char *, int, int);
```

```
extern void (*ptr_R_ResetConsole)();
extern void (*ptr_R_FlushConsole)();
extern void (*ptr_R_ClearerrConsole)();
extern void (*ptr_R_Busy)(int);
extern void (*ptr_R_CleanUp)(SA_TYPE, int, int);
extern int (*ptr_R_ShowFiles)(int, const char **, const char **,
                              const char *, Rboolean, const char *);
extern int (*ptr_R_ChooseFile)(int, char *, int);
extern int (*ptr_R_EditFile)(const char *);
extern void (*ptr_R_loadhistory)(SEXP, SEXP, SEXP);
extern void (*ptr_R_savehistory)(SEXP, SEXP, SEXP);
extern void (*ptr_R_addhistory)(SEXP, SEXP, SEXP);
// added in R 3.0.0
extern int (*ptr_R_EditFiles)(int, const char **, const char **, const char *);
extern SEXP (*ptr_do_selectlist)(SEXP, SEXP, SEXP);
extern SEXP (*ptr_do_dataentry)(SEXP, SEXP, SEXP);
extern SEXP (*ptr_do_dataviewer)(SEXP, SEXP, SEXP);
extern void (*ptr_R_ProcessEvents)();
```

which allow standard R callbacks to be redirected to your GUI. What these do is generally documented in the file src/unix/system.txt.

void R_ShowMessage (char *message)

[Function]

This should display the message, which may have multiple lines: it should be brought to the user's attention immediately.

void R_Busy (int which)

[Function]

This function invokes actions (such as change of cursor) when R embarks on an extended computation (which=1) and when such a state terminates (which=0).

These functions interact with a console.

R_ReadConsole prints the given prompt at the console and then does a fgets(3)-like operation, transferring up to buffer characters into the buffer buf. The last two bytes should be set to '" \n ' to preserve sanity. If hist is non-zero, then the line should be added to any command history which is being maintained. The return value is 0 is no input is available and >0 otherwise.

R_WriteConsoleEx writes the given buffer to the console, *otype* specifies the output type (regular output or warning/error). Call to R_WriteConsole(buf, buflen) is equivalent to R_WriteConsoleEx(buf, buflen, 0). To ensure backward compatibility of the callbacks, ptr_R_WriteConsoleEx is used only if ptr_R_WriteConsole is set to NULL. To ensure that stdout() and stderr() connections point to the console, set the corresponding files to NULL *via*

```
R_Outputfile = NULL;
R_Consolefile = NULL;
```

R_ResetConsole is called when the system is reset after an error. R_FlushConsole is called to flush any pending output to the system console. R_ClearerrConsole clears any errors associated with reading from the console.

int R_ShowFiles (int nfile, const char **file, const char **headers, const char *wtitle, Rboolean del, const char *pager)

This function is used to display the contents of files.

```
int R_ChooseFile (int new, char *buf, int len)
```

[Function]

Choose a file and return its name in buf of length len. Return value is 0 for success, > 0 otherwise.

```
int R_EditFile (const char *buf)
```

[Function]

Send a file to an editor window.

Send *nfile* files to an editor, with titles possibly to be used for the editor window(s).

```
SEXP R_loadhistory (SEXP, SEXP, SEXP); [Function]
SEXP R_savehistory (SEXP, SEXP, SEXP, SEXP); [Function]
SEXP R_addhistory (SEXP, SEXP, SEXP, SEXP); [Function]
```

.Internal functions for loadhistory, savehistory and timestamp.

If the console has no history mechanism these can be as simple as

```
SEXP R_loadhistory (SEXP call, SEXP op, SEXP args, SEXP env)
{
    errorcall(call, "loadhistory is not implemented");
    return R_NilValue;
}
SEXP R_savehistory (SEXP call, SEXP op , SEXP args, SEXP env)
{
    errorcall(call, "savehistory is not implemented");
    return R_NilValue;
}
SEXP R_addhistory (SEXP call, SEXP op , SEXP args, SEXP env)
{
    return R_NilValue;
}
```

The R_addhistory function should return silently if no history mechanism is present, as a user may be calling timestamp purely to write the time stamp to the console.

```
void R_Suicide (const char *message)
```

[Function]

This should abort R as rapidly as possible, displaying the message. A possible implementation is

```
void R_Suicide (const char *message)
{
    char pp[1024];
    snprintf(pp, 1024, "Fatal error: %s\n", s);
    R_ShowMessage(pp);
    R_CleanUp(SA_SUICIDE, 2, 0);
}
```

void R_CleanUp (SA_TYPE saveact, int status, int RunLast) [Function] This function invokes any actions which occur at system termination. It needs to be quite complex:

```
#include <Rinterface.h>
#include <Rembedded.h>
                          /* for Rf_KillAllDevices */
void R_CleanUp (SA_TYPE saveact, int status, int RunLast)
    if(saveact == SA_DEFAULT) saveact = SaveAction;
    if(saveact == SA_SAVEASK) {
       /* ask what to do and set saveact */
    switch (saveact) {
    case SA_SAVE:
        if(runLast) R_dot_Last();
        if(R_DirtyImage) R_SaveGlobalEnv();
        /* save the console history in R_HistoryFile */
        break;
    case SA_NOSAVE:
       if(runLast) R_dot_Last();
        break;
    case SA_SUICIDE:
   default:
        break;
   }
   R_RunExitFinalizers();
    /* clean up after the editor e.g. CleanEd() */
   R_CleanTempDir();
   /* close all the graphics devices */
    if(saveact != SA_SUICIDE) Rf_KillAllDevices();
    fpu_setup(FALSE);
   exit(status);
}
```

These callbacks should never be changed in a running R session (and hence cannot be called from an extension package).

```
SEXP R_dataentry (SEXP, SEXP, SEXP, SEXP); [Function]
SEXP R_dataviewer (SEXP, SEXP, SEXP, SEXP); [Function]
SEXP R_selectlist (SEXP, SEXP, SEXP, SEXP); [Function]
External functions for dataentry (and edit on matrices and data frames). View and
```

.External functions for dataentry (and edit on matrices and data frames), View and select.list. These can be changed if they are not currently in use.

8.1.3 Registering symbols

An application embedding R needs a different way of registering symbols because it is not a dynamic library loaded by R as would be the case with a package. Therefore R reserves a special DllInfo entry for the embedding application such that it can register symbols to be

used with .C, .Call etc. This entry can be obtained by calling getEmbeddingDllInfo, so a typical use is

```
DllInfo *info = R_getEmbeddingDllInfo();
R_registerRoutines(info, cMethods, callMethods, NULL, NULL);
```

The native routines defined by cMethods and callMethods should be present in the embedding application. See Section 5.4 [Registering native routines], page 102, for details on registering symbols in general.

8.1.4 Meshing event loops

One of the most difficult issues in interfacing R to a front-end is the handling of event loops, at least if a single thread is used. R uses events and timers for

- Running X11 windows such as the graphics device and data editor, and interacting with them (e.g., using locator()).
- Supporting Tcl/Tk events for the **tcltk** package (for at least the X11 version of Tk).
- Preparing input.
- Timing operations, for example for profiling R code and Sys.sleep().
- Interrupts, where permitted.

Specifically, the Unix-alike command-line version of R runs separate event loops for

- Preparing input at the console command-line, in file src/unix/sys-unix.c.
- Waiting for a response from a socket in the internal functions underlying FTP and HTTP transfers in download.file() and for direct socket access, in files src/modules/internet/nanoftp.c, src/modules/internet/nanohttp.c and src/modules/internet/Rsock.c
- Mouse and window events when displaying the X11-based dataentry window, in file src/modules/X11/dataentry.c. This is regarded as modal, and no other events are serviced whilst it is active.

There is a protocol for adding event handlers to the first two types of event loops, using types and functions declared in the header R_ext/eventloop.h and described in comments in file src/unix/sys-std.c. It is possible to add (or remove) an input handler for events on a particular file descriptor, or to set a polling interval (via R_wait_usec) and a function to be called periodically via R_PolledEvents: the polling mechanism is used by the tcltk package.

It is not intended that these facilities are used by packages, but if they are needed exceptionally, the package should ensure that it cleans up and removes its handlers when its namespace is unloaded.

An alternative front-end needs both to make provision for other R events whilst waiting for input, and to ensure that it is not frozen out during events of the second type. This is not handled very well in the existing examples. The GNOME front-end ran a private handler for polled events by setting

whilst it is waiting for console input. This obviously handles events for Gtk windows (such as the graphics device in the **gtkDevice** package), but not X11 events (such as the X11() device) or for other event handlers that might have been registered with R. It does not attempt to keep itself alive whilst R is waiting on sockets. The ability to add a polled handler as R_timeout_handler is used by the **tcltk** package.

8.1.5 Threading issues

Embedded R is designed to be run in the main thread, and all the testing is done in that context. There is a potential issue with the stack-checking mechanism where threads are involved. This uses two variables declared in Rinterface.h (if CSTACK_DEFNS is defined) as

```
extern uintptr_t R_CStackLimit; /* C stack limit */
extern uintptr_t R_CStackStart; /* Initial stack address */
```

Note that uintptr_t is a C99 type for which a substitute is defined in R, so your code needs to define HAVE_UINTPTR_T appropriately.

These will be set³ when Rf_initialize_R is called, to values appropriate to the main thread. Stack-checking can be disabled by setting R_CStackLimit = (uintptr_t)-1 immediately after Rf_initialize_R is called, but it is better to if possible set appropriate values. (What these are and how to determine them are OS-specific, and the stack size limit may differ for secondary threads. If you have a choice of stack size, at least 10Mb is recommended.)

You may also want to consider how signals are handled: R sets signal handlers for several signals, including SIGINT, SIGSEGV, SIGPIPE, SIGUSR1 and SIGUSR2, but these can all be suppressed by setting the variable R_SignalHandlers (declared in Rinterface.h) to 0.

Note that these variables must not be changed by an R **package**: a package should not calling R internals which makes use of the stack-checking mechanism on a secondary thread.

8.2 Embedding R under Windows

All Windows interfaces to R call entry points in the DLL R.dll, directly or indirectly. Simpler applications may find it easier to use the indirect route *via* (D)COM.

8.2.1 Using (D)COM

(D)COM is a standard Windows mechanism used for communication between Windows applications. One application (here R) is run as COM server which offers services to clients, here the front-end calling application. The services are described in a 'Type Library' and are (more or less) language-independent, so the calling application can be written in C or C++ or Visual Basic or Perl or Python and so on. The 'D' in (D)COM refers to 'distributed', as the client and server can be running on different machines.

The basic R distribution is not a (D)COM server, but two addons are currently available that interface directly with R and provide a (D)COM server:

• There is a (D)COM server called StatConnector written by Thomas Baier available *via* http://sunsite.univie.ac.at/rcom/, which works with R packages to support transfer of data to and from R and remote execution of R commands, as well as embedding of an R graphics window.

Recent versions have usage restrictions.

• Another (D)COM server, RDCOMServer, may be available from Omegahat, http://www.omegahat.net/. Its philosophy is discussed in http://www.omegahat.net/RDCOMServer/Docs/Paradigm.html and is very different from the purpose of this section.

8.2.2 Calling R.dll directly

The R DLL is mainly written in C and has _cdecl entry points. Calling it directly will be tricky except from C code (or C++ with a little care).

There is a version of the Unix-alike interface calling

```
int Rf_initEmbeddedR(int ac, char **av);
```

³ at least on platforms where the values are available, that is having getrlimit and on Linux or having sysctl supporting KERN_USRSTACK, including FreeBSD and OS X.

```
void Rf_endEmbeddedR(int fatal);
```

which is an entry point in R.dll. Examples of its use (and a suitable Makefile.win) can be found in the tests/Embedding directory of the sources. You may need to ensure that R_HOME/bin is in your PATH so the R DLLs are found.

Examples of calling R.dll directly are provided in the directory src/gnuwin32/front-ends, including a simple command-line front end rtest.c whose code is

```
#define Win32
#include <windows.h>
#include <stdio.h>
#include <Rversion.h>
#define LibExtern __declspec(dllimport) extern
#include <Rembedded.h>
#include <R_ext/RStartup.h>
/* for askok and askyesnocancel */
#include <graphapp.h>
/* for signal-handling code */
#include <psignal.h>
/* simple input, simple output */
/* This version blocks all events: a real one needs to call ProcessEvents
   frequently. See rterm.c and ../system.c for one approach using
   a separate thread for input.
int myReadConsole(const char *prompt, char *buf, int len, int addtohistory)
    fputs(prompt, stdout);
    fflush(stdout);
    if(fgets(buf, len, stdin)) return 1; else return 0;
void myWriteConsole(const char *buf, int len)
    printf("%s", buf);
}
void myCallBack(void)
    /* called during i/o, eval, graphics in ProcessEvents */
void myBusy(int which)
    /* set a busy cursor ... if which = 1, unset if which = 0 */
static void my_onintr(int sig) { UserBreak = 1; }
int main (int argc, char **argv)
    structRstart rp;
    Rstart Rp = &rp;
    char Rversion[25], *RHome;
    sprintf(Rversion, "%s.%s", R_MAJOR, R_MINOR);
    if(strcmp(getDLLVersion(), Rversion) != 0) {
        fprintf(stderr, "Error: R.DLL version does not match\n");
        exit(1);
    R_setStartTime();
    R_DefParams(Rp);
```

```
if((RHome = get_R_HOME()) == NULL) {
        fprintf(stderr, "R_HOME must be set in the environment or Registry\n");
        exit(1);
    }
    Rp->rhome = RHome;
    Rp->home = getRUser();
    Rp->CharacterMode = LinkDLL;
    Rp->ReadConsole = myReadConsole;
    Rp->WriteConsole = myWriteConsole;
    Rp->CallBack = myCallBack;
    Rp->ShowMessage = askok;
    Rp->YesNoCancel = askyesnocancel;
    Rp->Busy = myBusy;
                               /* Default is FALSE */
    Rp->R_Quiet = TRUE;
    Rp->R_Interactive = FALSE; /* Default is TRUE */
    Rp->RestoreAction = SA_RESTORE;
    Rp->SaveAction = SA_NOSAVE;
    R_SetParams(Rp);
    R_set_command_line_arguments(argc, argv);
    FlushConsoleInputBuffer(GetStdHandle(STD_INPUT_HANDLE));
    signal(SIGBREAK, my_onintr);
    GA_initapp(0, 0);
    readconsolecfg();
    setup_Rmainloop();
#ifdef SIMPLE_CASE
    run_Rmainloop();
#else
    R_ReplDLLinit();
    while(R_ReplDLLdo1() > 0) {
/* add user actions here if desired */
/* only get here on EOF (not q()) */
#endif
    Rf_endEmbeddedR(0);
    return 0:
```

The ideas are

- Check that the front-end and the linked R.dll match other front-ends may allow a looser match.
- Find and set the R home directory and the user's home directory. The former may be available from the Windows Registry: it will be in HKEY_LOCAL_MACHINE\Software\R-core\R\InstallPath from an administrative install and HKEY_CURRENT_USER\Software\R-core\R\InstallPath otherwise, if selected during installation (as it is by default).
- Define startup conditions and callbacks *via* the Rstart structure. R_DefParams sets the defaults, and R_SetParams sets updated values.
- Record the command-line arguments used by R_set_command_line_arguments for use by the R function commandArgs().
- Set up the signal handler and the basic user interface.
- Run the main R loop, possibly with our actions intermeshed.
- Arrange to clean up.

An underlying theme is the need to keep the GUI 'alive', and this has not been done in this example. The R callback R_ProcessEvents needs to be called frequently to ensure that Windows events in R windows are handled expeditiously. Conversely, R needs to allow the GUI

code (which is running in the same process) to update itself as needed – two ways are provided to allow this:

• R_ProcessEvents calls the callback registered by Rp->callback. A version of this is used to run package Tcl/Tk for tcltk under Windows, for the code is

```
void R_ProcessEvents(void)
{
    while (peekevent()) doevent(); /* Windows events for GraphApp */
    if (UserBreak) { UserBreak = FALSE; onintr(); }
    R_CallBackHook();
    if(R_tcldo) R_tcldo();
}
```

• The mainloop can be split up to allow the calling application to take some action after each line of input has been dealt with: see the alternative code below #ifdef SIMPLE_CASE.

It may be that no R GraphApp windows need to be considered, although these include pagers, the windows() graphics device, the R data and script editors and various popups such as choose.file() and select.list(). It would be possible to replace all of these, but it seems easier to allow GraphApp to handle most of them.

It is possible to run R in a GUI in a single thread (as RGui.exe shows) but it will normally be easier⁴ to use multiple threads.

Note that R's own front ends use a stack size of 10Mb, whereas MinGW executables default to 2Mb, and Visual C++ ones to 1Mb. The latter stack sizes are too small for a number of R applications, so general-purpose front-ends should use a larger stack size.

8.2.3 Finding R_HOME

Both applications which embed R and those which use a system call to invoke R (as Rscript.exe, Rterm.exe or R.exe) need to be able to find the R bin directory. The simplest way to do so is the ask the user to set an environment variable R_HOME and use that, but naive users may be flummoxed as to how to do so or what value to use.

The R for Windows installers have for a long time allowed the value of R_HOME to be recorded in the Windows Registry: this is optional but selected by default. Where it is recorded has changed over the years to allow for multiple versions of R to be installed at once, and to allow 32- and 64-bit versions of R to be installed on the same machine.

The basic Registry location is Software\R-core\R. For an administrative install this is under HKEY_LOCAL_MACHINE and on a 64-bit OS HKEY_LOCAL_MACHINE\Software\R-core\R is by default redirected for a 32-bit application, so a 32-bit application will see the information for the last 32-bit install, and a 64-bit application that for the last 64-bit install. For a personal install, the information is under HKEY_CURRENT_USER\Software\R-core\R which is seen by both 32-bit and 64-bit applications and so records the last install of either architecture. To circumvent this, there are locations Software\R-core\R32 and Software\R-core\R64 which always refer to one architecture.

When R is installed and recording is not disabled then two string values are written at that location for keys InstallPath and Current Version, and these keys are removed when R is uninstalled. To allow information about other installed versions to be retained, there is also a key named something like 3.0.0 or 3.0.0 patched or 3.1.0 Pre-release with a value for InstallPath.

So a comprehensive algorithm to search for R_HOME is something like

⁴ An attempt to use only threads in the late 1990s failed to work correctly under Windows 95, the predominant version of Windows at that time.

- Decide which of personal or administrative installs should have precedence. There are arguments both ways: we find that with roaming profiles that HKEY_CURRENT_USER\Software often gets reverted to an earlier version. Do the following for one or both of HKEY_CURRENT_USER and HKEY_LOCAL_MACHINE.
- If the desired architecture is known, look in Software\R-core\R32 or Software\R-core\R64, and if that does not exist or the architecture is immaterial, in Software\R-core\R.
- If key InstallPath exists then this is R_HOME (recorded using backslashes). If it does not, look for version-specific keys like 2.11.0 alpha, pick the latest (which is of itself a complicated algorithm as 2.11.0 patched > 2.11.0 > 2.11.0 alpha > 2.8.1) and use its value for InstallPath.

Prior to R 2.12.0 R.dll and the various front-end executables were in R_HOME\bin, but they are now in R_HOME\bin\i386 or R_HOME\bin\x64. So you may need to arrange to look first in the architecture-specific subdirectory and then in R_HOME\bin.

Function and variable index

~	\packageMaintainer 73
.C99	\packageTitle
.Call	\pkg 67
.External	\preformatted66
.Fortran99	\RdOpts72
.Last.lib41	\references 62
.onAttach	\renewcommand
.onDetach	\R70
.onLoad	\S3method61
.onUnload41	\samp
.Random.seed	\section65
	\seealso
1	\Sexpr
	\source
	\sQuote
\acronym	\strong
\alias 59	-
\arguments 61	\tabular
\author	\title
\bold66	\url
\cite67	\usage
\code66	\value 61
\command	\var 67
\concept	\verb
\cr	
\deqn69	A
\describe	\mathbf{A}
\description	allocVector
\details	AUTHORS
\dfn	
\dontrun	
\dontshow	В
\dots	bessel_i141
\dQuote	
\email	bessel_j
\emph	bessel_k
\enc	bessel_y141
\enumerate	beta
\env	BLAS_LIBS
·	browser
\eqn	
\examples	
	\mathbf{C}
\figure	\mathbf{C}
\file	C Calloc
\file	_
\file	Calloc
\file 67 \format 64 \href 67 \if 71	Calloc 13 CAR 12 CDR 12
\file 67 \format 64 \href 67 \if 71 \ifelse 71	Calloc 13 CAR 12 CDR 12 cgmin 144
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15,56
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140
\file 67 \format 64 \href 67 \if 71 \if 87 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15,56
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63	Calloc 138 CAR 121 CDR 122 cgmin 144 choose 146 CITATION 15, 56 COPYRIGHTS 6, 18
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15, 56 COPYRIGHTS 6, 18
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68	Calloc 138 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15, 50 COPYRIGHTS 6, 15 D debug 8
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60	Calloc 13 CAR 121 CDR 122 cgmin 144 choose 146 CITATION 15, 56 COPYRIGHTS 6, 15 D debug 87 debugger 86
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60 \name 59	Calloc 133 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15, 50 COPYRIGHTS 6, 19 D debug 8 debugger 8 defineVar 118
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60 \name 59 \newcommand 73	Calloc 13 CAR 121 CDR 122 cgmin 144 choose 146 CITATION 15,56 COPYRIGHTS 6,18 D debug 87 debugger 86 defineVar 118 digamma 140
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60 \name 59 \newcommand 73 \note 62	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15,56 COPYRIGHTS 6,18 D debug 8 debugger 86 defineVar 118 digamma 140 dump.frames 86
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60 \name 59 \newcommand 73 \note 62 \option 67	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15,50 COPYRIGHTS 6,15 D debug 8 debugger 8 defineVar 11 digamma 14 dump.frames 8 duplicate 115
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60 \name 59 \newcommand 73 \note 62 \option 67 \out 71	Calloc 133 CAR 121 CDR 122 cgmin 144 choose 140 CITATION 15,56 COPYRIGHTS 6,18 D debug 87 debugger 86 defineVar 118 digamma 140 dump.frames 86 duplicate 119 dyn.load 101
\file 67 \format 64 \href 67 \if 71 \ifelse 71 \itemize 68 \kbd 66 \keyword 63 \ldots 70 \link 68 \method 60 \name 59 \newcommand 73 \note 62 \option 67 \out 71 \packageAuthor 73	Calloc 13 CAR 12 CDR 12 cgmin 14 choose 140 CITATION 15,56 COPYRIGHTS 6,18 D debug 87 debugger 86 defineVar 118 digamma 140 dump.frames 86 duplicate 115

\mathbf{E}	P
exp_rand	pentagamma
export	PKG_CFLAGS 106
exportClasses	PKG_CPPFLAGS
exportClassPattern	PKG_CXXFLAGS
exportMethods 44 exportPattern 39, 44	PKG_FCFLAGS
exportrattern	PKG_FFLAGS
	PKG_LIBS
\mathbf{F}	PKG_OBJCFLAGS 106 PKG_OBJCXXFLAGS 106
FALSE	prompt
findVar	PROTECT
FLIBS	PROTECT_WITH_INDEX
Free	psigamma
	PutRNGstate
\mathbf{G}	
gammafn	\mathbf{R}
gctorture	R CMD build
getAttrib	R CMD build
getCharCE	R CMD config
GetRNGstate	R CMD Rd2pdf
	R CMD Rdconv
I	R CMD SHLIB
import	R CMD Stangle74
importClassesFrom	R CMD Sweave74
importFrom 40	R_alloc
importMethodsFrom	R_allocLD
install	R_FINITE
ISNA	R_GetCCallable 105 R_GetCurrentSrcref 130
ISNAN 123, 137	R_GetSrcFilename
	R_INLINE
L	R_IsNaN
LAPACK_LIBS	R_LIBRARY_DIR
lbeta	R_NegInf
lbfgsb	R_PACKAGE_DIR
1choose	R_PACKAGE_NAME
lgammafn	R_ParseVector
library.dynam 12, 101	R_PosInf 137 R_PreserveObject 113
	R_RegisterCCallable
M	R_registerRoutines
	R_ReleaseObject113
M_E	R_Srcref
M_PI	R_Version
mkCharCE	Rdqagi
mkCharLen	Rdqags
mkCharLenCE	Realloc
	recover
N	REprintf
	REPROTECT
NA_REAL	REvprintf
NEWS.Rd	Rprintf
nmmin 143 norm_rand 136	Rprof
потш_тани	Rprofmem
	Rvprintf
O	
OBJECTS 21, 106	

\mathbf{S}	translateCharUTF8
S_alloc	trigamma
S_realloc	TRUE
S3method	
SAFE_FFLAGS	
samin	
seed_in	\mathbf{U}
seed_out	
setAttrib	undebug
setVar	unif_rand
summaryRprof	UNPROTECT
system	UNPROTECT_PTR
system.time	untracemem 79
system2	useDynLib41
\mathbf{T}	
tetragamma	\mathbf{V}
trace	·
traceback85	vmaxget
tracemem	vmaxset
translatoChar 129	1/19

Concept index

•	\mathbf{F}
.install_extras file37.Rbuildignore file33.Rinstignore file14	Figures in documentation 69 finalizer 130 Finding variables 118
	\mathbf{G}
\linkS4class	
	Gamma function140Garbage collection112Generic functions152
\mathbf{A}	Generic functions
Allocating storage	Н
	handling character data
B	Handling lists 117 Handling R objects in C 111
Bessel functions	
Beta function	т
Building binary packages	I
	IEEE special values 123, 137 INDEX file 12
\mathbf{C}	Indices
	Inspecting R objects when debugging 97
C++ code, interfacing	integration
Checking packages	Interfacing C++ code
citation	Interrupts
Classes 117 cleanup file 3	
conditionals	$\mathbf L$
configure file	LICENCE file
Copying objects 119 CRAN 3	LICENSE file
Creating packages	Lists and tables in documentation
Creating shared objects	
Cross-references in documentation	${f M}$
cumulative nazard	Marking text in documentation 66
D	Mathematics in documentation
D	Memory allocation from C
Debugging 95 DESCRIPTION file 4	Method functions
Details of R types	Missing values
Distribution functions from C	
Documentation, writing	${f N}$
dynamic pages	namespaces
	news
E	Numerical analysis subroutines from C
	rumerical derivatives
Editing Rd files 75 encoding 74	0
Error handling from C	O
Error handling from FORTRAN	OpenMP 23, 148 Operating system access 99
Evaluating R expressions from C	optimization

Concept index 170

P	${f T}$
Package builder33Package structure3Package subdirectories12	tarballs
Packages 2 Parsing R code from C 129 Platform-specific documentation 71 Printing from C 137 Printing from FORTRAN 137	U user-defined macros
Processing Rd format 74 Profiling 76, 78, 79	\mathbf{V}
R Random numbers in C	Version information from C. 148 vignettes. 35 Visibility. 149
Random numbers in FORTRAN	\mathbf{W}
\mathbf{S}	weak reference
Setting variables 118 Sort functions from C 145	${f Z}$
Sweave	Zero-finding